

CMPE 110: Computer Architecture

Week 2

Performance

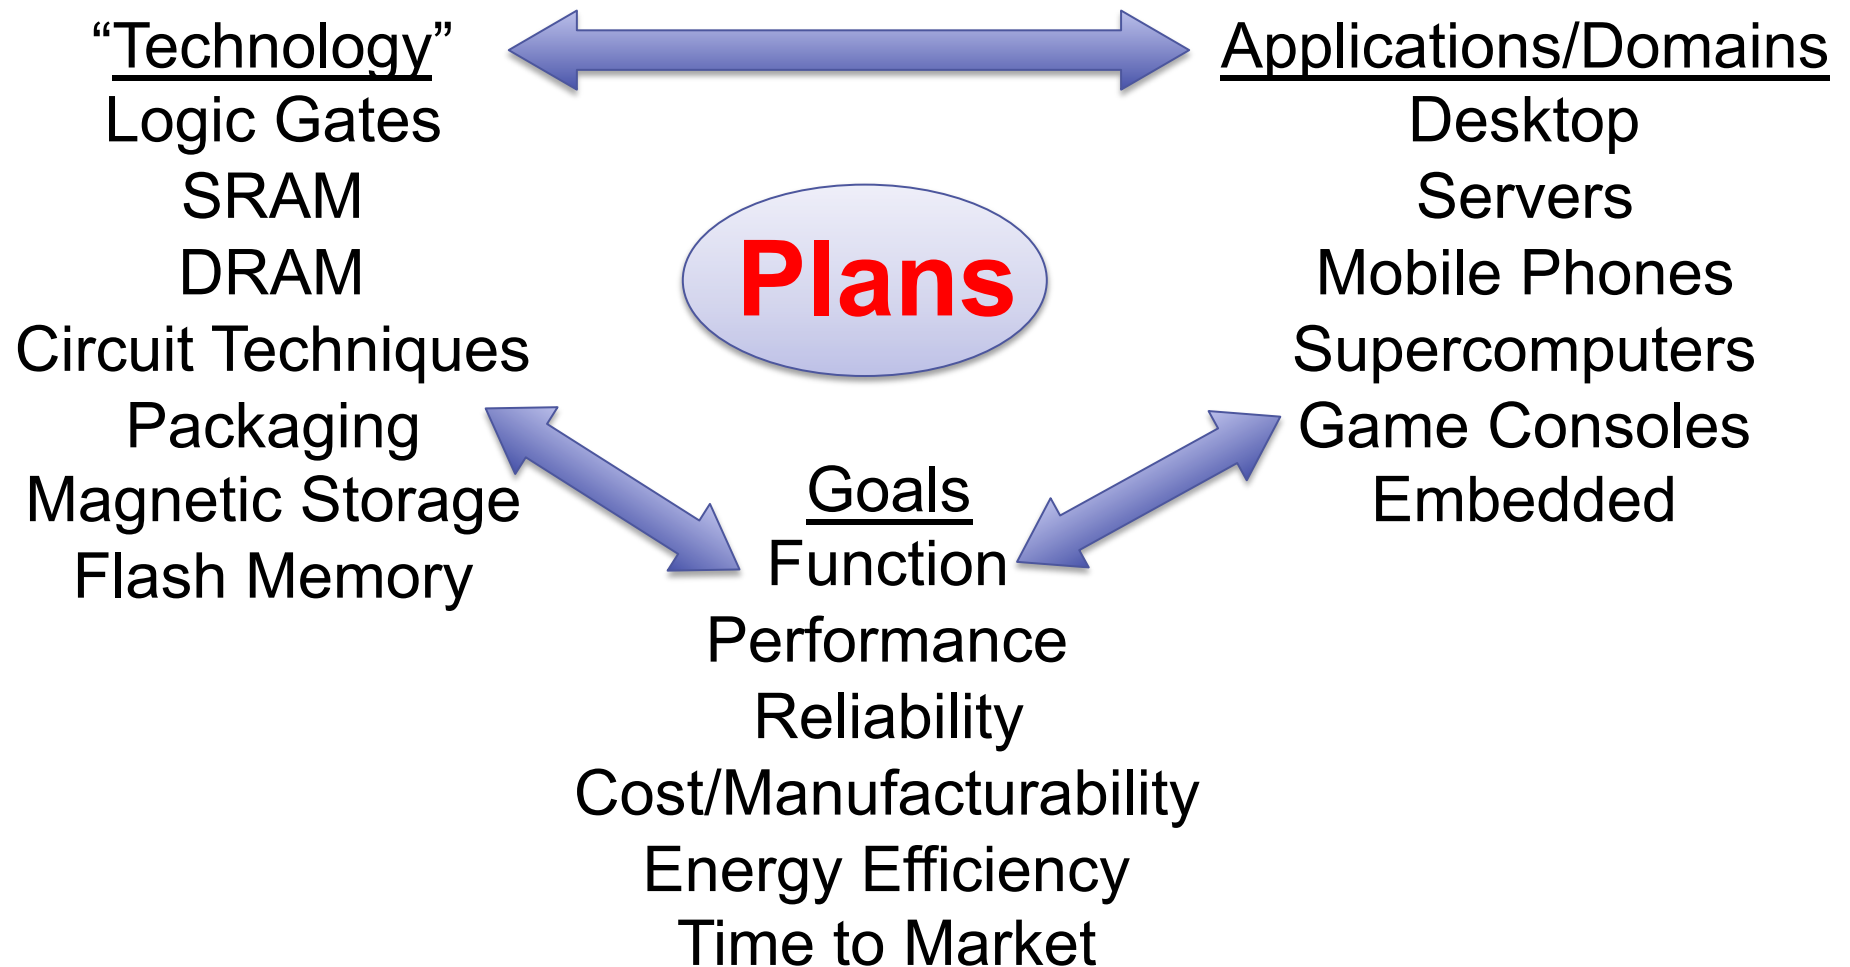
Jishen Zhao (<http://users.soe.ucsc.edu/~jzhao/>)

[Adapted in part from Jose Renau, Mary Jane Irwin, Joe Devietti, Onur Mutlu, and others]

Reminders

- TA office hours
 - Rebecca Rashkin: Mon 4-5pm, E2-480
 - Xin Li: TBD
 - Narendra Kumar Govinda Raju: Mon 5:30-6:30pm, BE-312B
 - Aziz (Abdulazaz) Albalawi: Thu 5-6pm, BE-312B
- Discussion sessions
 - Section 1: Tue 08:30 - 09:35am in Kresge Clrm 327, TA: Aziz
 - Section 2: Wed 04:00 - 05:05pm in Kresge Clrm 327, TA: Rebecca
 - Section 3: Wed 05:20 - 06:25pm in Kresge Clrm 327, TA: Narendra
 - Section 4: TBA, TA: Xin

Review: What is computer architecture?



Today: Performance

- Metrics
 - Latency and throughput
 - Speedup
 - Averaging
- CPU Performance



Performance Metrics

Performance: Latency vs. Throughput

- **Latency (execution time)**: time to finish a fixed task
- **Throughput (bandwidth)**: number of tasks per unit time
 - Different: exploit parallelism for throughput, not latency



- Choose definition of performance that matches your goals
 - Scientific program? latency.
 - Web server? throughput.

Examples

- How to measure the performance of moving people for 10 miles **round trip**?
 - Car: capacity = 5, speed = 60 miles/hour
 - Bus: capacity = 60, speed = 20 miles/hour
 - **Latency**: how long does ~ take to move one person for 20 miles?
 - **Throughput**: how many people can ~ move per hour?



Answer:

- **Latency**: **car = 20 min**, bus = 60 min
- **Throughput**: car = 15 PPH, **bus = 60 PPH**

Examples

- Fastest way to send 10TB of data from US to UK?
FTP, SMB, Rsync / Robocopy, other?

FedEx

Used FedEx overnight to deliver the drive

Overnight
EXPRESS Shipping

Even 1 Gbps data transfer takes days!



Amazon Does This...

Available Internet Connection	Theoretical Min. Number of Days to Transfer 1TB at 80% Network Utilization	When to Consider AWS Import/Export?
T1 (1.544Mbps)	82 days	100GB or more
10Mbps	13 days	600GB or more
T3 (44.736Mbps)	3 days	2TB or more
100Mbps	1 to 2 days	5TB or more
1000Mbps	Less than 1 day	60TB or more



[Amazon Web Services](#) » [AWS Import/Export](#) » AWS Import/Export Calculator

Operation Type		Import to S3 ▼
Location	AWS Region	US Standard Region ▼
AWS Import/Export Data Load	Total Terabytes to Load	1 TB
	Number of Devices	1
	Wipe Device After Import	No ▼
Estimated Transfer Speed	Average File Size*	1 MB
	Interface Type	eSATA ▼
	Transfer Speed**	22.51 MB/sec

What we learned

Measuring performance

Latency & throughput



Comparing Performance - Speedup

- Speedup of A over B
 - $X = \text{Latency}(B) / \text{Latency}(A)$ (divide by the faster)
 - $X = \text{Throughput}(A) / \text{Throughput}(B)$ (divide by the slower)
- A is X% faster than B if
 - $X = ((\text{Latency}(B) / \text{Latency}(A)) - 1) * 100$
 - $X = ((\text{Throughput}(A) / \text{Throughput}(B)) - 1) * 100$
 - $\text{Latency}(A) = \text{Latency}(B) / (1 + (X/100))$
 - $\text{Throughput}(A) = \text{Throughput}(B) * (1 + (X/100))$
- Car/bus example
 - Latency?
 - Throughput?
 - See next slide...

Car/bus example

- **Latency:** car = **20 min**, bus = 60 min
- **Throughput:** car = 15 PPH, **bus = 60 PPH**

Speedup?

- Latency:
 - Speedup of car over bus is 3
 - Car is 200% faster than bus
- Throughput:
 - Speedup of bus over car is 4
 - Bus is 300% faster than car

Comparing Performance - Speedup

- Program A runs for 200 cycles
- Program B runs for 350 cycles

What is “cycle”?
Execution time * clock frequency
i.e., second * Hz

- Speedup of A over B?
 - Speedup = $350/200 = 1.75$
 - As a percentage: $(1.75 - 1) * 100 = 75\%$ (Program A runs 75% faster than program B)
- If program C is 50% faster than A, how many cycles does C run for?
 - 133 cycles



Note

- Speedup of A over B
 - $X = \text{Latency}(B) / \text{Latency}(A)$
 - $X = \text{Throughput}(A) / \text{Throughput}(B)$

What if $X < 1$?

-- means A is slower than B

Speedup and % Increase and Decrease

- Program A runs for 200 cycles
- Program B runs for 350 cycles
- Percent increase and decrease are **not the same**.
 - % increase of cycles: $((350 - 200)/200) * 100 = 75\%$
 - % decrease of cycles: $((350 - 200)/350) * 100 = 42.3\%$

What we learned

Comparing performance

Speedup

Performance metrics

Latency, throughput, speedup

Averaging performance

Mean (Average) Performance Numbers

- **Arithmetic:** $(1/N) * \sum_{P=1..N} \text{Latency}(P)$
 - For units that are proportional to time (e.g., latency)
- **Harmonic:** $N / \sum_{P=1..N} 1/\text{Throughput}(P)$
 - For units that are inversely proportional to time (e.g., throughput)
- You can add latencies, but not throughputs
 - $\text{Latency}(P1+P2, A) = \text{Latency}(P1, A) + \text{Latency}(P2, A)$
 - $\text{Throughput}(P1+P2, A) \neq \text{Throughput}(P1, A) + \text{Throughput}(P2, A)$
- **Geometric:** $N\sqrt[N]{\prod_{P=1..N} \text{Speedup}(P)}$
 - For unitless quantities (e.g., speedup ratios)

For Example...

1 mile @ 30 miles/hour + 1 mile @ 90 miles/hour

- You drive two miles
 - 30 miles per hour for the first mile
 - 90 miles per hour for the second mile
- Question: what was your average speed?
 - Hint: the answer is not 60 miles per hour
 - Why?



Answer: 45 miles/hour

- You drive two miles
 - 30 miles per hour for the first mile
 - 90 miles per hour for the second mile
- Question: what was your average speed?
 - Hint: the answer is not 60 miles per hour
 - 0.03333 hours per mile for 1 mile
 - 0.01111 hours per mile for 1 mile
 - 0.04444 hours for 2 miles
 - = 45 miles per hour
 - $\neq (30 + 90) / 2$

What we learned

Averaging performance

Arithmetic mean for latency

Harmonic mean for throughput

Geometric mean for speedup



CPU Performance

How to evaluate

Latency, throughput, and speedup

CPU Performance Equation

- Latency = seconds / program =
 - $(\text{insns} / \text{program}) * (\text{cycles} / \text{insn}) * (\text{seconds} / \text{cycle})$
 - **Insns / program**: insn count
 - Impacted by program, compiler, ISA
 - **Cycles / insn**: **CPI**
 - Impacted by program, compiler, ISA, **micro-arch**
 - **Seconds / cycle**: **clock period (Hz)**
 - Impacted by micro-arch, technology
- For low latency (better performance) minimize all three
 - Difficult: often pull against one another
 - Example we have seen: RISC vs. CISC ISAs
 - ± RISC: low CPI/clock period, high insn count
 - ± CISC: low insn count, high CPI/clock period