

CMPE 110: Computer Architecture

Week 2

ISA II

Jishen Zhao (<http://users.soe.ucsc.edu/~jzhao/>)

[Adapted in part from Jose Renau, Mary Jane Irwin, Joe Devietti, Onur Mutlu, and others]

Reminder

- Homework 1 will be posted on course Google site by 4:30pm today
 - Performance and ISA
 - Due by midnight Oct. 17
- Midterm 1
 - Questions reduced to four
 - Cover everything up until last Friday class

Q1	12	
Q2	12	
Q3	18	
Q4	18	
Total	60	



This exam is closed book and closed notes. Personal calculators (four-function calculators only) *are* allowed. Show your work on the attached sheets (front and back) and insert your answer in the space(s) provided. **Please provide details on how you reach a result.** Ask for extra paper sheets if necessary.

You have 70 minutes to complete the exam. This exam is worth 60 points. This exam counts for 15% of your course grade.

Reminder

- No office hour on Wednesday (Oct. 5)
- Xin's TA office hour announced
 - Time: Thu 9 - 10 am
 - Office: BE-312B

Review: ISA Overview



- What is an ISA?
 - An ISA includes a **specification** of the set of opcodes and the native commands implemented by a particular processor. <Wikipedia>
- Program execution model
 - **High-level programming language** (C/C++/Java/C#/...)
 - **Assembly language**
 - Human-readable representation
 - **Machine language**
 - 1s and 0s (often displayed in "hex")
- Instruction execution model

Review: An ADD Documentation

ADD[S] : Addition

ADD will add two values.

Operand 1 is a **register**, operand 2 can be a register, **shifted register**, or an immediate value (which may be **shifted**).

If the S bit is set (**ADDS**), the N and Z flags are set according to the result, and the C and V flags are set as follows: **C** if the result generated a carry (unsigned overflow); **V** if the result generated a signed overflow.

ADD is useful for basic addition. Use **ADC** to perform addition with the Carry flag considered.

Syntax

```
ADD<suffix> <dest>, <op 1>, <op 2>
```

Function

```
dest = op_1 + op_2
```

Technical

The instruction bit pattern is as follows:

31 - 28	27	26	25	24 - 21	20	19 - 16	15 - 12	11 - 0
condition	0	0	I	0 1 0 0	S	op_1	dest	op_2/shift

Note: If the I bit is zero, and bits 4 and 7 are both **one** (with bits 5,6 zero), the instruction is **UMULL**, not ADD.

Review: ISA vs. Performance

- Latency = seconds / program =
 - $(\text{insns} / \text{program}) * (\text{cycles} / \text{insn}) * (\text{seconds} / \text{cycle})$
 - Insns / program: **insn count**
 - Impacted by program, compiler, ISA
 - Cycles / insn: **CPI**
 - Impacted by program, compiler, ISA, micro-arch
 - Seconds / cycle: clock period (Hz)
 - Impacted by micro-arch, technology
- How does ISA affect insn count and CPI?
 - Example: CISC vs. RISC

CISC vs. RISC

- **CISC** (Complex Instruction Set Computing) **ISAs**
 - Examples: Intel/AMD x86, IBM System/360, Intel 8051
 - Big heavyweight instructions (lots of work per instruction)
 - + Low "insns/program"
 - Higher "cycles/insn" and "seconds/cycle"
- **RISC** (Reduced Instruction Set Computer) **ISAs**
 - Examples: MIPS, ARM, Alpha
 - Simple instructions only: minimalist approach to an ISA
 - + Low "cycles/insn" and "seconds/cycle"
 - Higher "insn/program", but hopefully not as much

Today: ISA / Pipelining

- ISA design goals
 - Programmability
 - Performance/implementability
 - Compatibility
- Aspects of ISAs
- Pipelining

ISA Design Goals

What Makes a Good ISA?

- **Programmability**
 - Easy to express programs efficiently?
- **Performance/Implementability**
 - Easy to design high-performance implementations?
 - Easy to design low-power implementations?
 - Easy to design low-cost implementations?
- **Compatibility**
 - Easy to maintain as languages, programs, and technology evolve?
 - x86 (IA32) generations: 8086, 286, 386, 486, Pentium, PentiumII, PentiumIII, Pentium4, Core2, Core i7, ...

Compatibility

- In many domains, ISA must remain compatible
 - IBM's 360/370 (the *first* "ISA family")
 - Another example: Intel's x86 and Microsoft Windows
 - x86 one of the worst designed ISAs EVER, but it survives
- **Backward compatibility**
 - New processors supporting old programs
 - **Hard to drop features**
 - Update software/OS to emulate dropped features (slow)
- **Forward (upward) compatibility**
 - Old processors supporting new programs
 - Include a "CPU ID" so the software can test for features
 - Add ISA hints by overloading no-ops (example: x86's PAUSE)
 - New firmware/software on old processors to emulate new insns

Translation and Virtual ISAs



- New compatibility interface: ISA + translation software
 - **Binary-translation** (static): transform static image, run native
 - **Emulation** (dynamic binary-translation):
unmodified image, interpret each dynamic insn
 - Typically optimized with just-in-time (JIT) compilation
 - Examples: DEC FX!32 (x86 on Alpha), Rosetta (PowerPC on x86)
 - Performance overheads reasonable (many advances over the years)
- **Virtual ISAs**: designed for translation, not direct execution
 - Target for high-level compiler (one per language)
 - Source for low-level translator (one per ISA)
 - Goals: Portability (abstract hardware nastiness), flexibility over time
 - Examples: Java Bytecodes, C# CLR (Common Language Runtime), NVIDIA's "PTX"

Ultimate Compatibility Trick

- Support old ISA by...
 - ...having a simple processor for that ISA somewhere in the system
 - How did PlayStation2 support PlayStation1 games?
 - Used PlayStation processor for I/O chip & **emulation**

What we learned


- What is ISA?
- Execution model:
 - Compilation
 - Assembly & machine language
- Instruction execution model
 - Registers, memory, PC
 - Instruction execution
- ISA design goals
 - Programmability
 - Performance/implementability
 - Compatibility



Aspects of ISAs


Aspects of ISAs

- Instruction length and format
 - Instruction encoding
- Where does data live?
 - Addressing modes
- Control transfers
 - How to find the next instruction
 - Branch
 - Jump



Length and Format

- **Length**
 - Fixed length
 - Most common is 32 bits
 - + Simple implementation (next PC often just PC+4)
 - Code density: 32 bits to increment a register by 1
 - Variable length
 - + Code density
 - x86 averages 3 bytes (ranges from 1 to 16)
 - Complex fetch (where does next instruction begin?)
 - Compromise: two lengths
 - E.g., MIPS16 or ARM's Thumb (16 bits)
- **Encoding**
 - A few simple encodings simplify decoder
 - Machine code (1s and 0s) <-> assembly




Example Instruction Encodings

- MIPS


- Fixed length
- 32-bits, 3 formats, simple encoding

add R1, R2, R3

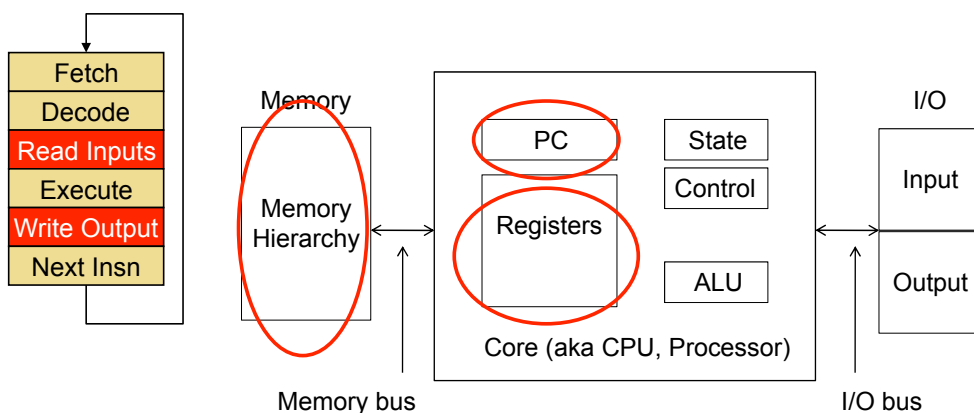


- x86


- Variable length encoding (1 to 15 bytes)



Where Does Data Live?



Where Does Data Live?



- **Registers** (e.g., R0, R1, **F0**)
 - "short term memory" `ADD R1, R2, R3`
 - Faster than memory, quite handy
 - Named directly in instructions
- **Memory** (e.g., (R3), **#20(R5)**)
 - "longer term memory" `ADD R1, R2, (R3)`
 - Accessed via "addressing modes"
 - Address to read or write calculated by instruction
- "Immediates" (e.g., #36, #7)
 - Values spelled out as bits in instructions
 - Input only

Memory Addressing

