

CMPE 110: Computer Architecture

Week 4

Pipelining III

Jishen Zhao (<http://users.soe.ucsc.edu/~jzhao/>)

[Adapted in part from Jose Renau, Mary Jane Irwin, Joe Devietti, Onur Mutlu, and others]


Reminder

- Quiz 1 posted
 - Due on Wednesday (Oct. 12) midnight on eCommons
 - Questions also posted on class Google website

Memory Addressing

- **Addressing mode:** way of specifying address
- Examples
 - **Displacement (based):** address = $[R2 + \text{immed}]$, e.g., $\#20(R2)$
 - **Index-base:** address = $[R2 + R3]$
 - **Memory-indirect:** address = $[\text{mem}[R2]]$
 - **Auto-increment:** address = $[R2]$, $R2 = R2 + 1$
 - **Auto-indexing:** address = $[R2 + \text{immed}]$, $R2 = R2 + \text{immed}$
 - **Scaled (scale-indexed):** address = $[R2 + R3 * \text{immed1} + \text{immed2}]$
 - **PC-relative:** address = $[PC + \text{imm}]$


Control Transfers



- Default next-PC = $PC + \text{sizeof}(\text{current insn})$
 - Branches and jumps can change that
- **Computing targets:** where to jump to
 - For all branches and jumps
 - **PC-relative:** e.g., J offset for branches and jumps with function
 - **Absolute:** e.g., J L3 for function calls
 - **Register indirect:** e.g., JR R5 for returns, switches & dynamic calls
- **Testing conditions:** whether to jump or not
 - Use registers & separate branch insns (MIPS)
e.g., `bnez R2, target`


```
L3:
    addu R7, R4, R3
    lw   R7, (R7)
    addu R8, R5, R3
    J    L3
    bne R3, R6, -4
```

Pipelining



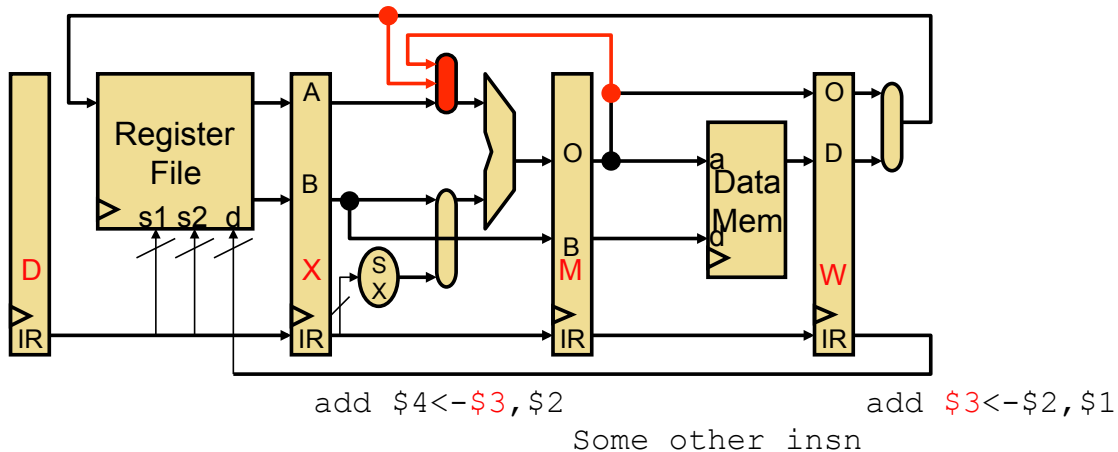
- Single-cycle datapaths vs. pipelined datapath
 - Basic pipelining: F, D, X, M, W
 - **Base CPI = 1**
 - Pipeline diagram (table)
- Data hazards
 - **Stalling**
 - Bypassing (forwarding)
- Structural hazards
 - **Stalling**
 - Add more hardware resources
- Multi-cycle operations
- Control hazards
 - Branch prediction

Review: Bypassing (forwarding)



- **Bypassing**
 - Reading a value from an intermediate (micro-architectural) source
 - Not waiting until it is available from primary source
 - Here, we are bypassing the register file
 - Also called **forwarding**
 - This example is an **MX** bypass

Review: Bypassing (forwarding)



- What about this combination?
 - Add another bypass path and MUX (multiplexor) input
 - This one is a **WX** bypass

Review: Pipeline Diagrams with Bypassing

- If bypass exists, “from”/“to” stages execute in same cycle

- Example: MX bypass

	1	2	3	4	5	6	7	8	9	10
add r1 <- r2, r3	F	D	X	M	W					
sub r2 <- -r1, r4		F	D	X	M	W				

- Example: WX bypass


	1	2	3	4	5	6	7	8	9	10
add r1 <- r2, r3	F	D	X	M	W					
ld r1, [r7+4]		F	D	X	M	W				
sub r2 <- -r1, r4			F	D	D	X	M	W		

- Example: WM bypass

	1	2	3	4	5	6	7	8	9	10
add r1 <- r2, r3	F	D	X	M	W					
sw r1, [r4+8]		F	D	X	M	W				

- Can you think of a code example that uses the WM bypass?

Review: Multi-cycle operations



- Multiplier itself is often pipelined, what does this mean?
 - Product/multiplicand register/ALUs/latches replicated
 - Can start different multiply operations in consecutive cycles
 - **But still takes 4 cycles to generate output value**

Review: Structural Hazard


- What about...
 - Two instructions trying to write register file in same cycle?
 - Structural hazard!
- Must prevent:

	1	2	3	4	5	6	7	8	9
mul \$4<-\$3,\$5	F	D	P0	P1	P2	P3	W		
addi \$6<-\$1,1		F	D	X	M	W			
add \$5<-\$6,\$10			F	D	X	M	W		

- Solution? stall the subsequent instruction

	1	2	3	4	5	6	7	8	9
mul \$4<-\$3,\$5	F	D	P0	P1	P2	P3	W		
addi \$6<-\$1,1		F	D	X	M	W			
add \$5<-\$6,\$10			F	D	d*	X	M	W	

Today



- Single-cycle datapaths vs. pipelined datapath
 - Basic pipelining: F, D, X, M, W
 - Base CPI = 1
 - Pipeline diagram (table)
- Data hazards
 - **Stalling**
 - Bypassing (forwarding)
- Structural hazards
 - **Stalling**
 - Add more hardware resources
- Multi-cycle operations
- Control hazards
 - Branch prediction

More Multiplier Nasties

- What about...
 - Mis-ordered writes to the same register
 - Software thinks `add` gets `$4` from `addi`, actually gets it from `mul`

	1	2	3	4	5	6	7	8	9
<code>mul \$4,\$3,\$5</code>	F	D	P0	P1	P2	P3	W		
<code>addi \$4,\$1,1</code>		F	D	X	M	W			
...									
...									
<code>add \$10,\$4,\$6</code>									

- Common? Not for a 4-cycle multiply with 5-stage pipeline
 - More common with deeper pipelines
 - In any case, must be correct

Corrected Pipeline Diagram

- With the correct stall logic
 - Prevent mis-ordered writes to the same register
 - Why two cycles of delay?

	1	2	3	4	5	6	7	8	9
mul \$4 , \$3, \$5	F	D	P0	P1	P2	P3	W		
addi \$4 , \$1, 1		F	D	d*	d*	X	M	W	
...									
...									
add \$10, \$4 , \$6									

- Multi-cycle operations complicate pipeline logic**

Pipelined Functional Units

- Almost all multi-cycle functional units are pipelined
 - Each operation takes N cycles
 - But can start initiate a new (independent) operation every cycle
 - Requires internal latching and some hardware replication
- + A cheaper way to improve throughput than multiple non-pipelined units


	1	2	3	4	5	6	7	8	9	10	11
mul f0, f1, f2	F	D	E0	E1	E2	E3	W				
mul f3, f4, f5		F	D	E0	E1	E2	E3	W			

- One exception: int/FP divide: difficult to pipeline and not worth it

	1	2	3	4	5	6	7	8	9	10	11
div f0, f1, f2	F	D	E	E	E	E	W				
div f3, f4, f5		F	D	s*	s*	s*	E	E	E	E	W


- s*** = stall for structural hazard

Where are we?



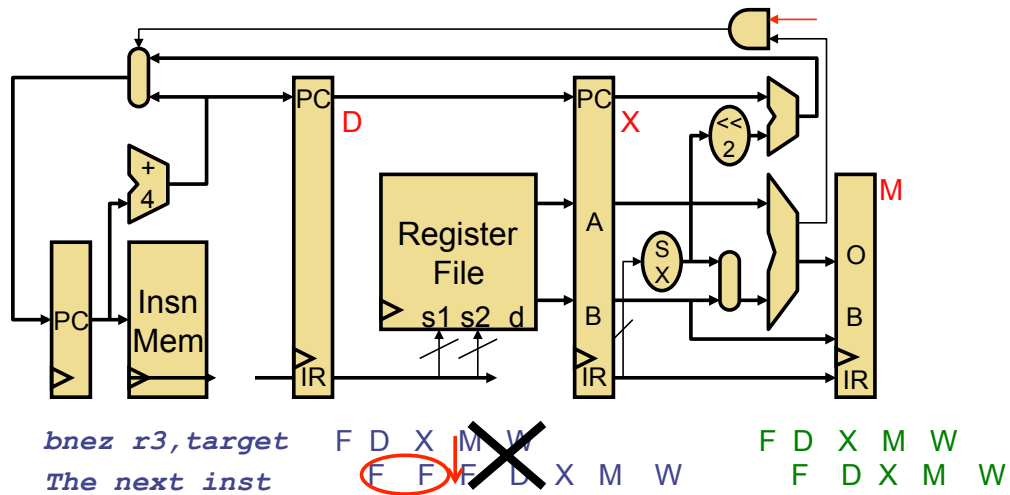
- Single-cycle datapaths vs. pipelined datapath
 - Basic pipelining: F, D, X, M, W
 - Base CPI = 1
 - Pipeline diagram (table)
- Data hazards
 - **Stalling**
 - Bypassing (forwarding)
- Structural hazards
 - **Stalling**
 - Add more hardware resources
- Multi-cycle operations
- Control hazards
 - Branch prediction

`bnez r3, target`



Control Dependences and Branch Prediction

What About Branches?

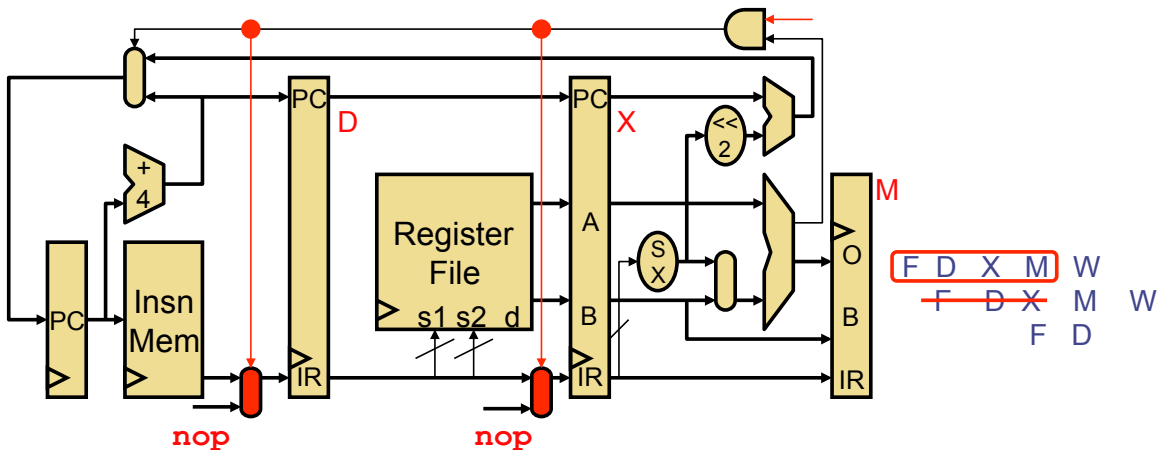


- Could just stall to wait for branch outcome (two-cycle penalty)
- **Branch speculation:** speculatively fetch some instruction before branch outcome is known
 - Default: assume "**not-taken**" (at fetch, can't tell it's a branch)

CMPE 110: Computer Architecture | Prof. Jishen Zhao | Week 4

17

Branch Recovery



- **Branch recovery**: what to do when branch is actually taken
 - Insns that will be written into D and X are wrong
 - **Flush them**, i.e., replace them with **nops**
- + They haven't written permanent state yet (regfile, DMem)
 - Two cycle penalty for taken branches

CMPE 110: Computer Architecture | Prof. Jishen Zhao | Week 4

18

Branch Speculation and Recovery

Speculation:

	1	2	3	4	5	6	7	8	9
addi r1,1→r3	F	D	X	M	W				
bnez r3,targ		F	D	X	M	W			
st r6→[r7+4]			F	D	X	M	W		
mul r8,r9→r10				F	D	P0	P1	P2	...

speculative

- **Mis-speculation recovery:** what to do on wrong guess
 - Not too painful in a short, in-order pipeline
 - Branch resolves in X
 - + Younger insns (in F, D) haven't changed permanent state
 - **On next cycle, flush** insns in D and X → **2 cycle overhead**

Recovery:

	1	2	3	4	5	6	7	8	9
addi r1,1→r3	F	D	X	M	W				
bnez r3,targ		F	D	X	M	W			
st r6→[r7+4]			F	D	--	--	--		
mul r8,r9→r10				F	--	--	--	--	
targ:add r4,r5→r4					F	D	X	M	W

CMPE 110: Computer Architecture | Prof. Jishen Zhao | Week 4

19

Branch Performance

- Back of the envelope calculation
 - **Branch: 20%**, load: 20%, store: 10%, other: 50%
 - Speculation: not-taken
 - But actually, **75% of branches are taken**
- CPI = ? (assuming $CPI_{base} = 1$)
- $CPI = 1 + 20\% * 75\% * 2 = 1.3$
- **Branches cause 30% slowdown**
 - Worse with deeper pipelines (higher misprediction penalty)
- Can we do better than assuming branch is not taken?

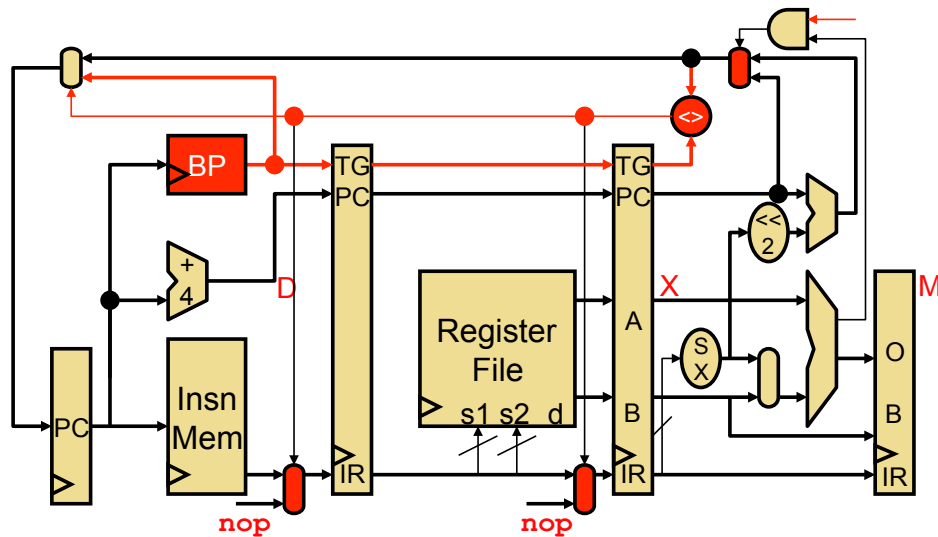
Big Idea: Speculative Execution

- **Speculation:** “risky transactions on chance of profit”
- **Speculative execution**
 - Execute before all parameters known with certainty
 - **Correct speculation**
 - + Avoid stall, improve performance
 - **Incorrect speculation (mis-speculation)**
 - Must abort/flush/squash incorrect insns
 - Must undo incorrect changes (recover pre-speculation state)
- **“Control speculation”:**
 - speculation aimed at control hazards

Control Speculation Mechanics

- Guess branch target, start fetching at guessed position
 - Doing nothing is implicitly guessing target is PC+length_in_bytes (e.g., PC+4 for 32-bit instructions)
 - Can actively guess other targets: **dynamic branch prediction**
- Execute branch to verify (check) guess
 - Correct speculation? keep going
 - Mis-speculation? Flush mis-speculated insns
 - Hopefully haven’t modified permanent state (Regfile, DMem)
 - + Happens naturally in in-order 5-stage pipeline

Dynamic Branch Prediction



- **Dynamic branch prediction:** hardware guesses outcome
 - Start fetching from guessed address
 - Flush on **mis-prediction**