

# CMPE 110: Computer Architecture

## Week 4


### Pipelining IV

Jishen Zhao (<http://users.soe.ucsc.edu/~jzhao/>)

[Adapted in part from Jose Renau, Mary Jane Irwin, Joe Devietti, Onur Mutlu, and others]

## Pipelining

---



- Single-cycle datapaths vs. pipelined datapath
  - Basic pipelining: F, D, X, M, W
  - Base CPI = 1
  - Pipeline diagram (table)
- Data hazards
  - **Stalling**
  - Bypassing (forwarding)
- Structural hazards
  - **Stalling**
  - Add more hardware resources
- Multi-cycle operations
- Control hazards
  - Branch prediction

## Review: Branch Speculation and Recovery

**Speculation:**

	1	2	3	4	5	6	7	8	9
addi r1,1→r3	F	D	X	M	W				
bnez r3,targ		F	D	X	M	W			
st r6→[r7+4]			<b>F</b>	<b>D</b>	X	M	W		
mul r8,r9→r10				<b>F</b>	D	P0	P1	P2	...

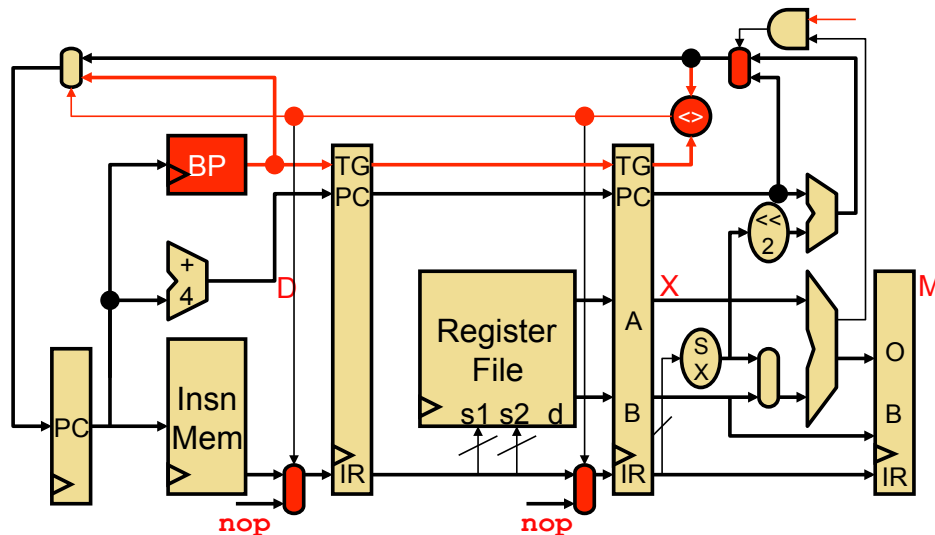
speculative

- **Mis-speculation recovery:** what to do on wrong guess
  - Not too painful in a short, in-order pipeline
  - Branch resolves in X
  - + Younger insns (in F, D) haven't changed permanent state
  - **On next cycle, flush** insns in D and X → **2 cycle overhead**

**Recovery:**

	1	2	3	4	5	6	7	8	9
addi r1,1→r3	F	D	X	M	W				
bnez r3,targ		F	D	<b>X</b>	M	W			
<del>st r6→[r7+4]</del>			<b>F</b>	<b>D</b>	--	--	--		
<del>mul r8,r9→r10</del>				<b>F</b>	--	--	--	--	
targ:add r4,r5→r4					<b>F</b>	D	X	M	W

## Review: Dynamic Branch Prediction




- **Dynamic branch prediction:** hardware guesses outcome
  - Start fetching from guessed address
  - Flush on **mis-prediction**

## When to perform branch prediction?

### ❑ During Decode

- Look at instruction opcode to determine branch instructions
- Can calculate next PC from instruction (for PC-relative branches)
- One cycle “mis-fetch” penalty **even if branch predictor is correct**

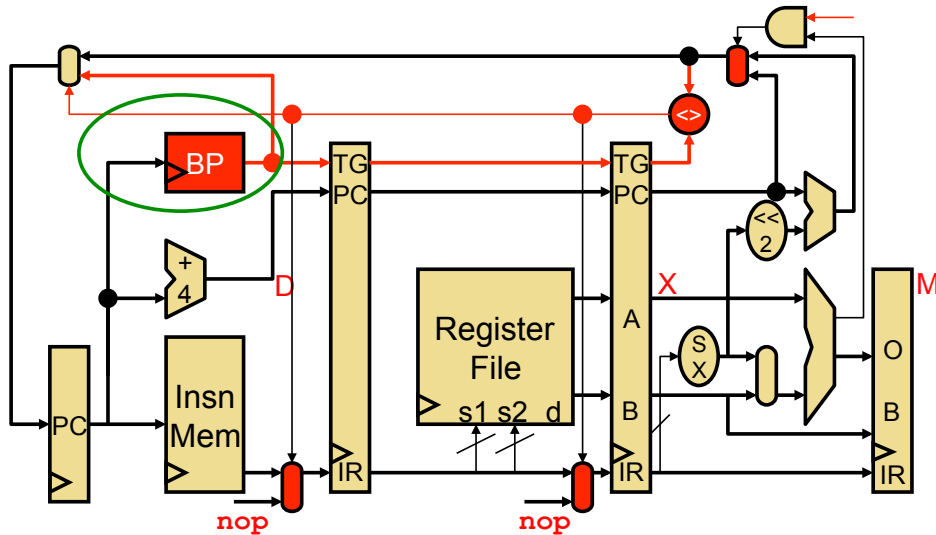


### ❑ During Fetch?

- How do we do that?
- Branch predictor locates at F stage

## Branch predictor


## Where is branch predictor (BP)?



- **Dynamic branch prediction:** hardware guesses outcome
  - Start fetching from guessed address
  - Flush on **mis-prediction**

## What's inside a branch predictor?

- BTB & branch direction predictor during fetch



- Step #1: is it a branch? BTB
- Step #2: is the branch taken or not taken? BHT
- Step #3: if the branch is taken, where does it go? BTB, RAS

## One-bit history-based branch prediction

- **Branch history table (BHT):**  
simplest direction predictor
  - PC indexes table of bits (0 = N, 1 = T)
  - Essentially: branch will go same way it went last time
  - Problem: **inner loop branch** below
 

```
for (i=0; i<100; i++)
    for (j=0; j<3; j++)
        // loop body
```

    - Two “built-in” mis-predictions per inner loop iteration
    - Branch predictor “changes its mind too quickly”
    - **6 wrong predictions**

Time	State	Prediction	Outcome
1	N	N	T
2	T	T	T
3	T	T	T
4	T	T	N
5	N	N	T
6	T	T	T
7	T	T	T
8	T	T	N
9	N	N	T
10	T	T	T
11	T	T	T
12	T	T	N


## Two-bit history-based branch prediction

- **Two-bit saturating counters (2bc)**  
[Smith 1981]
  - Replace each single-bit prediction
    - (0,1,2,3) = (N,n,t,T)
  - Adds “hysteresis”
    - Force predictor to mis-predict twice before “changing its mind”
  - One misprediction each loop execution (rather than two)
    - + Fixes this pathology (which is not contrived, by the way)
    - Can we do even better?
    - **5 wrong predictions**

Time	State	Prediction	Outcome
1	N	N	T
2	n	N	T
3	t	T	T
4	T	T	N
5	t	T	T
6	T	T	T
7	T	T	T
8	T	T	N
9	n	T	T
10	t	T	T
11	T	T	T
12	T	T	N


## Summary: inside a branch predictor?

- BTB & branch direction predictor during fetch



- Step #1: is it a branch? BTB
- Step #2: is the branch taken or not taken? BHT
- Step #3: if the branch is taken, where does it go? BTB, RAS

## Summary of pipelining



- Single-cycle datapaths vs. pipelined datapath
  - Basic pipelining
- Data hazards
- Structural hazards
- Multi-cycle operations
- Control hazards
- Fine-grained multithreading

**Any other methods on handling data and control hazards?**

**Do something else:  
fine-grained multithreading**

## Fine-Grained Multithreading

---

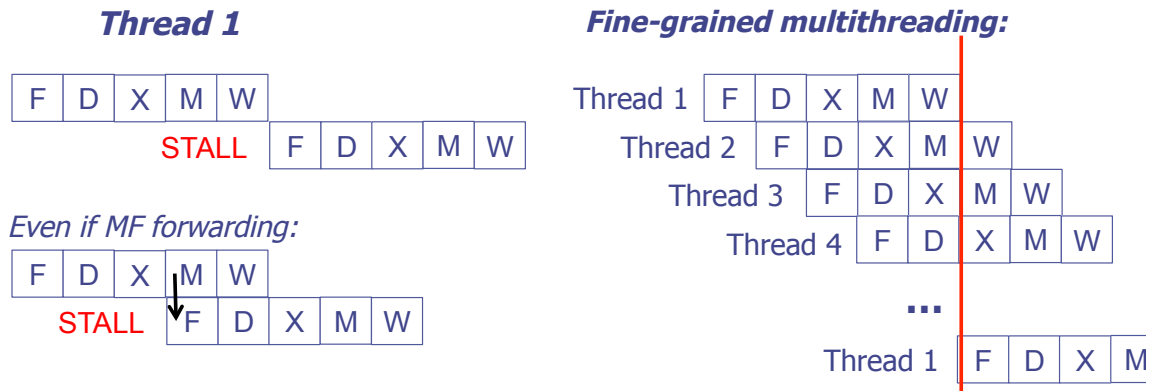
- Idea: Hardware has multiple thread contexts (hardware threads). Each cycle, fetch an instruction from a different thread. → the same pipeline, multiple threads
  - By the time the fetched branch instruction resolves, no instruction is fetched from the same thread
  - But...Branch instruction resolution latency can overlap with execution of other threads' instructions

# Fine-Grained Multithreading

*Thread 1*  
bne R3, R6, L3  
addi R1<- R1, #1  
L3: add R1<- R2, R4

*Thread 2*  
add R10<- R5, R8  
addi R10<- R10, #1

*Other threads...*  
Has no dependency




## Pros and cons

- **Advantages**
  - + No need for dependency checking between instructions (only one instruction in pipeline from a single thread)
  - + No need for branch prediction logic
  - + Otherwise-stall cycles used for executing useful instructions from different threads
  - + Improved system throughput and pipeline utilization
- **Disadvantages**
  - Extra hardware complexity: multiple hardware contexts (PCs, register files, ...), thread selection logic
  - Reduced single thread performance (one instruction fetched every N cycles from the same thread)
  - Resource contention between threads in caches and memory
  - Some dependency checking logic *between* threads remains (load/store)



# Hardware Multithreading

- **Not** the same as software multithreading!
- A **hardware thread** is a sequential stream of insns
  - could be a software thread or a single-threaded process



- **Hardware Multithreading (MT)**
  - Multiple hardware threads dynamically share a single pipeline
  - Replicate only per-thread structures: program counter & registers
  - Hardware interleaves instructions

# Hardware Multithreading

- Why use hw multithreading?
  - + **Multithreading improves utilization and throughput**
    - Single programs utilize <50% of pipeline (branch, cache miss)
    - allow insns from different hw threads in pipeline at once
  - **Multithreading does not improve single-thread performance**
    - Individual threads run as fast or even slower
  - **Coarse-grain MT**: switch on cache misses Why (will discuss later)?
  - **Simultaneous MT**: no explicit switching, fine-grain interleaving
    - Simultaneous multithreading (SMT): Issue multiple instructions from multiple threads in one cycle. The processor must be superscalar to do so.
    - Intel's "hyperthreading"

