

CMPE 110: Computer Architecture


Week 5

Out-of-Order Execution

Jishen Zhao (<http://users.soe.ucsc.edu/~jzhao/>)

[Adapted in part from Jose Renau, Mary Jane Irwin, Joe Devietti, Onur Mutlu, and others]


Review: Summary of pipelining



- Single-cycle datapaths vs. pipelined datapath
 - Basic pipelining
- Data hazards
- Structural hazards
- Multi-cycle operations
- Control hazards
- Fine-grained multithreading

Review: inside a branch predictor?

- BTB & branch direction predictor during fetch



- Step #1: is it a branch? BTB
- Step #2: is the branch taken or not taken? BHT
- Step #3: if the branch is taken, where does it go? BTB, RAS

Review: Uniprocessor Concurrency


- **Software "thread":** Independent flows of execution
 - "Per-thread" state
 - Context state: PC, registers
 - Stack (per-thread local variables)
 - "Shared" state: globals, heap, etc.
 - Threads generally share the same memory space
 - A process is like a thread, but with its own memory space
 - Java has thread support built in, C/C++ use the pthreads library
- Generally, system software (the O.S.) manages threads
 - "Thread scheduling", "context switching"
 - In single-core system, all threads share one processor
 - Hardware timer interrupt occasionally triggers O.S.
 - Quickly swapping threads gives illusion of concurrent execution
 - Much more in an operating systems course

Review: Uniprocessor Concurrency

- Programmer explicitly creates multiple threads
- A “thread switch” can occur at any time
 - Pre-emptive multithreading by OS
- Common uses:
 - Handling user interaction (GUI programming)
 - Handling I/O latency (send network message, wait for response)
 - **Expressing parallel work via Thread-Level Parallelism (TLP)**

Review: Hardware Multithreading


- **Not** the same as software multithreading!
- A **hardware thread** is a sequential stream of insns
 - could be a software thread or a single-threaded process




- **Hardware Multithreading (MT)**
 - Multiple hardware threads dynamically share a single pipeline
 - Replicate only per-thread structures: program counter & registers
 - Hardware interleaves instructions

Review: Hardware Multithreading

- Why use hw multithreading?
 - + **Multithreading improves utilization and throughput**
 - Single programs utilize <50% of pipeline (branch, cache miss)
 - allow insns from different hw threads in pipeline at once
 - **Multithreading does not improve single-thread performance**
 - Individual threads run as fast or even slower
 - **Coarse-grain MT**: switch on cache misses Why?
 - **Simultaneous MT**: no explicit switching, fine-grain interleaving
 - Intel's "hyperthreading"

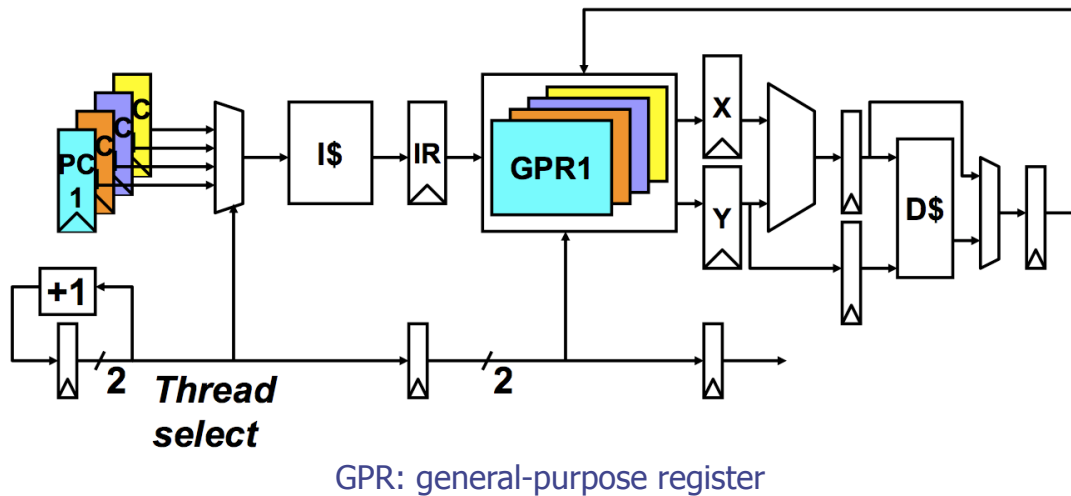


Reivew: Hardware Multithreading



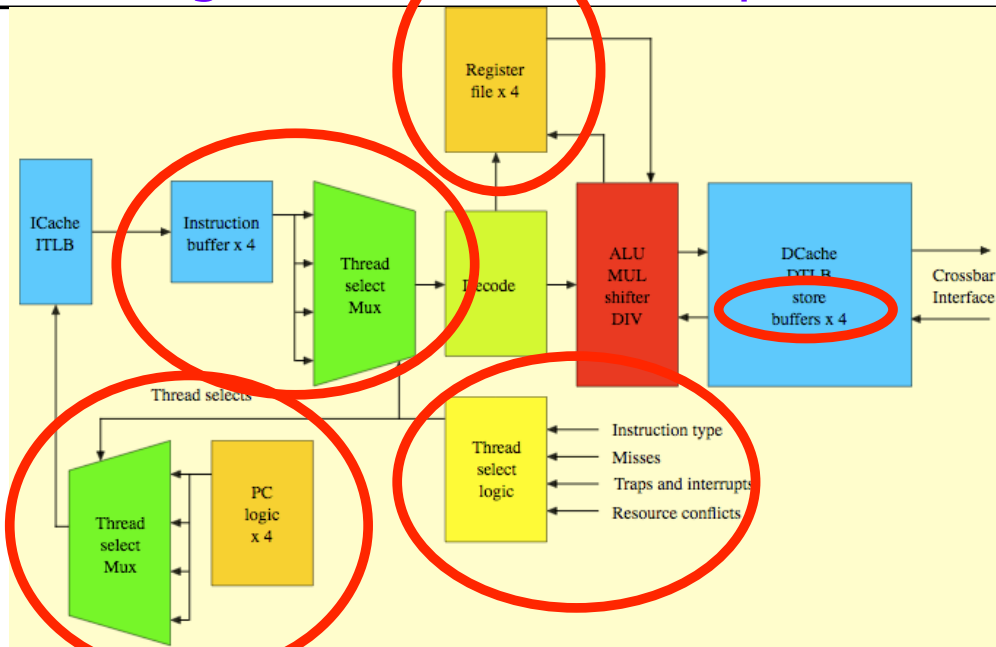
Multithreaded Pipeline Example

Need to store multiple "thread contexts"




Slide credit: Joel Emer

Sun Niagara Multithreaded Pipeline



Kongetira et al., "Niagara: A 32-Way Multithreaded Sparc Processor," IEEE Micro 2005.



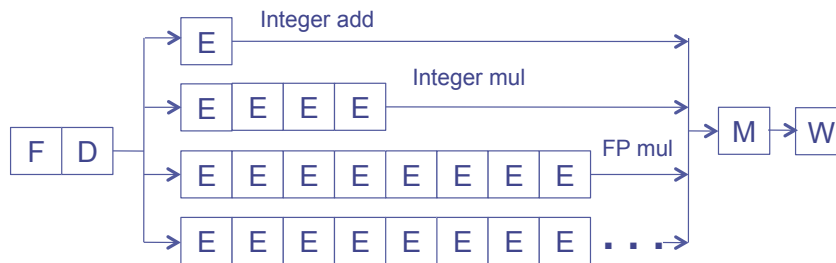
Out-of-order Execution

Preliminary: Dynamic Instruction Scheduling

- Hardware has knowledge of dynamic events on a per-instruction basis (i.e., at a very fine granularity)
 - Branch mispredictions
 - Load/store addresses
 - Cache misses
- Wouldn't it be nice if hardware did the scheduling of instructions?

An In-order Pipeline

F, D, **E**, M, W



- Problem: A true data dependency stalls dispatch of younger instructions into functional (execution) units
- Dispatch: Act of sending an instruction to an execution functional unit (i.e., in E stage)
- Load and store instructions can take multiple cycles in M stage

Can We Do Better?

- What do the following two pieces of code have in common (with respect to execution in the previous design)?

IMUL R3 \leftarrow R1, R2	LD R3 \leftarrow R1 (0) // M stage may take multiple cycles
ADD R3 \leftarrow R3, R1	ADD R3 \leftarrow R3, R1
ADD R1 \leftarrow R6, R7	ADD R1 \leftarrow R6, R7
IMUL R5 \leftarrow R6, R8	IMUL R5 \leftarrow R6, R8
ADD R7 \leftarrow R9, R9	ADD R7 \leftarrow R9, R9

- Answer: First ADD can stall the whole pipeline!
 - ADD cannot dispatch because its source registers unavailable
 - Later **independent** instructions cannot get executed
- How are the above code portions different?
 - Answer: Load latency is variable (unknown until runtime)
 - What does this affect? Think compiler vs. microarchitecture

Preventing Dispatch Stalls


- Multiple ways of doing it
- You have already seen at least one:
 - Fine-grained multithreadingMay see more later on:
 - Value prediction
 - Compile-time instruction scheduling/reordering (i.e., static instruction scheduling)
- What are the disadvantages of fine-grained multithreading?
- Any other way to prevent dispatch stalls?
 - Problem: in-order dispatch (scheduling, or execution)
 - Solution: out-of-order dispatch (scheduling, or execution)

Out-of-order Execution (Dynamic Scheduling)

- Idea: Move the dependent instructions out of the way of independent ones (s.t. independent ones can execute)
 - Resting areas for dependent instructions: **Reservation stations**
- Monitor the source “values” of each instruction in the reservation stations
- When all source “values” of an instruction are available, “fire” (i.e. dispatch) the instruction
- Benefit:
 - **Latency tolerance**: Allows independent instructions to execute and complete in the presence of a long latency operation


In-order vs. Out-of-order Dispatch

- In order dispatch + precise exceptions:



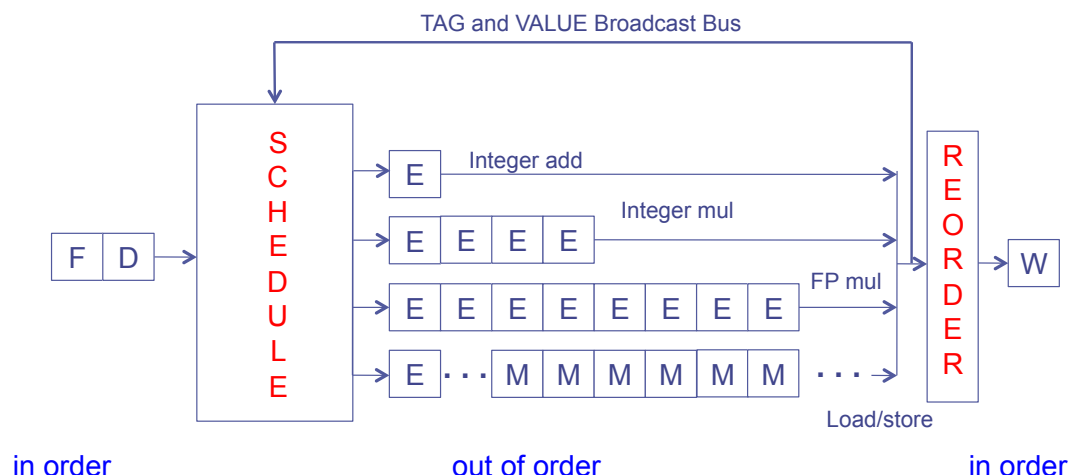
IMUL R3 \leftarrow R1, R2
 ADD R3 \leftarrow R3, R1
 ADD R1 \leftarrow R6, R7
 IMUL R5 \leftarrow R6, R8
 ADD R7 \leftarrow R3, R5

- Out-of-order dispatch + precise exceptions:



15 vs. 12 cycles

Two Humps in a Modern Pipeline



- Hump 1: Reservation stations (scheduling window)
- Hump 2: Reordering (reorder buffer, aka instruction window or active window)

Register Renaming

The processor only
has 4 registers

```
add R3<- R1, R2
addi R1<- R4, #100
```

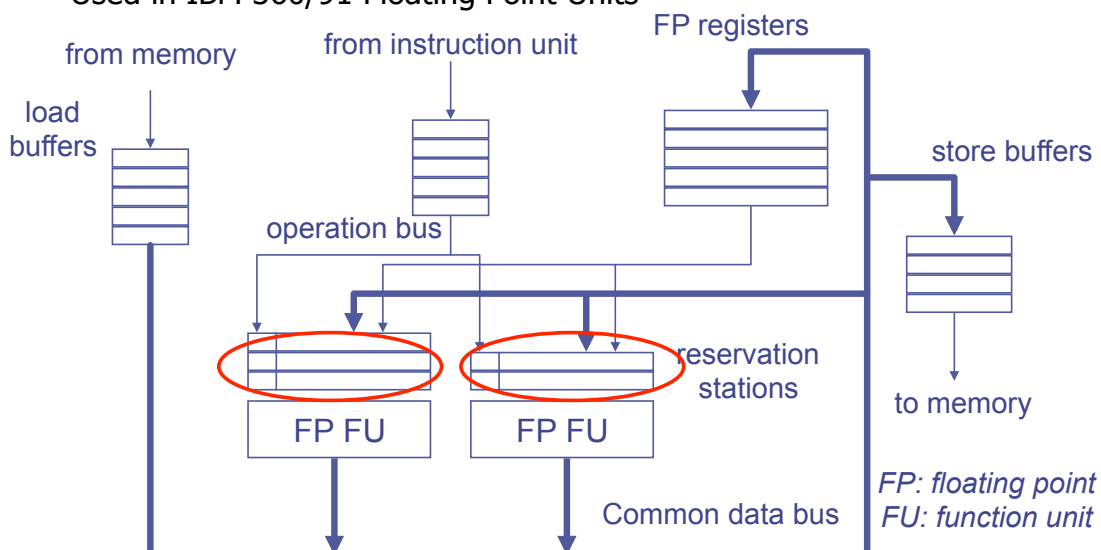
- Some dependencies are not true dependencies
 - The same register refers to values that have nothing to do with each other
 - They exist because not enough register ID's (i.e. names) in the ISA**
- True dependency:

```
addi R1<- R4, #100
add R3<- R1, R2
```
- Solution: The register ID is **renamed** to the reservation station entry that will hold the register's value
 - Register ID → Reservation Station entry ID
 - After renaming, Reservation Station entry ID is used to refer to the register
- This eliminates the dependency shown in our example
 - Achieves (approximates) the performance of having a large number of registers, even though ISA does not support that many

Register renaming

Out-of-order execution (with **register renaming**) invented by Robert Tomasulo in 1960's

- Used in IBM 360/91 Floating Point Units



Tomasulo's Algorithm

- OoO with register renaming invented by Robert Tomasulo
 - Used in IBM 360/91 Floating Point Units
 - Tomasulo, “**An Efficient Algorithm for Exploiting Multiple Arithmetic Units**,” IBM Journal of R&D, Jan. 1967.
- What is the major difference today?
 - **Precise exceptions**: IBM 360/91 did NOT have this
 - Patt, Hwu, Shebanow, “**HPS, a new microarchitecture: rationale and introduction**,” MICRO 1985.
 - Patt et al., “**Critical issues regarding HPS, a high performance microarchitecture**,” MICRO 1985.
- Variants are used in most high-performance processors
 - Initially in Intel Pentium Pro, AMD K5
 - Alpha 21264, MIPS R10000, IBM POWER5, IBM z196, Oracle UltraSPARC T4, ARM Cortex A15