

CMPE 110: Computer Architecture

Week 9

Virtual Memory

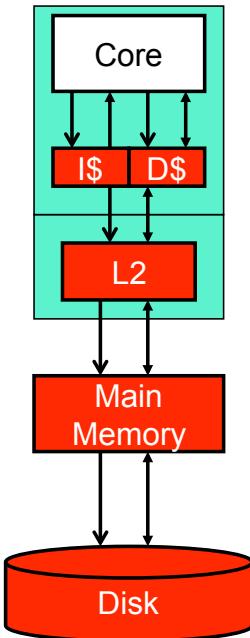
Jishen Zhao (<http://users.soe.ucsc.edu/~jzhao/>)

[Adapted in part from Jose Renau, Mary Jane Irwin, Joe Devietti, Onur Mutlu, and others]

Reminder

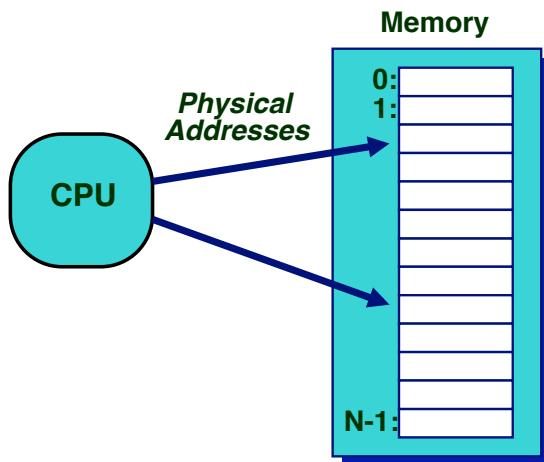
- Homework 4 will be posted on today

Today

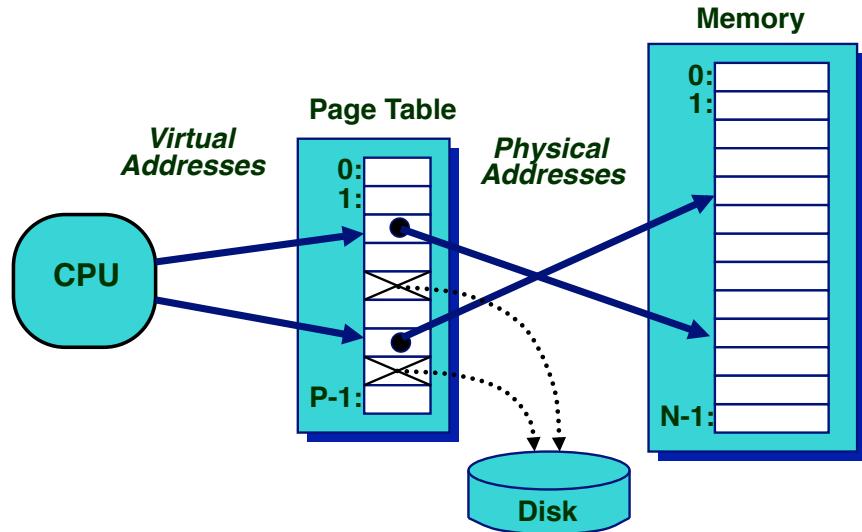


Virtualizing the memory hierarchy: An Infinite capacity

Review: If no virtual memory



Review: A System with Virtual Memory



- Address Translation: The hardware converts virtual addresses into physical addresses via an OS-managed lookup table (page table)

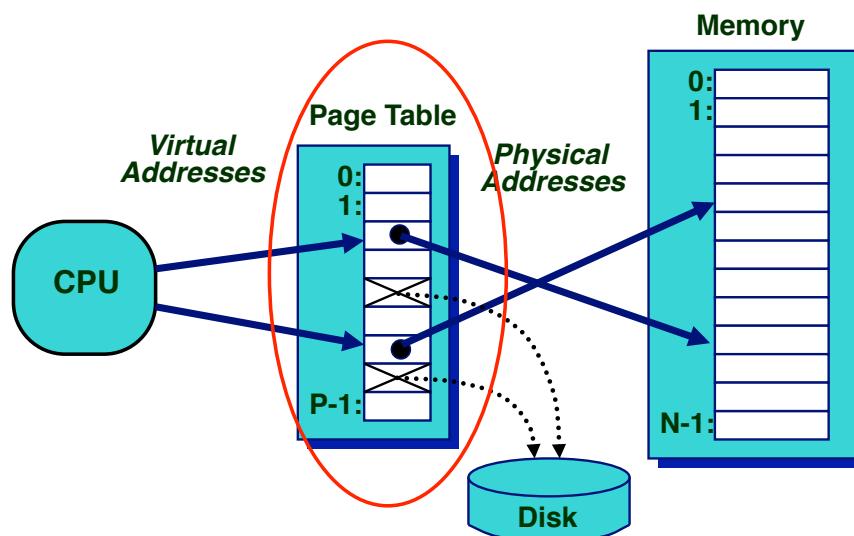
Virtual Memory

- Idea: Give the programmer the illusion of a large address space while having a small physical memory
 - So that the programmer does not worry about managing physical memory
- Programmer can assume he/she has “infinite” amount of physical memory
- Hardware and software cooperatively and automatically manage the physical memory space to provide the illusion
 - Illusion is maintained for each independent process

Today

- Virtual memory
 - “Infinite” memory, isolation, protection, inter-process communication
 - Virtual-physical address translation
 - Page tables
 - TLBs
 - Page faults
 - DMA
 - Virtual memory and cache
- Multicore

A logical view



Virtual Pages, Physical Frames

- Virtual address space divided into pages
- Physical address space divided into frames
- A virtual page is mapped to
 - A physical frame, if the page is in physical memory
 - A location in disk, otherwise
- If an accessed virtual page is not in memory, but on disk
 - Get a page fault
 - Virtual memory system brings the page into a physical frame and adjusts the mapping → this is called demand paging
- Page table is the table that stores the mapping of virtual pages to physical frames

Physical Memory as a Cache

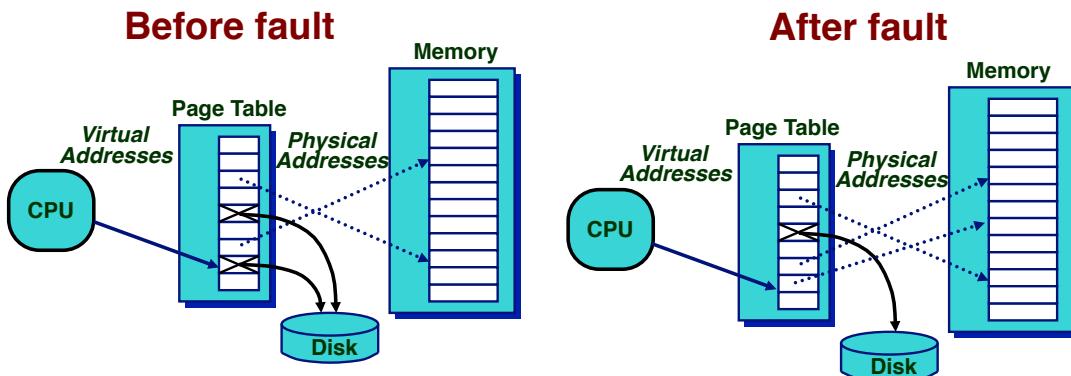
- In other words...
- Physical memory is a cache for pages stored on disk
 - In fact, it is a fully associative cache in modern systems (a virtual page can be mapped to any physical frame)
- Similar caching issues exist as we have covered earlier:
 - Placement: where and how to place/find a page in cache?
 - Replacement: what page to remove to make room in cache?
 - Granularity of management: large, small, uniform pages?
 - Write policy: what do we do about writes? Write back?

Supporting Virtual Memory

- Virtual memory requires both HW+SW support
 - Page Table is in memory
 - Can be **cached** in special hardware structures called Translation Lookaside Buffers (**TLBs**)
- The hardware component is called the **MMU** (memory management unit)
 - Includes Page Table Base Register(s), TLBs, page walkers
- It is the job of the software to leverage the MMU to
 - Populate page tables, decide what to replace in physical memory
 - Change the Page Table Register on context switch (to use the running thread's page table)
 - Handle page faults and ensure correct mapping

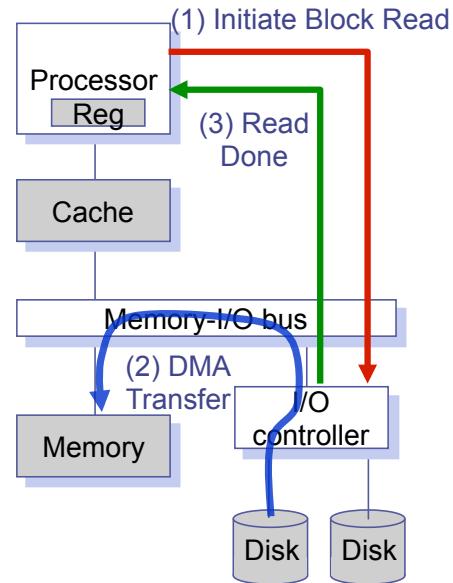
Page Fault ("A Miss in Physical Memory")

- If a page is not in physical memory but disk
 - Page table entry indicates virtual page not in memory
 - Access to such a page triggers a page fault exception
 - OS trap handler invoked to move data from disk into memory
 - Other processes can continue executing
 - OS has full control over placement



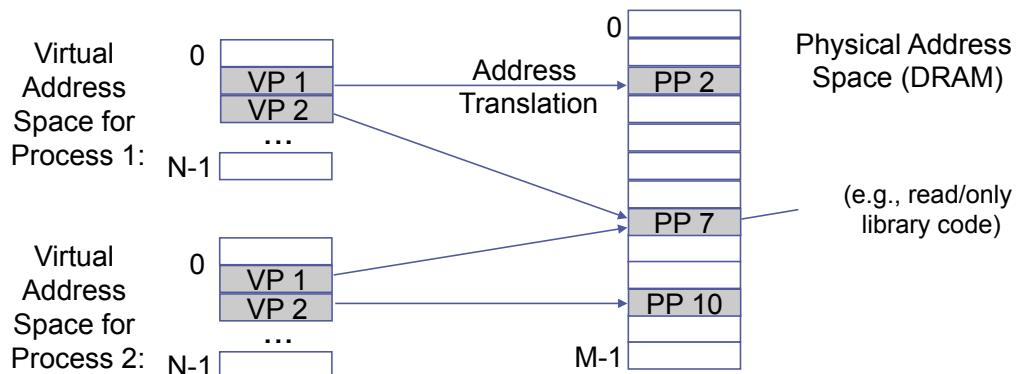
Servicing a Page Fault and DMA

- (1) Processor signals controller
 - Read block of length P starting at disk address X and store starting at memory address Y
- (2) Read occurs
 - Direct Memory Access (DMA)
 - Under control of I/O controller
- (3) Controller signals completion
 - Interrupt processor
 - OS resumes suspended process



Page Table is Per Process

- Each process has its own virtual address space
 - Full address space for each program
 - Simplifies memory allocation, sharing, linking and loading.



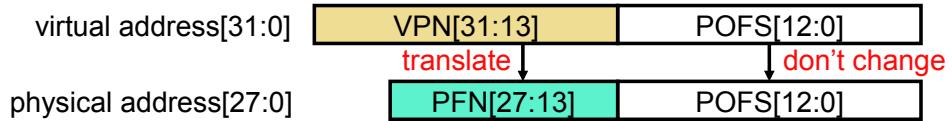
Virtual to Physical

Address translation

Address Translation

- How to obtain the physical address from a virtual address?
- Page size specified by the ISA
 - VAX: 512 bytes (by DEC in 1970s)
 - Today: 4KB, 8KB, 2GB, ... (small and large pages mixed together)
- Page Table contains an entry for each virtual page
 - Called Page Table Entry (PTE)
 - What is in a PTE?

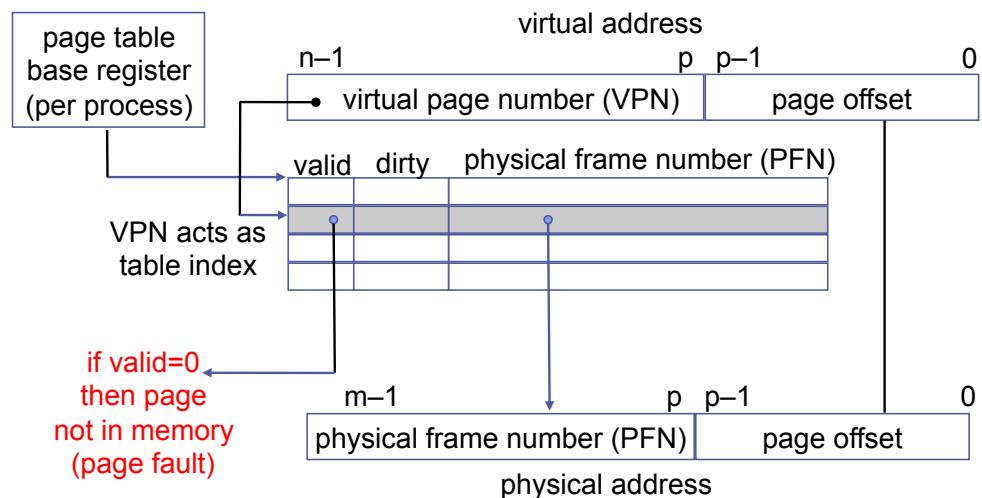
Address Translation (II)



- VA→PA mapping called **address translation**
 - Split VA into **virtual page number (VPN)** & **page offset (POFS)**
 - Translate VPN into **physical frame number (PFN)**
 - POFS is not translated
- Example above
 - 8KB page size
 - 32-bit virtual address
 - VPN 19 bits → 2^{19} virtual pages → 2^{19} PTEs in the page table (for each process)

What's in a page table

- Separate (set of) page table(s) per process
- VPN forms index into page table (points to a page table entry)
- Page Table Entry (PTE) provides information about page



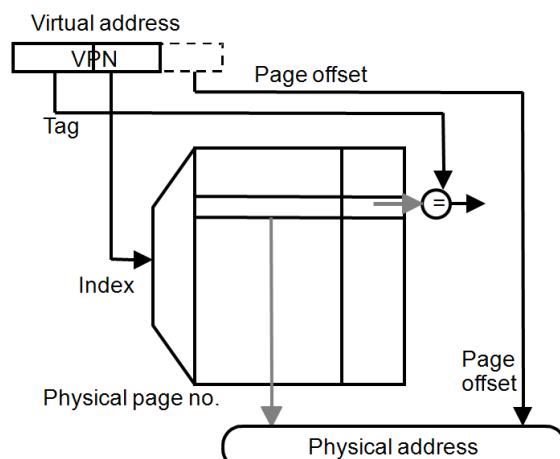
Virtual Memory Issue

- How fast is the address translation?
 - Cost at least one memory read
 - How can we make it fast?
- Idea: Use a hardware structure that caches PTEs →
 - Translation lookaside buffer (TLB)
 - Actually a cache, not a buffer

Speeding up Translation with a TLB

- Essentially a cache of recent address translations
 - Avoids going to the page table on every reference
- **Index** = lower bits of VPN (virtual page #)
- **Tag** = unused bits of VPN + process ID
- **Data** = a page-table entry
- **Status** = valid, dirty

The usual cache design choices (placement, replacement policy, multi-level, etc.) apply here too.



Handling TLB Misses

- The TLB is small; it cannot hold all PTEs
 - Some translations will inevitably miss in the TLB
 - Must access memory to find the appropriate PTE
 - Called **walking** the page directory/table
 - Large performance penalty
- Who handles TLB misses? Hardware or software?

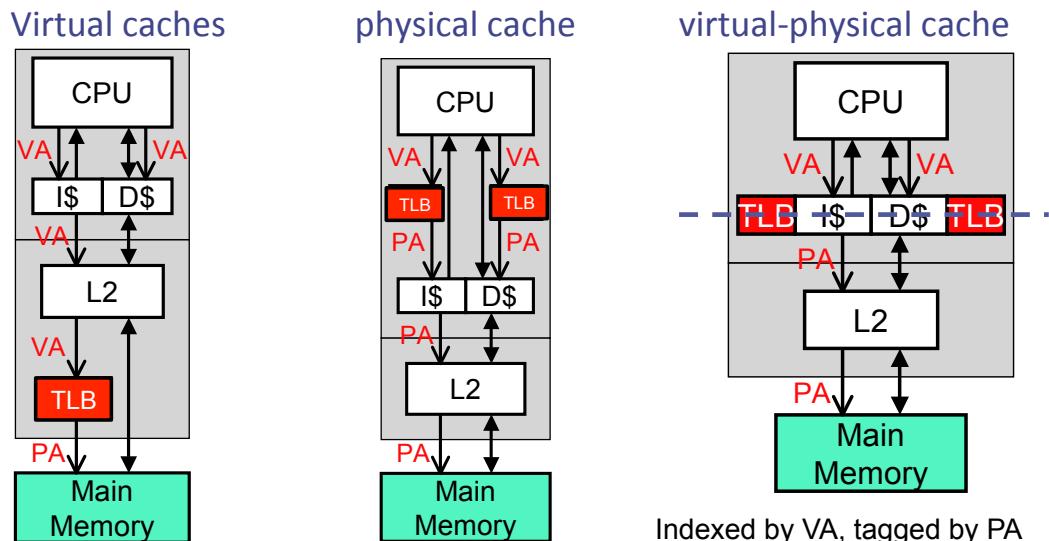
Handling TLB Misses (II)

- Approach #1. **Hardware-Managed** (e.g., x86)
 - The hardware does the **page walk**
 - The hardware fetches the PTE and inserts it into the TLB
 - If the TLB is full, the entry **replaces** another entry
 - Done transparently to system software
- Approach #2. **Software-Managed** (e.g., MIPS)
 - The hardware raises an exception
 - The operating system does the **page walk**
 - The operating system fetches the PTE
 - The operating system inserts/evicts entries in the TLB

Handling TLB Misses (III)

- Hardware-Managed TLB
 - + No exception on TLB miss. Instruction just stalls
 - + Independent instructions may continue
 - + No extra instructions/data brought into caches.
 - - Page table organization is etched into the system: OS has little flexibility in deciding these
- Software-Managed TLB
 - + The OS can define page table organization
 - + More sophisticated TLB replacement policies are possible
 - - Need to generate an exception → performance overhead due to pipeline flush, exception handler execution, extra instructions brought to caches

Cache-VM Interaction: Locations of TLB



What we learned

- OS virtualizes memory and I/O devices
- Virtual memory
 - “infinite” memory, isolation, protection, inter-process communication
 - Virtual-physical address translation
 - Page tables
 - TLBs
 - **It is a cache, not a buffer**
 - Page faults
 - DMA
 - Virtual memory and cache

What we learned

- Why do we need virtual memory?
- What is page table? Where is it?
- What is TLB? Is it a buffer or a cache?
- What is page fault? How is page fault handled? What is DMA?
- What is MMU? What does it do?
- What are virtual caches? What is physical cache? What is virtual-physical cache?
- Virtual-physical address translation. Example question on the next slide...