

CMPE 110: Computer Architecture

Week 10

Multicore

Jishen Zhao (<http://users.soe.ucsc.edu/~jzhao/>)

[Adapted in part from Onur Mutlu, Jose Renau, Mary Jane Irwin, Joe Devietti, and others]

Reminder

- Quiz 4 will be posted today
 - Due on Nov. 23 (Wed) 11:59pm

Review: Virtual memory


- OS virtualizes memory and I/O devices
- Virtual memory
 - “infinite” memory, isolation, protection, inter-process communication
 - Virtual-physical address translation
 - Page tables
 - TLBs
 - It is a cache, not a buffer
 - Page faults
 - DMA
 - Virtual memory and cache

3

Review: Virtual memory


- Why do need virtual memory?
- What is page table? Where is it?
- What is TLB? Is it a buffer or a cache?
- What is page fault? How is page fault handled? What is DMA?
- What is MMU? What does it do?
- What is virtual cache? What is physical cache? What is virtual-physical cache?
- Virtual-physical address translation. Example question on the next slide...

4



- Example above
 - What is page size in KB?
 - How many virtual pages can each process have?
 - VPN 19 bits $\rightarrow 2^{19}$ virtual pages
 - how many PTEs in total in maximum for 2 processes?
 - 2^{19} PTEs in the page table for each process $\rightarrow 2 * 2^{19}$ PTEs in total for two processes


5



Multicore & Multiprocessor Hardware

6

Roadmap Checkpoint




- Multicore vs. Hardware multithreading
- Cache coherence
- Memory consistency models

7

Multicore and Hardware Multithreading Implementation

8

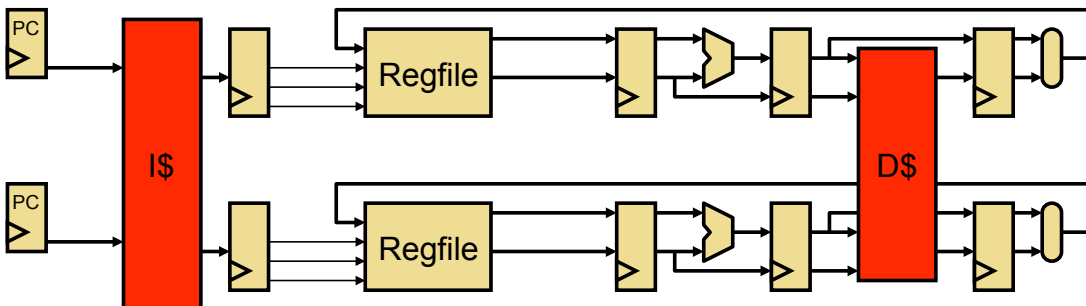
Recall: Hardware Multithreading



- **Hardware Multithreading (MT)**
 - Multiple threads dynamically share a single pipeline
 - Replicate only per-thread structures: program counter & registers
 - Hardware interleaves instructions
- + **Multithreading improves utilization and throughput**
 - Single programs utilize <50% of pipeline (branch, cache miss)
- **Multithreading does not improve single-thread performance**
 - Individual threads run as fast or even slower
- **Coarse-grain MT**: switch on cache misses Why?
- **Simultaneous MT**: no explicit switching, fine-grain interleaving

9


Simplest Multiprocessor



- Replicate entire processor pipeline!
 - Instead of replicating just register file & PC
 - Exception: share the caches (we'll address this bottleneck soon)
- Multiple threads execute
 - Shared memory programming model
 - Operations (loads and stores) are interleaved "at random"
 - Loads returns the value written by most recent store to location

10

Roadmap Checkpoint




- Multicore vs. Hardware multithreading
- Cache coherence
- Memory consistency models

11

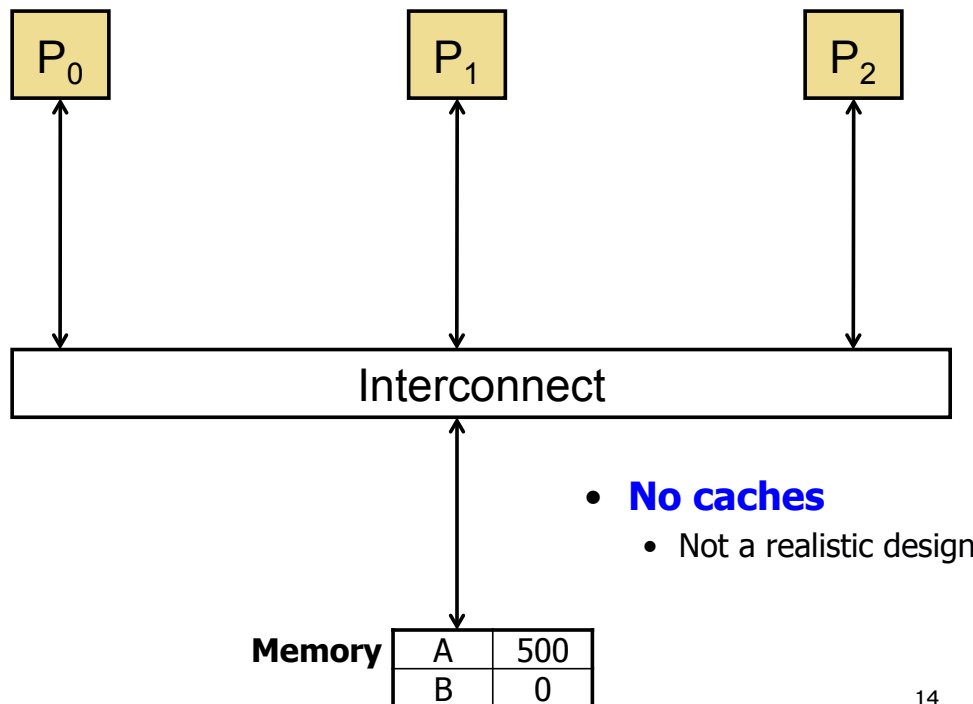
Cache coherence

12


No-Cache (Conceptual) Implementation




No-Cache (Conceptual) Implementation



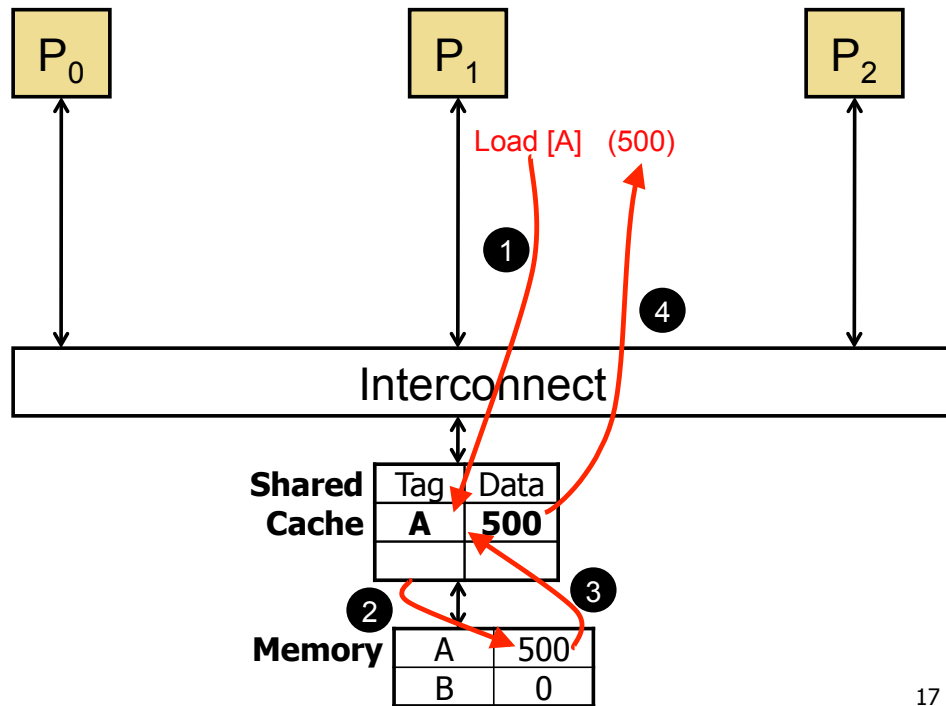
Shared Cache Implementation



Shared Cache Implementation

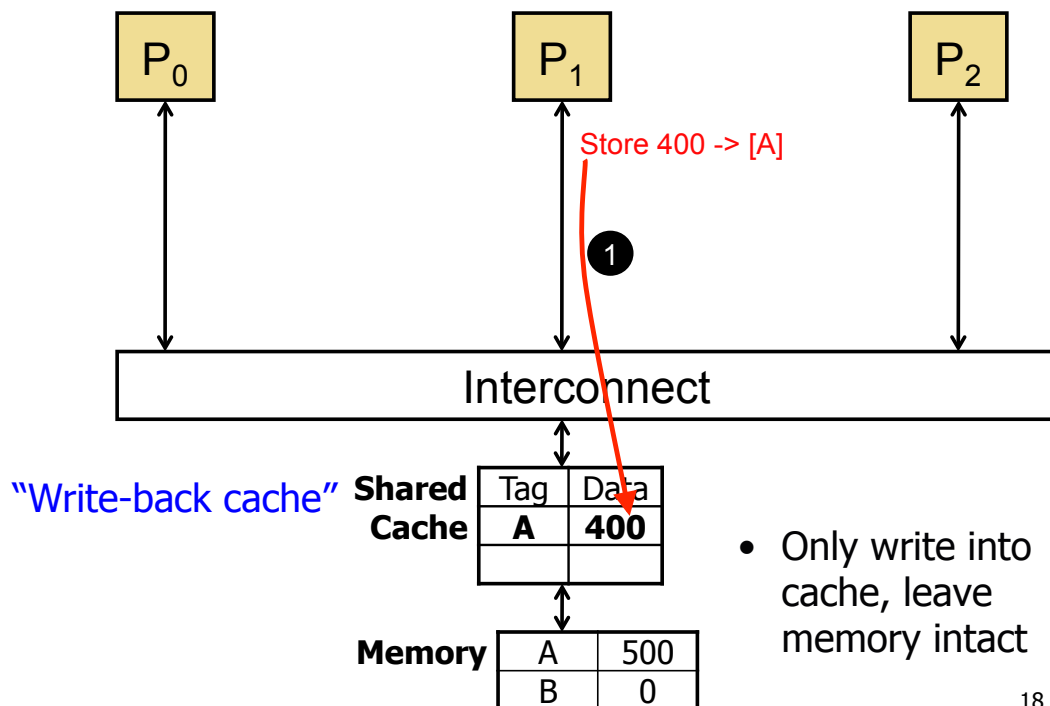


Shared Cache Implementation



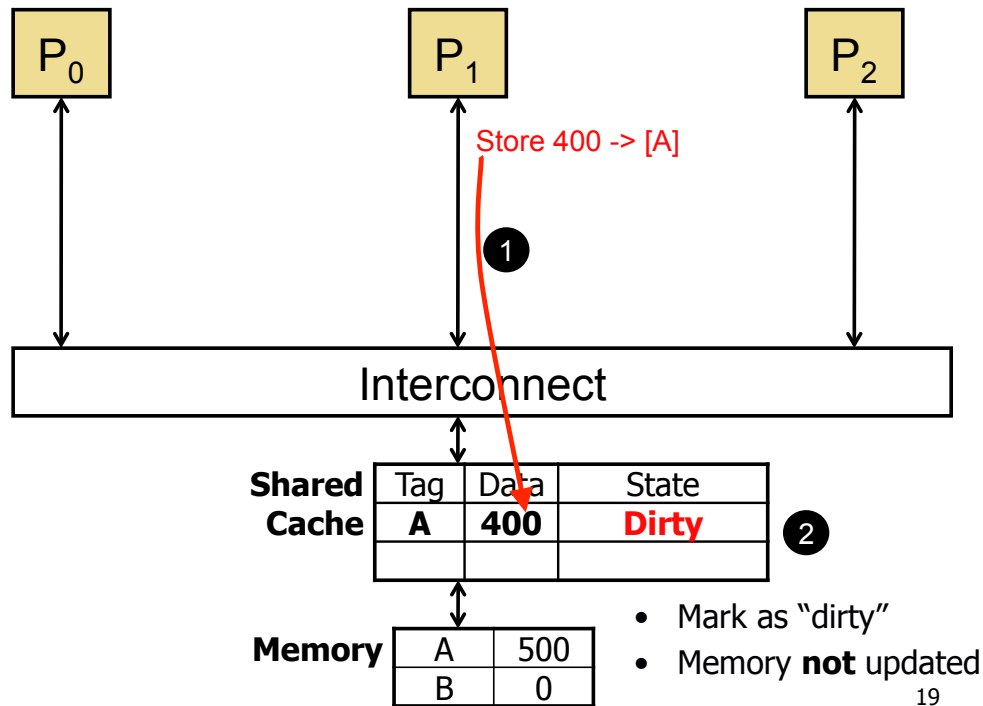
17

Shared Cache Implementation



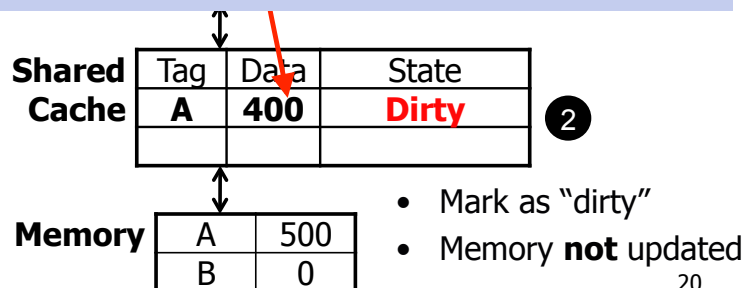
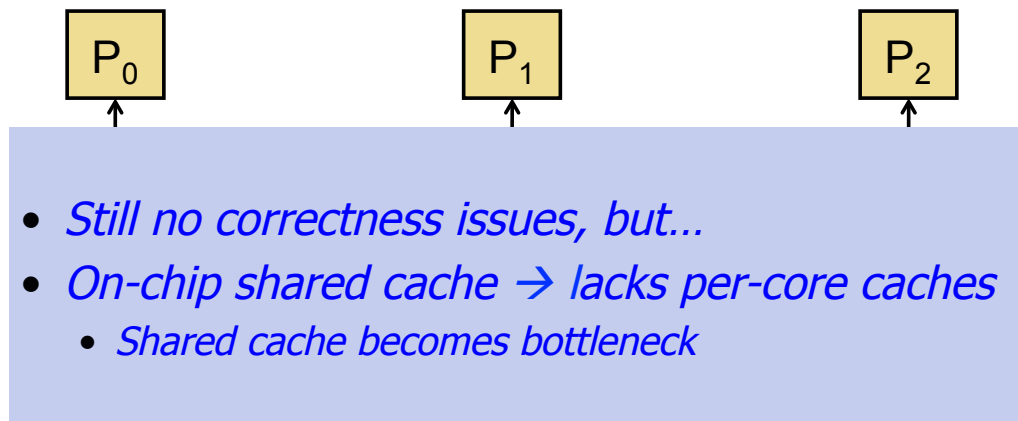
18

Shared Cache Implementation



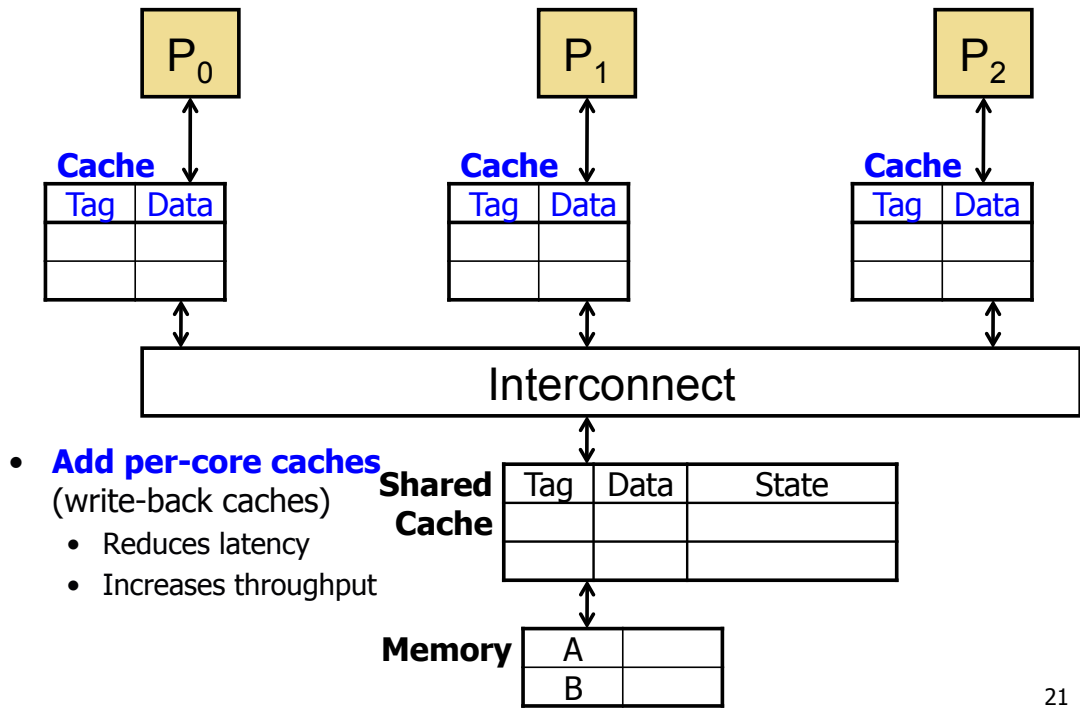
19

Shared Cache Implementation



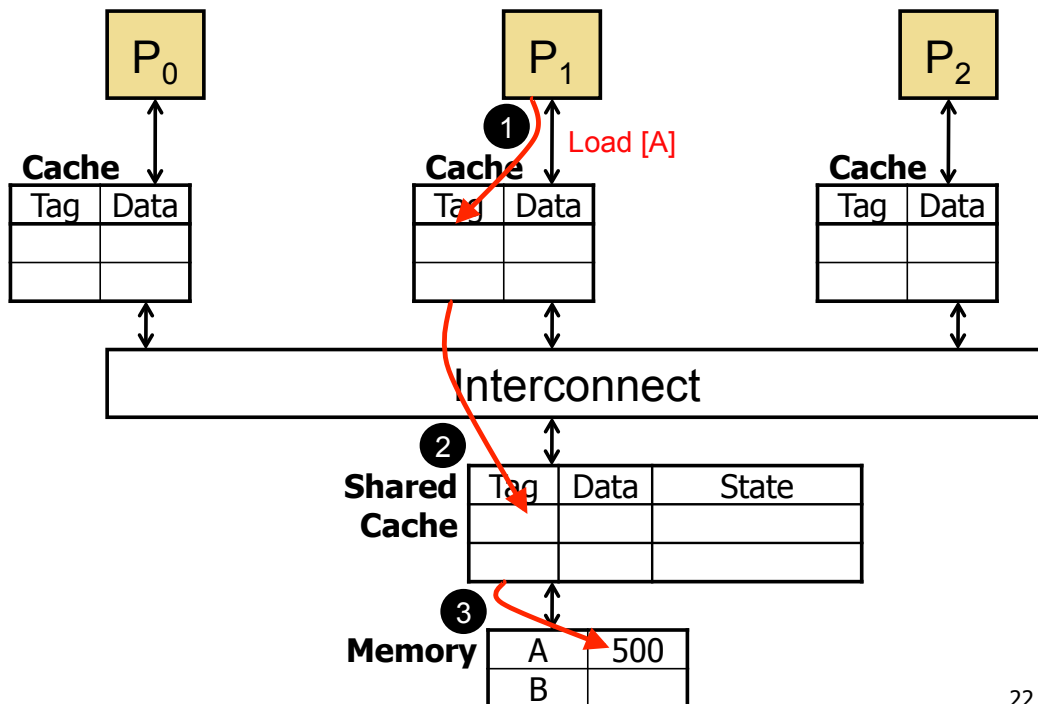
20

Adding Private Caches



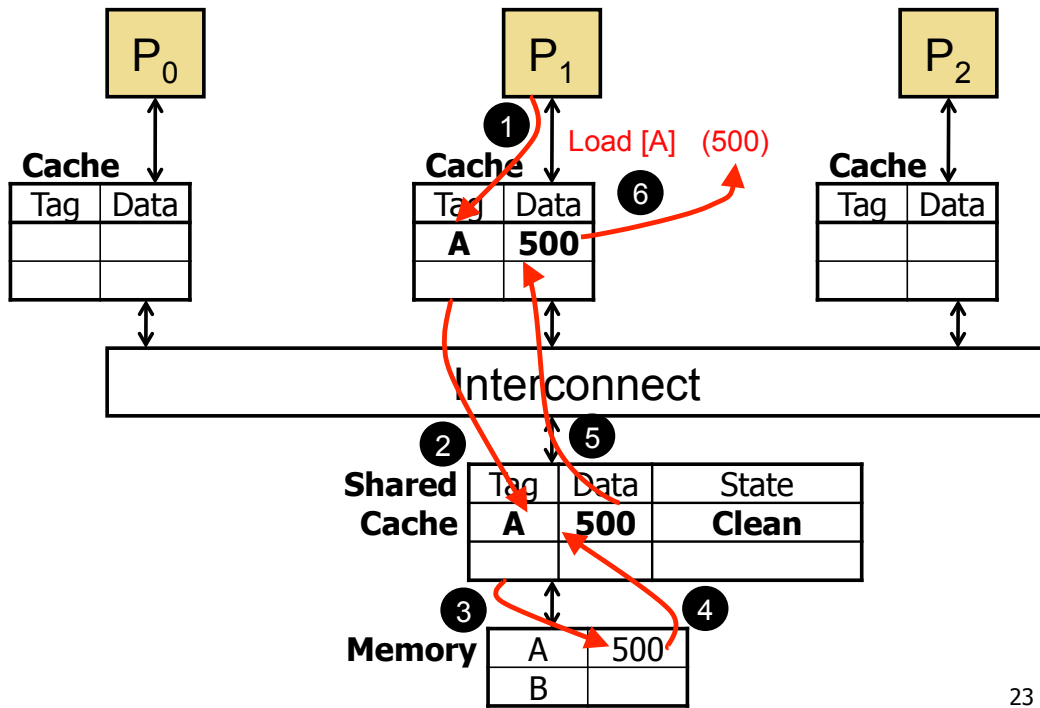
21

Adding Private Caches



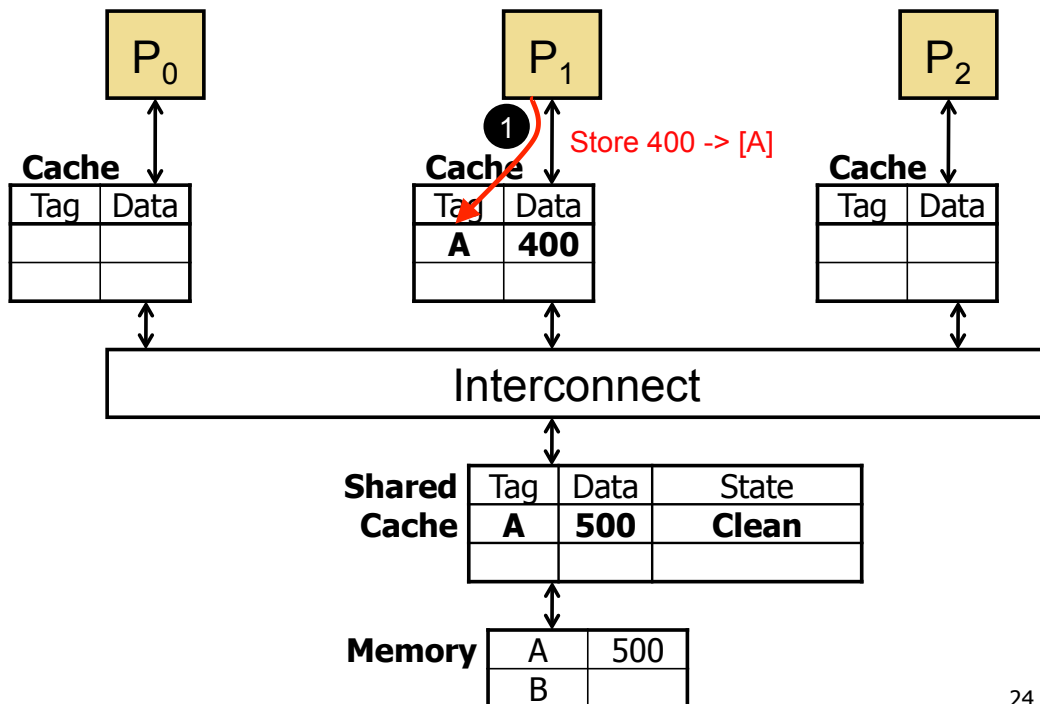
22

Adding Private Caches



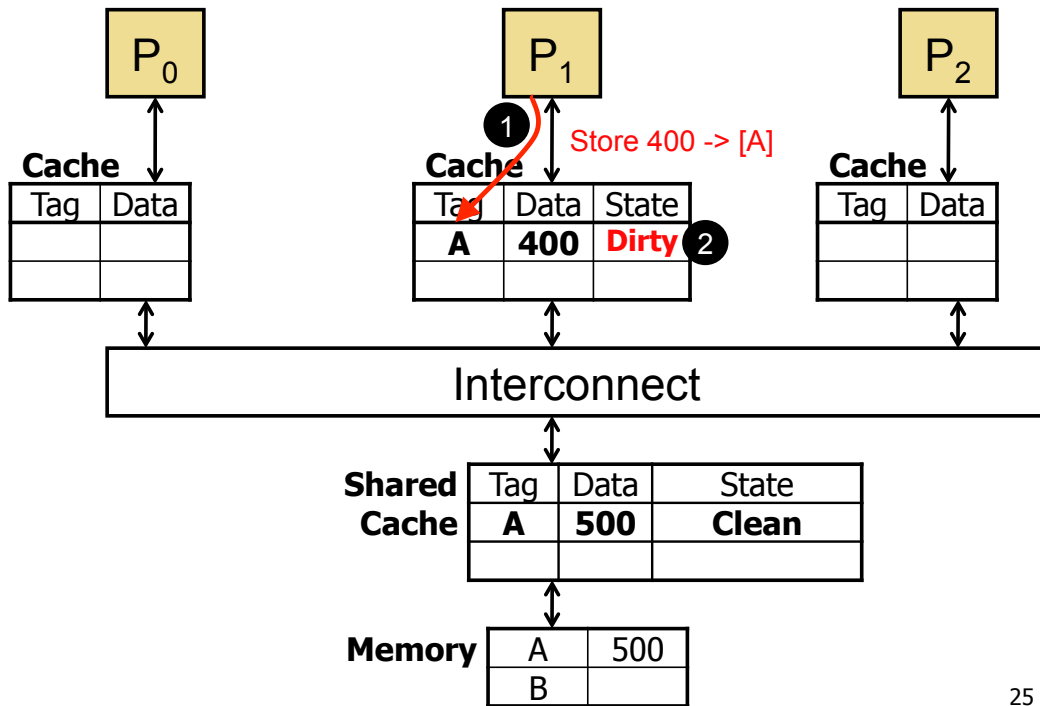
23

Adding Private Caches



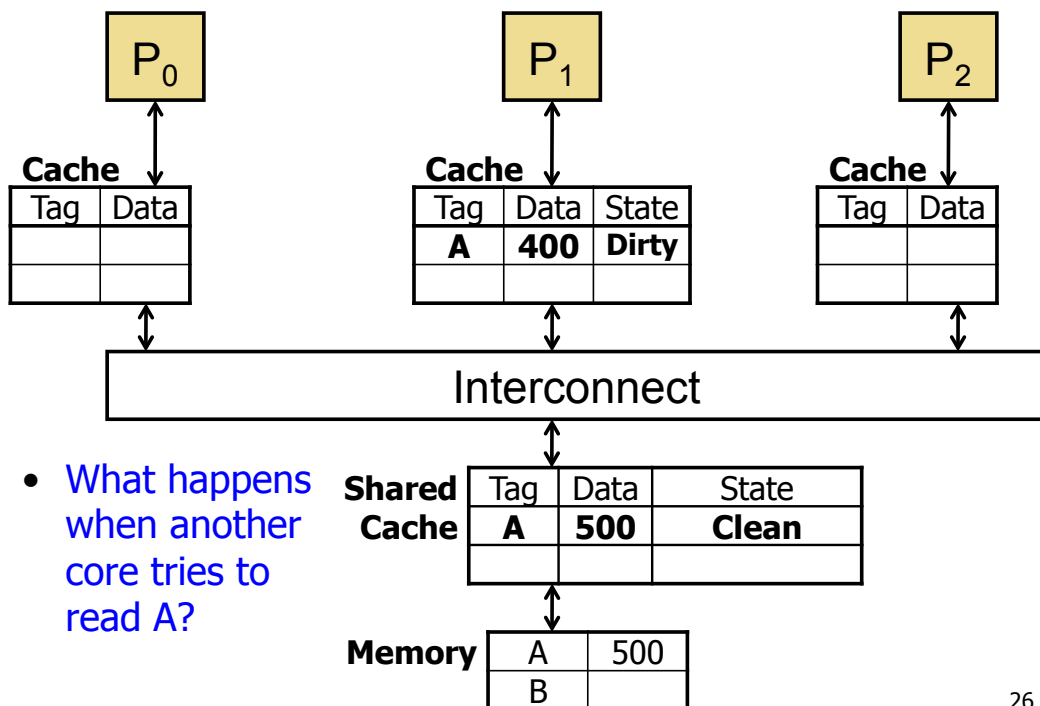
24

Adding Private Caches



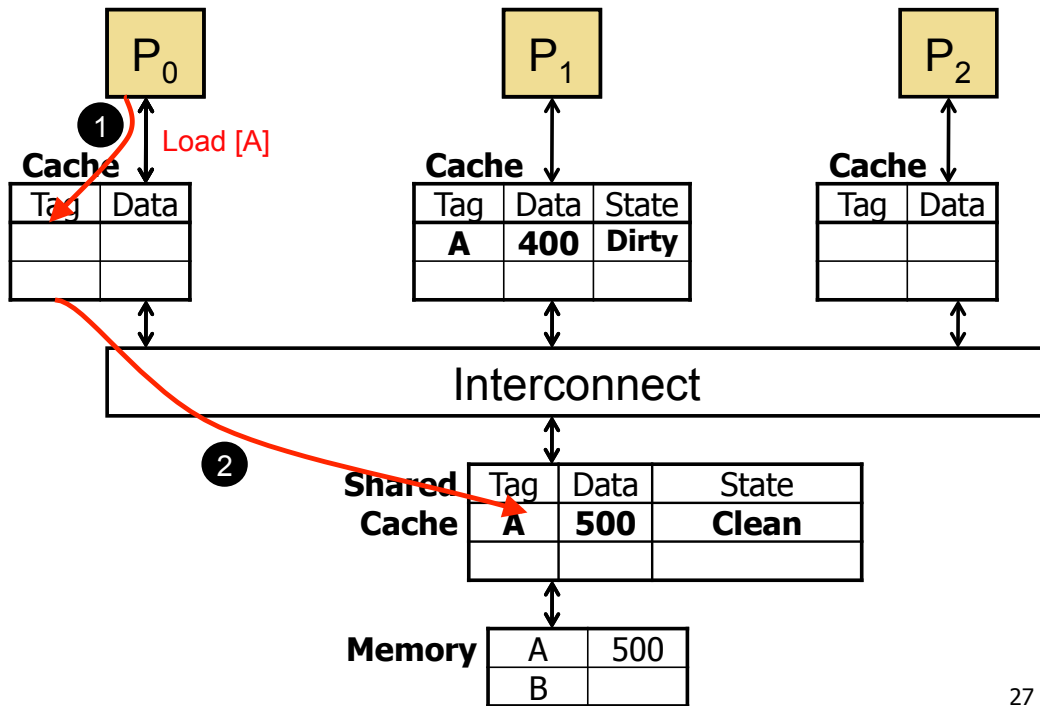
25

Private Cache Problem: Incoherence



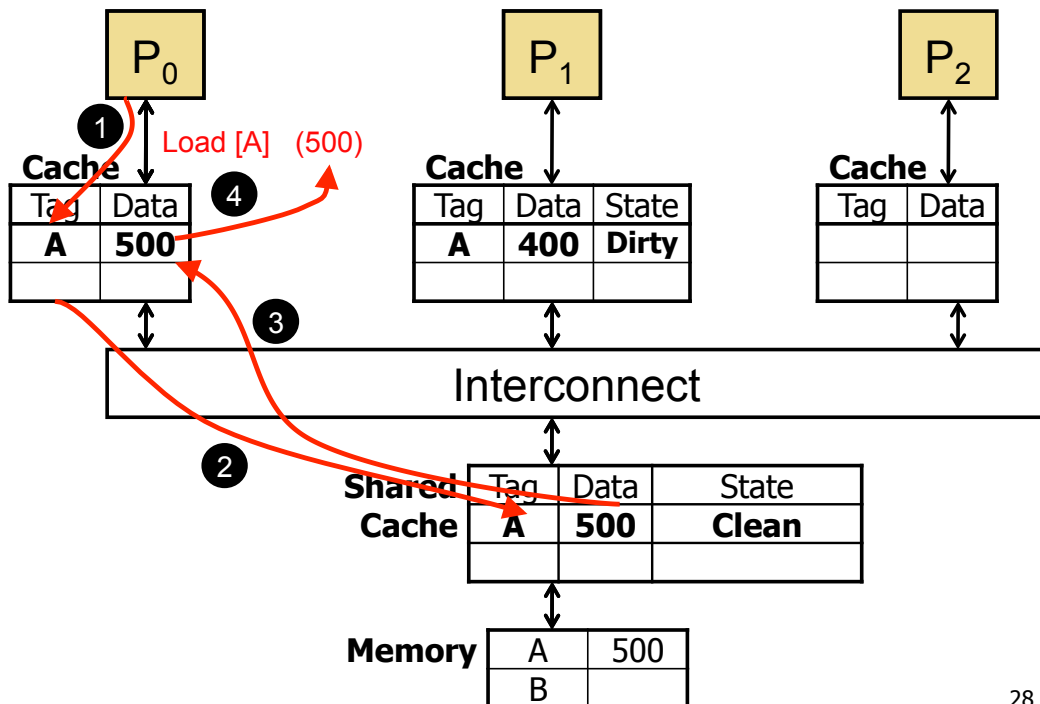
26

Private Cache Problem: Incoherence



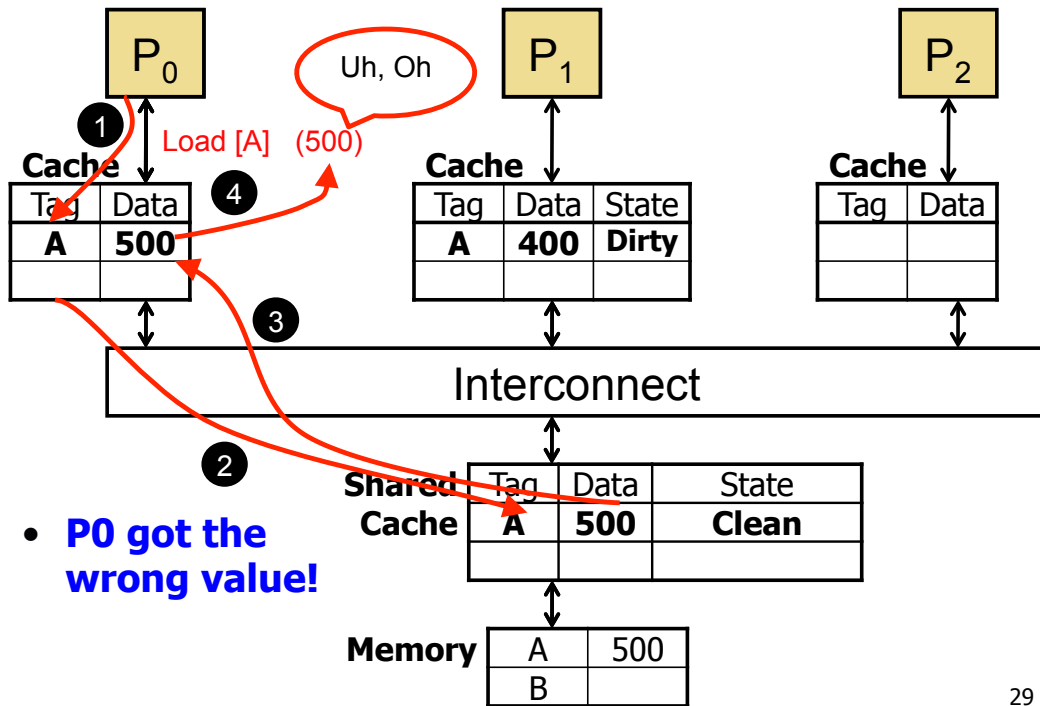
27

Private Cache Problem: Incoherence



28


Private Cache Problem: Incoherence



29

Cache Coherence: Who bears the brunt?

- Software
 - Caches are invisible to the programmer



EMM ROY

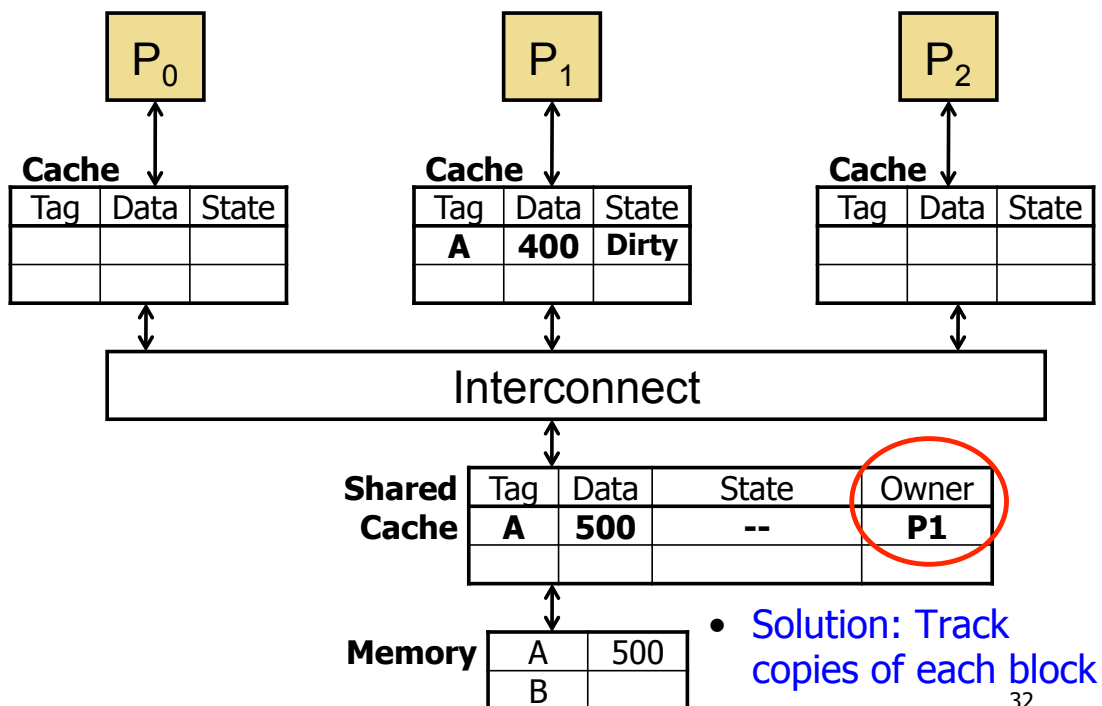
30

Cache Coherence: Who bears the brunt?

Let's do it in hardware

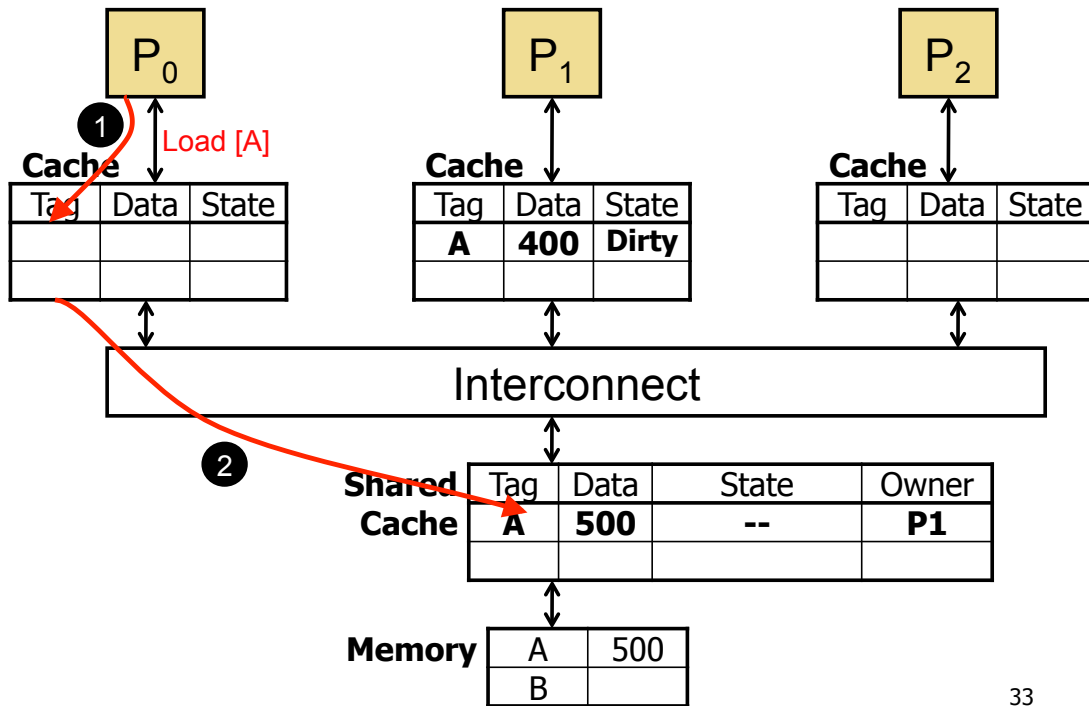
31

Rewind: Fix Problem by Tracking Sharers



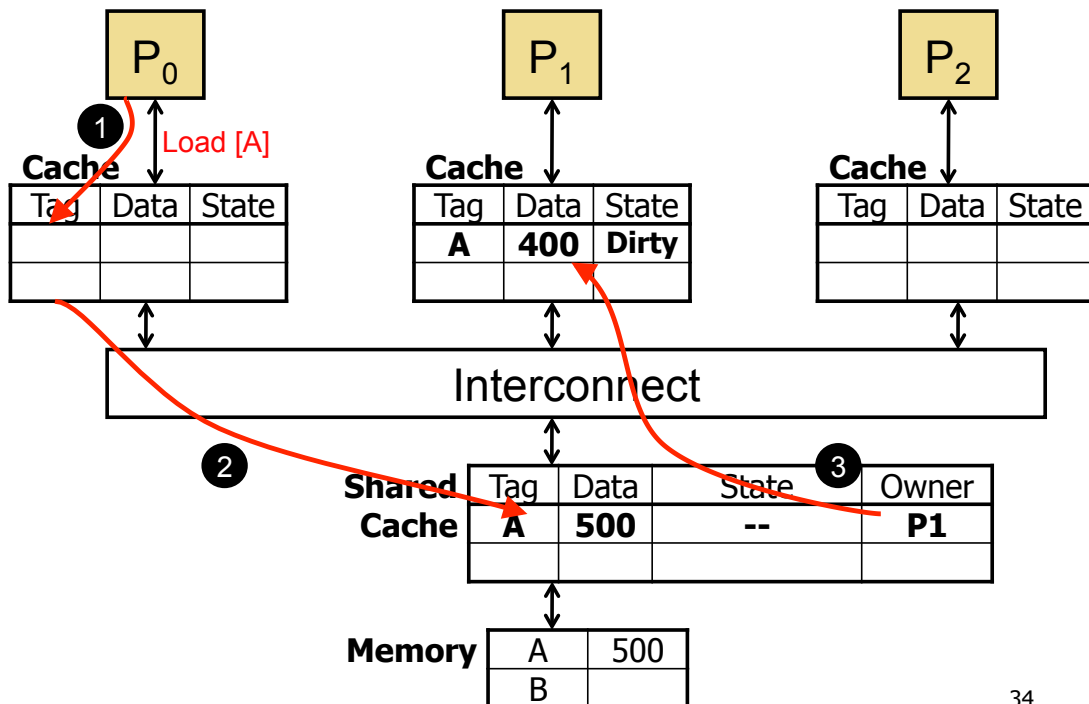
32

Use Tracking Information to "Invalidate"



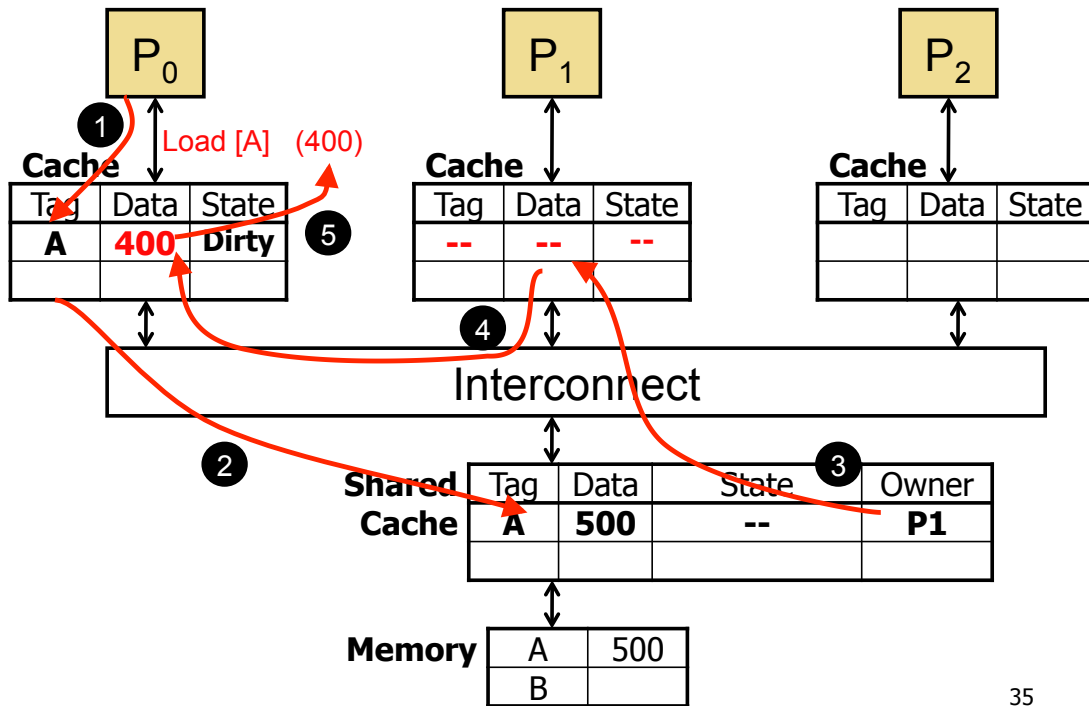
33

Use Tracking Information to "Invalidate"



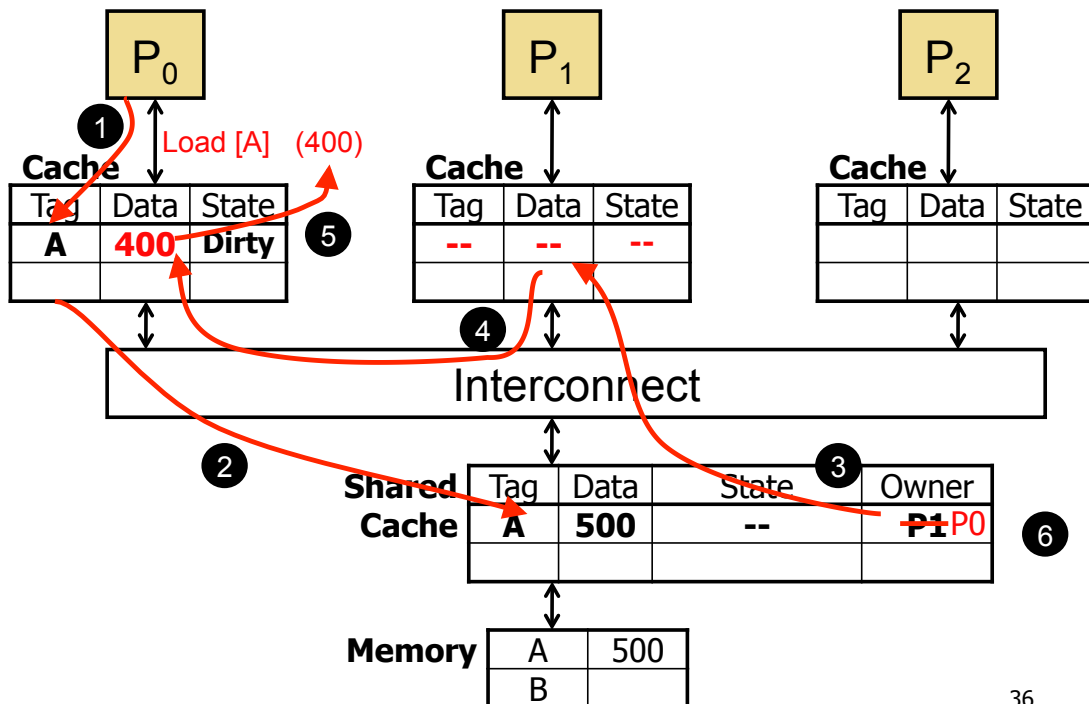
34

Use Tracking Information to "Invalidate"



35

Use Tracking Information to "Invalidate"




36

...This is "Valid/Invalid" Cache Coherence

- To enforce the shared memory invariant...
 - "Loads read the value written by the most recent store"
- Enforce the invariant...
 - **"At most one valid copy of the block"**
 - Simplest form is a **two-state "valid/invalid" protocol**
 - If a core wants a copy, must find and "invalidate" it
- On a cache miss, how is the valid copy found?
 - Option #1 **"Snooping"**: broadcast to all, whoever has it responds
 - Option #2: **"Directory"**: track sharers with separate structure

37

VI (MI) Coherence Protocol



- **VI (valid-invalid) protocol:**
 - Two states (per block in cache)
 - **V (valid)**: own the block
 - **I (invalid)**: don't own block
 - + Can implement with "valid bit"
- Protocol state transition (the left figure)
 - Summary
 - If anyone wants to read/write block
 - Give it up: transition to **I** state
 - Write-back if your own copy is dirty
- This is an **invalidate protocol**
- VI protocol is inefficient
 - Only one cached copy allowed in entire system
 - Multiple copies can't exist even if read-only
 - Not a problem in example
 - Big problem in reality

38

VI Protocol State Transition Table

State	<i>This Processor</i>		<i>Other Processor</i>	
	Load	Store	Load Miss	Store Miss
Invalid (I)	Load Miss → V	Store Miss → V	---	---
Valid (V)	Hit	Hit	Send Data → I	Send Data → I

- Rows are "states"
 - I vs V
- Columns are "events"
- Writeback events not shown
- Memory controller not shown
- **Memory sends data when no processor responds**

39

VI (MI) Coherence Inefficiency

- **VI (valid-invalid) protocol** is inefficient
 - Only one cached copy allowed in entire system
 - Multiple copies can't exist even if read-only
 - Not a problem in example
 - Big problem in reality

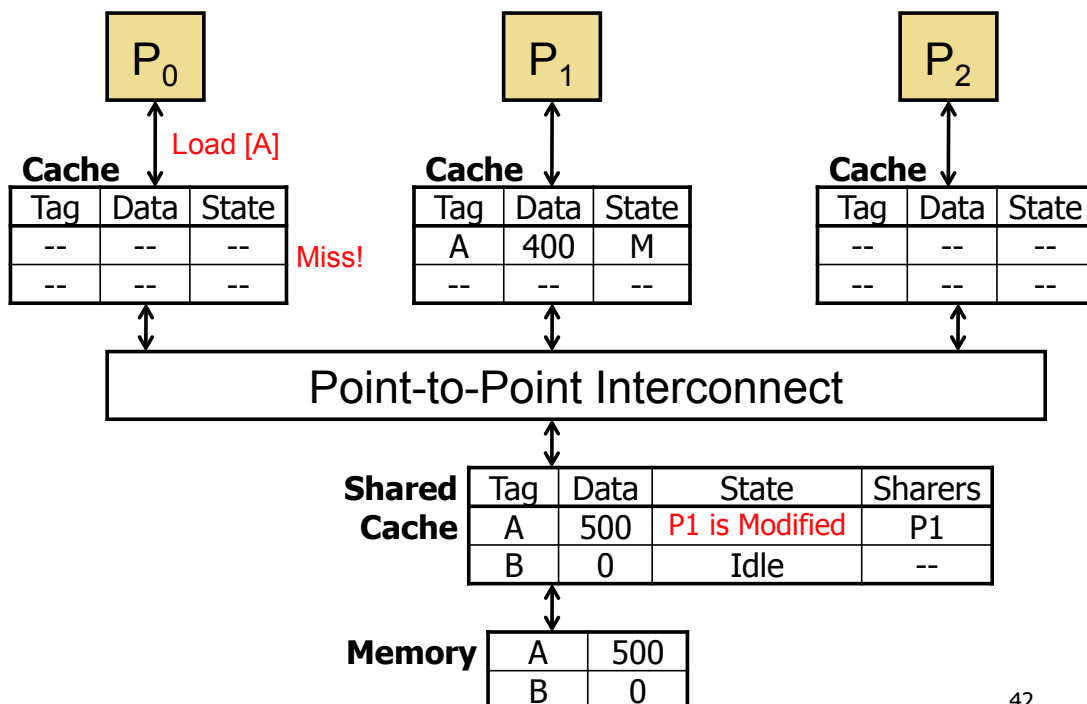
40

MSI Cache Coherence Protocol

- Solution: enforce the invariant...
 - **Multiple read-only copies** —OR—
 - **Single read/write copy**
- Track these MSI permissions (states) in per-core caches
 - **Modified (M):** read/write permission
 - **Shared (S):** read-only permission
 - **Invalid (I):** no permission
- Also track a **"Sharer" bit vector** in shared cache
 - One bit per core; tracks all shared copies of a block
 - Then, *invalidate all readers* when a write occurs
- Allows for many readers...
 - ...while still enforcing shared memory invariant ("Loads read the value written by the most recent store")

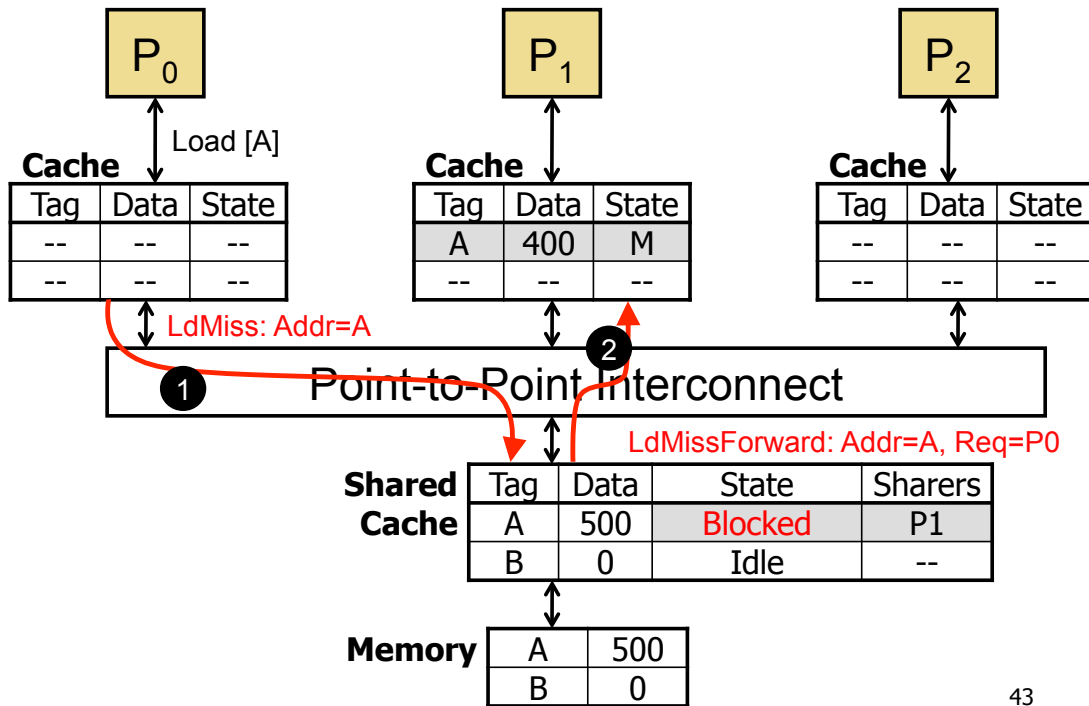
41

MSI Coherence Example: Step #1



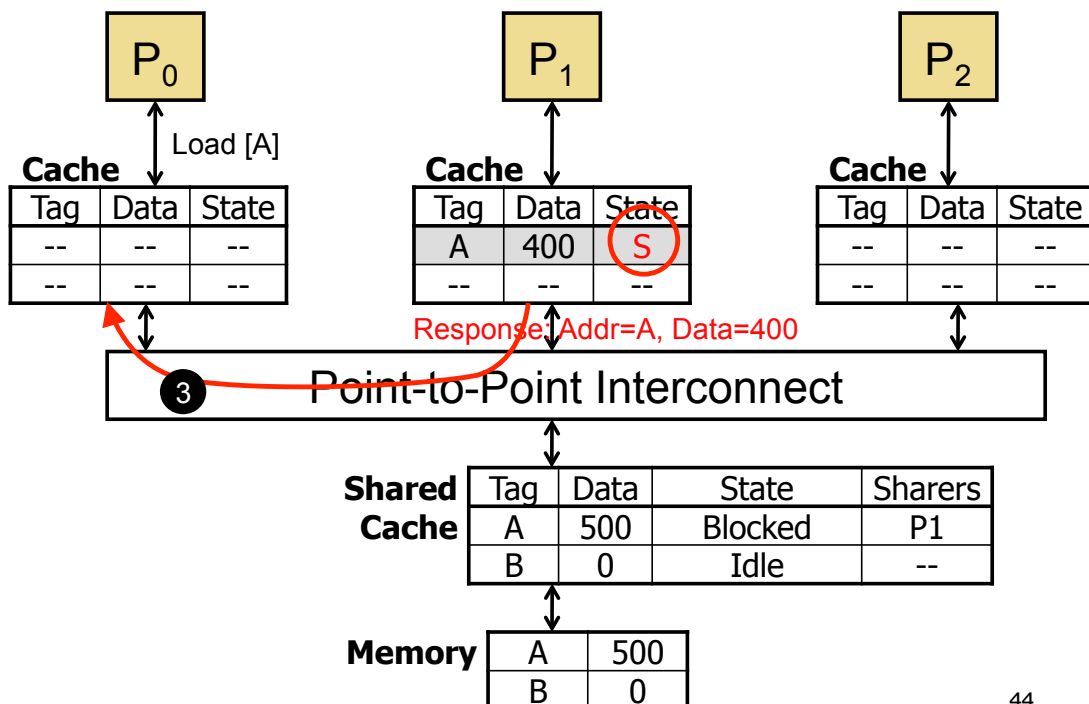
42

MSI Coherence Example: Step #2



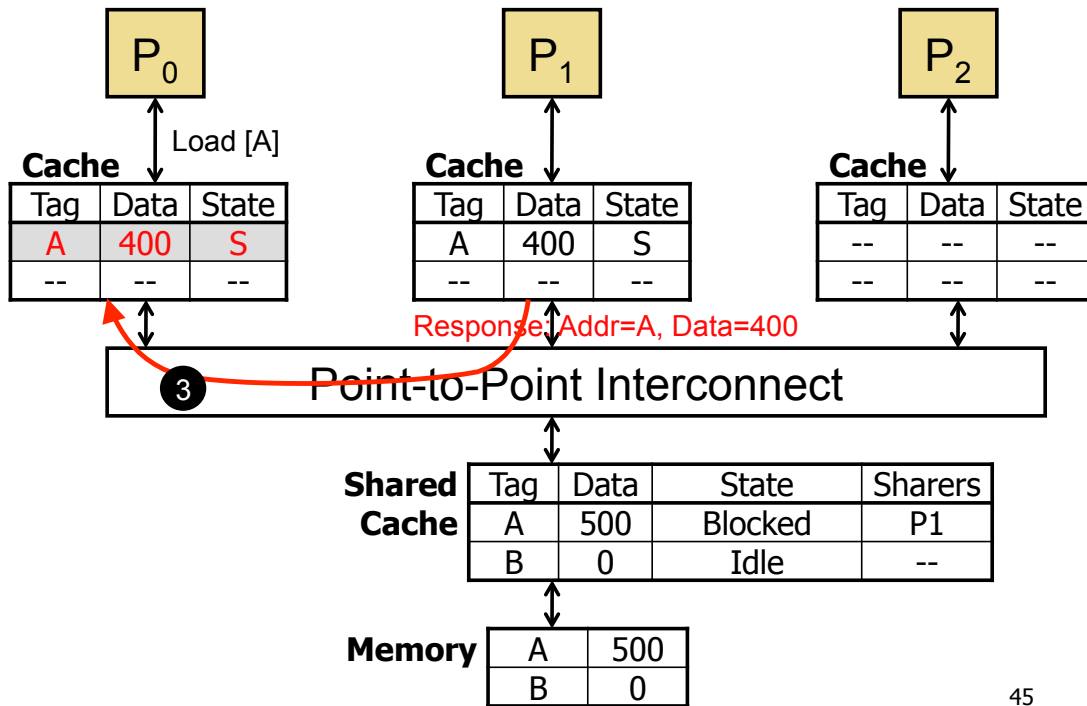
43

MSI Coherence Example: Step #3



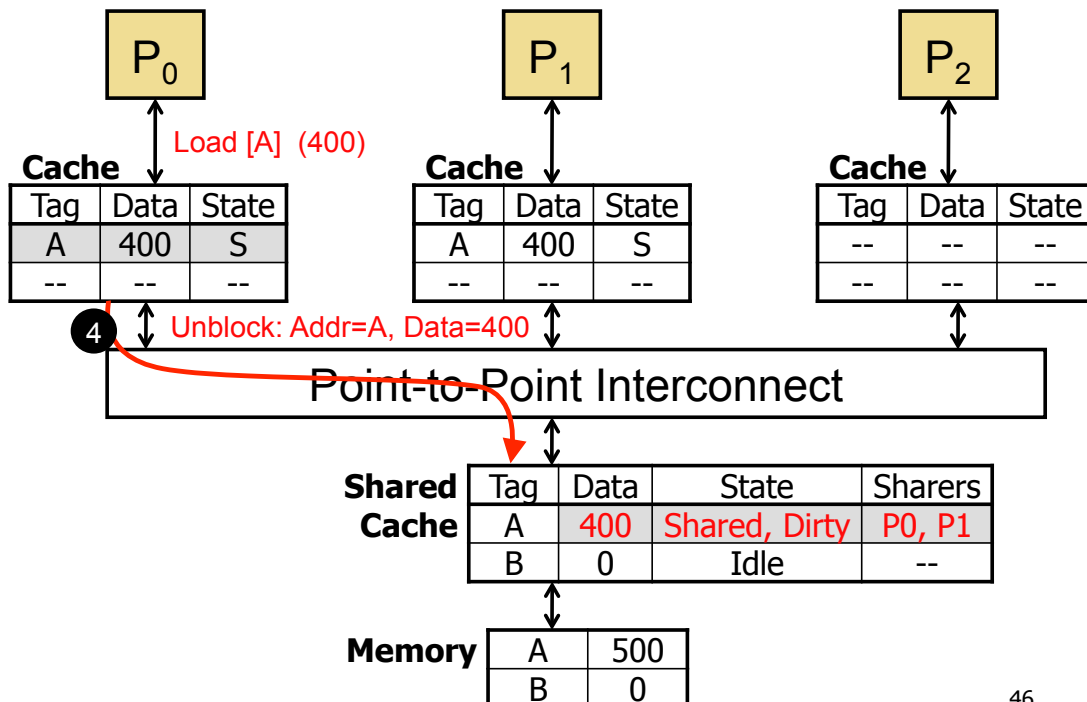
44

MSI Coherence Example: Step #4



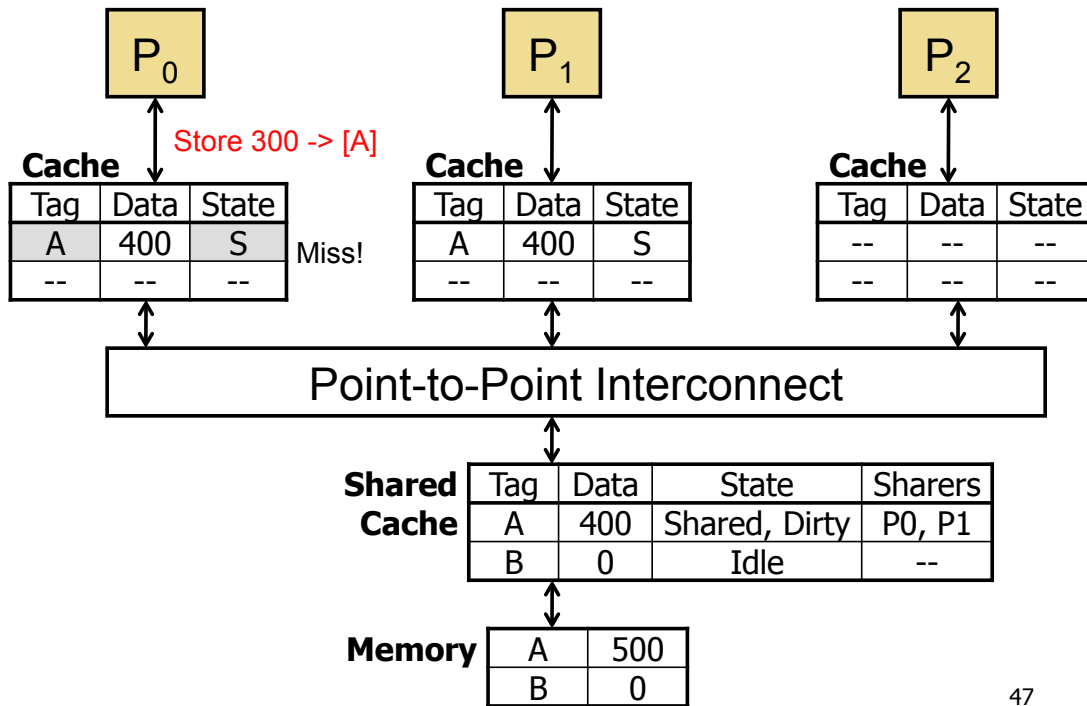
45

MSI Coherence Example: Step #5



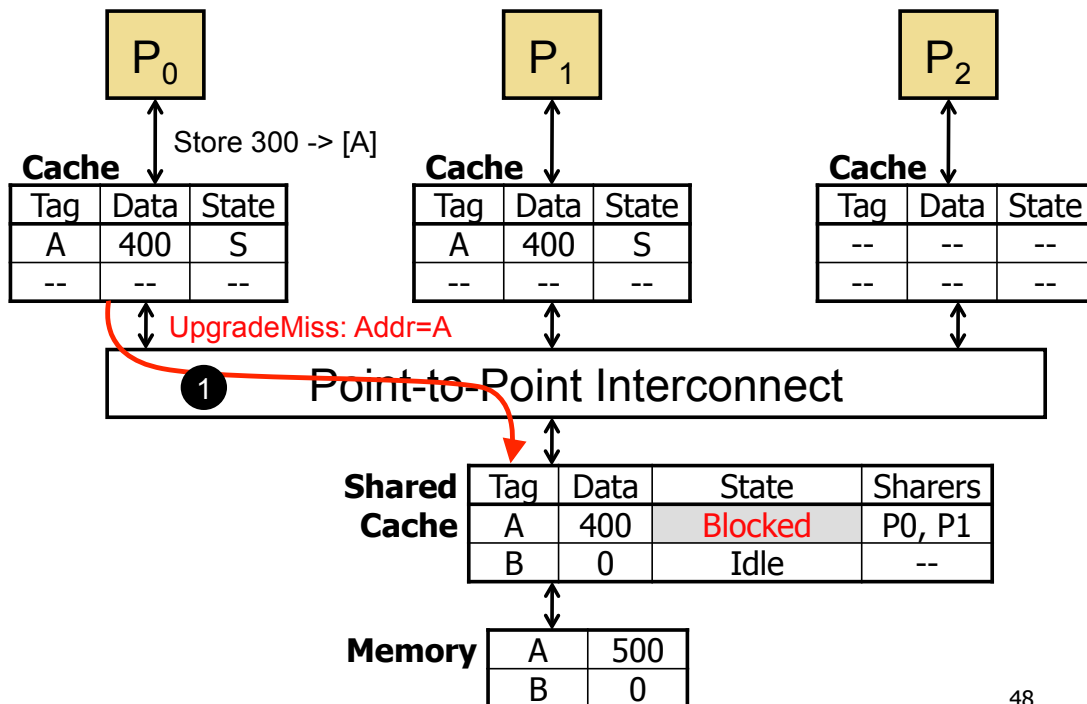
46

MSI Coherence Example: (store) Step #6



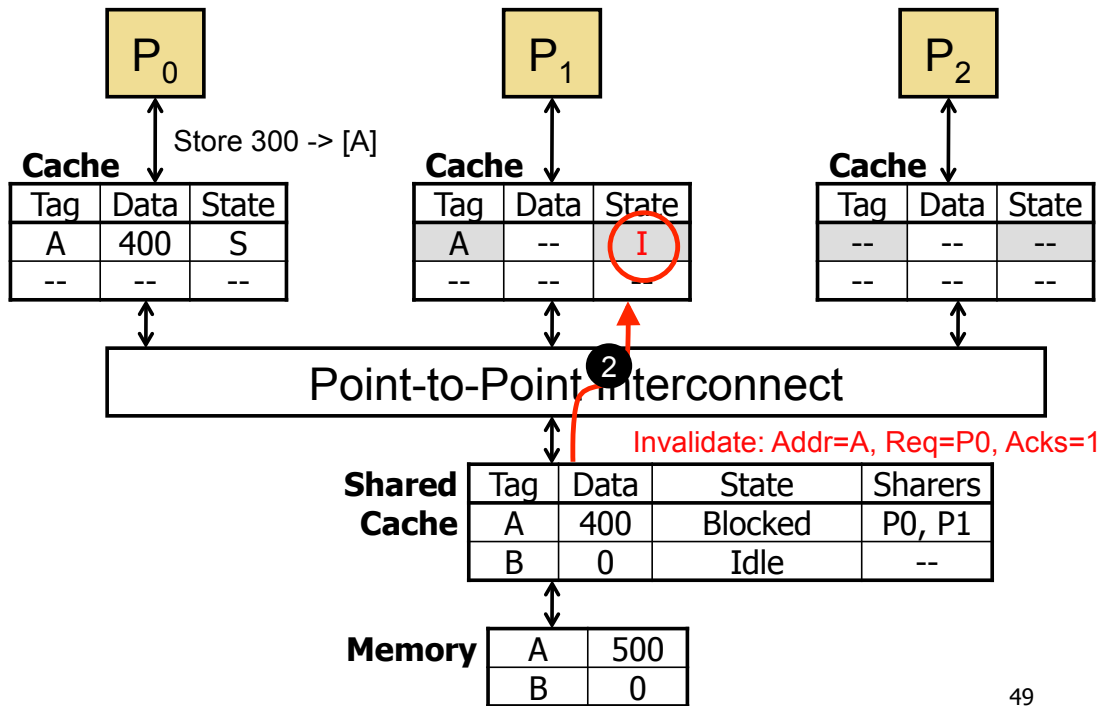
47

MSI Coherence Example: Step #7

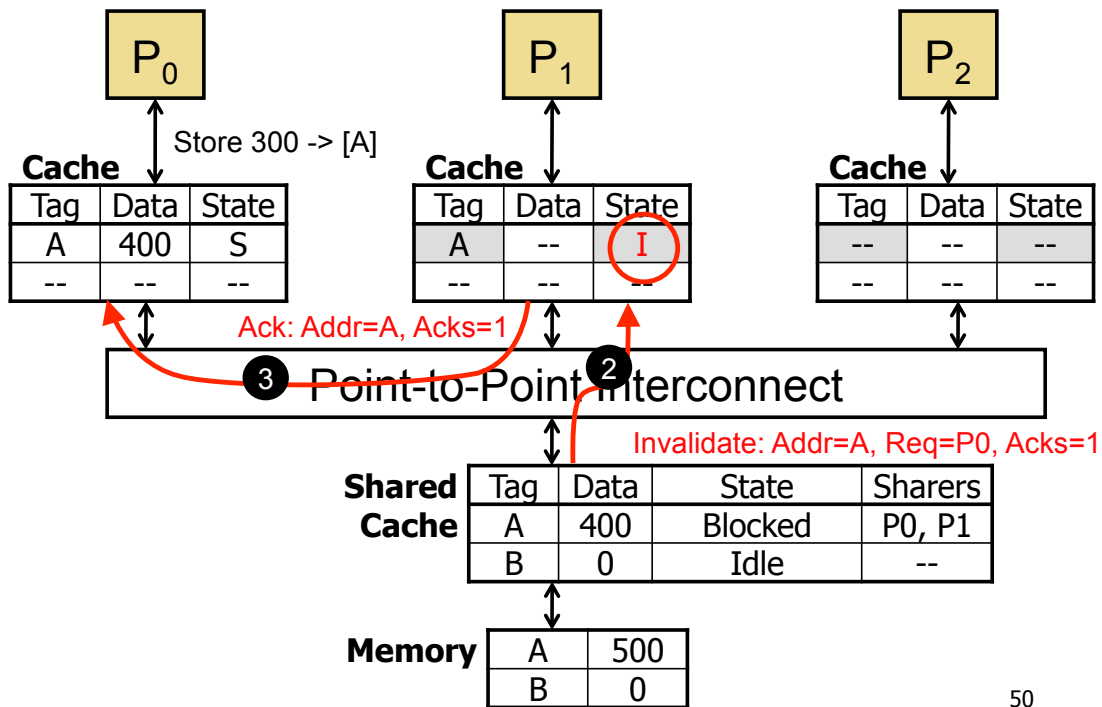


48

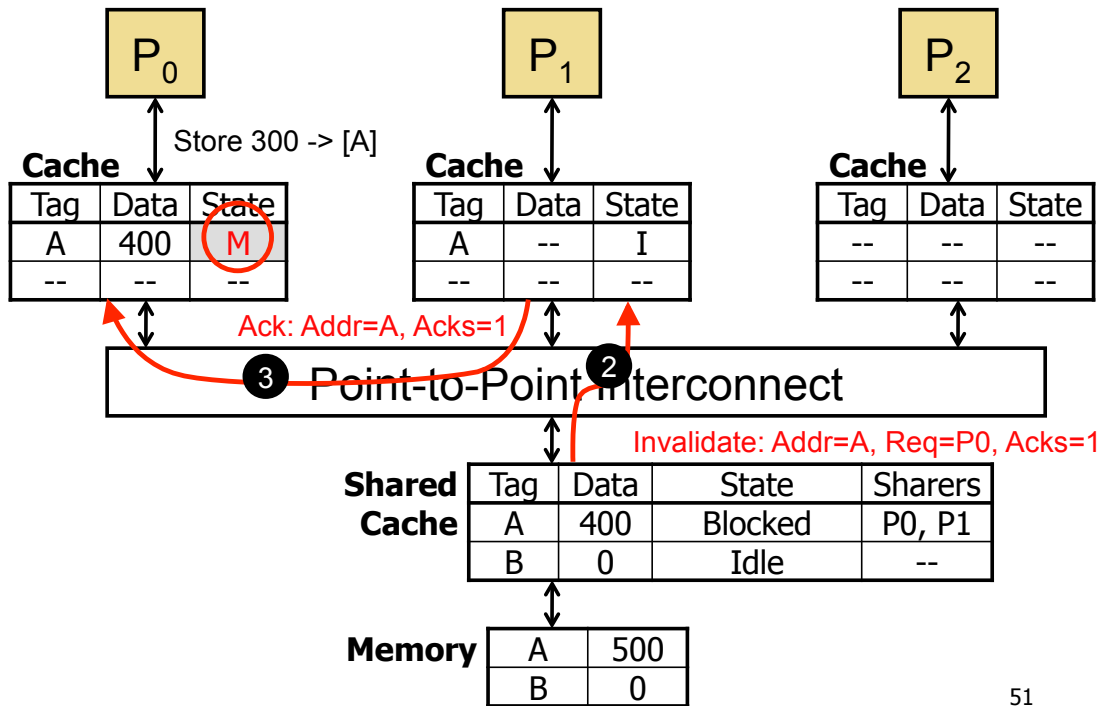
MSI Coherence Example: Step #8



MSI Coherence Example: Step #9

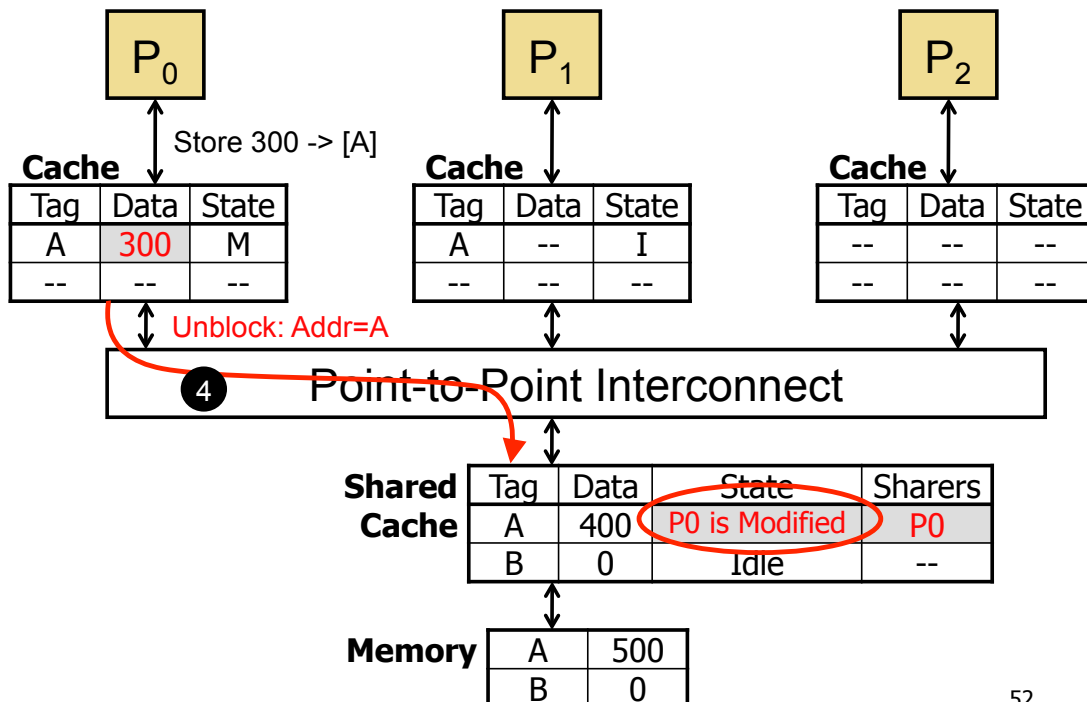


MSI Coherence Example: Step #10



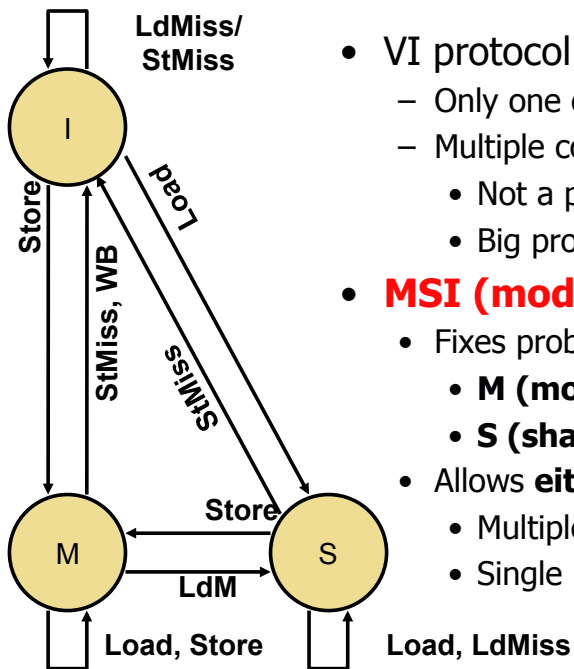
51

MSI Coherence Example: Step #11



52

MSI summary



- VI protocol is inefficient
 - Only one cached copy allowed in entire system
 - Multiple copies can't exist even if read-only
 - Not a problem in example
 - Big problem in reality
- **MSI (modified-shared-invalid)**
 - Fixes problem: splits "V" state into two states
 - **M (modified)**: local dirty copy
 - **S (shared)**: local clean copy
 - Allows **either**
 - Multiple read-only copies (S-state) **--OR--**
 - Single read/write copy (M-state)

53

MSI Protocol State Transition Table

State	This Processor		Other Processor	
	Load	Store	Load Miss	Store Miss
Invalid (I)	Load Miss → S	Store Miss → M	---	---
Shared (S)	Hit	Upgrade Miss → M	---	→ I
Modified (M)	Hit	Hit	Send Data → S	Send Data → I

- M → S transition also updates memory
- After which memory will respond (as all processors will be in S)

54

Classifying Misses: 3C Model

- Divide cache misses into three categories
 - **Compulsory (cold)**: never seen this address before
 - **Would miss even in infinite cache**
 - **Capacity**: miss caused because cache is too small
 - **Would miss even in fully associative cache**
 - Identify? Consecutive accesses to block separated by access to at least N other distinct blocks (N is number of frames in cache)
 - **Conflict**: miss caused because cache associativity is too low
 - Identify? **All other misses**
- One more case of cache misses
 - **(COHERENCE): MISS DUE TO EXTERNAL INVALIDATIONS**
 - **ONLY IN SHARED MEMORY MULTIPROCESSORS**