

CMPS 101

Homework Assignment 7

Solutions

- Let $N(n, h)$ denote the number of nodes at height $h \geq 0$ in an almost complete binary tree (ACBT) on $n \geq 0$ nodes. Prove that $N(n, h)$ satisfies the following double recurrence formula.

$$N(n, h) = \begin{cases} n - \lfloor n/2 \rfloor & \text{if } h = 0 \\ N(\lfloor n/2 \rfloor, h-1) & \text{if } h \geq 1 \end{cases}$$

(Hint: Let T be an ACBT with n nodes. First, show that T must have $n - \lfloor n/2 \rfloor$ leaves. Second, let $h \geq 1$ and consider the set of nodes in T of height h . Form a new ACBT T' by deleting all leaves from T . Argue that (1) T' has $\lfloor n/2 \rfloor$ nodes, and that (2) the nodes in T' of height $h-1$ are precisely the nodes in T of height h .)

Solution:

Let T be an ACBT on n nodes. If we index the nodes of T with integers in the range 1 to n , assigning 1 to the root, and descending level by level from left to right (just as in a heap), then we find that the parent of the node with index i has index $\lfloor i/2 \rfloor$. Therefore the parent of the rightmost leaf on the bottom level has index $\lfloor n/2 \rfloor$.

$$\underbrace{A_1 \ A_2 \ \cdots \ A_{\lfloor n/2 \rfloor}}_{\text{Internal Nodes}} \underbrace{A_{\lfloor n/2 \rfloor + 1} \ \cdots \ A_{n-1} \ A_n}_{\text{Leaves}}$$

This parent of the rightmost leaf must be the internal node of highest index. Thus T contains exactly $\lfloor n/2 \rfloor$ internal nodes, and hence $(\# \text{ of Leaves in } T) = n - \lfloor n/2 \rfloor$. But these leaves are precisely the nodes at height 0 in T , and so $N(n, 0) = n - \lfloor n/2 \rfloor$.

Now let $h \geq 1$ and consider the set of nodes in T of height h . Form a new ACBT T' by deleting all leaves from T . Observe that T' has $n - (n - \lfloor n/2 \rfloor) = \lfloor n/2 \rfloor$ nodes. Observe also that upon passing to T' , every internal node in T has had its height reduced by 1, since every descending path to a leaf was shortened by 1. Thus

$$\{ \text{Nodes in } T \text{ of height } h \} = \{ \text{Nodes in } T' \text{ of height } h-1 \},$$

and therefore

$$N(n, h) = N(\lfloor n/2 \rfloor, h-1),$$

as was claimed. ///

2. Let $N(n, h)$ be defined as in the previous problem. Prove that $N(n, h) \leq \left\lceil \frac{n}{2^{h+1}} \right\rceil$. (Hint: use the above recurrence formula together with (weak) induction on h starting at $h = 0$.)

Solution:

- I. Observe $n = \lfloor n/2 \rfloor + \lceil n/2 \rceil$ for any $n \geq 0$. Thus $N(n, 0) = n - \lfloor n/2 \rfloor = \lceil n/2 \rceil$, so the base case is satisfied.
- II. Let $h > 0$ and assume for any $m \geq 0$ that $N(m, h-1) \leq \left\lceil \frac{m}{2^h} \right\rceil$. Then

$$\begin{aligned}
 N(n, h) &= N(\lfloor n/2 \rfloor, h-1) && \text{by the recurrence formula for } N(n, h) \\
 &\leq \left\lceil \frac{\lfloor n/2 \rfloor}{2^h} \right\rceil && \text{by the induction hypothesis with } m = \lfloor n/2 \rfloor \\
 &\leq \left\lceil \frac{n/2}{2^h} \right\rceil && \text{since } \lfloor x \rfloor \leq x \text{ and } y \leq z \Rightarrow \lceil y \rceil \leq \lceil z \rceil \\
 &= \left\lceil \frac{n}{2^{h+1}} \right\rceil
 \end{aligned}$$

The result therefore holds for all $n \geq 0$ and $h \geq 0$.

///

3. 6.5-3 page 165

Write pseudocode for the procedures `HeapMinimum`, `HeapExtractMin`, `HeapDecreaseKey`, and `HeapInsert` that implement a min-priority queue with a min-heap.

Solution:

HeapMinimum(A) pre: $\text{heapSize}(A) \geq 1$

1. return $A[1]$

HeapExtractMin(A) pre: $\text{heapSize}(A) \geq 1$

1. $\text{min} = A[1]$
2. exchange $A[1]$ with $A[\text{heapSize}[A]]$
3. $\text{heapSize}[A] = \text{heapSize}[A] - 1$
4. `Heapify(A, 1)`
5. return min

HeapDecreaseKey(A, i, k) pre: $0 \leq i \leq \text{heapSize}[A]$

1. if $k < A[i]$
2. $A[i] = k$
3. while $i \geq 2$ and $A[\text{parent}(i)] > A[i]$
4. exchange $A[i]$ with $A[\text{parent}(i)]$
5. $i = \text{parent}(i)$

HeapInsert(A, k) pre: $\text{heapSize}[A] < \text{length}[A]$

1. $\text{heapSize}[A] = \text{heapSize}[A] + 1$
2. $A[\text{heapSize}[A]] = \infty$
3. `HeapDecreaseKey(A, heapSize[A], k)`

4. Let $G=(V, E)$ be a weighted directed graph and let $x \in V$. Suppose that after $\text{Initialize}(G, s)$ is executed, some sequence of calls to $\text{Relax}()$ causes $d[x]$ to be set to a finite value. Prove that G contains an s - x path of weight $d[x]$. (Hint: use induction on the number of calls to $\text{Relax}()$).

Proof: Let n denote the length of the relaxation sequence. If $n=0$, then the only d -value which is finite after Initialization is that of the source s . Indeed, G does contain an s - s path of weight $d[s]=0$, namely the trivial path. The base case is therefore verified.

Let $n > 0$, and assume for any vertex x , that if $d[x]$ achieves a finite value during a sequence of fewer than n relaxations, then there exists an s - x path in G of weight $d[x]$. Now let $y \in V$ and consider a sequence of n relaxations in which $d[y]$ becomes finite. An edge of the form (x, y) must have been relaxed during this sequence, for some vertex x . On that relaxation step, $d[y]$ was set to $d[x] + w(x, y)$. Since we suppose that this number is finite, $d[x]$ must have been finite before $\text{Relax}(x, y)$ was executed. Thus $d[x]$ became finite during a sequence of fewer than n relaxations, and by our induction hypothesis, there must exist an s - x path in G of weight $d[x]$. That path, followed by the edge (x, y) , constitutes an s - y path in G of weight $d[x] + w(x, y) = d[y]$. ///

5. 24.1-3 p. 654

Given a weighted directed graph $G=(V, E)$ with no negative-weight cycles, let m be the maximum over all vertices $x \in V$ of the minimum number of edges in a shortest path from the source $s \in V$ to x . (Here, the shortest path is by weight, not by the number of edges.) Suggest a simple change to the Bellman-Ford algorithm that allows it to terminate in $m+1$ passes, even if m is not known in advance.

Solution:

We reproduce Bellman-Ford here for reference:

BellmanFord(G, s)

1. $\text{Initialize}(G, s)$
2. for $i = 1$ to $|V| - 1$
3. for each edge $(u, v) \in E$
4. $\text{Relax}(u, v)$
5. for each edge $(u, v) \in E$
6. if $d[v] > d[u] + w(u, v)$
7. return false
8. return true

Note that we cannot simply alter the **for** statement on line 2 to say “for $i = 1$ to $m+1$ ”, since the value of m is not known ahead of time. Instead we modify Bellman-Ford so that loop 2-4 terminates as soon as one complete pass over the edge set results in no d -values being changed. Obviously no d -values will be changed by performing any further passes, so if we accept the correctness of Bellman-Ford (Lemma 24.2 and Theorem 24.4), the d and π values must be correct at that point. It remains

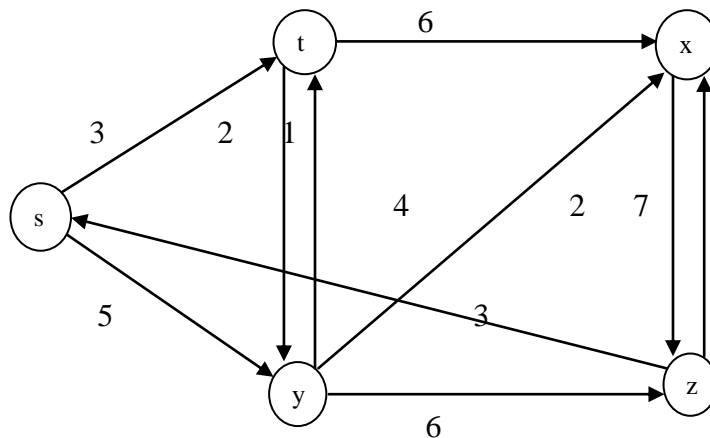
only to show that this rule causes loop 2-4 to terminate after $m+1$ passes. To prove this it is sufficient to show that the d -values are correct after exactly m passes. This follows from the path-relaxation property (Lemma 24.15) which says:

If $p = (v_0, v_1, \dots, v_k)$ is a shortest path from $s = v_0$ to v_k , and the edges of p are relaxed in the order (v_0, v_1) , (v_1, v_2) , \dots , (v_{k-1}, v_k) , then $d[v_k] = \delta(s, v_k)$. This property holds regardless of any other relaxation steps that occur, even if they are intermixed with relaxations of the edges of p .

Each of the edges (v_0, v_1) , (v_1, v_2) , \dots , (v_{k-1}, v_k) will be relaxed exactly once on each pass over the edge set, so k iterations of loop 2-4 suffice to correctly set the d -value of v_k . But by our definition of m , every vertex v (which is reachable from s) lies at the end of a shortest s - v path containing at most m edges, hence m iterations suffice to correctly set the d -values of all vertices in G . ///

6. 24.3-1 p. 662

Run Dijkstra's algorithm on the directed graph of Figure 24.2 p. 648 (pictured below), first using vertex s as the source and then using vertex z as the source. Show the d and π values and the vertices in set S after each iteration of the **while** loop.

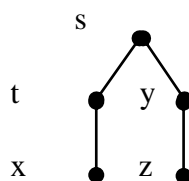


Solution:

With s as source:

	d/π -values after i -th iteration of while					
Vertex	0	1	2	3	4	5
s	0/ nil	0/ nil	0/ nil	0/ nil	0/ nil	0/ nil
t	∞ / nil	3/ s	3/ s	3/ s	3/ s	3/ s
x	∞ / nil	∞ / nil	9/ t	9/ t	9/ t	9/ t
y	∞ / nil	5/ s	5/ s	5/ s	5/ s	5/ s
z	∞ / nil	∞ / nil	∞ / nil	11/ y	11/ y	11/ y
Set S	\emptyset	{s}	{s, t}	{s, t, y}	{s, t, y, x}	{s, t, y, x, z}

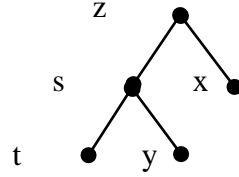
Predecessor subgraph:



With z as source:

	d/π -values after i -th iteration of while					
Vertex	0	1	2	3	4	5
s	∞/nil	3/z	3/z	3/z	3/z	3/z
t	∞/nil	∞/nil	6/s	6/s	6/s	6/s
x	∞/nil	7/z	7/z	7/z	7/z	7/z
y	∞/nil	∞/nil	8/s	8/s	8/s	8/s
z	0/nil	0/nil	0/nil	0/nil	0/nil	0/nil
Set S	\emptyset	{z}	{z, s}	{z, s, t}	{z, s, t, x}	{z, s, t, x, y}

Predecessor subgraph:



7. 24.3-6 p. 663

We are given a directed graph $G=(V,E)$ on which each edge $(u,v) \in E$ has an associated value $r(u,v)$, which is a real number in the range $0 \leq r(u,v) \leq 1$ that represents the reliability of a communication channel from vertex u to vertex v . We interpret $r(u,v)$ as the probability that the channel from u to v will not fail, and we assume that these probabilities are independent. Give an efficient algorithm to find the most reliable path between two given vertices.

Solution:

Let p be a directed x - y path consisting of vertices: $x=v_0, v_1, v_2, \dots, v_k=y$. Since the probabilities associated with each edge are independent, the probability that no edge along p fails is given by

$r(p) = \prod_{i=1}^k r(v_{i-1}, v_i) = r(v_0, v_1) \cdot r(v_1, v_2) \cdots r(v_{k-1}, v_k)$. The most reliable x - y path which we seek, is the

one that maximizes this quantity $r(p)$. Dijkstra's algorithm can be used to find this path by carefully defining an appropriate weight function on edges. Given $(u,v) \in E$, define $w(u,v) = -\log(r(u,v))$, where the \log function can have any base greater than 1. Since $0 \leq r(u,v) \leq 1$ we have $-\infty \leq \log(r(u,v)) \leq 0$, and hence $0 \leq w(u,v) \leq \infty$. Edge weights are therefore non-negative (and some may be infinite.) Running Dijkstra's algorithm on the source x will determine an x - y path which minimizes the quantity

$$\begin{aligned}
 w(p) &= \sum_{i=1}^k w(v_{i-1}, v_i) \\
 &= \sum_{i=1}^k -\log(r(v_{i-1}, v_i)) \\
 &= -\sum_{i=1}^k \log(r(v_{i-1}, v_i)) \\
 &= -\log\left(\prod_{i=1}^k r(v_{i-1}, v_i)\right) \\
 &= -\log(r(p)).
 \end{aligned}$$

But then p must maximize the quantity $\log(r(p))$, and since \log is an increasing function, the path p also maximizes $r(p)$ as required. The following algorithm determines the most reliable directed x - y path in G .

Max-Reliable(G, x, y, r)

1. for each $(u, v) \in E(G)$
2. $w(u, v) = -\log(r(u, v))$
3. Dijkstra(G, w, x)
4. PrintPath(G, x, y)

///