

Assignment 4 Design Document

CMPS 111, Spring 2016

1 Goals

The goal of this assignment is to use the FUSE file system framework to implement a simple FAT-based file system in user space. This assignment will be run in user space and therefore, will give us experience in extending operating system functionality with a user program.


Extending an operating system in such a way presents unique challenges, in that we must implement what is normally kernel logic without access to kernel level operations.

It is important that we learn how the operating system should interact with user space, since it restricts user programs to prevent them from crashing the machine or messing with any other programs that are in the kernel space, which teaches us about data protection.

2 Assumptions

We will be using a FAT (File Allocation Table) based file system and integrating it into FUSE; a user-space file system interface. For the disk structure, one can assume that the disk will be partitioned into 1 KB blocks. The file allocation table will occupy a number of these blocks. Since each block pointer in the FAT is 4 bytes, each block in the FAT can hold 256 block pointers, therefore we will need the number of blocks in the file systems divided by 256 FAT blocks. For the directories we will assume that each directory entry has a length of 64 bytes. The directory entries will carry the information specified in the assignment, but the only flag each directory entry will have to contain is the directory flag, which is a single bit that will indicate if the directory entry is for a regular file or a directory.

This is how we believe the fuse system works.



We are also assuming that if the amount of bytes to be read in `read()` goes over the amount actually present in the file, it is not an error and it will fill up the rest of the buffer with zeros.

3 Design

The first step to making the file system would be to create a disk image as specified in the assignment. It will contain the superblock, the file allocation table, and data blocks.

Superblock

The disk is divided into 1KB blocks and contains N blocks in total. The 0th block is reserved for the superblock, which will contain; the magic number, which is defined to be 0x2345beef, in the 0th word index, the total number of blocks in the file system (N) in the 1st word index, the number of blocks in the file allocation table (k or $1 + \frac{N}{256}$) in the 2nd word index, the block size which will be 1024 bytes for our file system in word index 3, and the starting block of the root directory in the last word index of 4. We will fill the rest of the super block with empty space.

File Allocation Table

The file allocation table will follow the superblock up until the kth block with k equaling $1 + \frac{N}{256}$. The first step to filling in the file allocation table would be to fill all the blocks that are already in use by the superblock and allocation table with a -1, indicating that they are invalid and can't be used. The root directory will be marked with a -2, which indicates that this block is the end of a file, and the root block is the one and only block needed for an empty file system. All the other blocks will be marked with zero because they are free. As information is written into the FAT the blocks will either point to other blocks that continue the file until it reaches an end of file, a -2. Any blocks that are in the file allocation table but aren't being used will be set to a -1.

Fuse:

To traverse through the file system we will be overriding Fuse operations with are own. We need to be able to use basic functionality and the following operations are the ones we need to implement to be able use them. Below are the main operations we are overriding and in the pseudocode we list the operations they are used for.

Operations we need to implement:

- Getattr
- Mkdir
- Unlink
- Rmdir
- Rename
- Open
- Read
- Write
- Release
- Readdir
- Opendir
- Main
- Create

4 Pseudocode

CreateDiskImage

Disk Image that will be created in the underlying file system. Split up into superblock, file allocation table, and file data (includes directory entries).

CreateDiskImage(imageName, nBlocks):

```
magicNumber = 0x2345beef
blockSize = 1024
nBlocksInFileAllocationTable = (nBlocks/(blockSize/4))+1
rootBlock = nBlocksInFileAllocationTable+1
fileDescriptor = Open imageName as write only, create if doesn't exist
//Creation of superblock
write(fileDescriptor, address of magicNumber, 4) //Write magic number 0x2345beef
write(fileDescriptor, address of nBlocks, 4) //Write number of Blocks to fileDescriptor
write(fileDescriptor, address of nBlocksInFileAllocationTable, 4)
write(fileDescriptor, address of blockSize, 4) //Write block size of 1024 to fileDescriptor\
write(fileDescriptor, address of rootBlock, 4)
//fill out rest of superblock
emptyBuffer = empty array of size blockSize-(4*5)
write(fileDescriptor, emptyBuffer, blockSize-20)
// Fill in file allocation table
allocationTable = array of size nBlocks
for i from 0 to rootBlock
    allocationTable[i] = -1 // mark blocks in table and superblock as invalid
allocationTable[rootBlock] = -2 // root block is occupied by root directory
for i from rootBlock+1 to nBlocks
    allocationTable[i] = 0 //mark all other blocks as free
write(fileDescriptor, allocationTable, nBlocks)
//Fill in rest of last block in file allocation table
buffer = array of size ((nBlocksInFileAllocationTable * blockSize) - nBlocks)
for i from 0 to length of buffer
    buffer[i] = -1 //mark rest of blocks invalid, since they don't exist
write(fileDescriptor, buffer, size of buffer)
// Fill out rest of disk image with empty data
buffer = empty array of size blockSize
for i from rootBlock to nBlocks
    write(fileDescriptor, emptyBuffer, blockSize)
```

Fuse Implementations

These are the following functions used in the implementation of fuse

Global Variables:

```
int diskImage
uint32_t nBlocks
uint32_t nBlocksInFileAllocationTable
uint32_t blockSize
int rootBlock
int * allocationTable
char * prevPath
```

Main(argc, argv):

```
if (argc < 2)
    error("Insufficient argument")
diskImage = open(argv[1], READ-WRITE)
if (open of diskImage failed)
    error("Unable to open: "+argv[1])
buffer = array of size 20 bytes (5 ints)
bytesRead=read(diskImage, buffer, 20)
if (bytesRead != 20 OR buffer[0] != 0x2345beef)
    error("Disk Image is not of proper format")
nBlocks = buffer[1]
nBlocksInFileAllocationTable = buffer[2]
blockSize = buffer[3]
rootBlock = buffer[4]
```

Struct DirectoryEntry:

```
Char filename [24]
Size_t creationTime
Size_t modificationTime
Size_t accessTime
Uint32_t fileLength
Int startBlock
Int flags
Int unused
```

Intermediate Functions

These following functions are used in the rest of the functions to traverse and do several operations in the file structure

FindOpenBlock()

Looks for an open space within the blocks. Checks to make sure there's enough space for a file to be written to, and returns the open block

FindOpenBlock():

```
    block = rootBlock
    while (allocationTable[block] != 0)
        block = block + 1
    If (block < nBlocks)
        byteArray [1024]
        for every entry in byteArray
            byteArray[count] = 0;
        result=lseek(diskImage, block * blockSize, SET)
        result=write(diskImage, buffer, blockSize)
        return block
    else error("no empty blocks")
```

TraverseDirectoryTree

Goes through each of the blocks to find the block and file data specified by path

TraverseDirectoryTree(path):

```
    if((length of string is 1 and is equal to "/") == 0)
        if(offset isn't null)
            return root directory
    block = rootBlock
    path = trim leading slash from path
    node = leading node of path
    directoryEntry = null
    while (path isn't null) //when path is null, we have found the start block of designated file
        offset = lseek(diskImage, blockSize * block, SET)
        buffer = array of size blockSize
```

```

result=read(diskImage, buffer, blockSize)
if (result < 0)
    error("Reading from disk image failed")
    return null
i=0
while (i < blockSize) //find the given path node in the current block
    directoryEntry=GetDirectoryEntry(address of buffer[i])
    if (directoryEntry exists AND node equals directoryEntry.fileName)
        if (directoryEntry.flags == DIRECTORY)
            block = directoryEntry.startBlock
            break loop
        else
            error("Not a directory")
            return null
    i = i + 64
if (i >= blockSize AND allocationTable[block] > 0) //reached end of this block, but
                                                    there's another
    block = allocationTable[block]
else if (i < blockSize)
    path = trim leading slash from path, if no slash set to null
    node = leading node of path
else
    error("Path not found")
    return null
return directoryEntry

```

WriteAllocationTabletoDisk()
Writes the allocation table to disk

```

WriteAllocationTableToDisk():
    offset=lseek(diskImage, blockSize, SEEK_SET)
    result=write(diskImage, allocationTable, nBlocks*4)
    if (result < 0)
        error("Writing to disk image failed")

```

Directory Entry Functions

These following functions are used to read and write directory entries

//will receive a sequence of bytes and organize them into a DirectoryEntry structure

GetDirectoryEntry(array):

```
    dirEntry = new DirectoryEntry
    //read filename string into structure
    for i from 0 to 24
        dirEntry.fileName[i] = array[i]
    dirEntry.creationTime = *(size_t *)array[24]
    dirEntry.modificationTime = *(size_t *)array[32]
    dirEntry.accessTime = *(size_t *)array[40]
    dirEntry.fileLength = *(uint32_t *)array[48]
    dirEntry.startBlock = *(int *)array[52]
    dirEntry.flags = *(int *)array[56]
    dirEntry.unused = *(int *)array[60]
    return dirEntry
```

WriteDirectoryEntryToBuffer

Writes the directory entry struct to the buffer

WriteDirectoryEntryToBuffer(dirEntry, byteBuffer):

```
    intBuffer = (int *)byteBuffer
    longBuffer = (size_t *)byteBuffer
    uintBuffer = (uint32_t *)byteBuffer
    for i from 0 to 24
        byteBuffer[i] = dirEntry.fileName[i]
    longBuffer[3] = dirEntry.creationTime
    longBuffer[4] = dirEntry.modificationTime
    longBuffer[5] = dirEntry.accessTime
    uintBuffer[12] = dirEntry.fileLength
    intBuffer[13] = dirEntry.startBlock
    intBuffer[14] = dirEntry.flags
    intBuffer[15] = dirEntry.unused
```

removeDirectoryEntry

Removes the directory entry specified by path

removeDirectoryEntry (Path):

```
    check = traverseDirectoryTree(path)
    Previous = offset in block of check's directory entry
    Next = offset in block of entry after check
    currentEntry = check
    currentBlock = block containing currentEntry's directory entry
    done = false
    Buffer = array of size blockSize
    read currentBlock into memory
    while (not done)
        tempBuf = array of size 64
        if (next < blockSize)
            nextDirectoryEntry = getDirectoryEntry(address of buffer[next])
            if (nextDirectoryEntry exists)
                Write directory entry to tempBuf
            else
                Fill tempBuf with zeros
                done = true
            write tempBuf to block at offset previous
            previous = next
            next = next + 64
        else if (allocationTable[currentBlock] > 0)
            next = 0
            nextBlock = allocationTable[currentBlock]
            nextbuffer = array of size blockSize
            read nextBlock's disk block into nextbuffer
            nextDirectoryEntry = getDirectoryEntry(address of nextbuffer[next]);
            if(nextDirectoryEntry exists)
                Write directory entry to tempBuf
            else
                Fill tempBuf with zeros
                done = 1
            write tempBuf to current block at offset previous
            previous = 0
            next = next + 64
            currentBlock = nextBlock
```


buffer = nextbuffer

Directory Functions

These following functions are used to be able to traverse and read through directories

mkdir()

Creates a new directory by creating a new directory entry and writing to the disk used for mkdir operation.

mkdir(path):

```
origLoc = 0
directoryEntry = TraverseDirectoryTree(path except trailing node) //get directory that will
                                                                    contain our new directory

block = newdirectoryEntry.startBlock
while (allocationTable[block] > 0) //get to block where we can write
    block = allocationTable[block]
offset = lseek(diskImage, blockSize * block, SEEK_SET)
buffer = array of size blockSize
result = read(diskImage, buffer, blockSize)
if (result < 0)
    error "can't read diskImage"
read(diskImage, buffer, blockSize)
count = 0
while (buffer[count] isn't empty AND count < blockSize) //get to free space in block
    count = count + 64
if (count >= blockSize) //there is no free space in block, make a new one
    allocationTable[block] = findOpenBlock()
    block = allocationTable[block]
    allocationTable[block] = -2
    buffer = empty buffer of size blockSize
    count = 0

dirEntry = new DirectoryEntry
dirEntry.fileName = node
dirEntry.creationTime = time of now
dirEntry.modificationTime = dirEntry.creationTime
dirEntry.accessTime = dirEntry.creationTime
dirEntry.fileLength = 0
dirEntry.startBlock = findOpenBlock()
```

```

    dirEntry.flags = DIRECTORY
    allocationTable[dirEntry.startBlock] = -2
    WriteAllocationTableToDisk()
    WriteDirectoryEntryToBuffer(dirEntry, address of buffer[count])
    offset = lseek(diskImage, blockSize * block, SEEK_SET)
    result = write(diskImage, buffer, blockSize)
    If the result is less than zero
        Throw a "can't read diskImage" error
        return -1
newDir.fileLength = new.fileLength + 64
newDir.modificationTime = time of now
newDir.accessTime = newDir.modificationTime

If (origLoc > 0)
    tempBuf = DirectoryEntry
    writeDirectoryEntryToBuffer(newDir, tempBuf)
    offset = lseek(diskImage, origLoc, SEEK_SET)
    result = write(diskImage, tempBuf, DirectoryEntry)
return 0

```

opendir()

Opens directory and shouldn't have to check for permissions. Used for the cd operation

```

opendir(path, fi):
    Print "opendir"
    If (path matches "..")
        Path = prevpath
    Int origloc
    Directoryentry check = traverseDirectoryTree(path, &origLoc);
    If (check is null AND check.flags does not equal IS_DIRECTORY)
        Return -ENOENT
    Check.accesstime = time of now;
    If (origloc > 0)
        Char tempBuf
        writeDirectoryEntryToBuffer(check, tempbuf)
        Offset = lseek(diskImage, origLoc, SEEK_SET);
        Result = write(diskImage, tempBuf, sizeof(struct DirectoryEntry));
    Prevpath = dirname(path);
    Return 0

```

readdir()

Reads directory entries and used for the ls operation

```
readdir(path,buf,filler,offset,fi)
    If (“.” == path) path = prevPath
    DirectoryEntry * check = TraverseDirectoryTree(path)
    If (check == NULL || check.flags != Directory)
        Return -enoent
    filler(buf, ".", NULL, 0);
    filler(buf, "..", NULL, 0);

    block = check.startBlock
    done = 0
    While (done == 0)
        offset = lseek(diskImage, blockSize * block, SET)
        buf = array of size blockSize
        result=read(diskImage, buf, blockSize)
        if (result < 0)
            error("Reading from disk image failed")
            return EROENT

        i=0
        while (i < blockSize) //find the given path node in the current block
            directoryEntry=GetDirectoryEntry(address of buf[i])
            filler(buf, directoryEntry.fileName, NULL, 0);
            i = i + 64
        if (i >= blockSize AND allocationTable[block] > 0) //reached end of this block, but
                                                                there's another
            block = allocationTable[block]
        Else done = 1;

    return 0
```

rmdir()

Removes directory and all files in it and updates the path. Used to rm operation for directories

```
rmdir(path):
    DirectoryEntry* check = TraverseDirectoryTree(path)
    filename = path of names of files in this directory
```

```

block = check.startBlock
While (done == 0)
    offset = lseek(diskImage, blockSize * block, SET)
    buffer = array of size blockSize
    result=read(diskImage, buffer, blockSize)
    if (result < 0)
        error("Reading from disk image failed")
        return ENOENT
    i=0
    Done = 0
    while (i < blockSize) //find the given path node in the current block
        directoryEntry=GetDirectoryEntry(address of buffer[i])
        if(directoryEntry.flags == DIRECTORY)
            rmdir(path+directoryEntry.filename)
        Else
            unlink(path+directoryEntry.filename)
        i = i + 64
    if (i >= blockSize AND allocationTable[block] > 0) //reached end of this block, but
                                                                    there's another
        block = allocationTable[block]
    Else done = 1;
removeDirectoryEntry(check, check.startBlock)
WriteAllocationTableToDisk()

```

File Functions

These following functions are used to be able to do all the operations to traverse and manipulate files

create()

Creates and opens a file used for any creating file operations

```
*****
create(path, mode, fi)
    origLoc= 0
    DirectoryEntry * check = TraverseDirectoryTree(path)
    If (check != NULL)
        Return open(path,fi)

    directoryEntry = TraverseDirectoryTree(path except trailing node) //get directory that will
                                                                    contain our new directory

    block = directoryEntry.startBlock
    while (allocationTable[block] > 0) //get to block where we can write
        block = allocationTable[block]
    offset = lseek(diskImage, blockSize * block, SET)
    buffer = array of size blockSize
    read(diskImage, buffer, blockSize)
    i=0
    while (buffer[i] isn't empty AND i < blockSize) //get to free space in block
        i = i + 64
    if (i >= blockSize) //there is no free space in block, make a new one
        allocationTable[block] = FindOpenBlock()
        block = allocationTable[block]
        allocationTable[block] = -2
        buffer = empty buffer of size blockSize
        i = 0
    dirEntry = new DirectoryEntry
    dirEntry.fileName = node
    dirEntry.creationTime = TODO: get time of now
    dirEntry.modificationTime = dirEntry.creationTime
    dirEntry.accessTime = dirEntry.creationTime
    dirEntry.fileLength = 0
    dirEntry.startBlock = FindOpenBlock()
    dirEntry.flags = 0
```

```

allocationTable[dirEntry.startBlock] = -2
WriteAllocationTableToDisk()
WriteDirectoryEntryToBuffer(dirEntry, address of buffer[i])
offset = lseek(diskImage, blockSize * block, SET)
write(diskImage, buffer, blockSize)
WriteAllocationTableToDisk()
Return open(path,fi)

```

open()

File open operation

open(path, fi):

```

//Set the open bit and then write the directory entry
origLoc = 0
DirectoryEntry * check = TraverseDirectoryTree(path)
If (check == NULL || check.flags == IS_DIRECTORY)
    return -enoent
Check.flags = OPEN //Set the open flag on
check.accessTime = null
if(origLoc > 0)
    tempBuf = size of Directory entry
    writeDirectoryEntrytoBuffer(check, tempBuf)
    offset=lseek(diskImage,origLoc,SEEK_SET)
    result=write(diskImage,tempBuf,
Return 0

```

release()

Releases an open file

release(path, fi):

```

//Set the open bit and then write the directory entry
origLoc = 0
DirectoryEntry * check = TraverseDirectoryTree(path)
If (check == NULL || Check.flags != OPEN)
    return -enoent
Check.flags = CLOSE //Set the open flag on
writeDirectoryEntrytoBuffer(check, check.startblock)
Return 0

```

read()

Reads data from an open file and returns exactly the numbers of bytes requested except on an end of file or error

```
read (path, buffer, size, offset, fi):
    origLoc = 0
    DirectoryEntry * check = TraverseDirectoryTree(path)
    If (check == NULL || Check.flags != IS_OPEN)
        for (int count = 0; count < size; count++) buf[count] = 0;
        return -enoent
    If (offset > check .fileLength)
        for (int count = 0; count < size; count++) buf[count] = 0;
        return 0;
    bytesRead = 0
    currentBlock = check.startBlock
    offsetCount = offset
    while (offsetCount > blockSize AND allocationTable[currentBlock] > 0)
        offsetCount = offsetCount - blockSize
        currentblock = allocationTable[currentBlock]
    if (offsetCount > blockSize)
        for (int count = 0; count < size; count++) buf[count] = 0;
        return 0
    while (bytesRead < size AND bytesRead+offset < check.fileLength)
        oft=lseek(diskImage, currentBlock * blockSize+offsetCount, SEEK_SET)
        result=read(diskImage, buffer+bytesRead, blockSize-offsetCount)
        offsetCount = 0
        bytesRead = bytesRead + result
    if(currentBlock == -2) break
        currentBlock = allocationTable[currentBlock]
    If (bytesRead < size)
        For (int count = bytesRead; count < size; count++) buf[count] = 0;
    Return bytesRead
```

write()

Writes data to open file and returns the exact number of bytes requested except on an error

write (path, buffer, size, offset, fi):

DirectoryEntry * check = TraverseDirectoryTree(path)

If (check == NULL || Check.flags != OPEN)

Return -enoent //If write breaks write for case for file doesn't exist

bytesWritten = 0

currentBlock = check.startBlock

offsetCount = offset

While (offsetCount > blockSize)

offsetCount = offsetCount - blockSize

allocationTable[block] = FindOpenBlock()

currentBlock = allocationTable[block]

allocationTable[block] = -2

while (bytesWritten < size AND bytesWritten+offset < blockSize)

result=lseek(diskImage, currentBlock * blockSize+offsetCount, SET)

result=write(diskImage, buffer+bytesWritten, blockSize-offsetCount)

offsetCount = 0

bytesWritten += result

allocationTable[block] = FindOpenBlock()

currentBlock = allocationTable[block]

allocationTable[block] = -2

If (bytesWritten < size)

Return error //Error for max size

WriteAllocationTableToDisk()

Return bytesWritten

getattr()

Gets the attributes of the file

getattr(const char* path, struct stat *st):

DirectoryEntry *file = TraverseDirectoryTree(path)

If (file == NULL)

Return -ENOENT

st.st_atime = file.accessTime


```

st.st_ctime = file.creationTime
st.st_mtime = file.modificationTime
st.st_size = file.fileLength
st.st_mode = DEFFILEMODE
if(file.flags is IS_DIRECTORY) //checks to see if flag is set to directory or not
    st.st_mode |= S_IFDIR
else
    st.st_mode |= S_IFREG
st.st_nlink = 1
return 0

```

rename()

Renames a file

```

rename(path, newName)
    origLoc= 0
    DirectoryEntry* check = TraverseDirectoryTree(path)
    If (check == NULL)
        Return -ERONENT
    check.fileName = newName
    //write directory entry to disk

```

unlink()

Removes(unlinks) a file Used for the rm functions to remove files

```

unlink(path):
    DirectoryEntry* check = TraverseDirectoryTree(path)
    Int nextblock = 0;
    Int currentblock = check.startblock
    While (nextblock != -2)
        nextblock = allocationTable[currentblock]
        allocationTable[currentblock] = 0
        Currentblock = nextblock
    removeDirectoryEntry(check, check.startBlock)
    WriteAllocationTableToDisk()

```

