

Introduction to Operating Systems

Midterm Examination

Name:_____ E-mail:_____

Winter 2014

Part I: Multiple Choice – Mark your answer (2 points each)

1. Multiprogramming is:
 - (a) Multiple processes running at different times.
 - (b) Multiple processes residing in memory and running more or less at the same time.
 - (c) Multiple processes running at exactly the same time.
2. Which has the fastest access times:
 - (a) Cache.
 - (b) Registers.
 - (c) Main memory.
3. Which are all system calls:
 - (a) `chmod()`, `chdir()`, `read()`, and `main()`
 - (b) `waitpid()`, `lsdir()`, `mount()`, and `open()`
 - (c) `fork()`, `exit()`, `lseek()`, and `kill()`
4. Which is true about processes and threads:
 - (a) Threads in a process share the same stack.
 - (b) Threads in a process share the same file descriptors.
 - (c) Threads in a process share the same register values.
5. A race condition is when:
 - (a) Multiple processes are all trying to finish first.
 - (b) A process runs too slow.
 - (c) The correctness of the code depends upon the timing of the execution.
6. Mutual exclusion is when:
 - (a) A set of processes are prevented from running because higher priority processes keep getting in ahead of them.
 - (b) A set of processes is waiting for an event that only another process in the set can cause.
 - (c) A set of processes are prevented from simultaneously accessing a shared data structure.
7. *Semaphores, locks and condition variables, and monitors* support two distinct operations:
 - (a) Parallelization and event notification.
 - (b) Synchronization and event notification.
 - (c) Synchronization and parallelization.
8. Which is true:
 - (a) Round-robin has the highest turn-around time.
 - (b) Shortest remaining time next is the fairest scheduling algorithm.
 - (c) Priority-based scheduling is the most general.

9. The difference between preemptive and non-preemptive scheduling is:
 - (a) Whether or not a ready process can be involuntarily terminated from the ready state.
 - (b) Whether or not a running process is involuntarily removed from the running state.
 - (c) Whether or not a blocked process can have its resources involuntarily taken away from it.
10. The four conditions that must hold in order for deadlock to occur are:
 - (a) Mutual exclusion, hold-and-wait, circular wait, and no preemption.
 - (b) No preemption, no starvation, circular wait, and mutual exclusion.
 - (c) Hold-and-wait, circular wait, no starvation and mutual exclusion.
11. Four strategies for dealing with deadlock are:
 - (a) Conservative resource allocation through the Banker's algorithm, spooling, two-phase locking, and mutual exclusion.
 - (b) Do nothing, kill one of the processes, pre-emptible resources, and spooling.
 - (c) Two-phase locking, resource ordering, resource discovery, and do nothing.
12. Which is false:
 - (a) Memory compaction is time consuming.
 - (b) Degree of multiprogramming refers to the amount of time a process spends waiting for I/O.
 - (c) Swapping can result in external fragmentation.
13. The main advantage of paging is:
 - (a) That it doesn't require any extra hardware support.
 - (b) The efficiency with which data can be brought into memory.
 - (c) That only the current working set of the process need actually be in memory at any one time.
14. The main advantage of multilevel page tables is that they:
 - (a) Use page table memory efficiently.
 - (b) Reduce the complexity of the paging system.
 - (c) Speed up page table references.
15. Which are both Stack algorithms:
 - (a) Optimal and LRU.
 - (b) LRU and FIFO.
 - (c) FIFO and Optimal.

Part II: Short answer – Write a paragraph or two (no more) (5 points each)

1. Explain the difference between processes and threads. What are the advantages of using threads (rather than processes) in implementing a complex application?
2. List two different ways a higher-priority process could be given a larger fraction of the CPU time. Assume processes get to run for a fixed-length quantum: the only freedom you have is in choosing which process to run for the next quantum.
3. The five (basic) states that a process can be in are: *new*, *ready*, *blocked*, *running* and *terminated*. Describe the transitions among these states, and state when they occur.

Hint: The best way to do this is to draw a diagram.

4. Recall that a *monitor* is a method for controlling access to a critical section (like a semaphore, but a higher level abstraction). Briefly describe the monitor abstraction, and do not forget to consider the case when a process may have to wait *inside* the monitor (condition variables).

Hint: You may find it helpful to draw a picture.

5. Explain the differences between a *safe* state, an *unsafe* state, and a *deadlocked* state. For each of the following transitions, briefly explain what causes a system to make that transition, or if such a transition is impossible, why is impossible: *safe* → *unsafe*; *unsafe* → *deadlocked*; *safe* → *deadlocked*.
6. Using the reference string $\langle 2\ 1\ 0\ 3\ 2\ 1\ 4\ 2\ 1\ 0\ 3\ 4 \rangle$, fill in the two tables below (representing systems with *three* and *four* page frames, respectively) using the FIFO page replacement policy. Indicate when a page fault occurs, and note how many page faults occur. Do you notice any anomalies?

| | | | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |

| | | | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |

7. Using the same reference string as in the previous question, fill in the following two tables as before using the LRU page replacement policy. Indicate when a page fault occurs, and note how many page faults occur.

| | | | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |

| | | | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |

8. The LRU page replacement algorithm has the property that $\mathcal{M}(r, m) \subseteq \mathcal{M}(r, m + 1)$. What is “Belady’s anomaly” and why does an algorithm with this property preclude it from happening?
9. The *shortest job first* (SJF) scheduling CPU algorithm has a definite benefit, but also has a serious problem. The benefit is that it is provably optimal. In what way is it considered optimal? What is the serious problem?
10. Consider the *working set* model of virtual memory. In a multiprocess environment, what should be done if the working set of the resident processes cannot be kept in main memory?