

Deadlocks



But crises and deadlocks when they occur have at least this advantage, that they force us to think.

— *Jawaharlal Nehru*

Overview

- ❖ Resources
- ❖ Why do deadlocks occur?
- ❖ Dealing with deadlocks
 - Ignoring them: ostrich algorithm
 - Detecting & recovering from deadlock
 - Avoiding deadlock
 - Preventing deadlock

Resources

- ❖ Resource: something a process uses
 - Usually limited (at least somewhat)
- ❖ Examples of computer resources
 - Printers
 - Semaphores / locks
 - Memory
 - Tables (in a database)
- ❖ Processes need access to resources in reasonable order
- ❖ Two types of resources:
 - Preemptable resources: can be taken away from a process with no ill effects
 - Nonpreemptable resources: will cause the process to fail if taken away

Using resources

- ❖ Sequence of events required to use a resource
 - Request the resource
 - Use the resource
 - Release the resource
- ❖ Can't use the resource if request is denied
 - Requesting process has options
 - Block and wait for resource
 - Continue (if possible) without it: may be able to use an alternate resource
 - Process fails with error code
 - Some of these may be able to prevent deadlock...


When do deadlocks happen?

❖ Suppose

- Process 1 holds resource A and requests resource B
- Process 2 holds B and requests A
- Both can be blocked, with neither able to proceed

❖ Deadlocks occur when ...

- Processes are granted exclusive access to devices or software constructs (resources)
- Each deadlocked process needs a resource held by another deadlocked process



DEADLOCK!

What is a deadlock?

- ❖ Formal definition:

“A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.”
- ❖ Usually, the event is release of a currently held resource
- ❖ In deadlock, none of the processes can
 - Run
 - Release resources
 - Be awakened

Four conditions for deadlock

❖ Mutual exclusion

- Each resource is assigned to at most one process

❖ Hold and wait

- A process holding resources can request more resources

❖ No preemption


- Previously granted resources cannot be forcibly taken away

❖ Circular wait

- There must be a circular chain of 2 or more processes where each is waiting for a resource held by the next member of the chain



Resource allocation graphs



- ❖ Resource allocation modeled by directed graphs
- ❖ Example 1:
 - Resource R assigned to process A
- ❖ Example 2:
 - Process B is requesting / waiting for resource S
- ❖ Example 3:
 - Process C holds T, waiting for U
 - Process D holds U, waiting for T
 - C and D are in deadlock!


Dealing with deadlock

- ❖ How can the OS deal with deadlock?
 - Ignore the problem altogether!
 - Hopefully, it'll never happen...
 - Detect deadlock & recover from it
 - Dynamically avoid deadlock
 - Careful resource allocation
 - Prevent deadlock
 - Remove at least one of the four necessary conditions
- ❖ We'll explore these tradeoffs


Getting into deadlock

A

Acquire R
Acquire S
Release R
Release S




Acquire R




Acquire S

B

Acquire S
Acquire T
Release S
Release T




Acquire S




Acquire T

C

Acquire T
Acquire R
Release T
Release R



Acquire T



The Ostrich Algorithm

- ❖ Pretend there's no problem
- ❖ Reasonable if
 - Deadlocks occur very rarely
 - Cost of prevention is high
- ❖ UNIX and Windows take this approach
 - Resources (memory, CPU, disk space) are plentiful
 - Deadlocks over such resources rarely occur
 - Deadlocks typically handled by rebooting
- ❖ Trade off between convenience and correctness


Not getting into deadlock...

- ❖ Many situations may result in deadlock (but don't have to)
 - In previous example, A could release R before C requests R, resulting in no deadlock
 - Can we always get out of it this way?
- ❖ Find ways to:
 - Detect deadlock and reverse it
 - Stop it from happening in the first place

Detecting deadlocks using graphs

- ❖ Process holdings and requests in the table and in the graph (they're equivalent)
- ❖ Graph contains a cycle \Rightarrow deadlock!
 - Easy to pick out by looking at it (in this case)
 - Need to mechanically detect deadlock
- ❖ Not all processes are deadlocked (A, C, F not in deadlock)

Process	Holds	Wants
A	R	S
B		T
C		S
D	U	S,T
E	T	V
F	W	S
G	V	U



Deadlock detection algorithm

- ❖ General idea: try to find cycles in the resource allocation graph
- ❖ Algorithm: depth-first search at each node
 - Mark arcs as they're traversed
 - Build list of visited nodes
 - If node to be added is already on the list, a cycle exists!
- ❖ Cycle \Rightarrow deadlock

```
For each node N in the graph {  
  Set L = empty list  
  unmark all arcs  
  Traverse (N,L)  
}  
If no deadlock reported by now, there isn't  
any  
  
define Traverse (C,L) {  
  If C in L, report deadlock!  
  Add C to L  
  For each unmarked arc from C {  
    Mark the arc  
    Set A = arc destination  
    /* NOTE: L is a  
       local variable */  
    Traverse (A,L)  
  }  
}
```

Recovering from deadlock

❖ Recovery through preemption

- Take a resource from some other process
- Depends on nature of the resource and the process

❖ Recovery through rollback

- Checkpoint a process periodically
- Use saved state to restart the process if it's in deadlock
- May present a problem if the process affects lots of “external” things

❖ Recovery through killing processes

- Crudest but simplest way to break a deadlock: kill one of the processes in the deadlock cycle
- Other processes can get its resources
- Try to choose a process that can be rerun from the start
 - Pick one that hasn't run too far already

Preventing deadlock

- ❖ Deadlock can be completely prevented!
- ❖ Ensure that at least one of the conditions for deadlock never occurs
 - Mutual exclusion
 - Circular wait
 - Hold & wait
 - No preemption
- ❖ Not always possible...

Eliminating mutual exclusion


- ❖ Some devices (such as printer) can be spooled
 - Only the printer daemon uses printer resource
 - This eliminates deadlock for printer
- ❖ Not all devices can be spooled
- ❖ Principle:
 - Avoid assigning resource when not absolutely necessary
 - As few processes as possible actually claim the resource

Attacking “hold and wait”

- ❖ Require processes to request resources before starting
 - A process never has to wait for what it needs
- ❖ This can present problems
 - A process may not know required resources at start of run
 - This also ties up resources other processes could be using
 - Processes will tend to be conservative and request resources they might need
- ❖ Variation: a process must give up all resources before making a new request
 - Process is then granted all prior resources as well as the new ones
 - Problem: what if someone grabs the resources in the meantime—how can the process save its state?


Attacking “no preemption”

- ❖ This is not usually a viable option
- ❖ Consider a process given the printer
 - Halfway through its job, take away the printer
 - Confusion ensues!
- ❖ May work for some resources
 - Forcibly take away memory pages, suspending the process
 - Process may be able to resume with no ill effects



Attacking “circular wait”

- ❖ Assign an order to resources
- ❖ Always acquire resources in numerical order
 - Need not acquire them all at once!
- ❖ Circular wait is prevented
 - A process holding resource n can't wait for resource m if $m < n$
 - No way to complete a cycle!
 - Place processes above the highest resource they hold and below any they're requesting
 - All arrows point up!



What does FreeBSD do?

❖ What resources are at issue?

- Locks & semaphores: one holder at a time!
- Physical resources: not typically a concern
 - There are millions of (interchangeable) pages
 - Limited resources (*e.g.*, printer) only used for a short time and then released

❖ Goal: prevent deadlock

- Mechanism must be low-cost (fast & efficient)
- Mechanism must be simple & flexible

❖ Two basic approaches

- Locks / semaphores: use resource ordering
 - Must acquire resources in “group” order
- Printers and other “limited” hardware resources: manage with a single process
 - No hold & wait: processes that want to use them queue up requests
 - Manager process is the only one to directly use the device

Deadlock prevention: summary

Condition	Prevented by
Mutual exclusion	Spool everything
Hold and wait	Request all resources initially
No preemption	Take resources away if there's not a complete set
Circular wait	Order resources numerically

Example: two-phase locking

❖ Phase One

- Process tries to lock all data it needs, one at a time
- If needed data found locked, start over
(no real work done in phase one)

❖ Phase Two


- Perform updates
- Release locks

❖ Note similarity to requesting all resources at once

❖ This is often used in databases

❖ What condition does this avoid?

Resource trajectories



Two process resource trajectories

Handling resources with multiple instances

- ❖ Previous algorithm only works if there's one instance of each resource
- ❖ If there are multiple instances of each resource, we need a different method
 - Track current usage and requests for each process
 - To detect deadlock, try to find a scenario where all processes can finish
 - If no such scenario exists, we have deadlock

Deadlock detection algorithm

	A	B	C	D
Avail	2	3	0	1

Hold

Process	A	B	C	D
1	0	3	0	0
2	1	0	1	1
3	0	2	1	0
4	2	2	3	0

Want

Process	A	B	C	D
1	3	2	1	0
2	2	2	0	0
3	3	5	3	1
4	0	4	1	1

```
current = avail;
for (j = 0; j < N; j++) {
    for (k=0; k<N; k++) {
        if (finished[k])
            continue;
        if (want[k] < current) {
            finished[k] = 1;
            current += hold[k];
            break;
        }
    }
    if (k==N) {
        printf "Deadlock!\n";
        // finished[k]==0 means process
        // is in the deadlock
        break;
    }
}
```

Note: want[j],hold[j],current,avail are arrays!

Safe and unsafe states

	Has	Max
A	3	9
B	2	4
C	2	7
Free: 3		

	Has	Max
A	3	9
B	4	4
C	2	7
Free: 1		

	Has	Max
A	3	9
B	0	-
C	2	7
Free: 5		

	Has	Max
A	3	9
B	0	-
C	7	7
Free: 0		

	Has	Max
A	3	9
B	0	-
C	0	-
Free: 7		

Demonstration that the first state is safe

	Has	Max
A	3	9
B	2	4
C	2	7
Free: 3		

	Has	Max
A	4	9
B	2	4
C	2	7
Free: 2		

	Has	Max
A	4	9
B	4	4
C	2	7
Free: 0		

	Has	Max
A	4	9
B	0	-
C	2	7
Free: 4		

Demonstration that the second state is unsafe

Banker's Algorithm for a single resource

Has Max		
A	0	6
B	0	5
C	0	4
D	0	7
Free: 10		

Any sequence finishes

Has Max		
A	1	6
B	1	5
C	2	4
D	4	7
Free: 2		

C,B,A,D finishes

Has Max		
A	1	6
B	2	5
C	2	4
D	4	7
Free: 1		

Deadlock (unsafe state)

- ❖ Bankers' algorithm: before granting a request, ensure that a sequence exists that will allow all processes to complete
 - Use previous methods to find such a sequence
 - If a sequence exists, allow the requests
 - If there's no such sequence, deny the request
- ❖ Can be slow: must be done on each request!

Banker's Algorithm with multiple resources

	Process	Tape drives	Plotters	Scanners	DVD-ROMs
A	3	0	1	1	
B	0	1	0	0	
C	1	1	1	0	
D	1	1	0	1	
E	0	0	0	0	

Resources assigned

	Process	Tape drives	Plotters	Scanners	DVD-ROMs
A	1	1	0	0	
B	0	1	1	2	
C	3	1	0	0	
D	0	0	1	0	
E	2	1	1	0	

Resources still needed

Total resources:	6	3	4	2
Already claimed:	5	3	2	2
Available:	1	0	2	0

Starvation

- ❖ Algorithm to allocate a resource
 - Give the resource to the shortest job first
- ❖ Works great for multiple short jobs in a system
- ❖ May cause long jobs to be postponed indefinitely
 - Even though not blocked
- ❖ Solution
 - First-come, first-serve policy
- ❖ Starvation can lead to deadlock
 - Process starved for resources can be holding resources
 - If those resources aren't used and released in a timely fashion, shortage could lead to deadlock

Livelock

- ❖ Sometimes, processes can still run, but not make progress
- ❖ Example: two processes want to use resources A and B
 - P0 gets A, P1 gets B
 - Each realizes that a deadlock will occur if they proceed as planned!
 - P0 drops A, P1 drops B
 - P0 gets B, P1 gets A
 - Same problem as before
 - This can go on for a very long time...
- ❖ Real-world example: Ethernet transmission collisions
 - If there's a "collision" on the wire, wait and try again
 - Multiple processes waited the exact same amount of time...