

Memory Management



Memory is merely the process of tuning into vibrations that have been left behind in space and time.

- Michio Kushi, *Spiritual Journey: Michio Kushi's Guide to Endless Self-Realization and Freedom* (with Edward Esko)

Memory management

- ❖ Basic memory management
- ❖ Swapping
- ❖ Virtual memory
- ❖ Page replacement algorithms
- ❖ Modeling page replacement algorithms
- ❖ Design issues for paging systems
- ❖ Implementation issues
- ❖ Segmentation

In an ideal world...

- ❖ The ideal world has memory that is
 - Very large
 - Very fast
 - Non-volatile (doesn't go away when power is turned off)
- ❖ The real world has memory that is:
 - Very large
 - Very fast
 - Affordable!

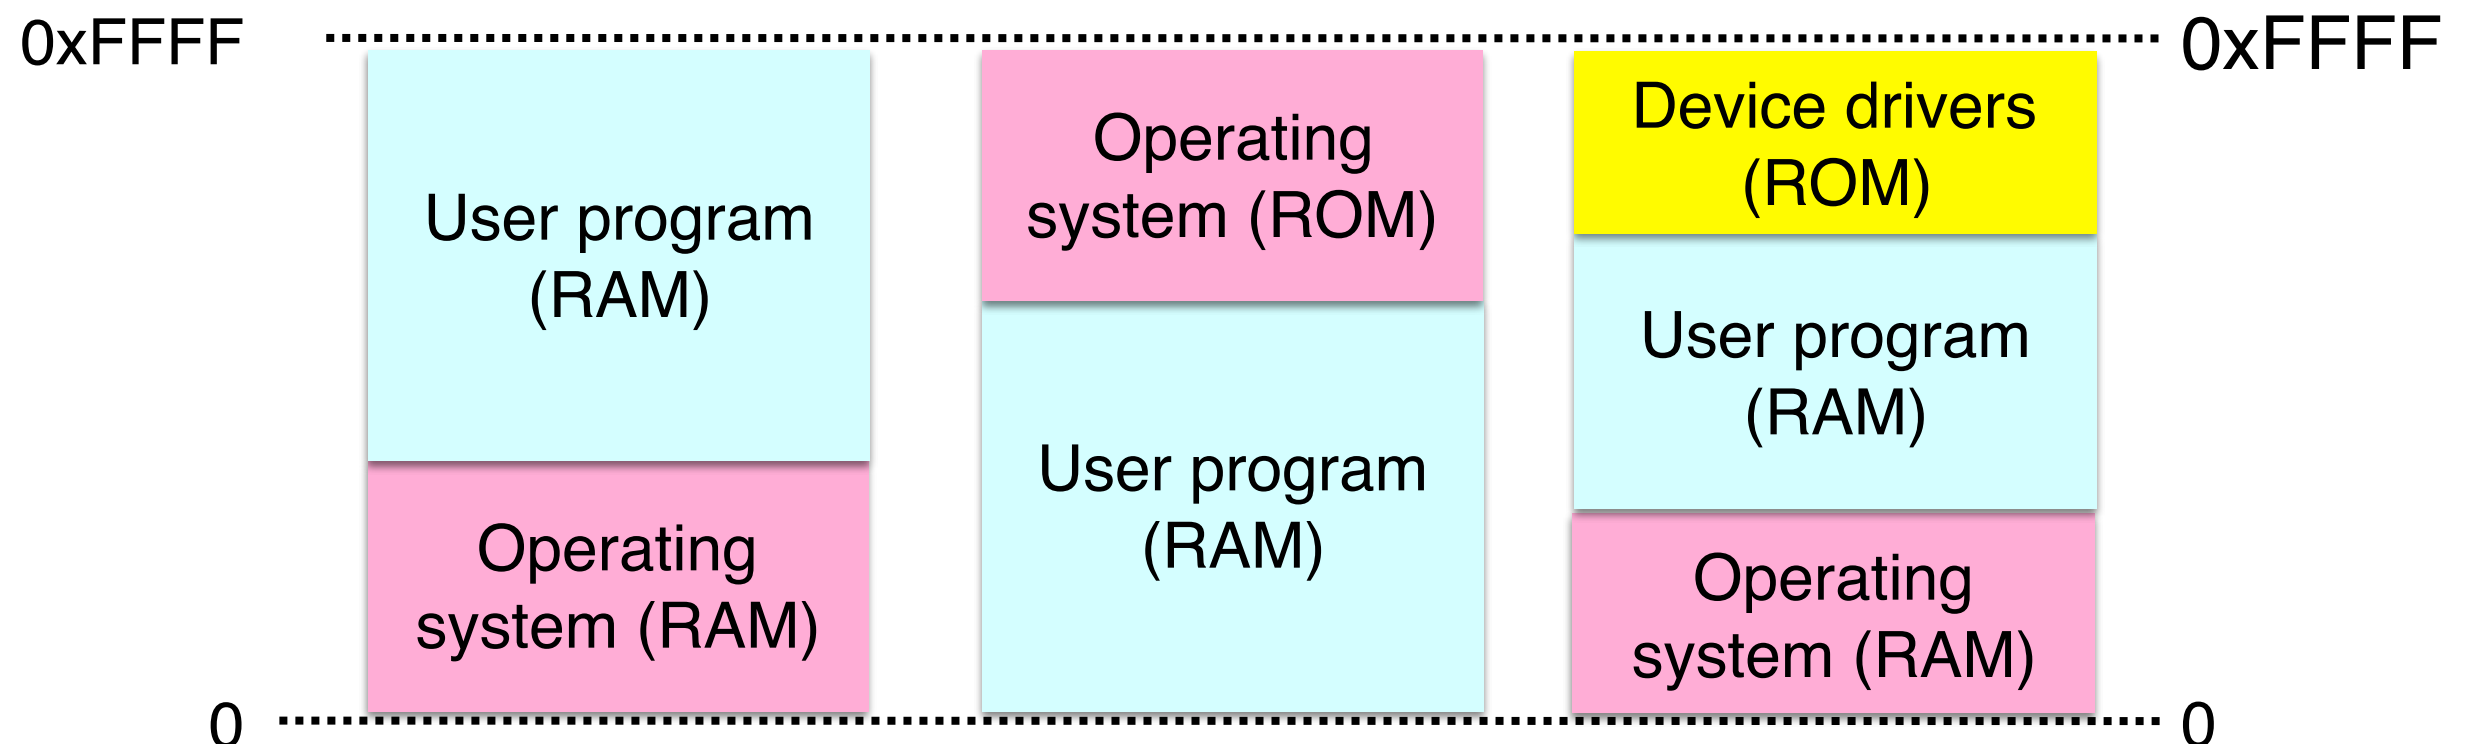
➡ Pick any two...
- ❖ Memory management goal: make the real world look as much like the ideal world as possible

Memory hierarchy

- ❖ What is the memory hierarchy?
 - Different levels of memory
 - Some are small & fast
 - Others are large & slow
- ❖ What levels are usually included?
 - Cache: small amount of fast, expensive memory
 - L1 (level 1) cache: usually on the CPU chip
 - L2: may be on or off chip
 - L3 cache: off-chip, made of SRAM
 - Main memory: medium-speed, medium price memory (DRAM)
 - Disk: many gigabytes of slow, cheap, non-volatile storage
- ❖ Memory manager handles the memory hierarchy

Basic memory management

- ❖ Components include
 - Operating system (perhaps with device drivers)
 - Single process
- ❖ Goal: lay these out in memory
 - Memory protection may not be an issue (only one program)
 - Flexibility may still be useful (allow OS changes, etc.)
- ❖ No swapping or paging




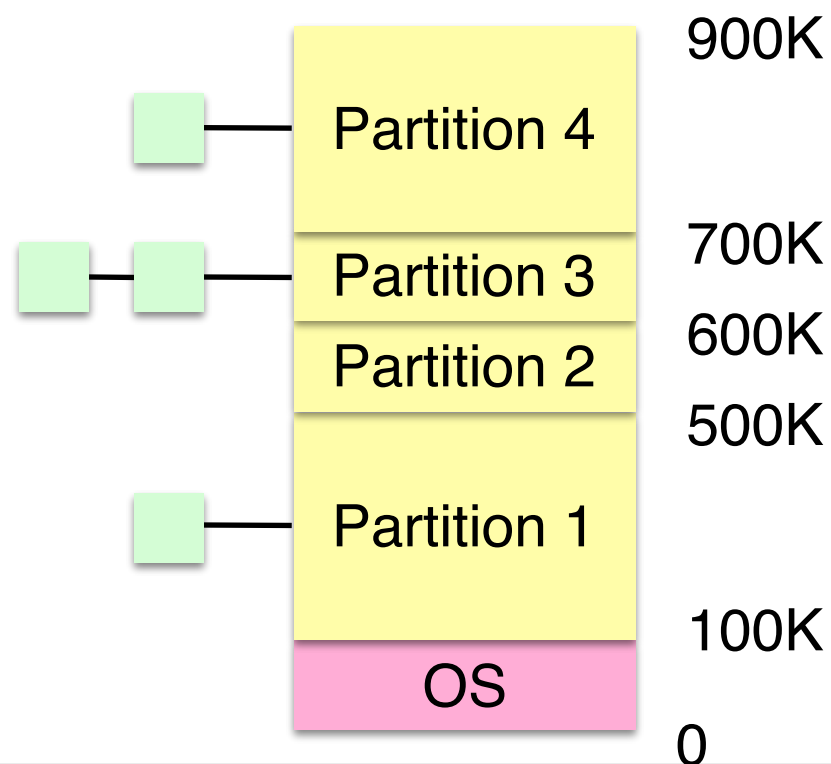
Fixed partitions: multiple programs

❖ Fixed memory partitions

- Divide memory into fixed spaces
- Assign a process to a space when it's free

❖ Mechanisms

- Separate input queues for each partition
- Single input queue: better ability to optimize CPU usage

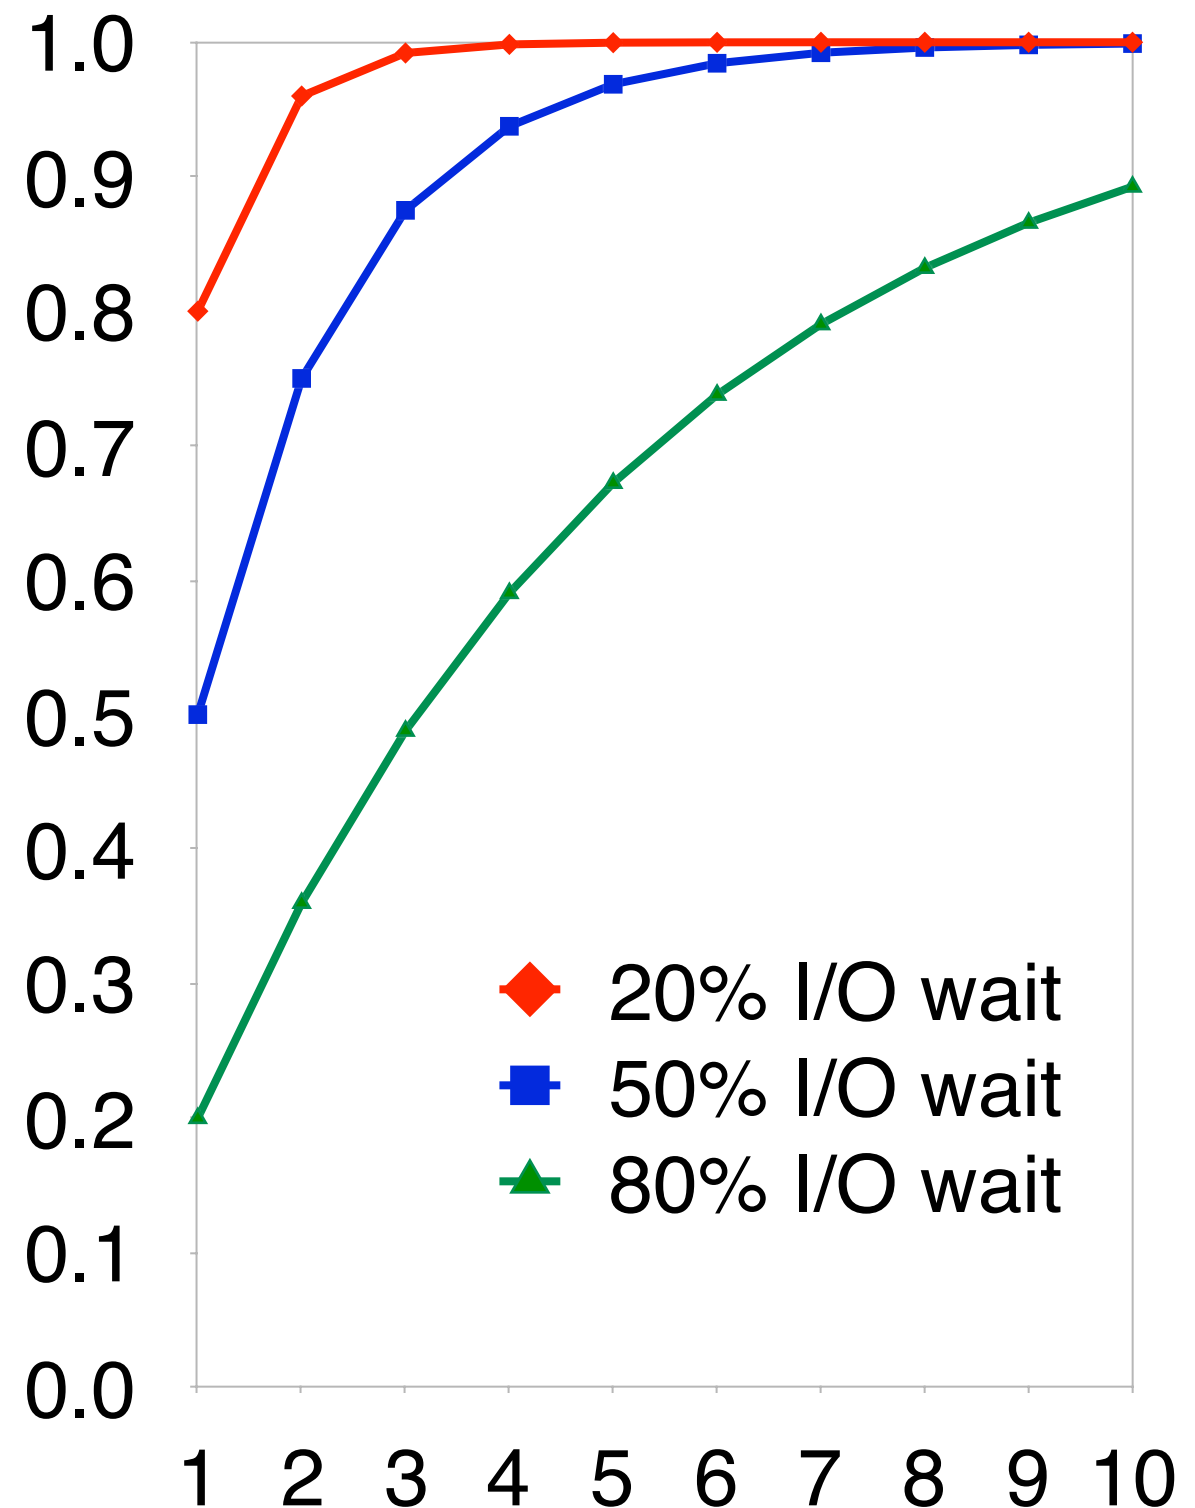


How many processes are enough?

- ❖ Several memory partitions (fixed or variable size)
- ❖ Lots of processes wanting to use the CPU
- ❖ Tradeoff
 - More processes utilize the CPU better
 - Fewer processes use less memory (cheaper!)
- ❖ How many processes do we need to keep the CPU fully utilized?
 - This will help determine how much memory we need
 - Is this still relevant with memory costing \$20/GB?

Modeling multiprogramming


- ❖ More I/O wait means less processor utilization
 - At 20% I/O wait, 3–4 processes fully utilize CPU
 - At 80% I/O wait, even 10 processes aren't enough
- ❖ This means that the OS should have more processes if they're I/O bound
- ❖ More processes → memory management & protection more important!



Multiprogrammed system performance

Job	Arrival time	CPU needed
1	10:00	4
2	10:10	3
3	10:15	2
4	10:20	2

	1	2	3	4
CPU idle	0.8	0.64	0.51	0.41
CPU busy	0.2	0.36	0.49	0.59
CPU/process	0.2	0.18	0.16	0.15



- ❖ Arrival and work requirements of 4 jobs
- ❖ CPU utilization for 1–4 jobs with 80% I/O wait
- ❖ Sequence of events as jobs arrive and finish
 - Numbers show amount of CPU time jobs get in each interval
 - More processes \Rightarrow better utilization, less time per process

Memory and multiprogramming

- ❖ Memory needs two things for multiprogramming
 - Relocation
 - Protection
- ❖ The OS cannot be certain where a program will be loaded in memory
 - Variables and procedures can't use absolute locations in memory
 - Several ways to guarantee this
- ❖ The OS must keep processes' memory separate
 - Protect a process from other processes reading or modifying its own memory
 - Protect a process from modifying its own memory in undesirable ways (such as writing to program code)

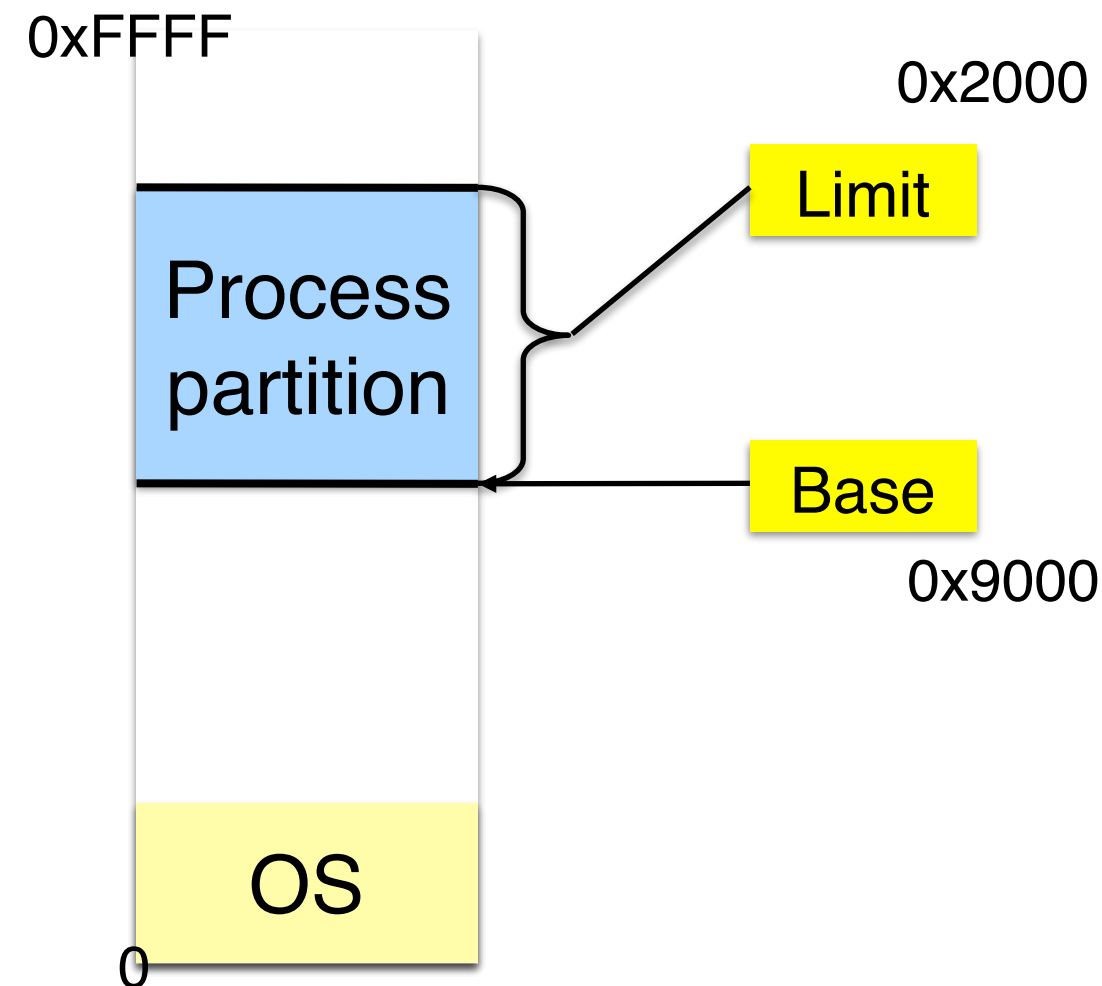
Base and limit registers

❖ Special CPU registers: base & limit

- Access to the registers limited to system mode
- Registers contain
 - Base: start of the process's memory partition
 - Limit: length of the process's memory partition

❖ Address generation

- Physical address: location in actual memory
- Logical address: location from the process's point of view
- Physical address = base + logical address
- Logical address larger than limit → error

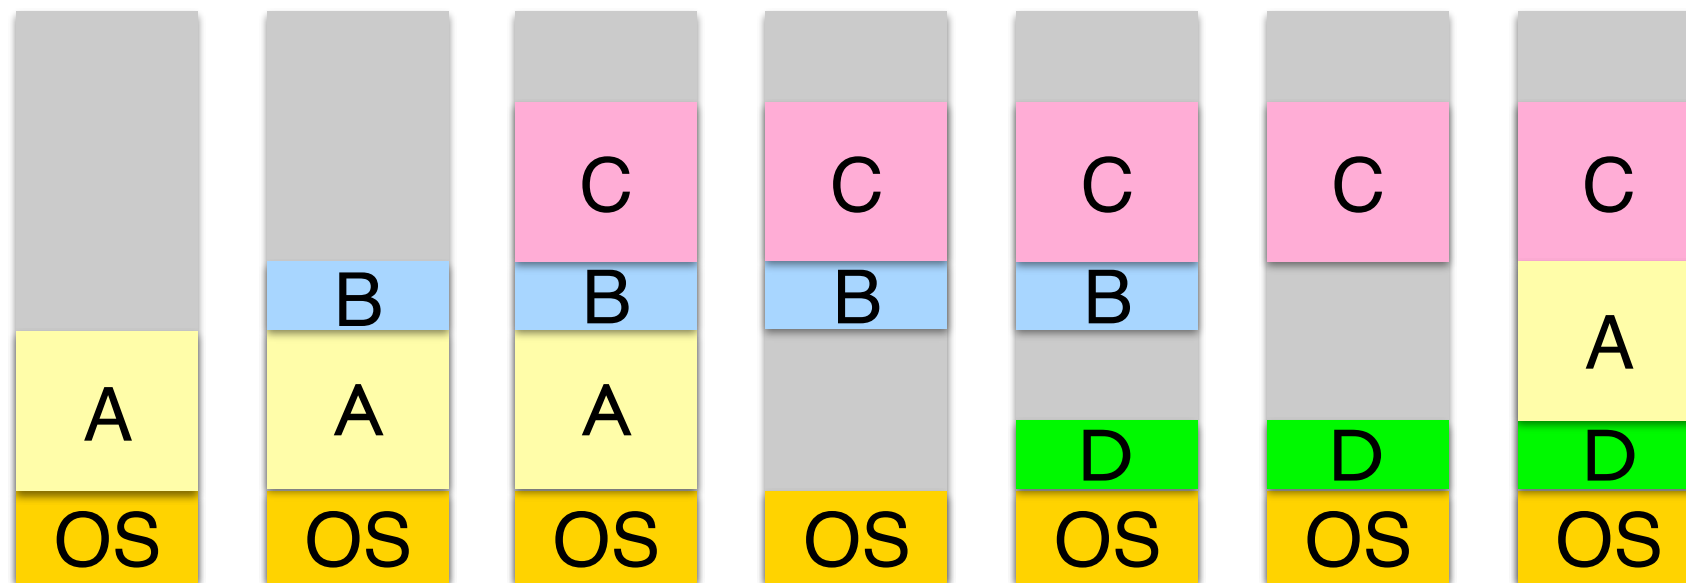


Logical address: 0x1204

Physical address:

$$0x1204 + 0x9000 = 0xa204$$


Swapping



- ❖ Memory allocation changes as
 - Processes come into memory
 - Processes leave memory
 - Swapped to disk
 - Complete execution
- ❖ Gray regions are unused memory

Swapping: leaving room to grow

- ❖ Need to allow for programs to grow
 - Allocate more memory for data
 - Larger stack
- ❖ Handled by allocating more space than is necessary at the start
 - Inefficient: wastes memory that's not currently in use
 - What if the process requests too much memory?




Tracking memory usage

- ❖ Operating system needs to track allocation state of memory
 - Regions that are available to hand out
 - Regions that are in use
 - Possibly what they're being used for
- ❖ Multiple approaches to doing this
 - Bitmap
 - Linked list
 - Buddy allocation
 - Slab allocation
- ❖ Techniques also used for tracking free / allocated storage space


Tracking memory usage: bitmaps

- ❖ Keep track of free / allocated memory regions with a bitmap
 - One bit in map corresponds to a fixed-size region of memory
 - Bitmap is a constant size for a given amount of memory regardless of how much is allocated at a particular time
- ❖ Chunk size determines efficiency
 - At 1 bit per 4KB chunk, we need just 256 bits (32 bytes) per MB of memory
 - For smaller chunks, we need more memory for the bitmap
 - Can be difficult to find large contiguous free areas in bitmap



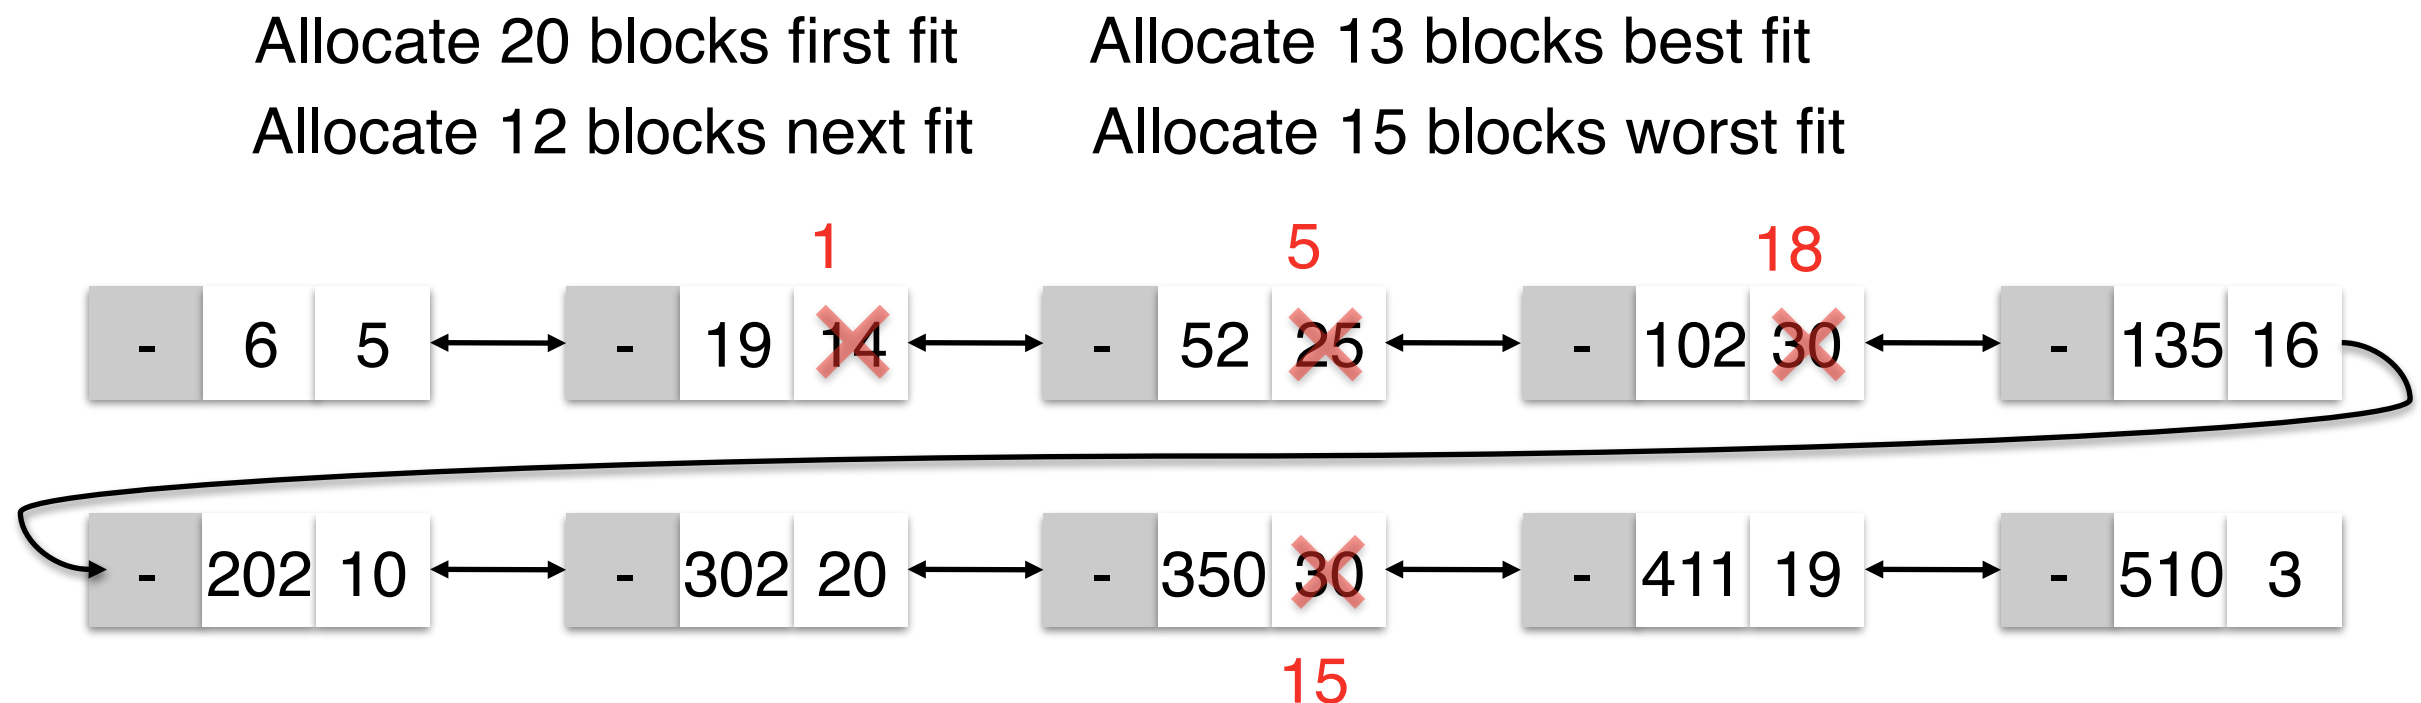
Tracking memory usage: linked lists

- ❖ Keep track of free / allocated memory regions with a linked list
 - Each entry in the list corresponds to a contiguous region of memory
 - Entry can indicate either allocated or free (and, optionally, owning process)
 - May have separate lists for free and allocated areas
- ❖ Efficient if chunks are large
 - Fixed-size representation for each region
 - More regions → more space needed for free lists




Allocating memory

- ❖ Search through region list to find a large enough space
- ❖ Suppose there are several choices: which one to use?
 - First fit: the first suitable hole on the list
 - Next fit: the first suitable after the previously allocated hole
 - Best fit: the smallest hole that is larger than the desired region (wastes least space?)
 - Worst fit: the largest available hole (leaves largest fragment)
- ❖ Option: maintain separate queues for different-size holes





Freeing memory

- ❖ Allocation structures must be updated when memory is freed
- ❖ Easy with bitmaps: just set the appropriate bits in the bitmap
- ❖ Linked lists: modify adjacent elements as needed
 - Merge adjacent free regions into a single region
 - May involve merging two regions with the just-freed area



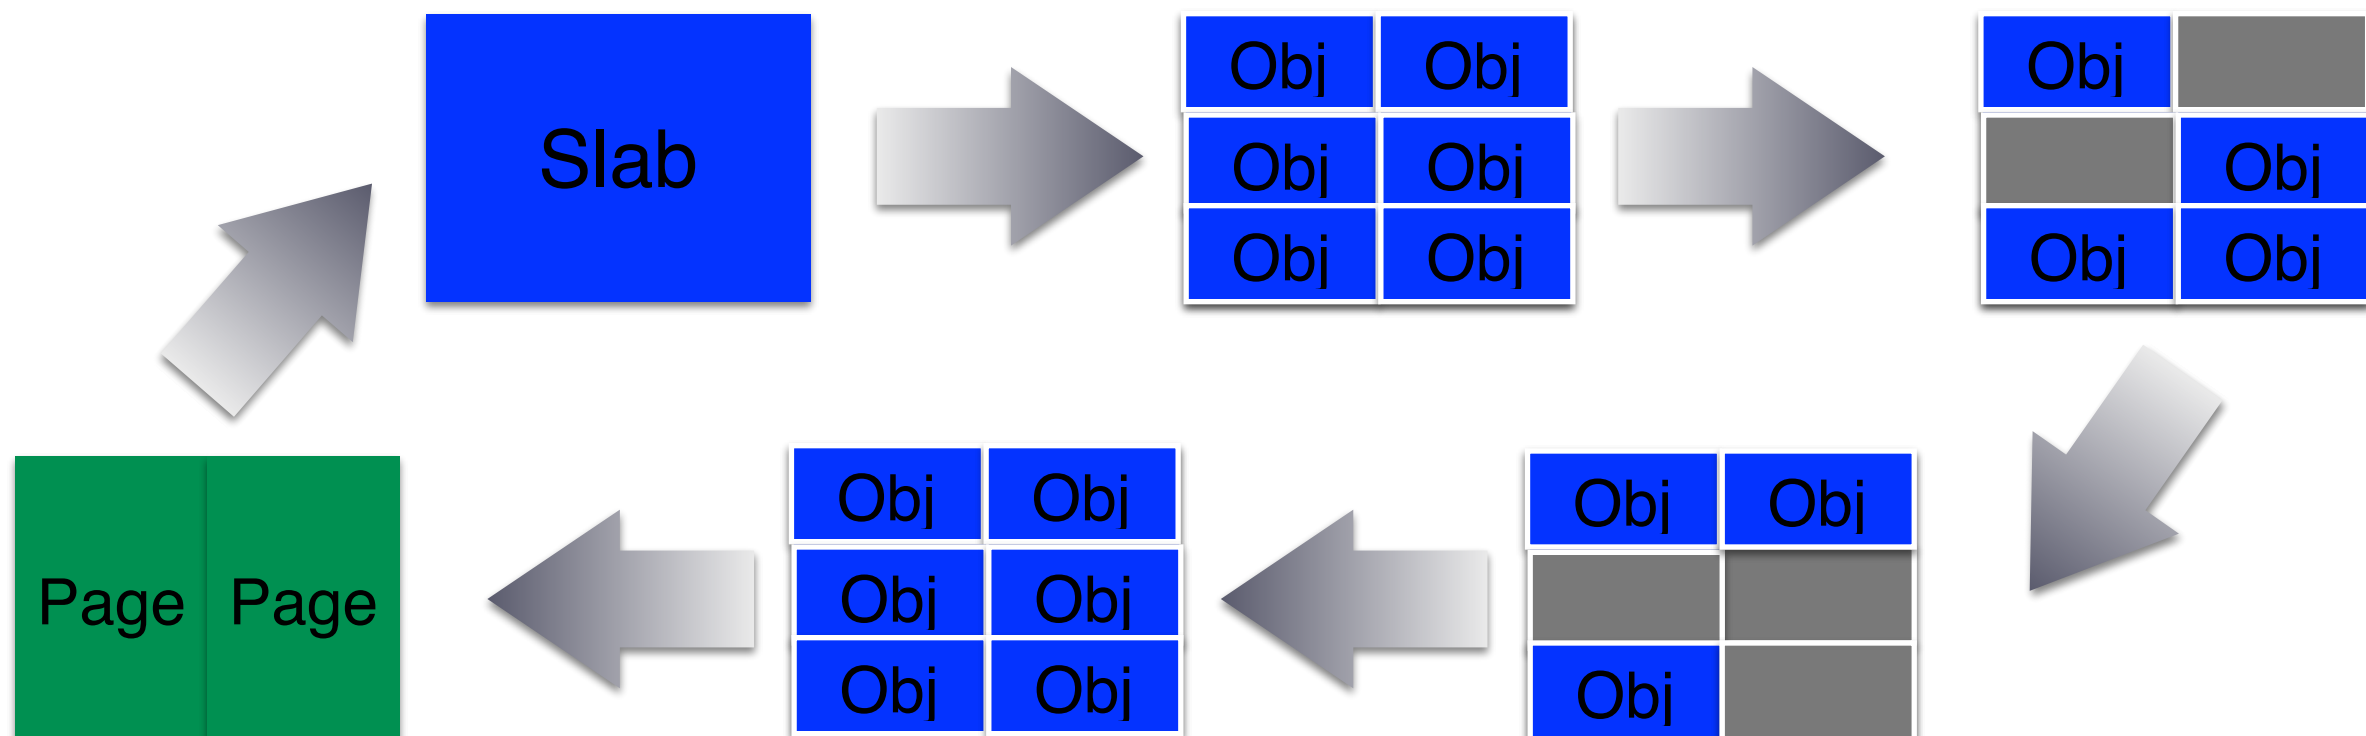
Buddy allocation

- ❖ Allocate memory in powers of two
 - Good for objects of varied sizes
- ❖ Split larger chunks to create two smaller chunks
- ❖ When chunk is freed, see if it can be combined with its buddy to rebuild a larger chunk
 - This is recursive...



Slab allocation

- ❖ Use slabs (each 1+ pages long) to allocate and manage a large number of equal-sized objects
 - Avoid internal fragmentation
 - Better performance: faster allocation / deallocation
- ❖ Carve slabs into objects
 - Free a slab only when all of its objects are freed
 - Keep track of free objects and hand out again when needed



Limitations of swapping

❖ Problems with swapping

- Process must fit into physical memory (impossible to run larger processes)
- Memory becomes fragmented
 - External fragmentation: lots of small free areas
 - Compaction needed to reassemble larger free areas
- Processes are either in memory or on disk: half and half doesn't do any good


❖ Overlays solved the first problem

- Bring in pieces of the process over time (typically data)
- Still doesn't solve the problem of fragmentation or partially resident processes

Virtual memory

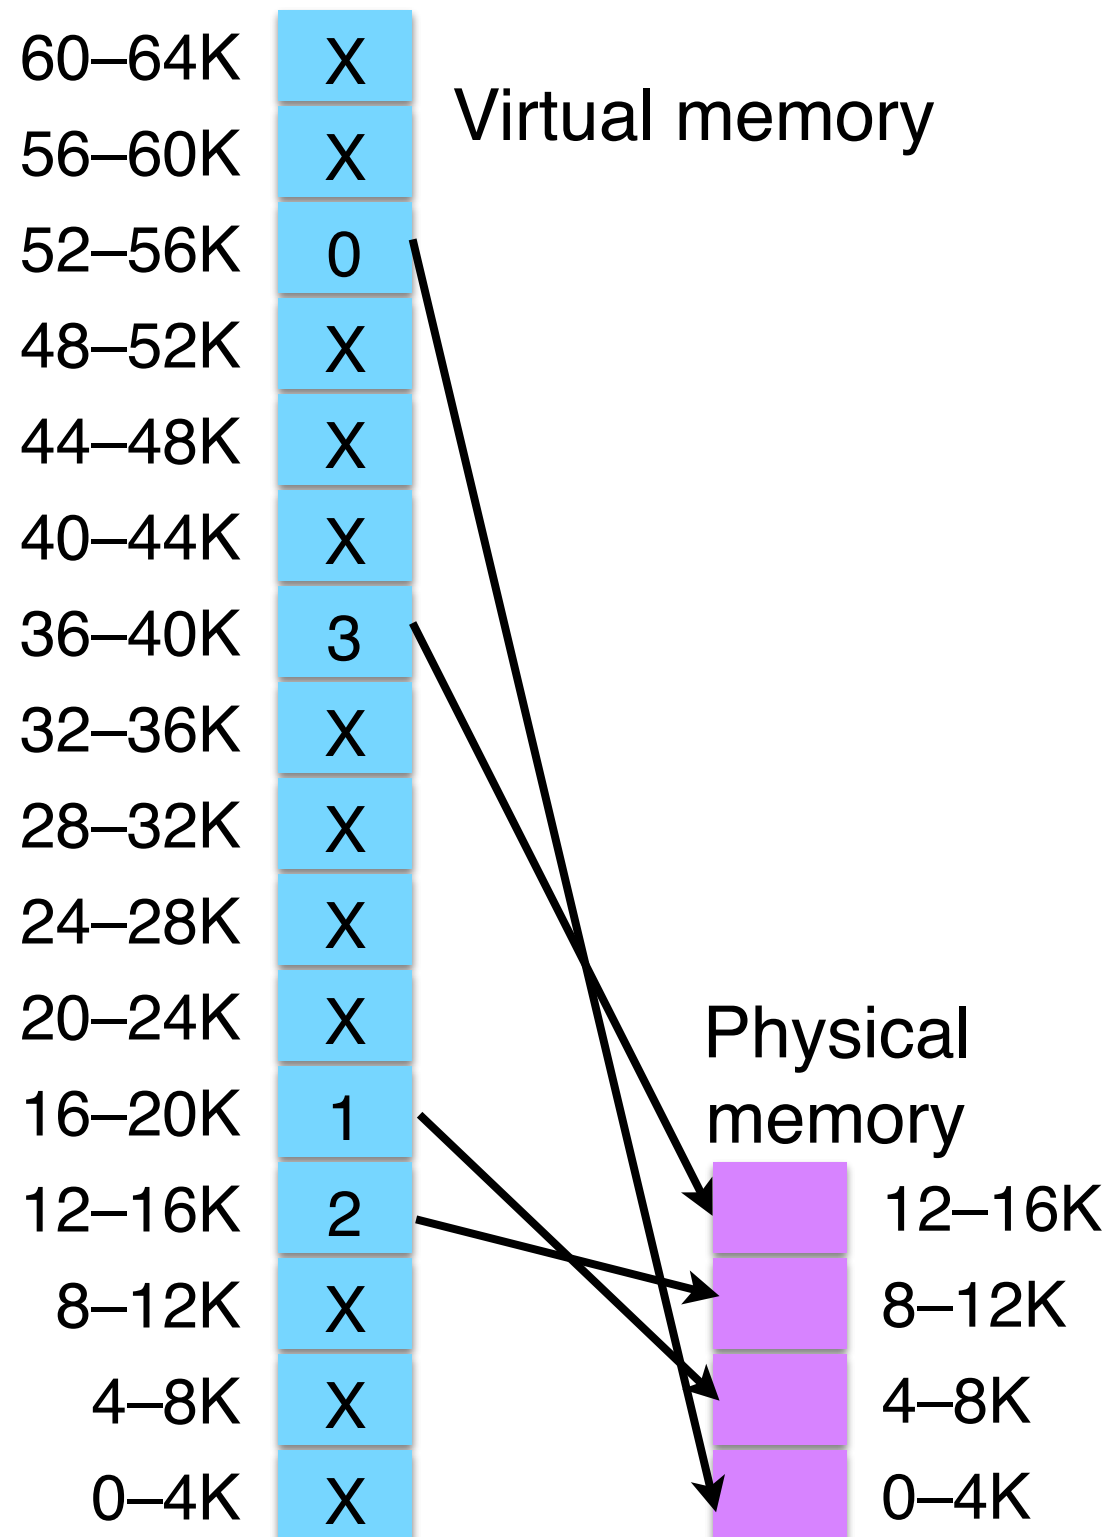
- ❖ Basic idea: allow the OS to hand out more memory than exists on the system
- ❖ Keep recently used stuff in physical memory
- ❖ Move less recently used stuff to disk
- ❖ Keep all of this hidden from processes
 - Processes still see an address space from 0–max_address
 - Movement of information to and from disk handled by the OS without process help
- ❖ Virtual memory (VM) especially helpful in multiprogrammed systems
 - CPU schedules process B while process A waits for its memory to be retrieved from disk

Virtual and physical addresses



- ❖ Program uses virtual addresses
 - Addresses local to the process
 - Hardware translates virtual address to physical address
- ❖ Translation done by the Memory Management Unit
 - Usually on the same chip as the CPU
 - ❑ Only physical addresses leave the CPU/MMU chip
- ❖ Physical memory indexed by physical addresses

Paging and page tables

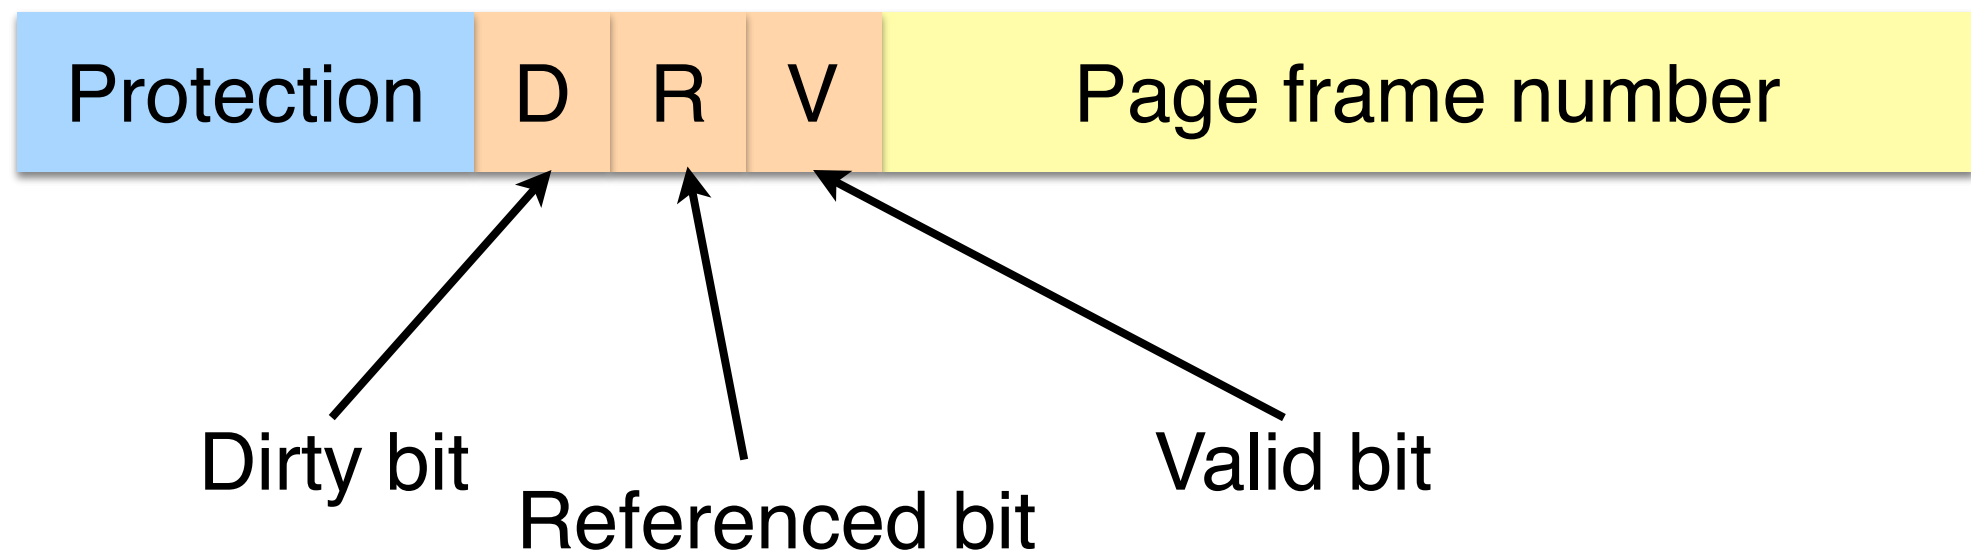


- ❖ Virtual addresses mapped to physical addresses
 - Unit of mapping is called a page
 - All addresses in the same virtual page are in the same physical page
 - Page table entry (PTE) contains translation for a single page
- ❖ Table translates virtual page number to physical page number
 - Not all virtual memory has a physical page
 - Not every physical page need be used
- ❖ Example:
 - 64 KB virtual memory
 - 16 KB physical memory

What's in a page table entry?

❖ Each entry in the page table contains

- Valid bit: set if this logical page number has a corresponding physical frame in memory
 - If not valid, remainder of PTE is irrelevant
- Page frame number: page in physical memory
- Referenced bit: set if data on the page has been accessed
- Dirty (modified) bit :set if data on the page has been modified
- Protection information



Mapping logical addresses to physical addresses


- ❖ Split address from CPU into two pieces
 - Page number (p)
 - Page offset (d)
- ❖ Page number
 - Index into page table
 - Page table contains base address of page in physical memory
- ❖ Page offset
 - Added to base address to get actual physical memory address
- ❖ Page size = 2^d bytes

Example:

- 4 KB (=4096 byte) pages
- 32 bit logical addresses

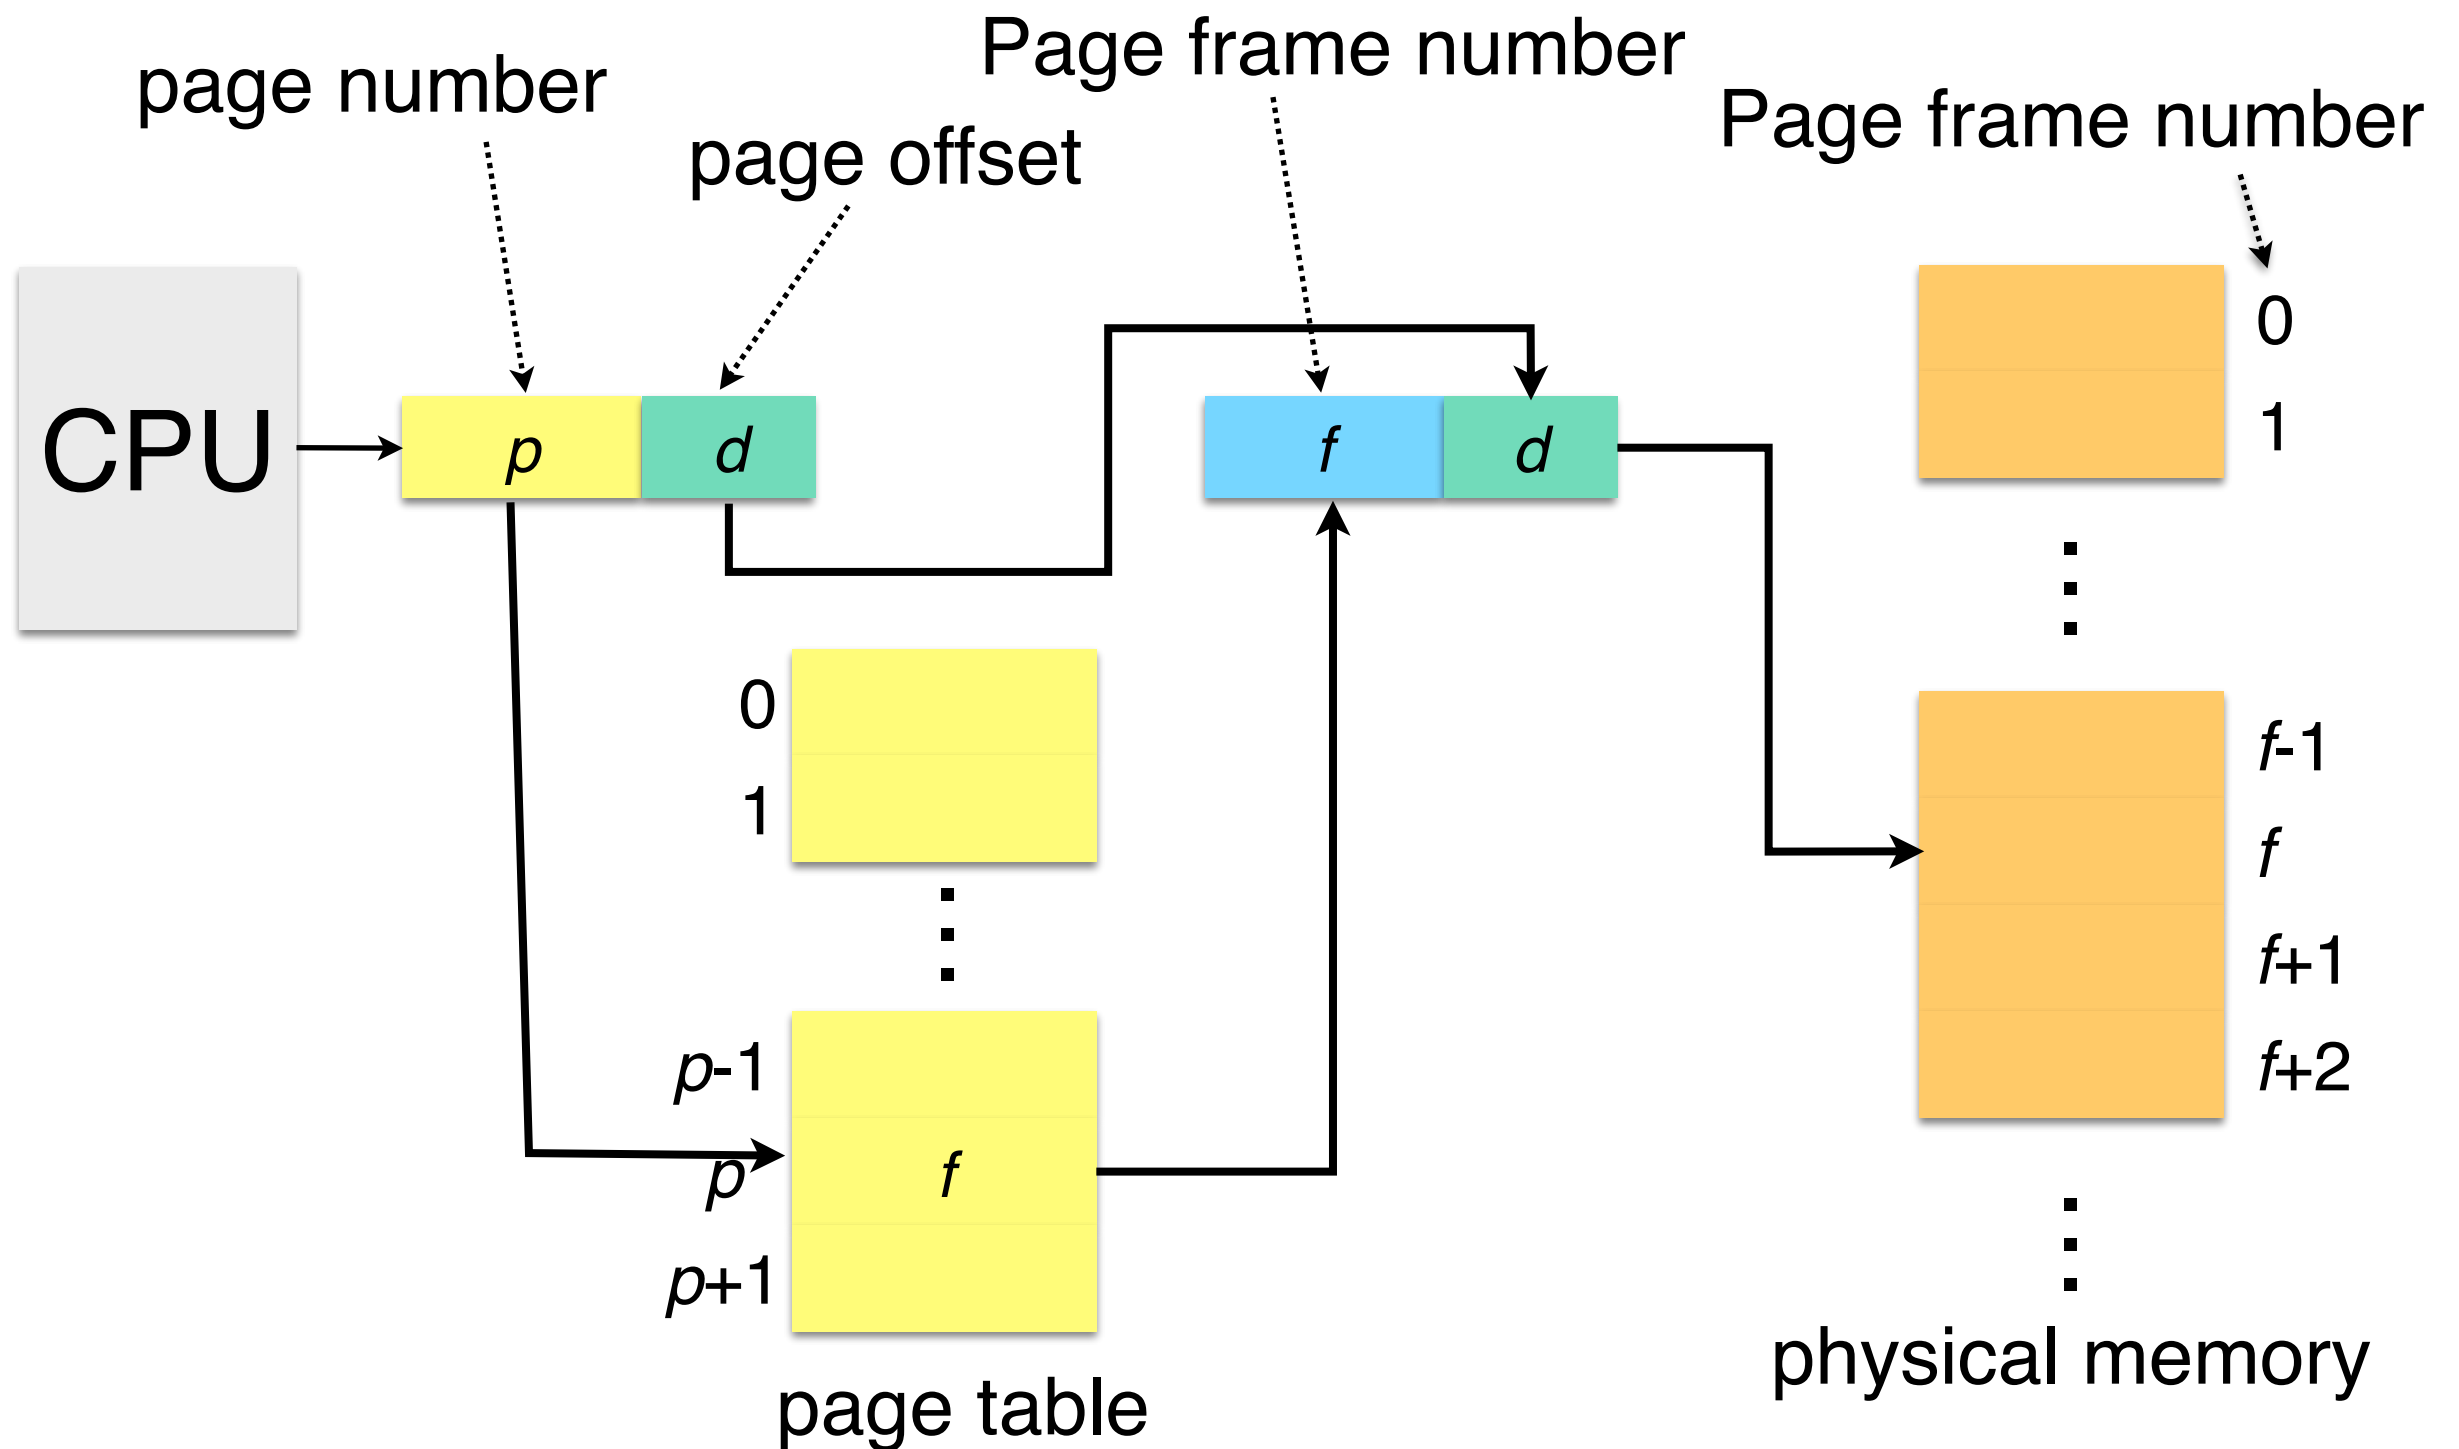
$$2^d = 4096 \longrightarrow d = 12$$

$$32 - 12 = 20 \text{ bits} \quad 12 \text{ bits}$$

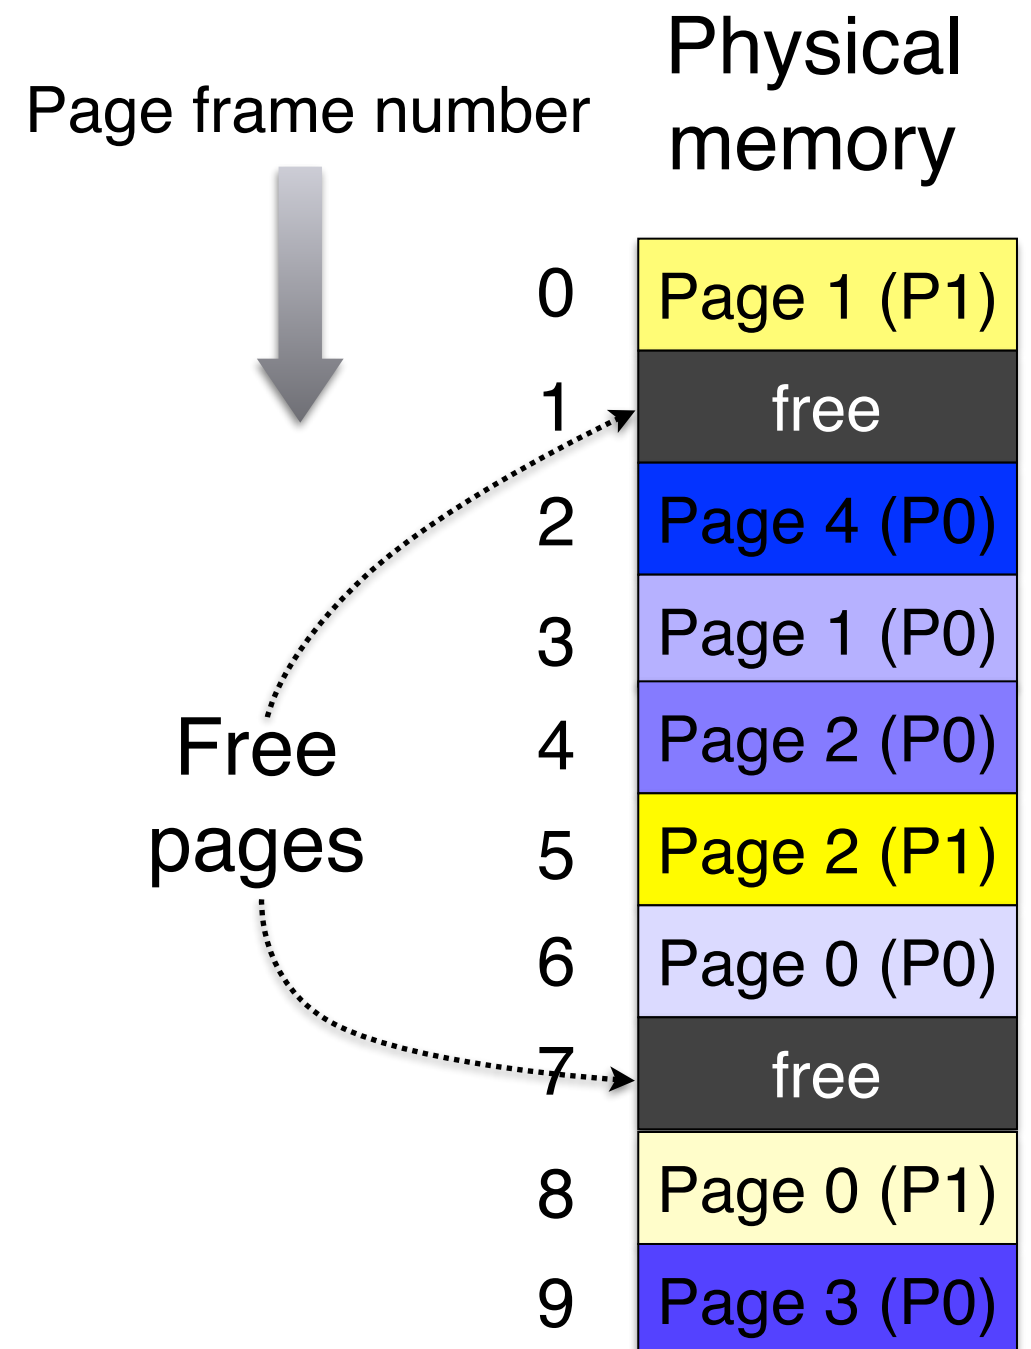
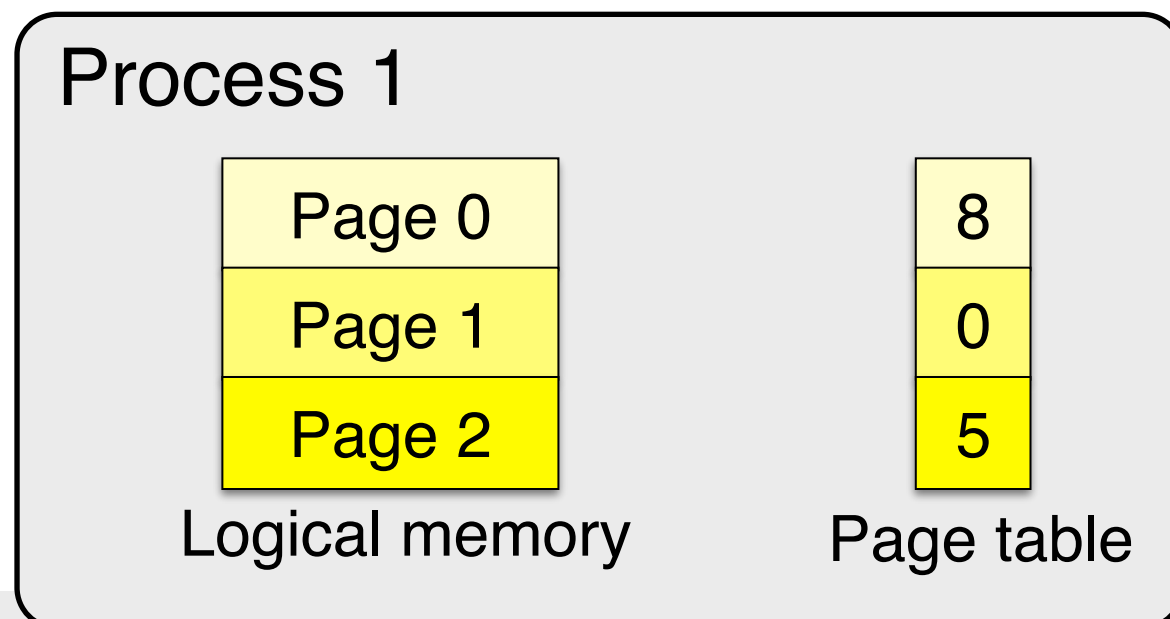
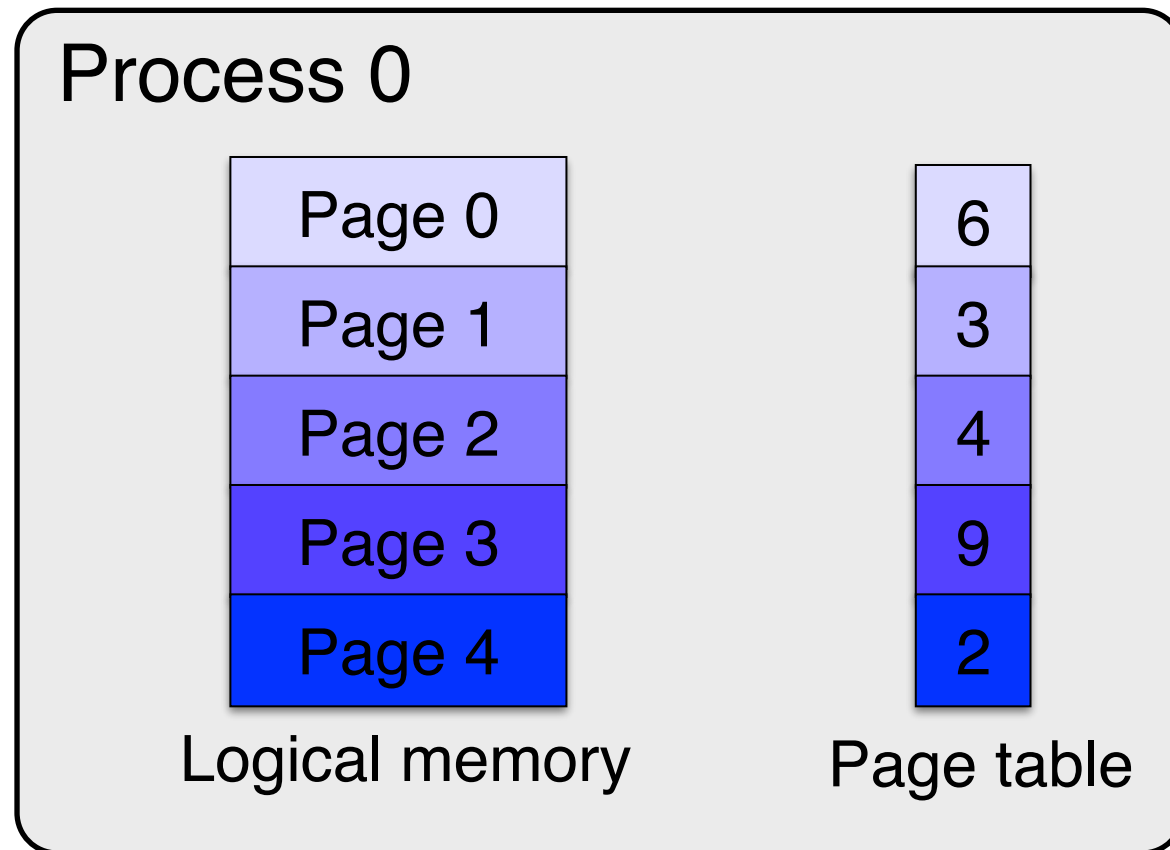


32 bit logical address

Address translation architecture

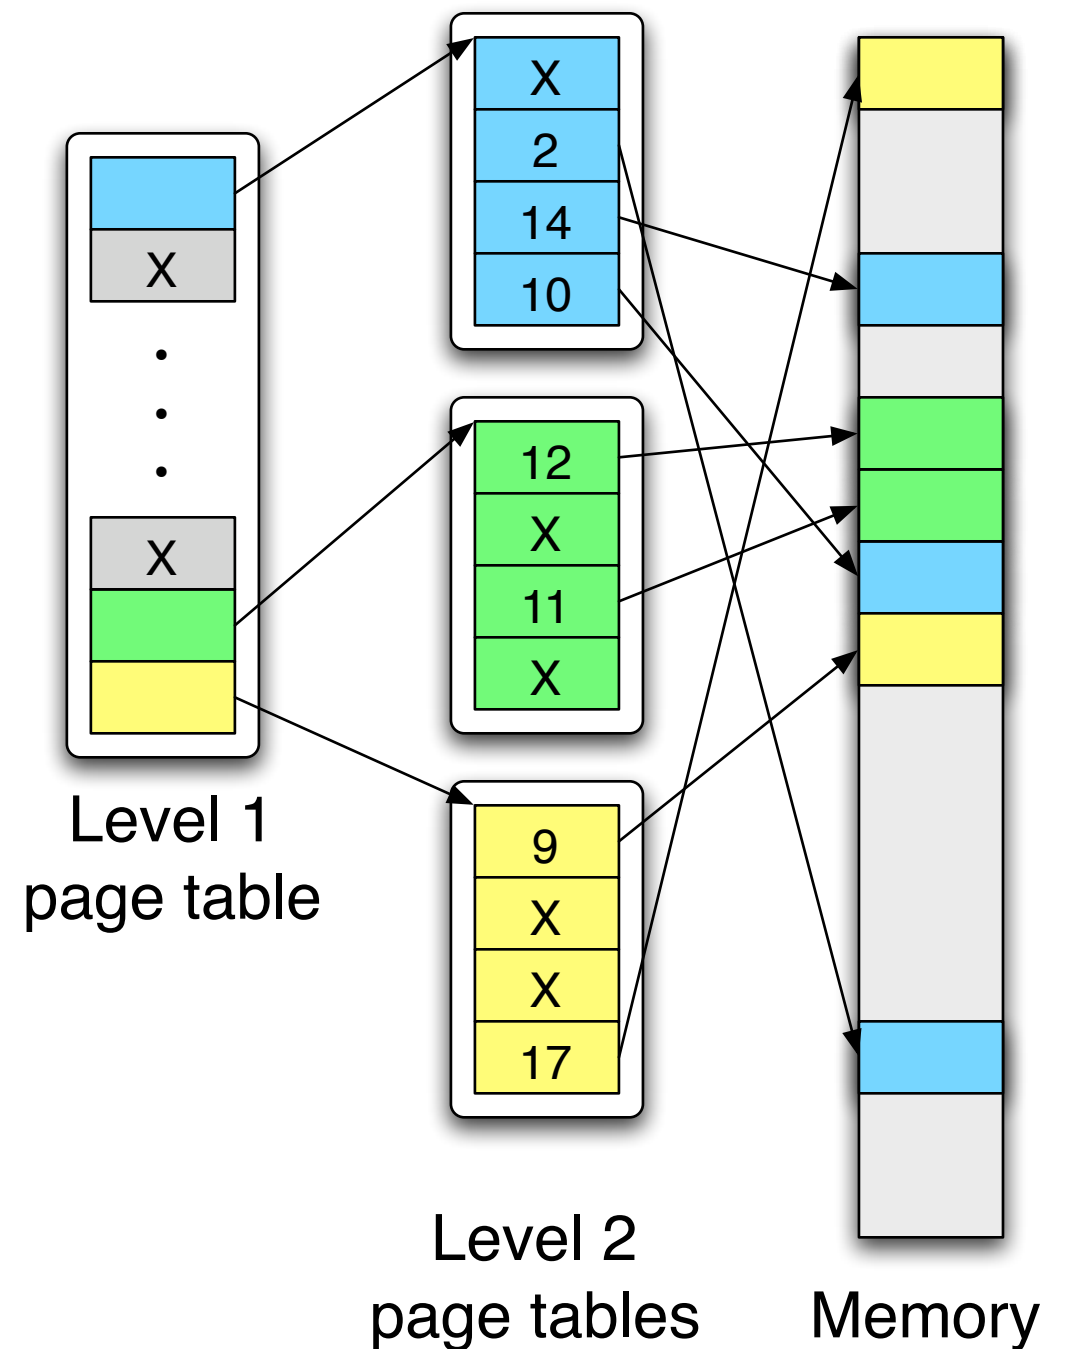


Memory & paging structures



Two-level page tables

- ❖ Problem: page tables can be too large
 - 2^{32} bytes in 4KB pages \rightarrow 1 million PTEs
 - Worse for 64-bit addressing!
- ❖ Solution: use multi-level page tables
 - “Page size” in first page table is large (megabytes)
 - PTE marked invalid in first page table needs no 2nd level page table
- ❖ 1st level page table has pointers to 2nd level page tables
- ❖ 2nd level page table has actual physical page numbers in it



More on two-level page tables

- ❖ Tradeoffs between 1st and 2nd level page table sizes
 - Total number of bits indexing 1st and 2nd level table is constant for a given page size and logical address length
 - Tradeoff between number of bits indexing 1st and number indexing 2nd level tables
 - More bits in 1st level: fine granularity at 2nd level
 - Fewer bits in 1st level: maybe less wasted space?
- ❖ All addresses in table are physical addresses
- ❖ Protection bits kept in 2nd level table

Two-level paging: example

❖ System characteristics


- 8 KB pages
- 32-bit logical address divided into 13 bit page offset, 19 bit page number

❖ Page number divided into:

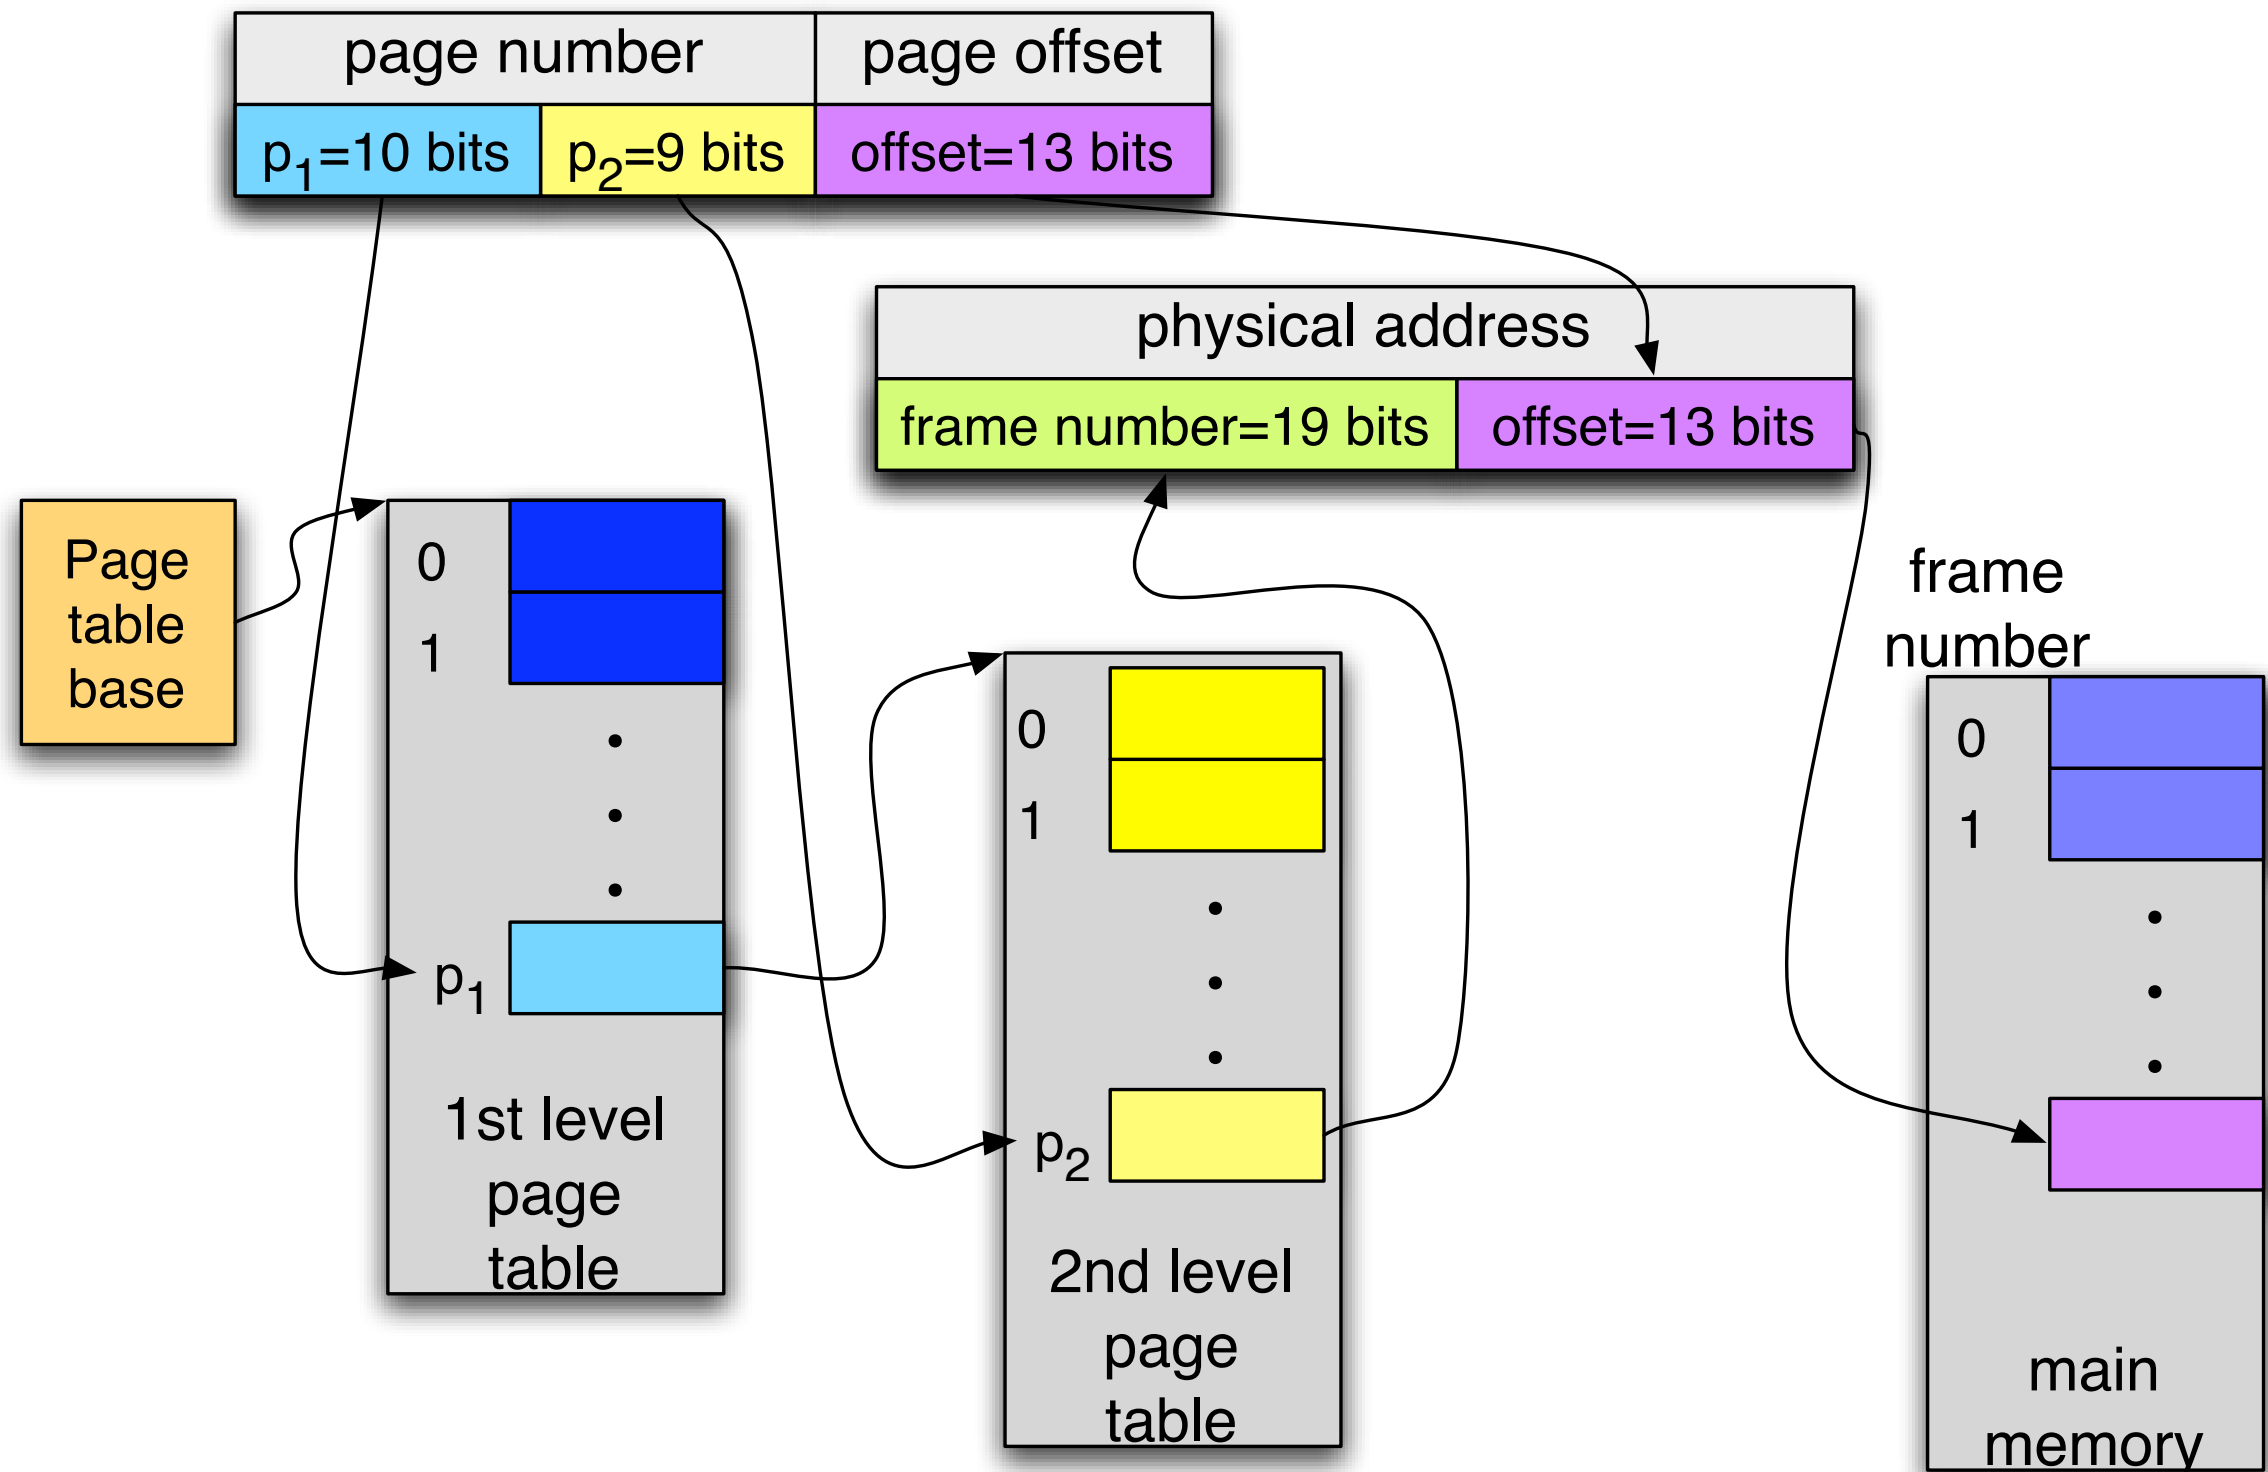
- 10 bit page number
- 9 bit page offset

❖ Logical address looks like this:

- p1 is an index into the 1st level page table
- p2 is an index into the 2nd level page table pointed to by p1



2-level address translation example



Implementing page tables in hardware

- ❖ Page table resides in main (physical) memory
- ❖ CPU uses special registers for paging
 - Page table base register (PTBR) points to the page table
 - Page table length register (PTLR) contains length of page table: restricts maximum legal logical address
- ❖ Translating an address requires two memory accesses
 - First access reads page table entry (PTE)
 - Second access reads the data / instruction from memory
- ❖ Reduce number of memory accesses
 - Can't avoid second access (we need the value from memory)
 - Eliminate first access by keeping a hardware cache (called a translation lookaside buffer or TLB) of recently used page table entries

Translation Lookaside Buffer (TLB)

- ❖ Search the TLB for the desired logical page number
 - Search entries in parallel
 - Use standard cache techniques
- ❖ If desired logical page number is found, get frame number from TLB
- ❖ If desired logical page number isn't found
 - Get frame number from page table in memory
 - Replace an entry in the TLB with the logical & physical page numbers from this reference

Logical page #	Physical frame #
8	3
invalid	
2	1
3	0
12	12
29	6
22	11
7	4

Example TLB

Handling TLB misses

- ❖ If PTE isn't found in TLB, OS needs to do the lookup in the page table
- ❖ Lookup can be done in hardware or software
- ❖ Hardware TLB replacement
 - CPU hardware does page table lookup
 - Can be faster than software
 - Less flexible than software, and more complex hardware
- ❖ Software TLB replacement
 - OS gets TLB exception
 - Exception handler does page table lookup & places the result into the TLB
 - Program continues after return from exception
 - Larger TLB (lower miss rate) can make this feasible

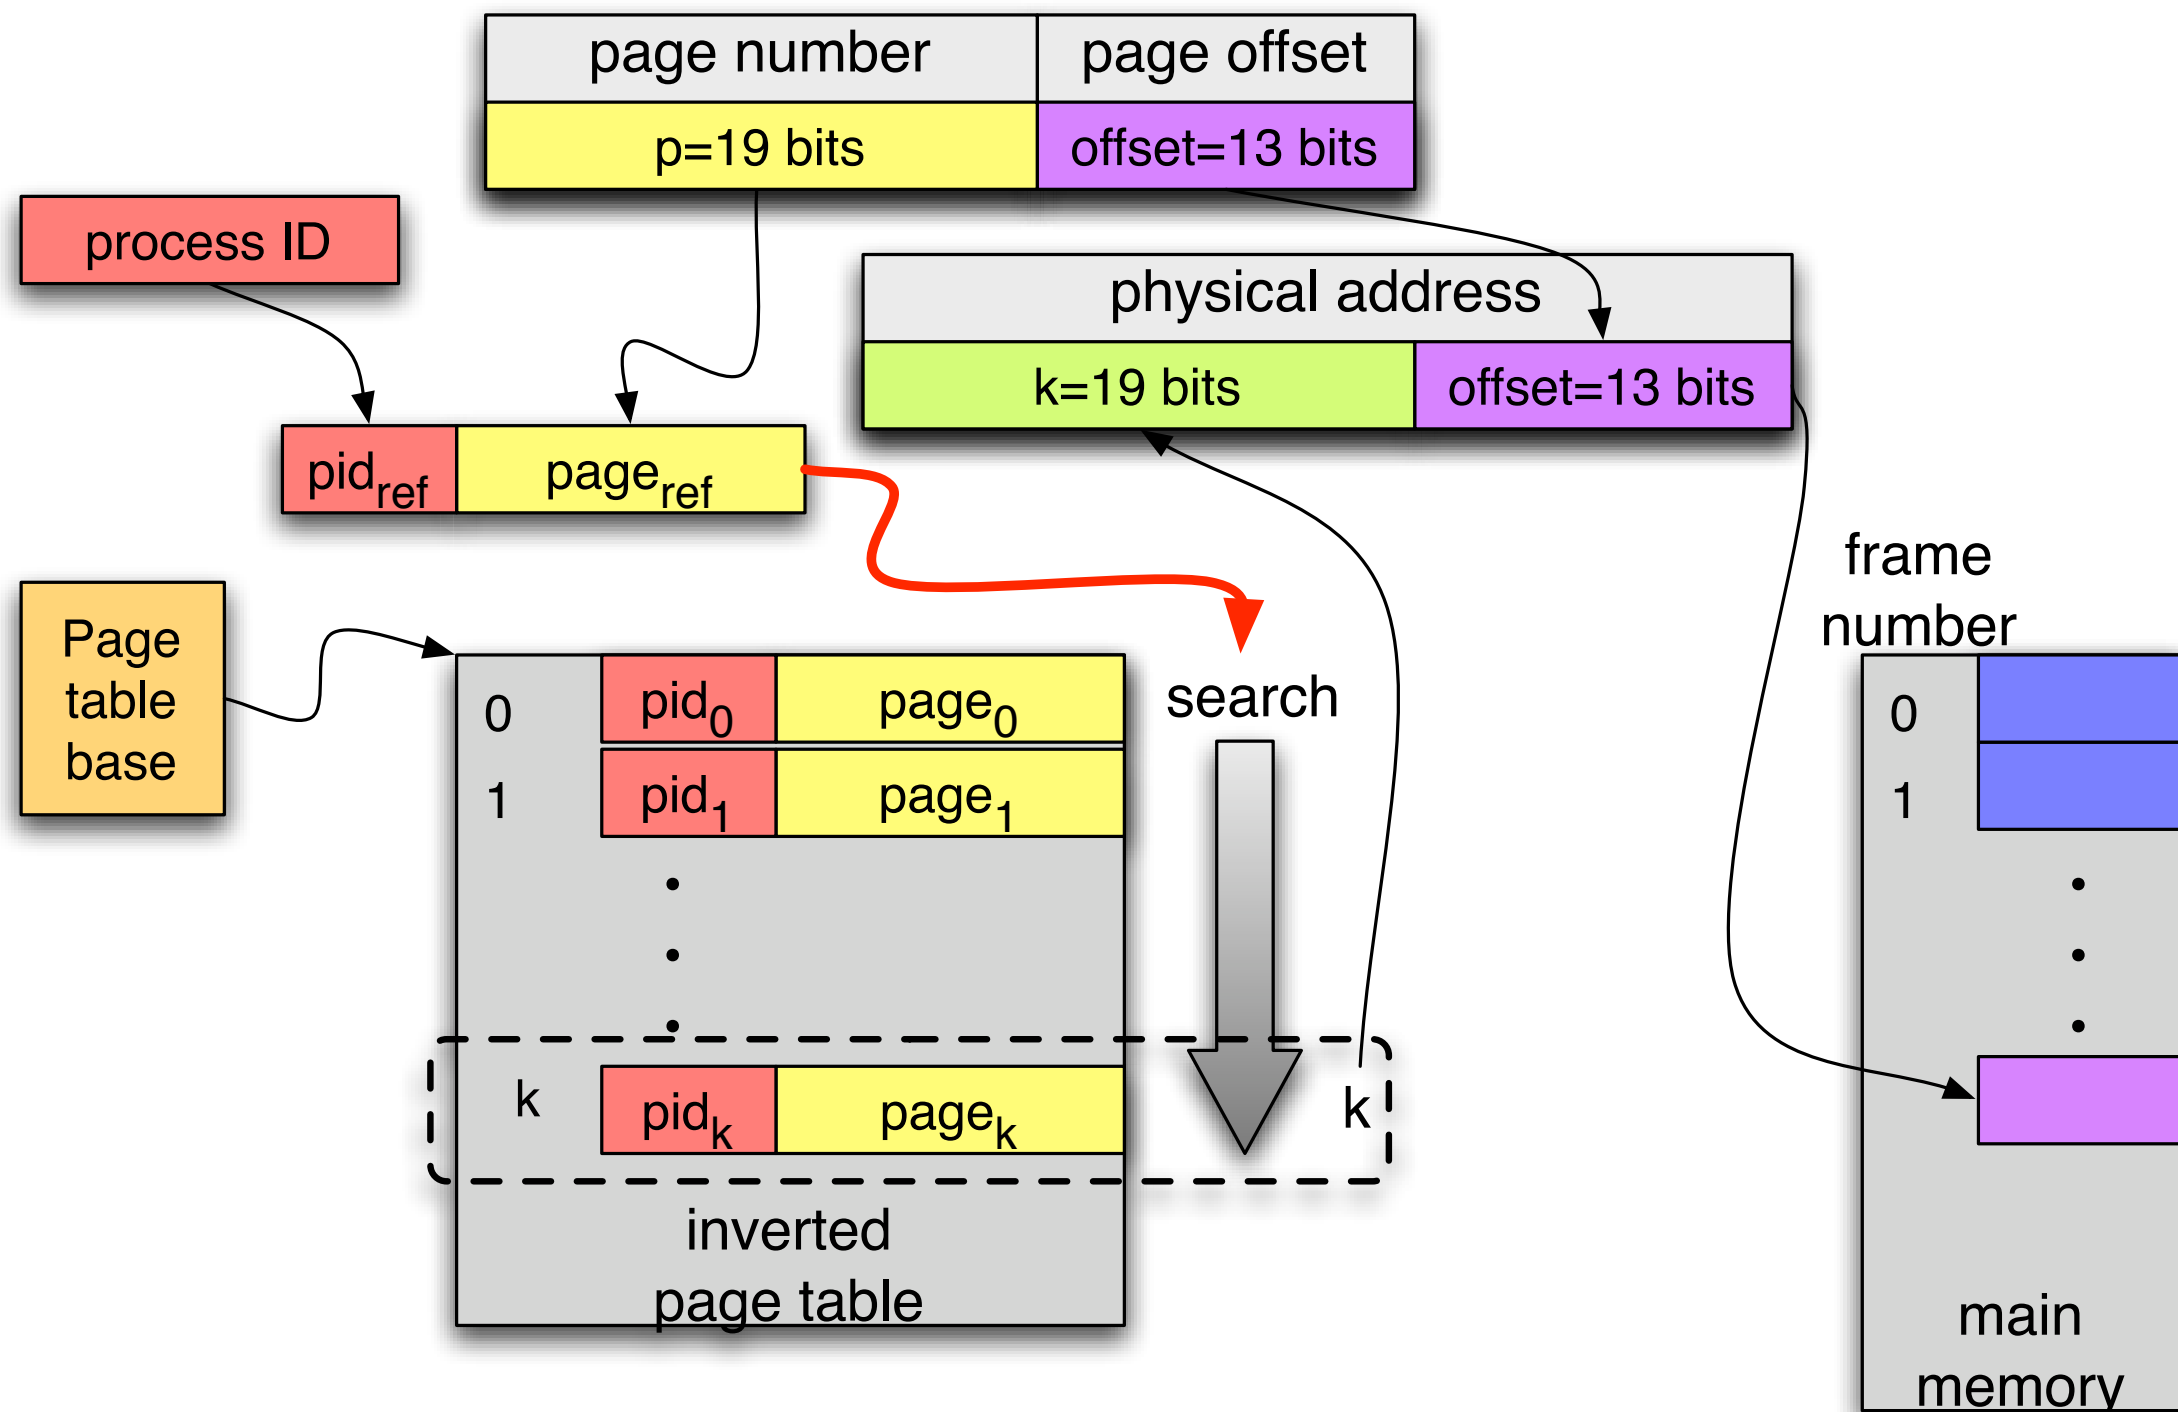
How long do memory accesses take?

- ❖ Assume the following times:
 - TLB lookup time = a (often zero—overlapped in CPU)
 - Memory access time = m
- ❖ Hit ratio (h) is percentage of time that a logical page number is found in the TLB
 - Larger TLB usually means higher h
 - TLB structure can affect h as well
- ❖ Effective access time (an average) is calculated as:
 - $EAT = (m + a)h + (m + m + a)(1-h)$
 - $EAT = a + (2-h)m$
- ❖ Interpretation
 - Reference always requires TLB lookup, 1 memory access
 - TLB misses also require an additional memory reference

Inverted page table

- ❖ Reduce page table size further: keep one entry for each frame in memory
 - Alternative: merge tables for pages in memory and on disk
- ❖ PTE contains
 - Virtual address pointing to this frame
 - Information about the process that owns this page
- ❖ Search page table by
 - Hashing the virtual page number and process ID
 - Starting at the entry corresponding to the hash result
 - Search until either the entry is found or a limit is reached
- ❖ Page frame number is index of PTE
- ❖ Improve performance by using more advanced hashing algorithms

Inverted page table architecture



One-to-one correspondence between page table entries and pages in memory