

# Processes and Threads



# Processes and threads

---

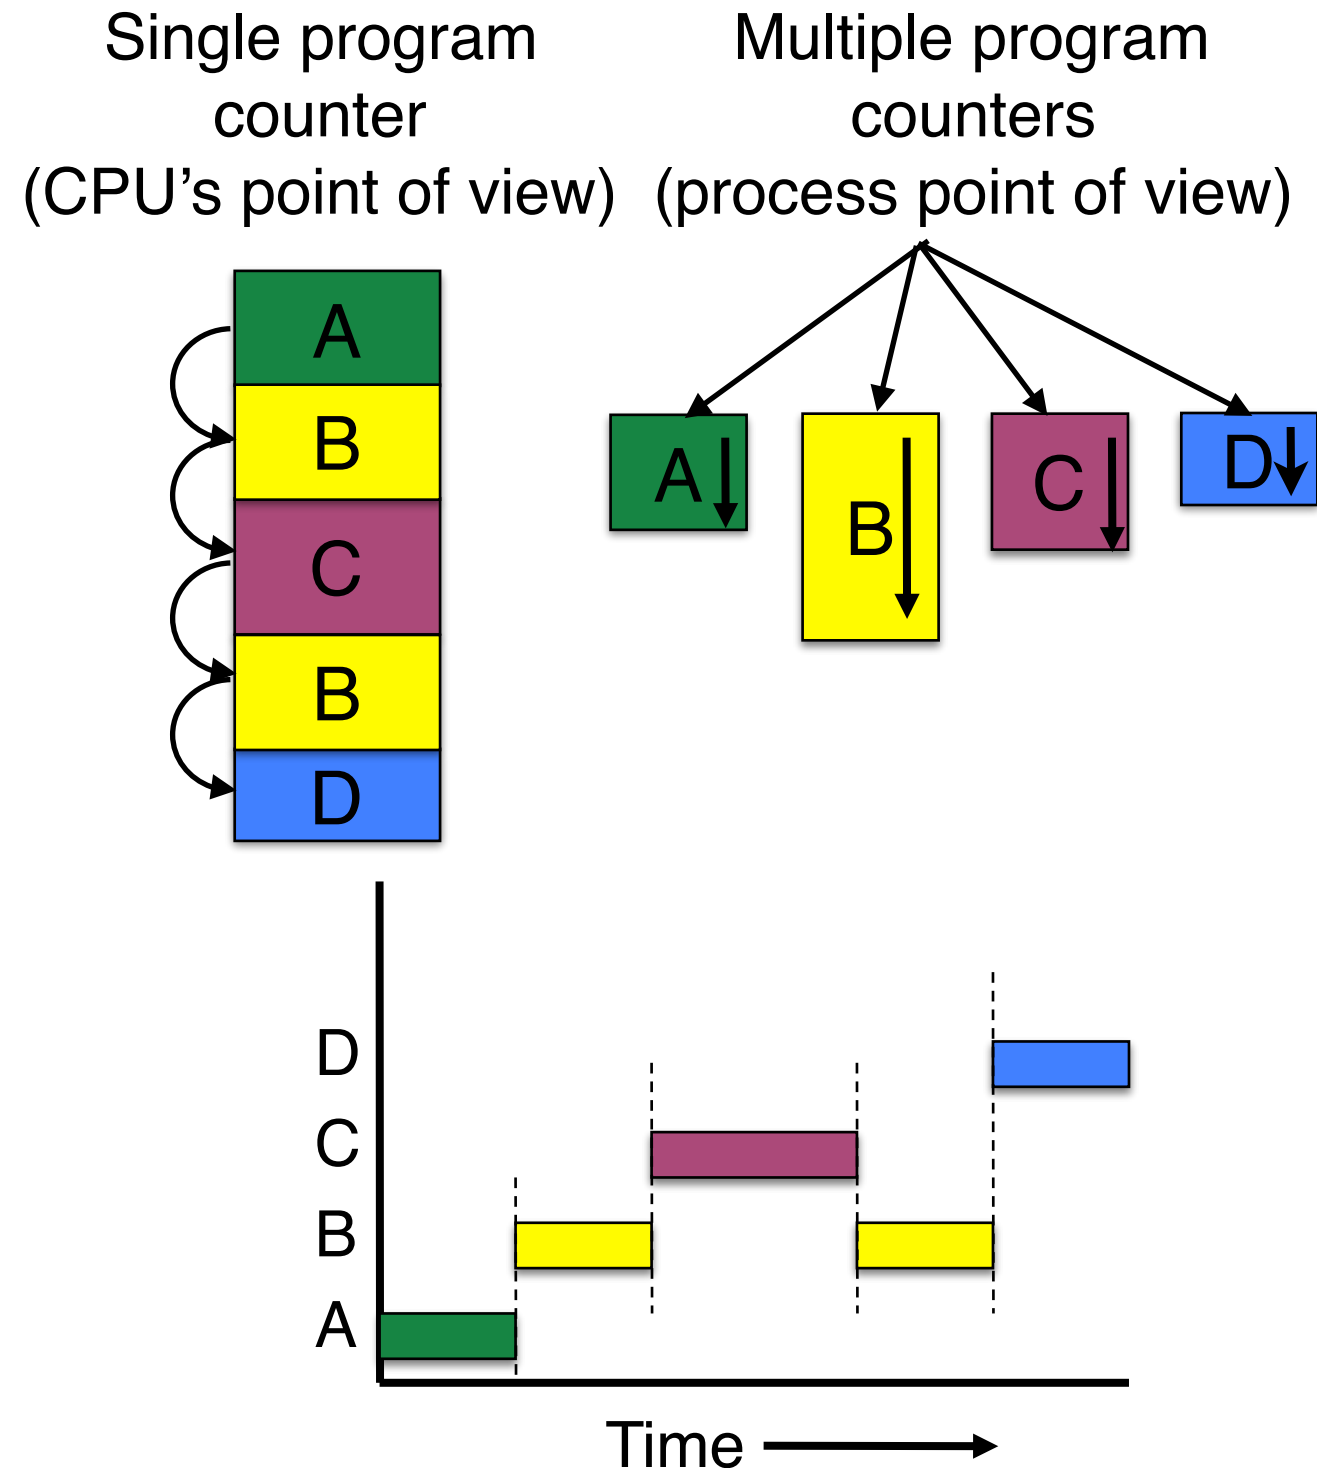
- ❖ Processes
- ❖ Threads
- ❖ Scheduling
- ❖ Interprocess communication (IPC)
- ❖ Classical IPC problems

# What is a process?

- ❖ Code, data, and stack
  - Usually (but not always) has its own address space
- ❖ Program state
  - CPU registers
  - Program counter (current location in the code)
  - Stack pointer
- ❖ Only one process can be running in a single CPU core at any given time!
  - Multi-core CPUs can support multiple processes

# The process model

- ❖ Multiprogramming of four programs
- ❖ Conceptual model
  - 4 independent processes
  - Processes run sequentially
- ❖ Only one program active at any instant!
  - That instant can be very short...
  - Only applies if there's a single CPU (with a single core) in the system



# When is a process created?

- ❖ Processes can be created in two ways
  - System initialization: one or more processes created when the OS starts up
  - Execution of a process creation system call: something explicitly asks for a new process
- ❖ System calls can come from
  - User request to create a new process (system call executed from user shell)
  - Already running processes
    - User programs
    - System daemons

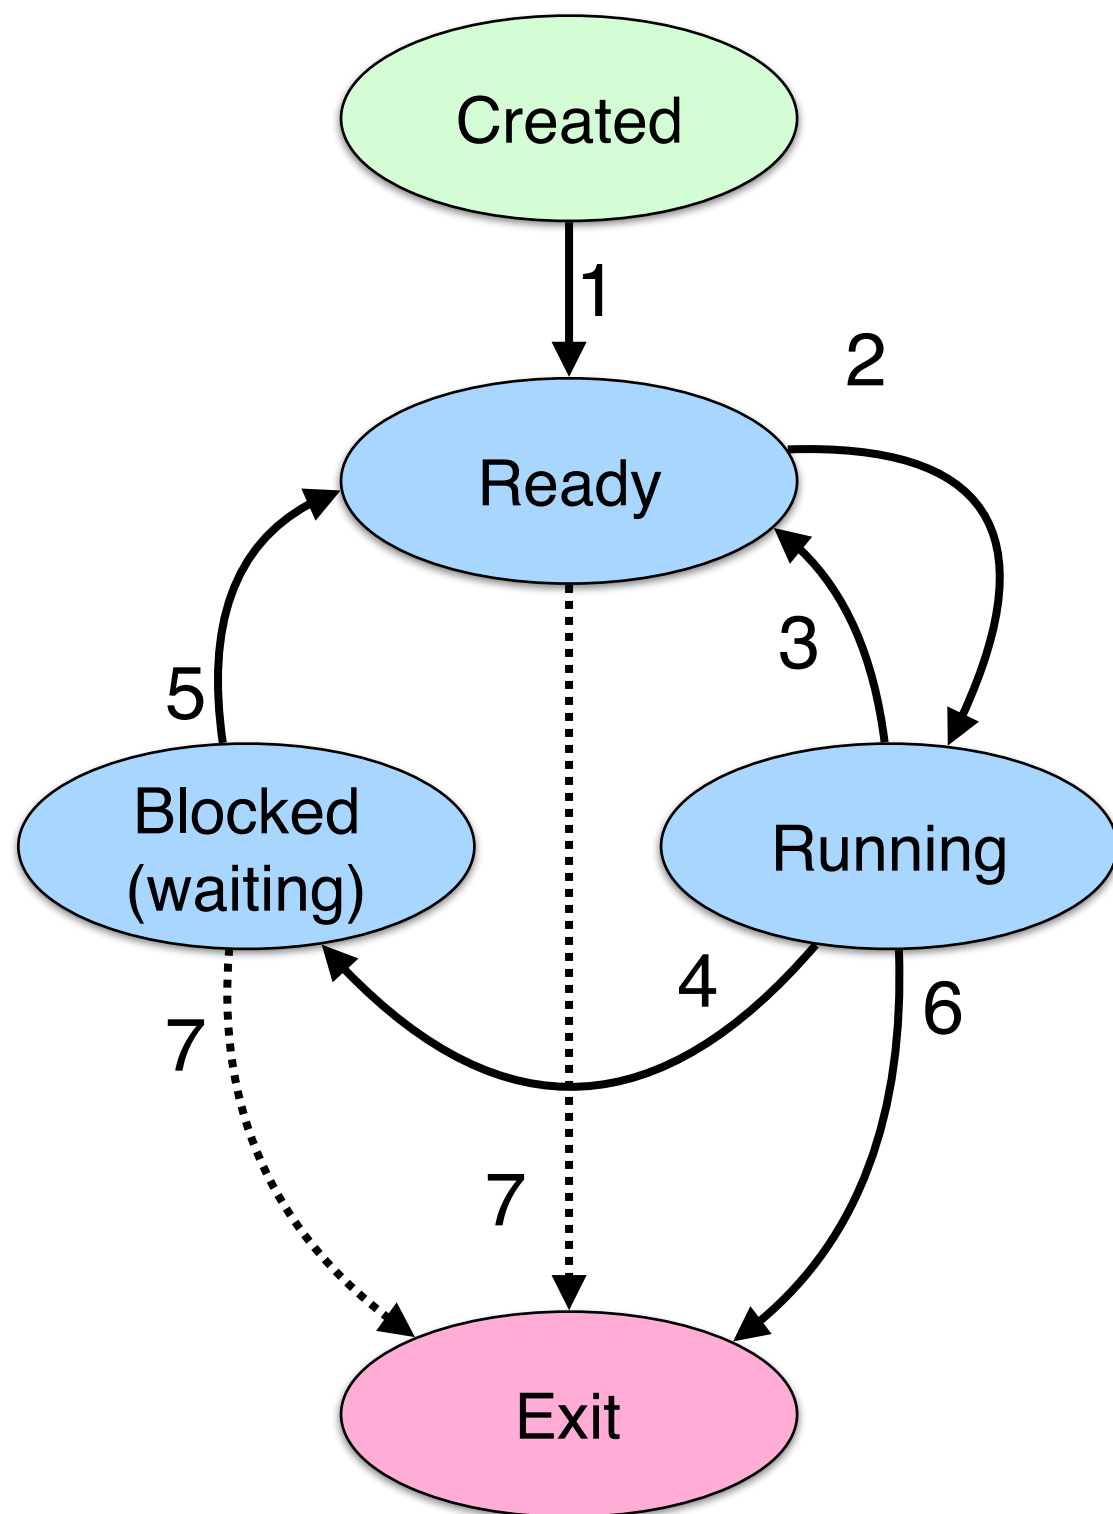
# When do processes end?

- ❖ Conditions that terminate processes can be
  - Voluntary
  - Involuntary
- ❖ Voluntary
  - Normal exit
  - Error exit
- ❖ Involuntary
  - Fatal error (only sort of involuntary)
  - Killed by another process

# Process hierarchies

- ❖ Parent creates a child process
  - Child processes can create their own children
- ❖ Forms a hierarchy
  - UNIX calls this a “process group”
  - If a process terminates, its children are “inherited” by the terminating process’s parent
- ❖ Windows has process groups
  - Multiple processes grouped together
  - One process is the “group leader”

# Process states



## ❖ Process in one of 5 states

- Created
- Ready
- Running
- Blocked
- Exit

## ❖ Transitions between states

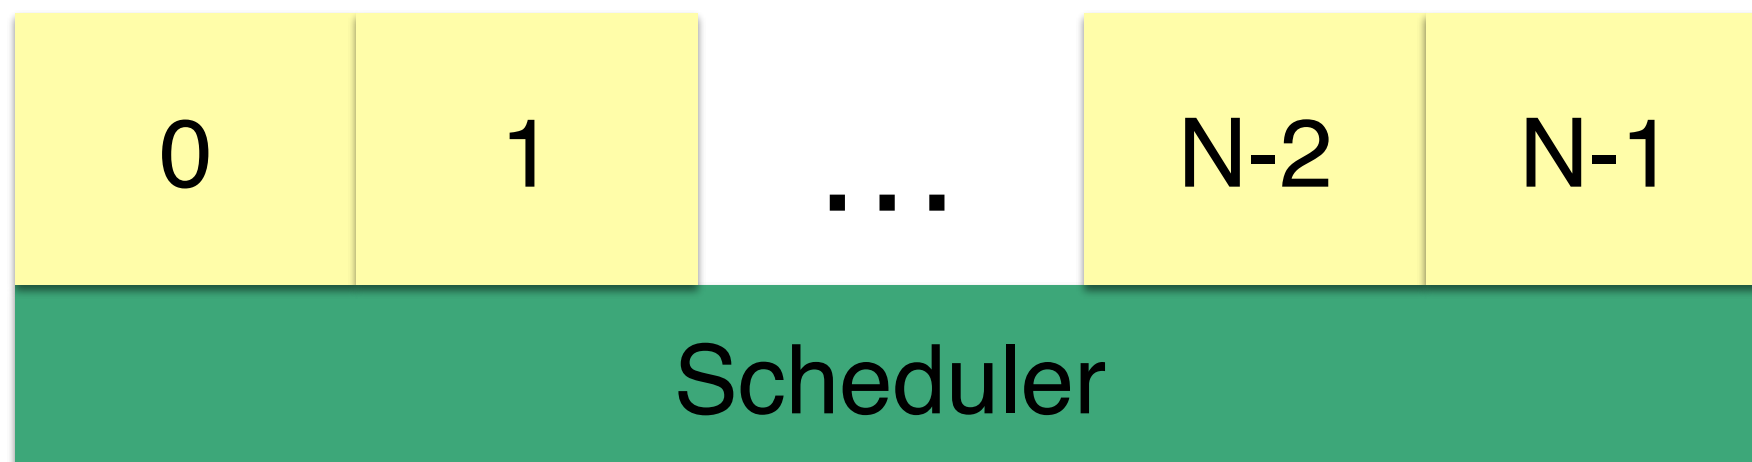
1. Process enters ready queue
2. Scheduler picks this process
3. Scheduler picks a different process
4. Process waits for an event such as I/O
5. Event occurs
6. Process exits
7. (Process ended by another process)



# Processes in the OS

- ❖ Two “layers” for processes
- ❖ Lowest layer of process-structured OS handles interrupts, scheduling
- ❖ Above that layer are sequential processes
  - Processes tracked in the process table
  - Each process has a process table entry

## Processes



# What's in a process table entry?

May be  
stored  
on stack

## Process management

Registers  
Program counter  
CPU status word  
Stack pointer  
Process state  
Priority / scheduling parameters  
Process ID  
Parent process ID  
Signals  
Process start time  
Total CPU usage

## File management

Root directory  
Working (current) directory  
File descriptors  
User ID  
Group ID

## Memory management


Pointers to text, data, stack  
- or -  
Pointer to page table

# What happens on a trap/interrupt?

1. Hardware saves program counter (on stack or in a special register)
2. Hardware loads new PC, identifies interrupt
3. Assembly language routine saves registers
4. Assembly language routine sets up stack
5. Assembly language calls C to run service routine
6. Service routine calls scheduler
7. Scheduler selects a process to run next (might be the one interrupted...)
8. Assembly language routine loads PC & registers for the selected process

# Threads: “processes” sharing memory

- ❖ Process  $\leftrightarrow$  address space
- ❖ Thread  $\leftrightarrow$  program counter / stream of instructions
- ❖ Two examples
  - Three processes, each with one thread
  - One process with three threads



# Process & thread information

## Per process items

Address space  
Open files  
Child processes  
Signals & handlers  
Accounting info  
*Global variables*

## Per thread items

Program counter  
Registers  
Stack & stack pointer  
State (local variables)


## Per thread items

Program counter  
Registers  
Stack & stack pointer  
State (local variables)

## Per thread items

Program counter  
Registers  
Stack & stack pointer  
State (local variables)

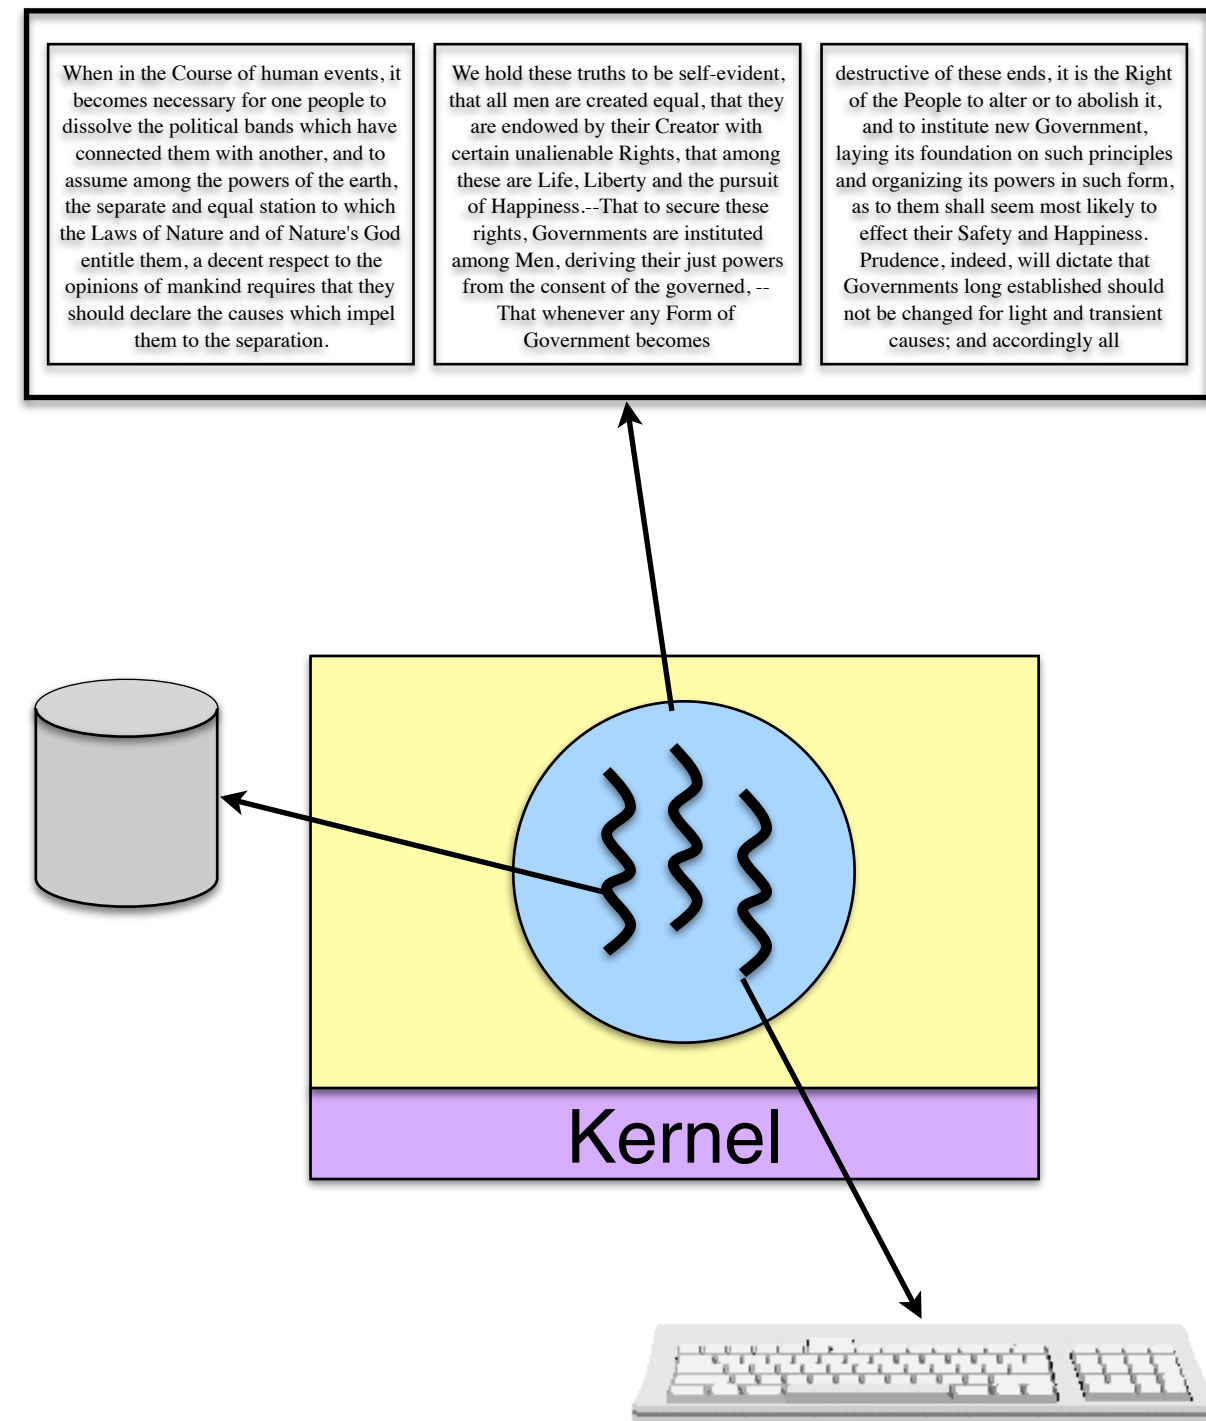
# Threads & stacks



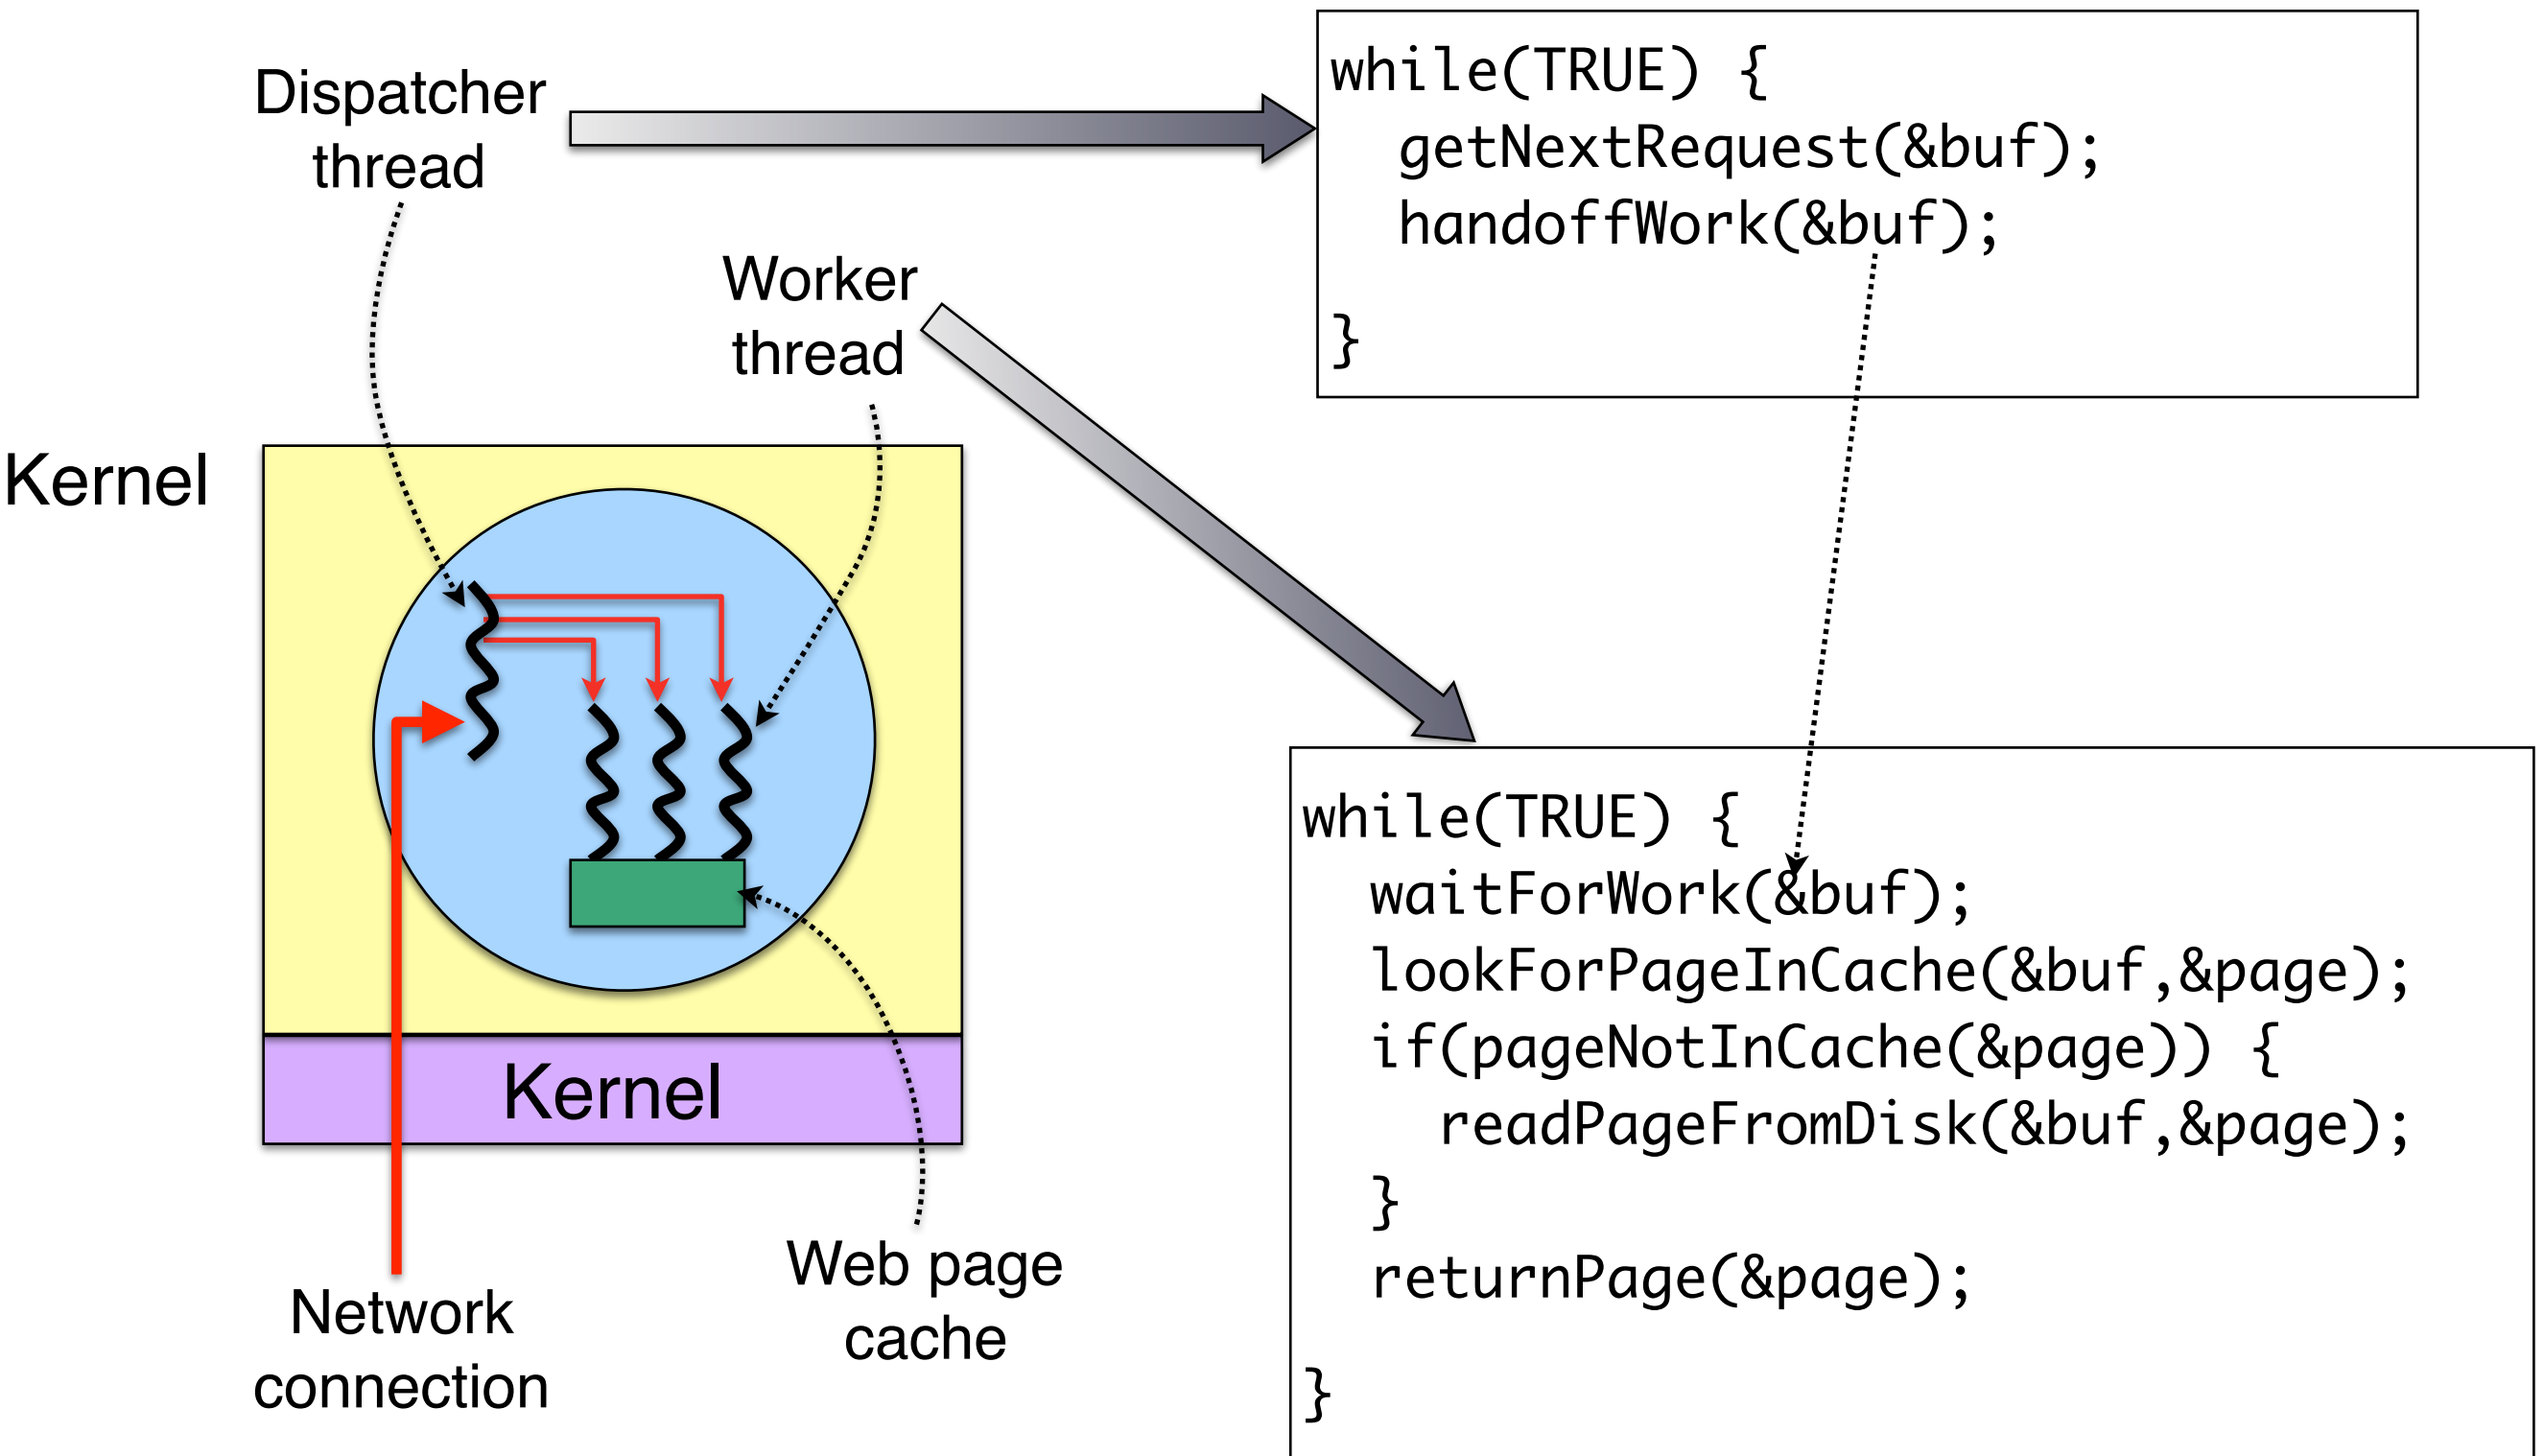
➡ Each thread has its own stack!

# Why use threads?

- ❖ Allow a single application to do many things at once
  - Simpler programming model
  - Less waiting
- ❖ Threads are faster to create or destroy
  - No separate address space
- ❖ Overlap computation and I/O
  - Could be done without threads, but it's harder
- ❖ Example: word processor
  - Thread to read from keyboard
  - Thread to format document
  - Thread to write to disk



# Multithreaded Web server





# Three ways to build a server

## ❖ Multithreaded server

- Parallelism
- Blocking system calls
- May use pop-up threads: create a new thread in response to an incoming message (rather than reusing a thread)

## ❖ Single-threaded process: slow, but easier to do

- No parallelism
- Blocking system calls


## ❖ Finite-state machine (event model)

- Each activity has its own state: states change when system calls complete or interrupts occur
- Parallelism
- Nonblocking system calls

# Issues with using threads

- ❖ May be tricky to convert single-threaded code to multithreaded code
- ❖ Re-entrant code
  - Code must function properly when multiple threads are using it simultaneously
  - Need to be careful when using static or global variables
    - Returned structures
    - Buffers
- ❖ Error management
  - What happens when just a single thread has an error?
  - Can't simply kill the process, since other threads might be running

# Implementing threads



## User-level threads

- + No need for kernel support
- May be slower than kernel threads
- Harder to do non-blocking I/O

## Kernel-level threads

- + More flexible scheduling
- + Non-blocking I/O
- Not (necessarily) portable

# POSIX threads

- ❖ Standard interface to threading library
- ❖ May be implemented in either user or kernel space
  - Some operating systems provide support for both!
- ❖ Allows thread-based programs to be portable

Thread call (Pthread_xx)	Description
create	Create a new thread
exit	Terminate the calling thread
join	Wait for a specific thread to exit
yield	Release the CPU, allowing another thread to run

# Processes & threads in Linux

- ❖ Supports POSIX standard
- ❖ Linux supports kernel-level threads (lightweight processes)
  - Share address space, file descriptors, etc.
  - Each has its own process descriptor in memory
- ❖ Linux processes (incl. lightweight) all have unique identifiers
  - Threads sharing address space are grouped into process groups
  - Identifier shared by the group is that of the leader
- ❖ Each process has its own 8KB region that stores
  - Kernel stack
    - Kernel has a small stack: about 4KB!
  - Low-level thread information
- ❖ Other information stored in a separate data structure
  - Memory allocated to the process
  - Open files
  - Signal information

# Processes & threads in FreeBSD

- ❖ Supports POSIX threads (as does Linux)
- ❖ Processes are heavyweight
  - Unique process ID
  - Individual address space
  - List of associated threads
- ❖ Threads are “variable-weight”
  - Must contain thread control block, thread kernel stack, and thread state
  - Lightest weight
    - Share *all* other process resources!
    - Thread library keeps track of user-level stacks
    - Might be useful for “thread pool” implementations
  - Heavier weight: created by `rfork()`
    - Share fewer resources
    - May have separate stack (likely) and data (possibly) spaces
    - Gets its own process ID
    - Still shares some address space, including global variables