

Scheduling

There cannot be a crisis next week. My schedule is already full.

— *Henry Kissinger*



Scheduling

- ❖ What is scheduling?
 - Goals
 - Mechanisms
- ❖ Scheduling on batch systems
- ❖ Scheduling on interactive systems
- ❖ Other kinds of scheduling
 - Real-time scheduling

Or you could do this ...

```

r1 01/08 05:50 FPS:30 Mem:00E788H800 CPU 3 1 1 1 1 1 1 1 ENTER
MENU 3);.....JukeSo=7FD9D428LX] MENU Task... Unname=7FD9AE28LX]
templeUS V3.12 templeUS V3.12

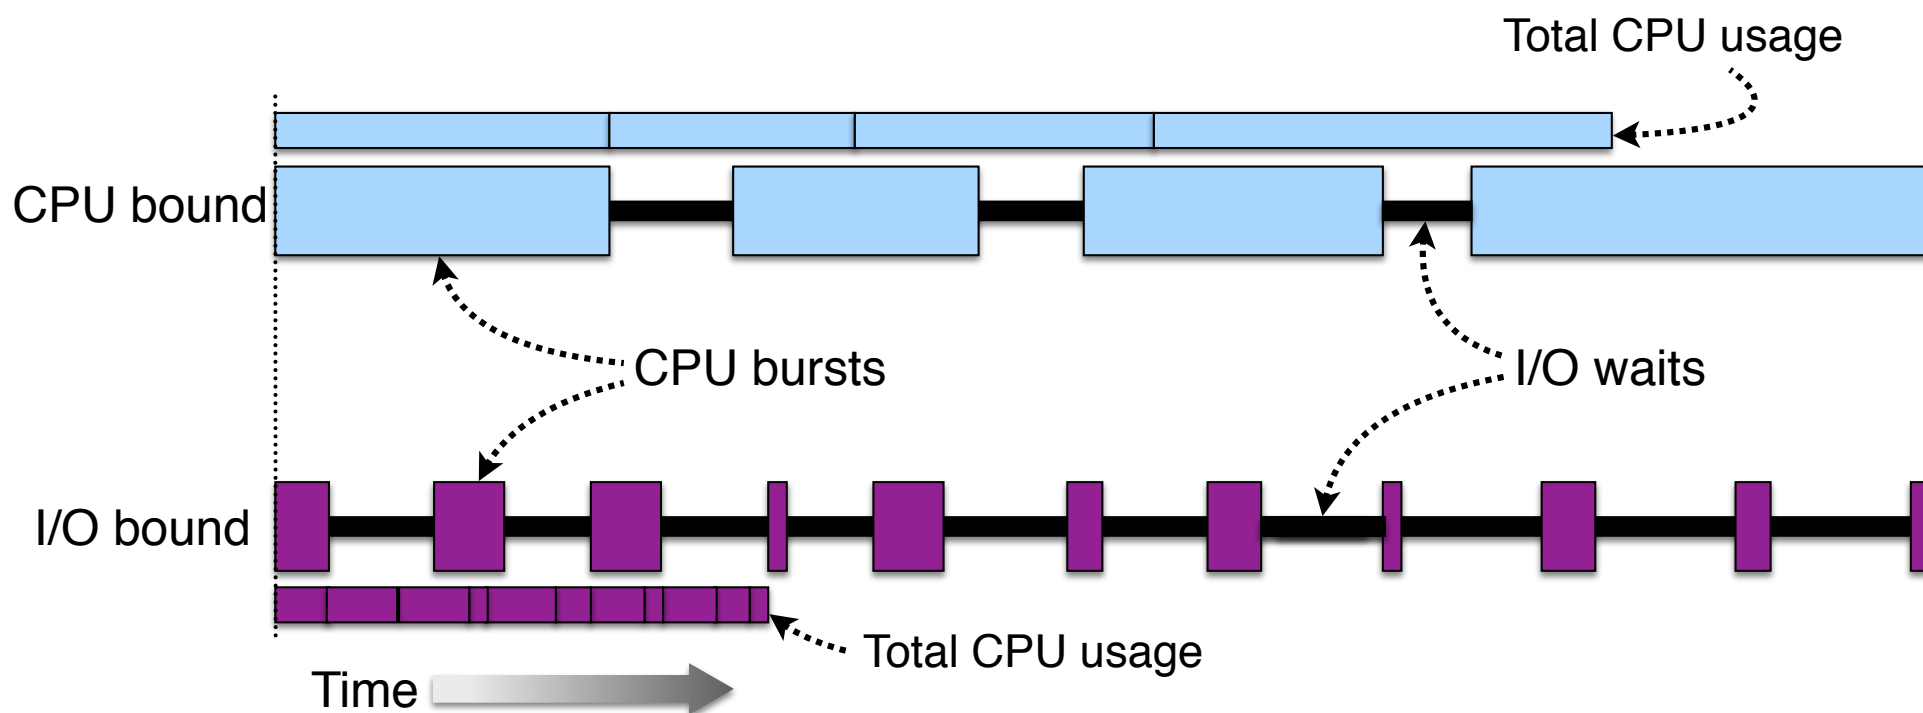
Public Domain Operating System
Help & Index, Quick Start: Cmd line

Directory of C:/Home
DATE TIME SIZE
01/08 04:59 00000 +
01/08 04:59 00000 Blog
01/08 04:59 00000 Budget
01/08 04:59 00000 Code
01/08 04:59 00000 Misc
01/08 04:59 00000 Videos
01/08 04:59 00000 Mb2
08/31 18:10 016E DoOnce.CPP.Z
01/06 10:44 0CDA HomeKeyPlugins.CPP.Z
11/01 19:24 0194 HomeLocalize.CPP.Z
12/03 08:41 00ED HomePkgs.CPP.Z
09/27 09:40 02DC HomeSys.CPP.Z
01/07 21:08 0272 Registry.CPP.Z
12/08 06:54 0742 OSCfg.CPP.Z
09/27 11:52 1E0E OSBistro.CPP.Z
01/07 22:48 211D OSHolySpirit.CPP.Z
07/26 06:26 02C6 OSDComputer.CPP.Z
01/08 05:08 11E4 OSMisc.CPP.Z
01/04 23:24 0025 Test.CPP.Z
01/08 05:00 4040 PersonalMenu.IX1.Z
01/07 21:57 29D4 PersonalNotes.IX1.Z
DFT:80,940LOC:8 Cores 3.900GHZ
C:/Home>Cd("::/Home");#include "DoOnce
";
0.000250s ans=0x00000000=0
C:/Home>JukeSong(53);

EOF

```

Why schedule processes?



- ❖ Bursts of CPU usage alternate with periods of I/O wait
- ❖ Some processes are CPU-bound: they don't make many I/O requests
- ❖ Other processes are I/O-bound and make many kernel requests

When are processes scheduled?

- ❖ At the time they enter the system
 - Common in batch systems
 - Two types of batch scheduling
 - Submission of a new job causes the scheduler to run
 - Scheduling only done when a job voluntarily gives up the CPU (i.e., while waiting for an I/O request)
- ❖ At relatively fixed intervals (clock interrupts)
 - Necessary for interactive systems
 - May also be used for batch systems
 - Scheduling algorithms at each interrupt, and picks the next process from the pool of “ready” processes

Scheduling goals

❖ All systems

- Fairness: give each process a fair share of the CPU
- Enforcement: ensure that the stated policy is carried out
- Balance: keep all parts of the system busy

❖ Batch systems

- Throughput: maximize jobs per unit time (hour)
- Turnaround time: minimize time users wait for jobs
- CPU utilization: keep the CPU as busy as possible

❖ Interactive systems

- Response time: respond quickly to users' requests
- Proportionality: meet users' expectations

❖ Real-time systems

- Meet deadlines: missing deadlines is a system failure!
- Predictability: same type of behavior for each time slice

Measuring scheduling performance

❖ Throughput

- Amount of work completed per second (minute, hour)
- Higher throughput usually means better utilized system

❖ Response time

- Response time is time from when a command is submitted until results are returned
- Can measure average, variance, minimum, maximum, ...
- May be more useful to measure time spent waiting
- Can also measure how often response time is faster than a given time (e.g., 100 ms): useful for real-time systems and servers

❖ Turnaround time

- Like response time, but for batch jobs (response is the completion of the process)

❖ Usually not possible to optimize for all metrics with a single scheduling algorithm

Interactive vs. batch scheduling

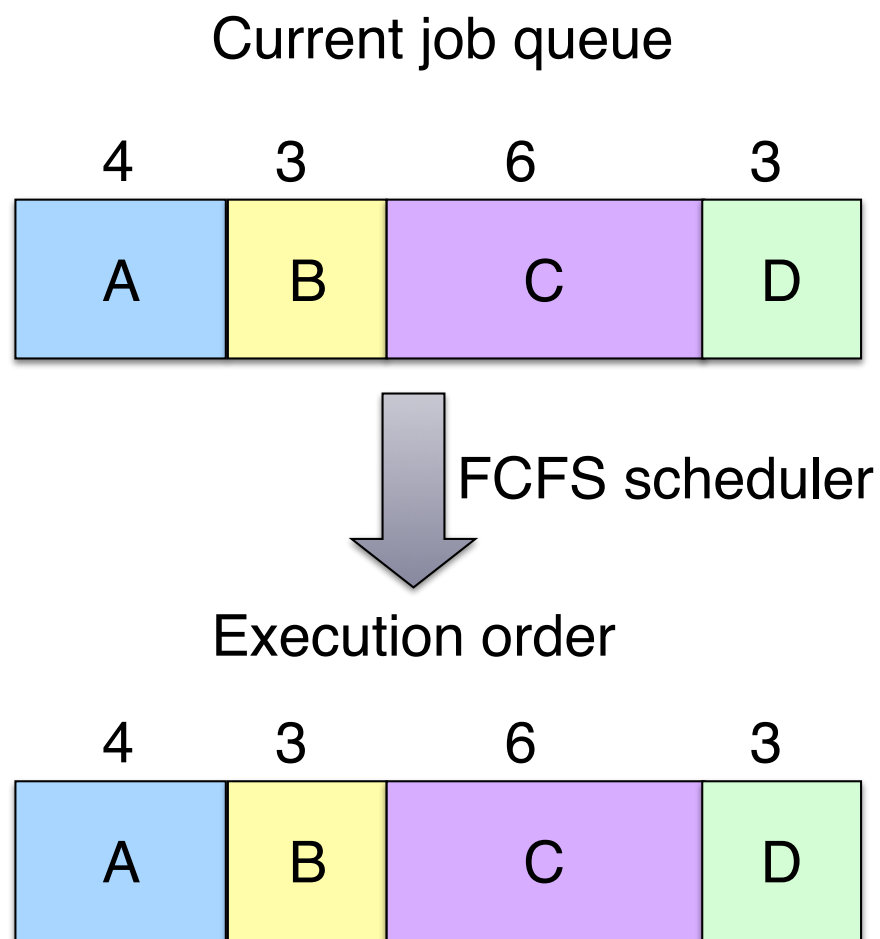
Batch

First-Come-First-Served (FCFS)
Shortest Job First (SJF)
Shortest Remaining Time First (SRTF)
Priority (non-preemptive)

Interactive

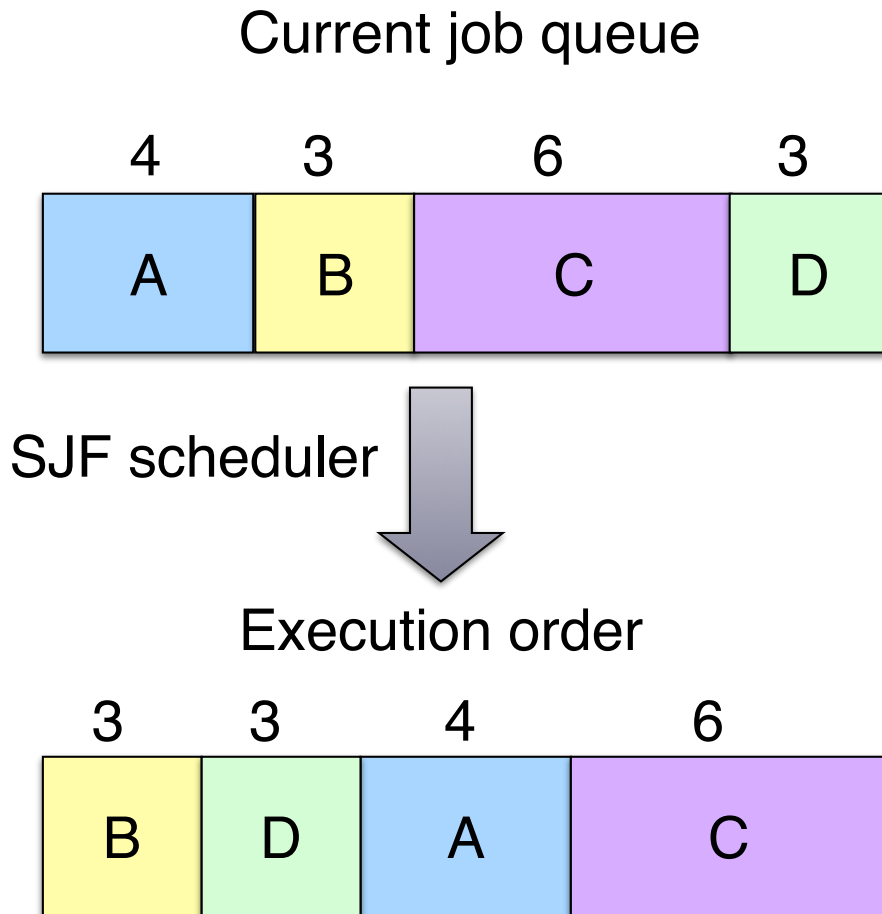
Round-Robin (RR)
Priority (preemptive)
Multi-level feedback queue
Lottery scheduling

First Come, First Served (FCFS)



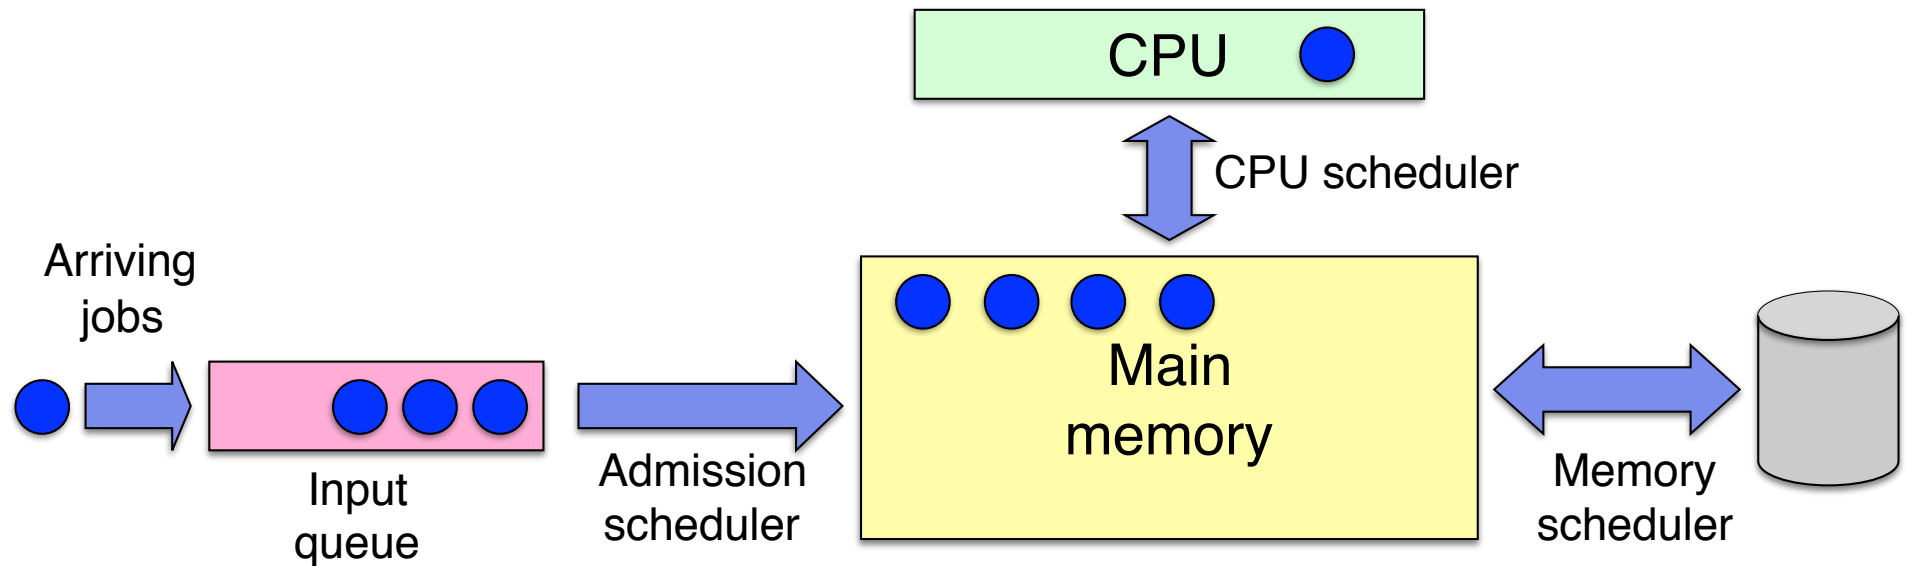
- ❖ Goal: do jobs in the order they arrive
 - Fair in the same way a bank teller line is fair
- ❖ Simple algorithm!
- ❖ Problem: long jobs delay every job after them
 - Many processes may wait for a single long job

Shortest Job First (SJF)



- ❖ Goal: do the shortest job first
 - Short jobs complete first
 - Long jobs delay every job after them
- ❖ Jobs sorted in increasing order of execution time
 - Ordering of ties doesn't matter
- ❖ Shortest Remaining Time First (SRTF): preemptive form of SJF
 - Re-evaluate when a new job is submitted
- ❖ Problem: how does the scheduler know how long a job will take?

Three-level scheduling



- ❖ Jobs held in input queue until moved into memory
 - Pick “complementary jobs”: small & large, CPU- & I/O-intensive
 - Jobs move into memory when admitted
- ❖ CPU scheduler picks next job to run
- ❖ Memory scheduler picks some jobs from main memory and moves them to disk if insufficient memory space

Round Robin (RR) scheduling


❖ Scheduling interactive processes

- Give each process a fixed time slot (quantum)
- Rotate through “ready” processes
- Each process makes some progress

❖ What’s a good quantum?

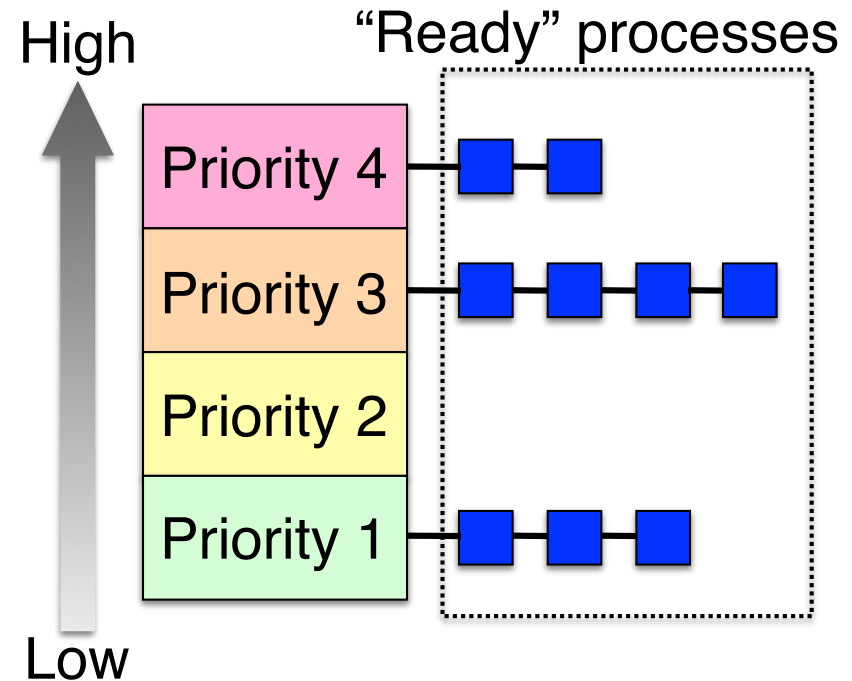
- Too short: many process switches hurt efficiency
- Too long: poor response to interactive requests
- Typical length: 10–100 ms

❖ “Strict” rotation: round robin



Priority scheduling

- ❖ Assign a priority to each process
 - “Ready” process with highest priority allowed to run
 - Running process may be interrupted after its quantum expires
- ❖ Priorities may be assigned dynamically
 - Reduced when a process uses CPU time
 - Increased when a process waits for I/O
- ❖ Often, processes grouped into multiple queues based on priority, and run round-robin per queue



Shortest process next

- ❖ Run the process that will finish the soonest
 - In interactive systems, job completion time is unknown!
- ❖ Guess at completion time based on previous runs
 - Update estimate each time the job is run
 - Estimate is a combination of previous estimate and most recent run time
- ❖ Not often used because round robin with priority works so well!


Lottery scheduling

- ❖ Give processes “tickets” for CPU time
 - More tickets \Rightarrow higher share of CPU
- ❖ Each quantum, pick a ticket at random
 - If there are n tickets, pick a number from 1 to n
 - Pseudo-random number is OK if it's a good RNG
 - Process holding the ticket gets to run for a quantum
 - This can be implemented efficiently without “real” tickets
 - Track range of tickets belonging to each process
- ❖ Over the long run, each process gets the CPU m/n of the time if the process has m of the n existing tickets
- ❖ Tickets can be transferred
 - Cooperating processes can exchange tickets
 - Clients can transfer tickets to a server so it can have a higher priority
 - Parent (shell) can transfer tickets to a child process

Scheduling in BSD4

- ❖ Quantum is 100 ms: longest that's OK for interactive scheduling
- ❖ Scheduler is based on multi-level feedback queues
 - Priority is based on two things
 - Resource requirements: blocked threads have higher priority when rescheduled
 - Previous CPU usage: CPU hogs have lower priority
- ❖ Each thread is placed into a run queue for its priority
 - Head of highest-priority run queue with a ready thread runs next

run queues



Calculating priority in BSD4


- ❖ Thread priority is set by:

$$pri = MIN + [estcpu/4] + 2 \times nice$$

- Values above *MAX* are set to *MAX*
 - *MIN*=160, *MAX*=223
 - *nice* is set by the user to manually lower thread priority
 - *estcpu* is an estimate of the number of “ready” processes in the CPU when the calculation is made
 - Has a bit of “memory” so it doesn’t change too quickly
 - *estcpu* is updated each clock tick
 - Higher numbers indicate lower priority: threads with lowest priority values are scheduled first
- ❖ Thread priority is set every 40 ms
 - ❖ Scheduling is more complex for multiprocessors...

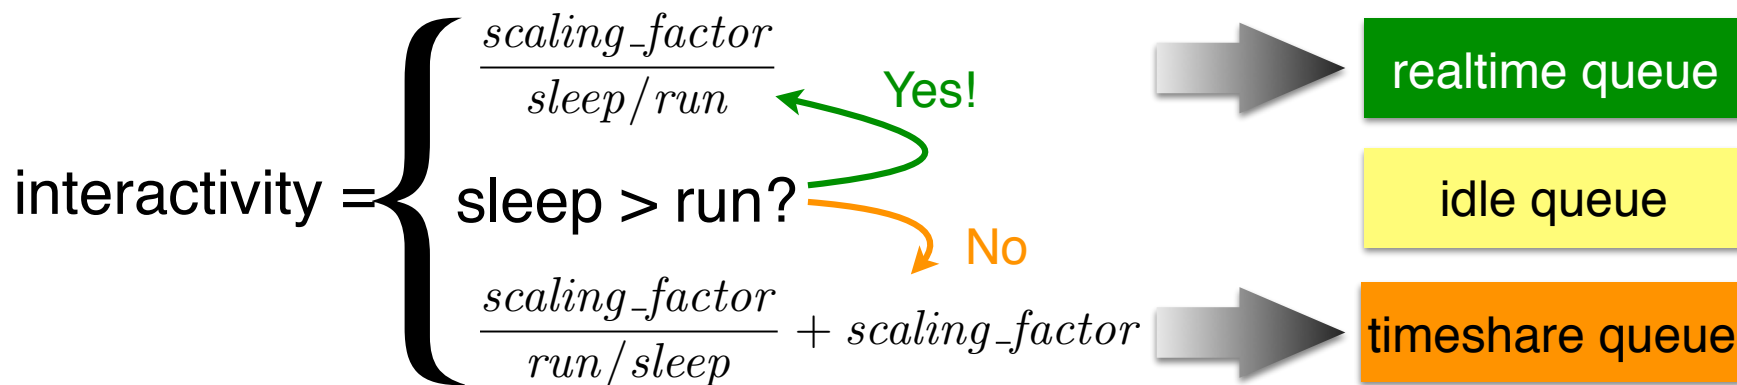
FreeBSD ULE scheduler

- ❖ Not an acronym: scheduler code is in `sched_ule.c`
- ❖ Divided into two parts
 - Low-level scheduler: runs on every context switch
 - Selects next runnable thread from the run queues
 - Uses bitmask to quickly find first non-empty queue
 - High-level scheduler: runs a few times per second
 - Sets thread priorities
 - Selects a processor on which to run a particular thread
- ❖ Run queues only contain runnable threads
 - Blocked threads placed on
 - Turnstile: blocked on short-term lock
 - Sleepqueue: blocked on medium/long-term lock
 - Runnable thread that consumes its quantum placed at the end of its run queue



FreeBSD ULE high-level scheduler

- ❖ High-level scheduler runs infrequently
 - Adjusts length of quantum
 - Adjusts process priority
- ❖ Prioritize interactive processes: use calculated “interactivity” score
- ❖ Three sets of queues per CPU
 - Idle queue: idle threads
 - Realtime queue: real-time & interactive threads
 - Timeshare queue: long-running (batch) threads
 - Uses “calendar queue” to handle lots of batch-type threads efficiently
 - Lower priority processes inserted further back in queue: take longer to bubble to front



Scheduling in Linux

❖ Three classes of processes

- Conventional
- Real-time (round-robin)
- Real-time (FIFO)

❖ Queue structure similar to BSD

- Two sets of queues (0–99 real-time, 100–139 conventional): active and expired
- Scheduler runs process in lowest-valued active queue
- Conventional threads placed in expired queue when their quantum is up
 - May be scheduled out before quantum expires: put back into active queue
- Real-time threads placed back into active queue

Linux: rescheduling processes


- ❖ Re-evaluate priorities when there are no threads on the active queue
- ❖ Thread priority is re-evaluated based on previous run time and sleep time
 - Threads that ran more and slept less are penalized with a worse priority
- ❖ Quantum is regenerated based on original (static) priority
 - Higher-priority threads are given longer quanta
 - Process may lose the CPU before quantum is up (if a higher-priority thread becomes available)

Further details in *Understanding Linux Kernel Internals* (3rd edition), Bovet & Cesati

Policy versus mechanism


- ❖ Separate what may be done from how it is done
 - Mechanism allows
 - Priorities to be assigned to processes
 - CPU to select processes with high priorities
 - Policy set by what priorities are assigned to processes
- ❖ Scheduling algorithm parameterized
 - Mechanism in the kernel
 - Priorities assigned in the kernel or by users
- ❖ Parameters may be set by user processes
 - Don't allow a user process to take over the system!
 - Allow a user process to voluntarily lower its own priority
 - Allow a user process to assign priority to its threads

Scheduling user-level threads



- ❖ Kernel picks a process to run next
- ❖ Run-time system (at user level) schedules threads
 - Run each thread for less than process quantum
 - Example
 - Processes get 40ms each
 - Threads get 10ms each
- ❖ Example schedule:
A1,A2,A3,A1,B1,B3,B2,B3
- ❖ Not possible:
A1,A2,B1,B2,A3,B3,A2,B1

Scheduling kernel-level threads



- ❖ Kernel schedules each thread
 - No restrictions on ordering
 - May be more difficult for each process to specify priorities
- ❖ Example schedule:
A1, A2, A3, A1,
B1, B3, B2, B3
- ❖ Also possible:
A1, A2, B1, B2,
A3, B3, A2, B1