

**Seongwoo Choi**  
**CMPS 111**  
**Winter 2017**  
**Prof. Darrell Long**  
**Week 5**  
**Notes**

**February 6th 2017 - February 8th 2017**

### Critical sections using semaphores

Define a class called Semaphore

- Class allows more complex implementations for semaphores
- Details hidden from processes

Code for individual process is simple.

### Implementing semaphores with blocking

Assume two (existing) operations:

sleep (): suspends current process

Wakeup (): allows process P to resume execution

### Semaphore is a class

Track value of semaphore

Keep a list of processes waiting for the semaphore

Operations still atomic

```
class Semaphore () {  
    int value;  
    ProcessList p1;  
    void down ();  
    void up ();  
};
```

### Semaphores for general synchronization

We want to execute B in P1 only after A executes in P0

Use a semaphore initialized to 0

Use up() to notify P1 at the appropriate time

This is called a rendezvous.

### Types of semaphores

Two different types of semaphores

Counting semaphores

Binary semaphores

Counting semaphore

Value can range over an unrestricted range

### Binary semaphore

Only two values possible

Value 1 means the semaphore is available

Value 0 means a process has acquired the semaphore

Maybe simpler to implement

Possible to implement one type using the other (midterm)

Deadly Embrace

Dijkstra was romantic. He did not write codes. He wrote only proofs.

### Monitors

A monitor is another kind of high-level synchronization primitive

One monitor has multiple entry points

Only one process may be in the monitor at any time

Enforces mutual exclusion: better at avoiding programming errors

### Monitors provided by high-level language

Variables belonging to monitor are protected from simultaneous access

Procedures in monitor are guaranteed to have mutual exclusion

### Monitor implementation

Language/compiler handles implementation

Can be implemented using semaphores

Xerox invented ethernet, video display, and mouse and Apple stole them

### Monitor Usage

This looks like C++ code, but it's not supported by C++

Provides the following features:

Variables foo, bar, and arr are accessible only by func1 & func2

Only one process can be executing in either func1 or func2 at any time

```
monitor mon {
    Int foo;
    Int bar;
    Double arr[100];
    Void func1(_) {
    }
    Void func2(_) {
    }
    Void mon() { // initialization code
    }
}
```

}

Set up process called semaphore process and give variables. I send a message and counter is bigger than 0 it executes something, but if it is zero, then wait.

If you have a monitor, then you do not want to use a monitor.

### Monitor semantics

Problem: P signals on condition variable X in monitor M, waking Q

Both can't be active in the monitor at the same time

Which one continues first?

### Mesa semantics

Signaling process (P) continues first

Q resumes when P leaves the monitor

Seems more logical: why suspend P when it signals?

### Hoare semantics

Awakened process (Q) continues first

P resumes when Q leaves the monitor

May be better: condition that Q wanted may no longer hold when P leaves the monitor

### Locks & condition variables

Monitors require native language support

Instead, provide monitor support using special

### Implementing locks with semaphores

Use mutex to ensure exclusion within the lock bounds

Use next

Peterson's algorithm was mentioned during the class.

Switch the order of something. Everything works

### Message passing

Synchronize by exchanging messages

Two primitives:

Send: send a message

Receive: receive a message

Both may specify a "channel" to use

Issue: how does the sender know the receiver got the message?

Issue: authentication.

### Barriers

Used for synchronizing multiple processes  
Processes wait at a “barrier” until all in the group arrive  
After all have arrived, all processes can proceed  
May be implemented using locks and condition variables.

### Implementing barriers using semaphores

Barrier b;  
B.bsem.value

### Deadlock and starvation

Deadlock: two or more processes are waiting indefinitely for an event that can only be caused by a waiting process

P0 gets A, needs B

P1 gets B, needs A

Each process waiting for the other to signal

### Starvation: indefinite blocking

Process is never removed from the semaphore queue in which it is suspended  
May be caused by ordering in queues (priority)

### Livelock

Sometimes, processes can still run, but not make progress

Example; two processes want to use resources A and B

P0 gets A, P1 gets B

Each realized that a deadlock will occur if then proceed as planned

P0 drops A, P1 drops B

Same problem as before

This can go on for a very long time

Real-world example: Ethernet transmission collisions

If there's a “collision” on the wire, wait and try again

Multiple processes waited the exact same amount of time...

### Classical synchronization problems

#### Bounded buffer problem

Goal: implement producer-consumer without busy waiting

#### Readers-writers problem

Shared variables

Int nreaders;

Semaphores mutex(1), writing(1);

Reader process

```

mutex.down();
Nreaders += 1;
If (nreaders == 1) // wait if
    Writing.down(); // 1st reader
mutex.up();
// Read some stuff
mutex.down()

```

### Dining Philosophers

N philosophers around a table  
 All are hungry  
 All like to think  
 N chopsticks available  
 1 between each pair of philosophers  
 Philosophers need two chopsticks to eat  
 Philosophers alternate between eating and thinking  
 Goal: coordinate use of chopsticks

#### Solution 1

Use a semaphore for each chopstick  
 A hungry philosopher  
 Gets the chopstick to his right  
 Gets

#### Solution 2

Gets lower, then higher numbered chopstick  
 Eats  
 Puts down the chopsticks  
 Potential problems?  
 Deadlock  
 Fairness

#### Shared variables

```

Const int n;
// initialize to 1
Semaphore chopstick[n];

```

#### Code for philosopher

```

Int i1, i2;
While (1) {
    If (i != (n-1)) {
        I1
    }
}

```

Dining philosophers with locks

Next time, we will talk about memory. Virtual memory

Memory

Deadlocks

But crises and deadlocks when they occur have at least this advantage, that they force us to think.

Jawaharlal Nehru

### Resources

-> Resource: something a process uses

Usually limited (at least somewhat)

-> Examples of computer resources

Printers

Semaphores / locks

Memory

Tables (in a database)

-> Processes need access to resources in reasonable order

-> Two types of resources:

Preemptable resources: can be taken away from a process with no ill effects

Non Preemptable resources: will cause the process to fail if taken away

Still millions of dollars are used in manufacturing tapes

The property of tape is sequential.

### Using resources

-> Sequence of events required to use a resource

Request the resource

Use the resource

Release the resource

-> Can't use the resource if request is denied

Requesting process has options

Block and wait for resource

Continue (if possible) without it: may be able to use an alternate resource

Process fails with error code

Some of these may be able to prevent deadlock

When do deadlocks happen?

Suppose

Process 1 holds resource A and resource B

Process 2 holds B and requests A

Both can be blocked, with neither able to proceed

Deadlocks occur when ...

Processes are granted

### What is a deadlock?

Formal definition:

“A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.”

Usually, the event is release of a currently held resource

In deadlock, none of the processes can

Run Release resources

Be awakened

### Four conditions for deadlock

Mutual exclusion

Each resource is assigned to at most one process

Hold and wait

A process holding resources can request more resources

No preemption

Previously granted resources cannot be forcibly taken away

Circular wait

There must be a circular chain of 2 or more processes where each is waiting for a resource held by the next member of the chain

### Resource allocation graphs

Resource allocation modeled by directed graphs

Example 1:

Resource R assigned to process A

Example 2:

Process B is requesting / waiting for resource S

Example 3:

Process C holds T, waiting for U

Process D holds U, waiting for T

C and D are in deadlock

### Dealing with deadlock

How can the OS deal with deadlock?

Ignore the problem altogether!

Hopefully, it'll never happen ...

Detect deadlock & recover from it

Dynamically avoid deadlock

Careful resource allocation

Prevent deadlock

Remove at least one of the four necessary conditions

We'll explore these tradeoffs

### Getting into deadlock

A	B	C
Acquire R Acquire S Release R Release S	Acquire S Acquire T Release S Release T	Acquire T Acquire R Release T Release R

Lives of the Others

The Ostrich Algorithm

Pretend there's no problem

Reasonable if

Deadlocks occur very rarely

Cost of prevention is high

UNIX and Windows take this approach

Resources (memor

Not getting into deadlock

Many situations may result in deadlock (but don't have to)

In previous example, A could release R before C requests R, resulting in no deadlock

Can we always get out of it this way?

Find ways to:

Detect deadlock and reverse it

Stop it from happening in the first place

### Detecting deadlocks using graphs

Process holdings and requests in the table and in the graph

Graph contains a cycle => deadlock

Easy to pick out by looking at it (in this case)

Need to mechanically detect deadlocks

### Deadlock detection algorithm

General idea: try to find cycles in the resource allocation graph

Algorithm: depth-first search at each code

Mark arcs as they're traversed

Build list of visited nodes

If node to be added is already on the list, a cycle exists!

Cycle => deadlock



## Recovering from deadlock

### ->Recovery through preemption

- Take a resource from some other process

- Depends on nature of the resource and the process

### ->Recovery through rollback

- Checkpoint a process periodically

- Use saved state to restart the process if it's in deadlock

- May present a problem if the process affects lots of "external" things

### ->Recovery through killing processes

- Crudest but simplest way to break a deadlock: kill one of the processes in the deadlock cycle

- Other processes can get its resources

- Try to choose a process that can be rerun from the start

- Pick one that hasn't run too far already

## Preventing Deadlock

Deadlock can be completely prevented

Ensure that at least one of the conditions for deadlock never occurs

Mutual exclusion

Circular wait

Hold & wait

## Eliminating mutual exclusion

Some devices (such as printer) can be spooled

- Only the printer daemon uses printer resource

- This eliminates deadlock for printer

Not all devices can be spooled

Principle:

- Avoid assigning resource when not absolutely necessary

- As few processes as possible actually claim the resource

Serialization

## Attacking "hold and wait"

Require processes to request resources before starting

A process never has to wait for what it needs

This can present problems

A process may not know required resources at start of run

This also ties up resources other processes could be using

Processes will tend to be conservative and request resources they might need

Variation: a process must give up all resources before making a new request

Process is then granted all prior resources as well as the new ones

Problem: what if someone grabs the resources in the meantime -- how can the process save its state?

### Attacking “no preemption”

This is not usually a viable option

Consider a process given the printer

Halfway through its jobs, take away the printer

Confusion ensues!

May work for some resources

Forcibly take away memory pages, suspending the process

Process may be able to resume with no ill effects

### Attacking “circular wait”

Assign an order to resources

Always acquire resources in numerical order

Need not acquire them all at once

Circular wait is prevented A process holding resource  $n$  can't wait for resource  $m$  if  $m < n$

No way to complete a cycle!

Place processes above the highest resource they hold and below any they're requesting

All arrows point up

### What does FreeBSD do?

What resources are at issue?

Locks and semaphores: one holder at a time

Physical resources: not typically

### Deadlock prevention: summary

Condition	Prevented by
Mutual exclusion	Spool everything
Hold and wait	Request all resources initially
No preemption	take resources away if there's not a complete set
Circular wait	Order resources numerically

### Example: two-phase locking

Phase One

Process tries to lock all data it needs, one at a time

If needed data found locked, start over (no real work done in phase one)

Phase Two

Banker's algorithm for a single resource

Banker's algorithm with multiple resources

Starvation

Algorithm to allocate a resource

Give the resource to the shortest job first

Works great for multiple short jobs in a system

May cause long jobs to be postponed indefinitely

Even though not blocked

Solution

First-come, first-serve policy

Starvation

### Livelock

-> Sometimes, processes can still run, but not make progress

-> Example: two processes want to use resources A and B

P0 gets A, P1 gets B

Each realizes that a deadlock will occur if they proceed as planned

P0 drops A, P1 drops B

P0 gets B, P1 gets A

Same problem as before

This can go on for a very long time...

-> Real-world example: Ethernet transmission collisions

If there's a "collision" on the wire, wait and try again

Multiple processes waited the exact same amount of time ...

### Exam

Lecture first One hour long

Short test

Bring a page of notes -> not from the slides (create your own notes)

Double sided

8 pt font

Deadlocks

Scheduling, IPC, Processes

Little memory (easy)