

# Relational Algebra

Instructor: Shel Finkelstein

*Reference:*

*A First Course in Database Systems,  
3<sup>rd</sup> edition, Chapter 2.4 – 2.6, plus Query Execution Plans*

# Important Notices

- Midterm and Midterm Answers have been posted on Piazza.
  - Midterm should be graded by the end of next week.
  - Grades will be posted to Canvas, and exam will be returned in class.
- Lab3 assignment was posted on Monday, Feb 6.
  - Due by Sunday, Feb 26, 11:59pm (3 weeks)
  - There will be Lab Sessions during all 3 weeks.
  - Lab3 has lots of parts and is worth 16 points, not 10 points.
- Gradiance 3 was posted on Tuesday, February 14
  - Due by Tuesday, February 21, 11:59pm.
  - 6 problems on Views, Transactions and Referential Integrity.

# What is a Data Model?

- A *data model* is a mathematical formalism that consists of three parts:
  1. A notation for describing and representing data (structure of the data)
  2. A set of operations for manipulating data.
  3. A set of constraints on the data.
- What is the associated query language for the relational data model?

# Two Query Languages

- Codd proposed two different query languages for the relational data model.
  - Relational Algebra
    - Queries are expressed as a sequence of operations on relations.
    - Procedural language.
  - Relational Calculus
    - Queries are expressed as formulas of first-order logic.
    - Declarative language.
- **Codd's Theorem:** The Relational Algebra query language has the same *expressive power* as the Relational Calculus query language.

# Procedural vs. Declarative Languages

- **Procedural program**
  - The program is specified as a sequence of operations to obtain the desired the outcome. I.e., *how* the outcome is to be obtained.
  - E.g., Java, C, ...
- **Declarative program**
  - The program specifies *what* is the expected outcome, and not *how* the outcome is to be obtained.
  - E.g., Scheme, Ocaml, ...

# SQL – Structured Query Language

- Is SQL a procedural or a declarative language?
  - SQL is usually described as declarative, but it's not fully declarative
  - However, relational database systems usually try to understand meaning of query, regardless of how query is expressed
    - There may be multiple equivalent ways to write a query
- SQL is the principal language used to describe and manipulate data stored in relational database systems.
  - Frequently pronounced as “Sequel”, but formally it's “Ess Cue El”
  - Not the same as Codd's Relational Algebra or Relational Calculus

# Some Properties of Good Database Query Languages and Database Systems

1. Physical database independence
    - Programmers should be able to write queries without understanding the mechanics of the physical layer
    - What was logical data independence?
  2. Highly expressive
    - Programmers should be able to formulate simple and complex queries using the language.
  3. Efficient execution
    - Systems should be able to compute answers to queries with “good” response time and throughput.
- Physical data independence is achieved by most query languages today.
  - Increased expressiveness may come at the expense of not-so-good performance on some complex queries

# Relational Algebra

- Relational Algebra: a query language for manipulating data in the relational data model.
  - Not used directly as a query language
- Internally, Relational Database Systems transform SQL queries into trees/graphs that are similar to relational algebra expressions.
  - Query analysis, transformation and optimization are performed based on these relational algebra expression-like representations.
  - Relational Databases use multi-sets/bags, but Relational Algebra is based on sets.
    - There are multi-set variations of Relational Algebra that permit duplicates, and that's more realistic for Relational Database ...
    - ... but we'll only discuss set-based Relational Algebra.



# Composition

- Each Relational Algebra operator is either a unary or a binary operator.
- A complex Relational Algebra expression is built up from basic ones by composing simpler expressions.
- This is similar to SQL queries and views.

# Relation Algebra Operators

- Queries in relational algebra are composed using basic operations or functions.
  - Selection (  $\sigma$  )
  - Projection (  $\pi$  )
  - Set-theoretic operations:
    - Union (  $\cup$  )
    - Set-difference (  $-$  )
    - Cross-product (  $\times$  )
    - Intersection (  $\cap$  )
  - Renaming (  $\rho$  )
  - Natural Join (  $\bowtie$  ), Theta-Join (  $\bowtie_{\theta}$  )
  - Division (  $/$  or  $\div$  )

# Relation Algebra Operators

- Codd proved that the relational algebra operators ( $\sigma$  ,  $\pi$  ,  $\times$  ,  $\cup$  ,  $-$  ) are independent of each other. That is, you can't define any of these operators using the others.
- However, there are other important operators that can be expressed using ( $\sigma$  ,  $\pi$  ,  $\times$  ,  $\cup$  ,  $-$  )
  - Theta Join, Join, Natural Join, Semi-Join
  - Set Intersection
  - Division
  - Outer Join (sections 5.2.7 and 6.3.8), which we'll discuss when we get to OLAP, On-Line Analytic Processing (section 10.6)

# Selection: $\sigma_{condition}(R)$

- Unary operation
  - Input: Relation with schema  $R(A_1, \dots, A_n)$
  - Output: Relation with attributes  $A_1, \dots, A_n$
  - Meaning: Takes a relation  $R$  and extracts only the rows from  $R$  that satisfy the *condition*
  - Condition is a logical combination (using AND, OR, NOT) of expressions of the form:  
 $\langle expr \rangle \langle op \rangle \langle expr \rangle$   
where  $\langle expr \rangle$  is an attribute name, a constant, a string, and  $op$  is one of  $(=, \leq, \geq, <, >, <>)$ 
    - E.g., “age > 20 OR height < 6”,
    - “name LIKE “Anne%” AND salary > 200000”
    - “NOT (age > 20 AND salary < 100000)”

# Example of $\sigma$

- $\sigma_{\text{rating} > 6}$  (Hotels)

Hotels

name	address	rating	capacity
Windsor	54 <sup>th</sup> ave	6.0	135
Astoria	5 <sup>th</sup> ave	8.0	231
BestInn	45 <sup>th</sup> st	6.7	28
ELodge	39 W st	5.6	45
ELodge	2nd E st	6.0	40

name	address	rating	capacity
Astoria	5 <sup>th</sup> ave	8.0	231
BestInn	45 <sup>th</sup> st	6.7	28

# Example of $\sigma$ with AND in Condition

- $\sigma_{\text{rating} > 6 \text{ AND capacity} > 50}(\text{Hotel})$

name	address	rating	capacity
Windsor	54 <sup>th</sup> ave	6.0	135
Astoria	5 <sup>th</sup> ave	8.0	231
BestInn	45 <sup>th</sup> st	6.7	28
ELodge	39 W st	5.6	45
ELodge	2nd E st	6.0	40

name	address	rating	capacity
Astoria	5 <sup>th</sup> ave	8.0	231

- Is  $\sigma_{C_1}(\sigma_{C_2}(R)) = \sigma_{C_1 \text{ AND } C_2}(R)$  ?
- Prove or give a counter-example.
- Is  $\sigma_{C_1}(\sigma_{C_2}(R)) = \sigma_{C_2}(\sigma_{C_1}(R))$  ?
- Prove or give a counter-example.

# Projection: $\pi_{\langle \textit{attribute list} \rangle}(\mathbf{R})$

- Unary operation
  - Input : Relation with schema  $R(A_1, \dots, A_n)$
  - Output: Relation with attributes in *attribute list*, which must be attributes of R
  - Meaning: For every tuple in relation R, output only the attributes appearing in *attribute list*
- May be duplicates; for Codd's Relational Algebra, duplicates are always eliminated (set-oriented semantics)
  - Reminder: For relational database, duplicates matter.
  - Why?

# Example of $\pi$

- $\pi_{\text{name, address}}(\text{Hotels})$

name	address
Windsor	54 <sup>th</sup> ave
Astoria	5 <sup>th</sup> ave
BestInn	45 <sup>th</sup> st
ELodge	39 W st
ELodge	2nd E st

- Suppose that name and address form the key of the Hotels relation. Is the cardinality of the output relation the same as the cardinality of Hotels? Why?



# Example of $\pi$

- $\pi_{\text{name}}(\text{Hotel})$

name
Windsor
Astoria
BestInn
ELodge

- Note that there are no duplicates.

# Set Union: $R \cup S$

- Binary operator
  - Input: Two relations R and S which must be **union-compatible**
    - They have the same arity, i.e., the same number of columns.
    - For every column i, the i'th column of R has the same type as the i'th column of S.
    - Note that field names are not used in defining union-compatibility.
      - We can think of relations R and S as being union-compatible if they are sets of records having the same record type.
  - Output: Relation that has the same type as R (or same type as S).
  - Meaning: The output consists of the **set** of all tuples in either R or S (or both)

# Example of $\cup$

Dell\_Desktops  $\cup$  HP\_Desktops

Dell\_Desktops

Harddisk	Speed	OS
20G	500Mhz	Windows
30G	1.0Ghz	Windows
20G	750Mhz	Linux

HP\_Desktops

Harddisk	Speed	OS
30G	1.2Ghz	Windows
20G	500Mhz	Windows

All tuples in R occurs in  $R \cup S$ .  
All tuples in S occurs in  $R \cup S$ .  
 $R \cup S$  contains tuples that either occur in R or S (or both).

Harddisk	Speed	OS
20G	500Mhz	Windows
30G	1.0Ghz	Windows
20G	750Mhz	Linux
30G	1.2Ghz	Windows

# Properties of $\cup$

Dell\_Desktops  $\cup$  HP\_Desktops

Dell\_Desktops

Harddisk	Speed	OS
20G	500Mhz	Windows
30G	1.0Ghz	Windows
20G	750Mhz	Linux

HP\_Desktops

Harddisk	Speed	OS
30G	1.2Ghz	Windows
20G	500Mhz	Windows

$R \cup S = S \cup R$  (commutativity)

$(R \cup S) \cup T = R \cup (S \cup T)$  (associativity)

Harddisk	Speed	OS
20G	500Mhz	Windows
30G	1.0Ghz	Windows
20G	750Mhz	Linux
30G	1.2Ghz	Windows

# Set Difference: $R - S$

- Binary operator.
  - Input: Two relations R and S which must be **union-compatible**
  - Output: Relation with the same type as R (or same type as S)
  - Meaning: Output consists of all tuples in R but not in S

# Example of -

- Dell\_Desktops - HP\_Desktops

Dell\_Desktops

Harddisk	Speed	OS
20G	500Mhz	Windows
30G	1.0Ghz	Windows
20G	750Mhz	Linux

HP\_Desktops

Harddisk	Speed	OS
30G	1.2Ghz	Windows
20G	500Mhz	Windows

Dell\_Desktops – HP\_Desktops

Harddisk	Speed	OS
30G	1.0Ghz	Windows
20G	750Mhz	Linux

# Properties of -

- HP\_Desktops – Dell\_Desktops

Dell\_Desktops

Harddisk	Speed	OS
20G	500Mhz	Windows
30G	1.0Ghz	Windows
20G	750Mhz	Linux

HP\_Desktops

Harddisk	Speed	OS
30G	1.2Ghz	Windows
20G	500Mhz	Windows

HP\_Desktops – Dell\_Desktops

Harddisk	Speed	OS
30G	1.2Ghz	Windows

Is it commutative?

Is it associative?

# Product: $R \times S$

- Binary operator
  - Input: Two relations  $R$  and  $S$ , where  $R$  has relation schema  $R(A_1, \dots, A_m)$  and  $S$  has relation schema  $S(B_1, \dots, B_n)$ .
  - Output: Relation of arity  $m+n$
  - Meaning:
$$R \times S = \{ (a_1, \dots, a_m, b_1, \dots, b_n) \mid (a_1, \dots, a_m) \in R \text{ and } (b_1, \dots, b_n) \in S \}.$$
    - Read “ $\mid$ ” as “such that”
    - Read “ $\in$ ” as “belongs to”



# Example and Properties of Product

R

A	B	C
$a_1$	$b_1$	$c_1$
$a_2$	$b_2$	$c_2$

S

D	E
$d_1$	$e_1$
$d_2$	$e_2$
$d_3$	$e_3$

$R \times S$

A	B	C	D	E
$a_1$	$b_1$	$c_1$	$d_1$	$e_1$
$a_1$	$b_1$	$c_1$	$d_2$	$e_2$
$a_1$	$b_1$	$c_1$	$d_3$	$e_3$
$a_2$	$b_2$	$c_2$	$d_1$	$e_1$
$a_2$	$b_2$	$c_2$	$d_2$	$e_2$
$a_2$	$b_2$	$c_2$	$d_3$	$e_3$

- Is it commutative?
- Is it associative?
- Is it distributive across  $\cup$ ? That is, does  $R \times (S \cup T) = (R \times S) \cup (R \times T)$ ?

# Product and Common Attributes

- What happens when we compute the Product of R and S if R and S contain common attributes, e.g., for R(A,B,C) and S(A,E)?

A.1	B	C	A.2	E
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>	d <sub>1</sub>	e <sub>1</sub>
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>	d <sub>2</sub>	e <sub>2</sub>
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>	d <sub>3</sub>	e <sub>3</sub>
a <sub>2</sub>	b <sub>2</sub>	c <sub>2</sub>	d <sub>1</sub>	e <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>	c <sub>2</sub>	d <sub>2</sub>	e <sub>2</sub>
a <sub>2</sub>	b <sub>2</sub>	c <sub>2</sub>	d <sub>3</sub>	e <sub>3</sub>

# Derived Operators

- So far, we have learned:
  - Selection
  - Projection
  - Product
  - Union
  - Difference
- Some other operators can be derived by composing the operators we have learned so far:
  - Theta-Join, Join, Natural Join, Semi-Join
  - Set Intersection
  - Division/Quotient
  - Outer Join (to be discussed when we get to OLAP)

# Theta-Join: $R \bowtie_{\theta} S$

- Binary operator
  - Input:  $R(A_1, \dots, A_m), S(B_1, \dots, B_n)$
  - Output: Relation consisting of all attributes  $A_1, \dots, A_m$  and all attributes  $B_1, \dots, B_n$ . Identical attributes in  $R$  and  $S$  are disambiguated with the relation names.
  - Meaning of  $\sigma_{\theta}(R \times S)$ : The  $\theta$ -Join outputs those tuples from  $R \times S$  that satisfy the condition  $\theta$ .
    - Compute  $R \times S$ , then keep only those tuples in  $R \times S$  that satisfy  $\theta$ .
    - Equivalent to writing  $\sigma_{\theta}(R \times S)$
- If  $\theta$  always evaluates to true, then  $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S) = R \times S$ .

# Example of Theta-Join

Enrollment(esid, ecid, grade)

Course(cid, cname, instructor-name)

Please give me an example to write on the board where ecid in Enrollment equals cid in Course.

- Joins involving equality predicates (usually just called Joins or Equi-Joins) are very common in database; other joins are less common.
  - Enrollment  $\bowtie_{\theta}$  Course, where  $\theta$  could be:  
“Enrollment.ecid = Course.cid”
- Could write any condition involving attributes of Enrollment and Course as  $\theta$ , just as with  $\sigma$ .

# Natural Join: $R \bowtie S$

- Often a query over two relations can be formulated using Natural Join.
- Binary operator:
  - Input: Two relations R and S where  $\{A_1, \dots, A_k\}$  is the set of common attributes (column names) between R and S.
  - Output: A relation where its attributes are  $\text{attr}(R) \cup \text{attr}(S)$ . In other words, the attributes consists of the attributes in  $R \times S$  without repeats of the common attributes  $\{A_1, \dots, A_k\}$

- Meaning:

$$R \bowtie S = \pi_{(\text{attr}(R) \cup \text{attr}(S))} (\sigma_{R.A_1=S.A_1 \text{ AND } R.A_2 = S.A_2 \text{ AND } \dots \text{ AND } R.A_k=S.A_k} (R \times S))$$

1. Compute  $R \times S$
2. Keep only those tuples in  $R \times S$  satisfying:  
 $R.A_1=S.A_1 \text{ AND } R.A_2 = S.A_2 \text{ AND } \dots \text{ AND } R.A_k=S.A_k$
3. Output is projection on the set of attributes in  $R \cup S$  (without repeats)

# Example of Natural Join

Enrollment(sid, cid, grade)

Course(cid, cname, instructor-name)

cid is the common attribute between the two relations.

- Want: Course-grade(sid, cid, grade, cname, instructor-name)
- $\pi_{(sid, cid, grade, cname, instructor-name)}(\text{Enrollment} \bowtie \text{Course})$
- What happens when R and S have no common attributes?
- What happens when R and S have only common attributes?

# Semi-Join: $R \bowtie S$

- Meaning:  $R \bowtie S = \pi_{\text{attr}(R)} (R \Join S)$ 
  1. Compute Natural Join of R and S
  2. Output is the projection on just the attributes of R
- Find all courses that have some enrollment:  
Course  $\bowtie$  Enrollment
- Find all faculty who are advising at least one student:  
Faculty  $\bowtie$  Student
- How does Semi-Join relate to EXISTS in SQL?



# Set Intersection: $R \cap S$

Find all desktops sold by both Dell and HP.

Dell\_Desktops  $\cap$  HP\_Desktops

Dell\_Desktops

Harddisk	Speed	OS
20G	500Mhz	Windows
30G	1.0Ghz	Windows
20G	750Mhz	Linux

HP\_Desktops

Harddisk	Speed	OS
30G	1.2Ghz	Windows
20G	500Mhz	Windows

Harddisk	Speed	OS
20G	500Mhz	Windows

# Intersect

- How would you write  $\text{Dell\_desktops} \cap \text{HP\_desktops}$  in SQL?

```
SELECT *  
FROM Dell_desktops  
INTERSECT  
SELECT *  
FROM HP_desktops;
```

- Intersection is a Derived Operator in Relational Algebra:

$$\begin{aligned} R \cap S &= R - (R - S) \\ &= S - (S - R) \end{aligned}$$

# Division: $R \div S$ (also written $R/S$ )

- Input: Two relations  $R$  and  $S$ , where:  
 $\text{attr}(S) \subset \text{attr}(R)$ , and  $\text{attr}(S)$  is non-empty.
- Output: Relation whose attributes are  $\text{attr}(R) - \text{attr}(S)$ .
- Example:  $R(A,B,C,D)$ ,  $S(B,D)$ .
- Meaning:  $R \div S = \{ (a, c) \mid \text{for all } (b,d) \in S, \text{ we have } (a,b,c,d) \in R \}$
- The quotient (or division)  $R \div S$  is the relation consisting of all tuples  $(a_1, \dots, a_{r-s})$  such that:  
For every tuple  $(b_1, \dots, b_s)$  in  $S$ , the tuple  $(a_1, \dots, a_{r-s}, b_1, \dots, b_s)$  is in  $R$

# Example of Division

Enrollment(sid, cid, grade)

Course(cid, cname, instructor-name)

- Find the sids of students who are enrolled in all courses

$\text{Enrollment} \div \pi_{\text{cid}}(\text{Course})$

- Find the sids of all students who are enrolled in all courses taught by “Ullman”

$\text{Enrollment} \div \pi_{\text{cid}}(\sigma_{\text{instructor-name}='Ullman'}(\text{Course}))$

# Example of Division

R

A	B	C
a1	b1	c1
a1	b2	c2
a2	b1	c1
a1	b3	c3
a4	b2	c2
a3	b2	c2
a4	b1	c1

S

B	C
b1	c1
b2	c2

$R \div S$

A
a1
a4

# Quotient (or Division) (cont'd)

- Can we express  $R \div S$  with basic operators (select, project, cross product, union, difference) ?
- Yes; see textbook

# Independence of Basic Operators

- Many interesting queries can be expressed using the five basic operators ( $\sigma$  ,  $\pi$  ,  $\chi$  ,  $\cup$  ,  $-$  )
- Can one of the five operators be derived by the other four operators?

## Theorem (Codd):

The five basic operators are independent of each other. In other words, for each relational operator  $o$ , there is no relational algebra expression that is built from the rest that defines  $o$ .

- $\chi$
- $\pi$
- $\sigma$
- $\cup$
- $-$

# Renaming: $\rho_{S(A_1, \dots, A_n)}(R)$

- To specify the attributes of a relational expression.
- Input: a relation, a relation symbol  $R$ , and a set of attributes  $\{B_1, \dots, B_n\}$
- Output: the same relation with name  $S$  and attributes  $A_1, \dots, A_n$ .
- Meaning: rename relation  $R$  to  $S$  with attributes  $A_1, \dots, A_n$ .



# Example

R

A	B	C
$a_1$	$b_1$	$c_1$
$a_2$	$b_2$	$c_2$

S

C	D
$d_1$	$e_1$
$d_2$	$e_2$
$d_3$	$e_3$

$R \times \rho_{T(X,D)} S$

A	B	C	X	D
$a_1$	$b_1$	$c_1$	$d_1$	$e_1$
$a_1$	$b_1$	$c_1$	$d_2$	$e_2$
$a_1$	$b_1$	$c_1$	$d_3$	$e_3$
$a_2$	$b_2$	$c_2$	$d_1$	$e_1$
$a_2$	$b_2$	$c_2$	$d_2$	$e_2$
$a_2$	$b_2$	$c_2$	$d_3$	$e_3$

# Renaming: $\rho_{S(A_1, \dots, A_n)}(R)$

- To specify the attributes of a relational expression.
- Input: a relation, a relation symbol  $R$ , and a set of attributes  $\{B_1, \dots, B_n\}$
- Output: the same relation with name  $S$  and attributes  $A_1, \dots, A_n$ .
- Meaning: Rename relation  $R$  to  $S$  with attributes  $A_1, \dots, A_n$ .

# More Complex Queries

- Relational operators can be composed to form more complex queries. We have already seen examples of this in SQL.

Enrollments(esid, ecid, grade)

Courses(cid, cname, instructor-name)

- Query 1: Find student id, grade and instructor for students whose grade was higher than 80 points in a course.

$$\sigma_{\text{grade} > 80} ( \pi_{\text{esid}, \text{grade}, \text{instructor-name}} ( \sigma_{\text{Enrollments.ecid} = \text{Courses.cid}} (\text{Enrollments x Courses}) ) )$$

# Query 2

Enrollments(esid, ecid, grade)

Courses(cid, cname, instructor-name)

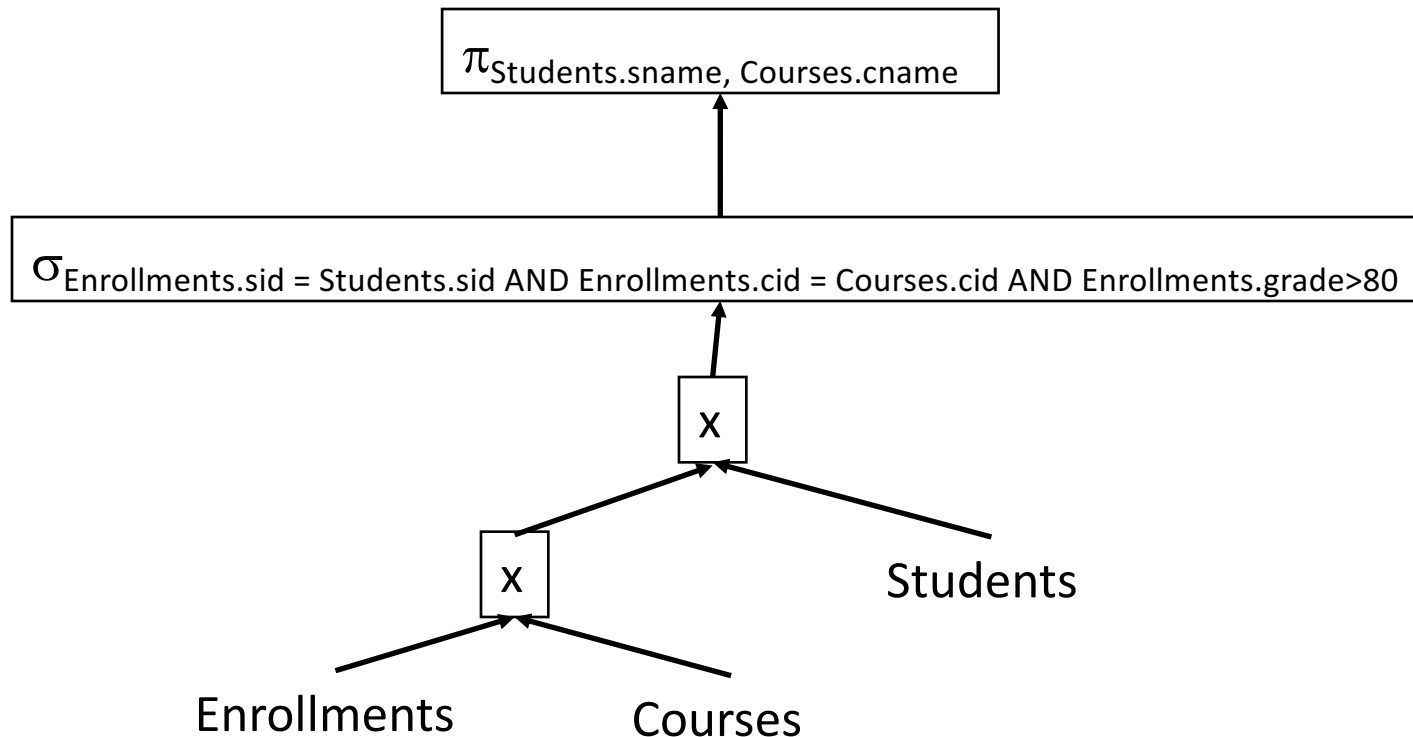
Students(sid, sname)

- Find the student name and course name where the student had a grade more than 80 points in a course.

$$\pi_{\text{Students.sname, Courses.cname}} \left( \begin{array}{l} \sigma_{\text{Enrollments.ecid} = \text{Courses.cid}} \quad (\text{Enrollments} \times \text{Courses} \times \text{Students}) \\ \text{AND Enrollment.esid} = \text{Students.sid} \\ \text{AND Grade} > 80 \end{array} \right)$$

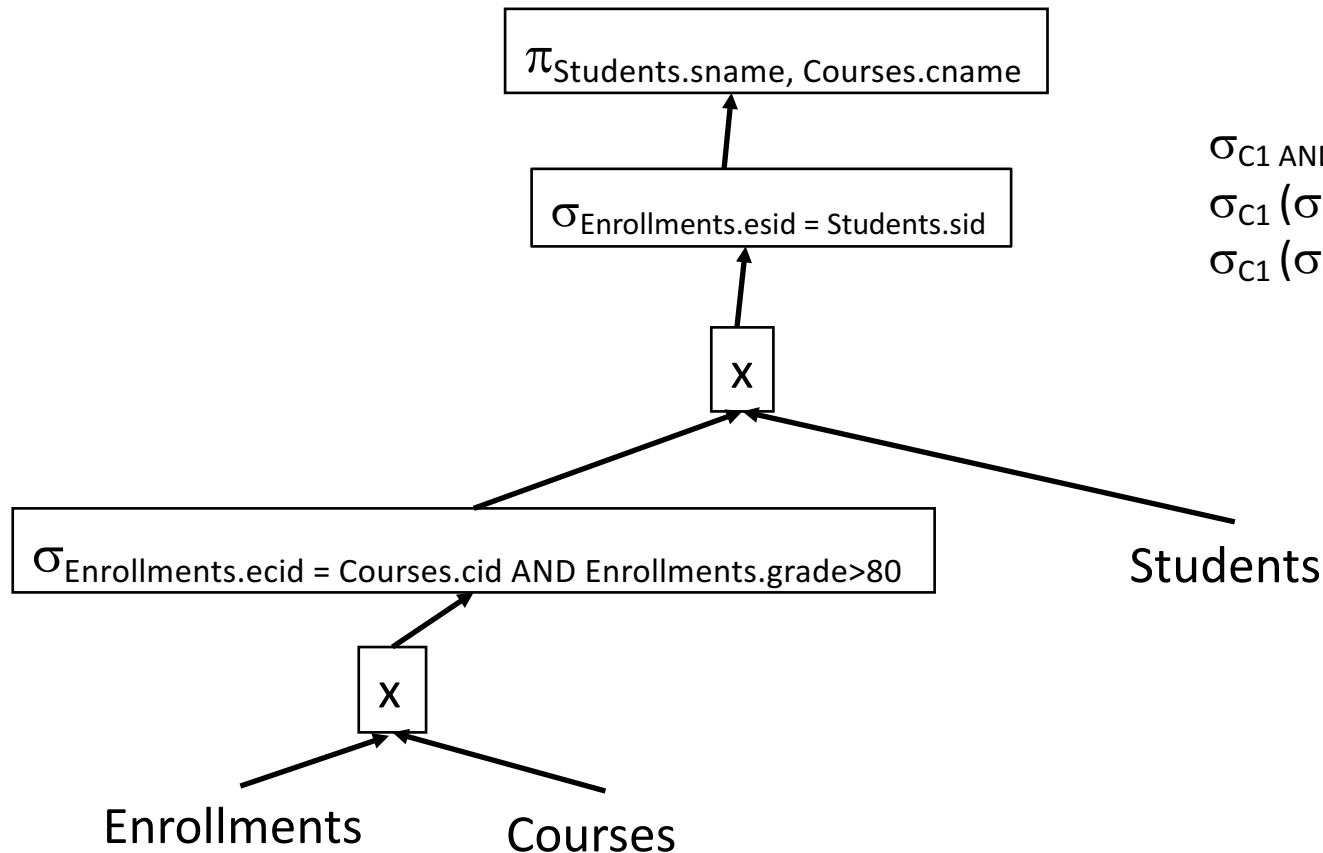
# An Execution Plan for Query 2

- Find the student name and course name where the student had a grade more than 80 points in a course.



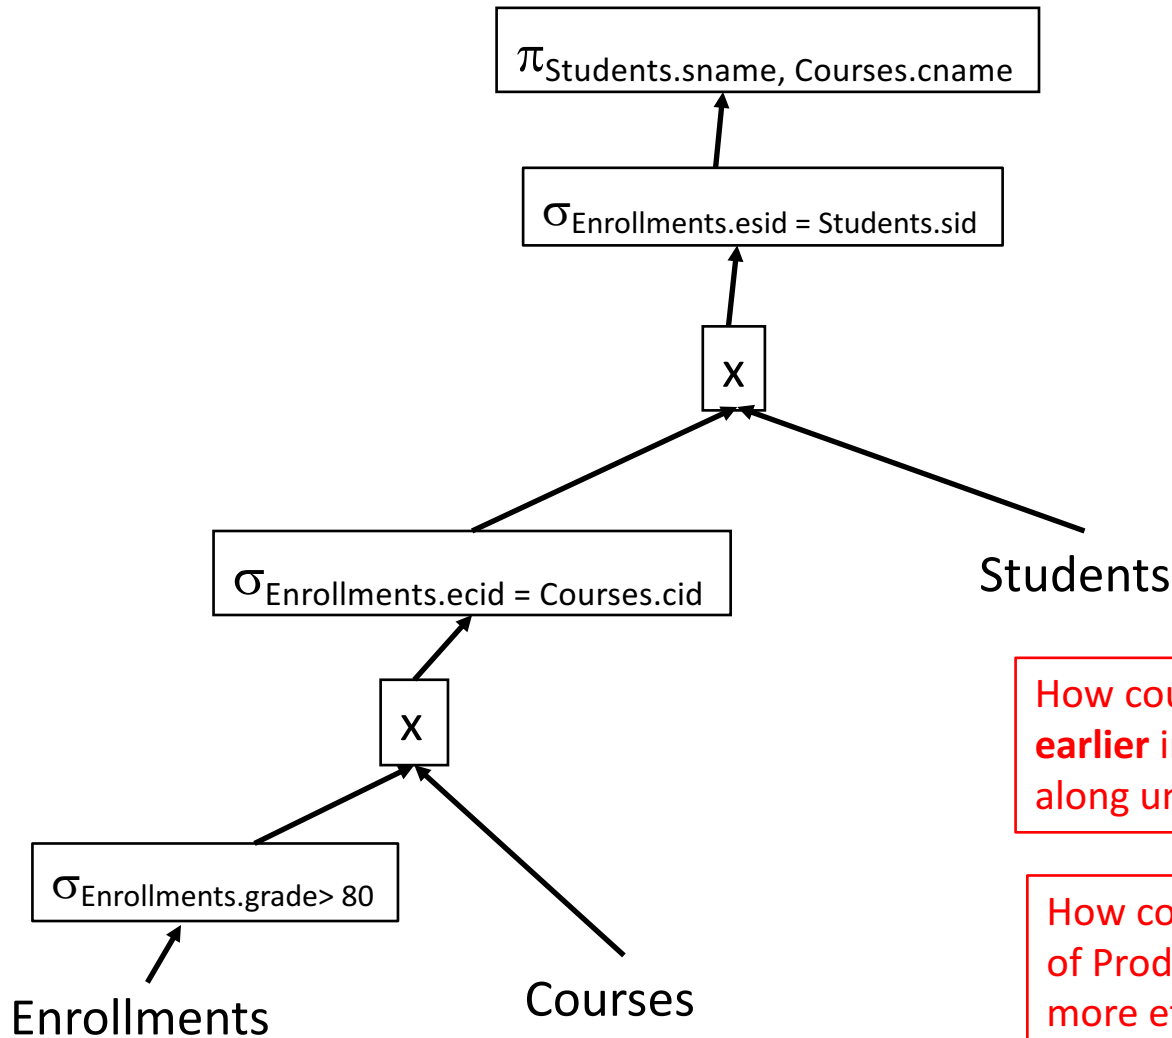
# Another Execution Plan for Query 2

- Find the student name and course name where the student had a grade more than 80 points in a course.



$$\begin{aligned}\sigma_{C1 \text{ AND } C2 \text{ AND } C3} (E \times C \times S) &= \\ \sigma_{C1} (\sigma_{C2 \text{ AND } C3} (E \times C \times S)) &= \\ \sigma_{C1} (\sigma_{C2 \text{ AND } C3} (E \times C) \times S) &= \end{aligned}$$

# A Third Execution Plan for Query 2



How could we do projections **earlier** in plan to avoid carrying along unnecessary attributes?

How could we do **Joins**, instead of Products to make plan more efficient?

# Query Transformations

- What were some of the query equivalences that we talked about earlier?
- What other query equivalences do you know about?



# Execution Plans

- When do you do SELECTION?
  - Predicate pushdown is always a good idea.
- How do you access each table?
  - Scan, index (which index), hash, ...
- What's the order in which you Join tables?
  - Join/Equi-join is common; avoid Cartesian product
  - But which table do you start with?
    - Predicates on indexed columns are often useful in picking first table, then next table, to join, ...
- What join method do you use for each join?
  - Nested loop join, merge join, hash-join, ...
- How much parallelism do you use?
  - How do you schedule tasks to hardware?
- Do you need to sort? If so, when do you sort?

# Query Optimization

- Comparing Execution Plans and finding a “good” (not necessarily best) plan
- Statistics that DBMS may keep to help calculate approximate query cost
  - Cardinality (number of rows) in table
  - Highest and lowest (non-null) value in column
  - Column cardinality (number of different values in column)
  - Number of appearances of the top 10 most frequent value in each column
  - Join cardinality between tables for particular equi-join
    - May be calculated, not stored; not well-defined if there are conditions (predicates) on the tables
  - Many other statistics are calculated approximately
- How frequently are stored statistics updated?
- Cost: CPU? I/O? Network? How do these get combined to compare plans?

# EXPLAIN Statement

- Shows information about query plan
  - Each DBMS that has EXPLAIN has its own variation
  - Try it with PostgreSQL
- You may want to try to rewrite query yourself to find better execution plan if Query Optimizer isn't smart enough to do so
- Should Optimizer take advice from users?

# Practice Homework 5

Sailors(sid, sname, rating, age) // sailor id, sailor name, rating, age

Boats(bid, bname, color) // boat id, boat name, color of boat

Reserves(sid, bid, day) // sailor id, boat id, date that sid reserved bid.

- Use **Relational Algebra** to write the following 8 queries.
  - How might you optimize execution of queries using ideas in this Lecture, per discussion in slides 44-47?
1. Find the names of sailors who reserved boat 103.
  2. Find the colors of boats reserved by Lubber.
  3. Find the names of sailors who reserved at least one boat.

## Practice Homework 5 (cont'd)

4. Find the names of sailors whose age  $> 20$  and have not reserved any boats.
5. Find the names of sailors who have reserved a red or a green boat.
6. Find the names of sailors who have reserved a red and a green boat.
7. Find the names of sailors who have reserved at least 2 different boats.
8. Find the names of sailors who have reserved exactly 2 different boats.