

Real SQL

Application Programming

Instructor: Shel Finkelstein

Reference:

*A First Course in Database Systems,
3rd edition, Chapter 9*

Important Notices

- Midterm and Midterm Answers are on Piazza.
 - Answers were reviewed in class on Wed, Feb 15.
 - Midterm grades (uncurved) were posted on Canvas.
 - Curve: $0.9 * X + 12$
 - Exam returned on Friday, Feb 24 and Monday, Feb 27.
- Lab3 Solution was posted on Monday, Feb 27.
- Lab4 assignment was posted on Monday, Feb 27.
 - Due by **Sunday, March 12, 11:59pm** (2 weeks).
 - Lab4 focusses on material in this Application Programming lecture (JDBC, Stored Procedures).
 - If you don't attend Lectures and Labs, you probably will find Lab4 difficult.
- Gradiance Assignment #4 is due by **Tuesday, March 7, 11:59pm**.

SQL in Real Programs

- We have seen only how SQL is used at a generic query interface --- an environment where we sit at a terminal and ask queries of (or modify) a database.
- Reality is almost always different!
 - Conventional programs written in C or Java, (or other languages) that interact with database using SQL.
 - Why?

Approaches

1. Code in a specialized language is stored in the database itself (e.g., **Stored Procedure** languages such as PSM and PL/SQL).
2. SQL statements are **embedded in a host language** (e.g., C).
3. **Connection tools/libraries** are used to allow a **conventional language** to access a database (e.g., CLI, JDBC).

Approach 1: Stored Procedures

- PSM, or “*persistent stored modules*,” allows us to store procedures as database schema elements.
- PSM = a mixture of conventional statements (if, while, etc.) and SQL.
- Lets us do things we cannot do in SQL alone.

Basic PSM Form

```
CREATE PROCEDURE <name> (  
    <parameter list> )  
    <optional local declarations>  
    <body>;
```

Basic PSM Form

```
CREATE FUNCTION <name> (  
    <parameter list> ) RETURNS <type>  
    <optional local declarations>  
    <body>;
```

Parameters in PSM

- Unlike the usual name-type pairs in languages like C, PSM uses mode-name-type triples, where the *mode* can be:
 - IN = procedure uses value, does not change value.
 - OUT = procedure changes, does not use.
 - INOUT = both.
- Function parameters must be of mode IN. Functions returns value, but must have no side-effects on parameters.

Example: Stored Procedure

- Let's write a procedure that takes two arguments b and p , and adds a tuple to **Sells(bar, beer, price)** that has bar = 'Joe's Bar', beer = b , and price = p .
- Used by Joe to add to his menu more easily.

The Procedure

CREATE PROCEDURE JoeMenu (

```
IN    b    CHAR(20),  
IN    p    REAL
```

Parameters are both
read-only, not changed

)

```
INSERT INTO Sells  
VALUES(' Joe' ' s Bar' , b, p);
```

The body ---
a single insertion

Invoking Procedures

- Use SQL/PSM statement CALL, with the name of the desired procedure and arguments.

- **Example:**

```
CALL JoeMenu('Moosedrool', 5.00);
```

- Functions may be used in SQL expressions wherever a value of their return type is appropriate.

Kinds of PSM statements – (1)

- RETURN <expression> sets the return value of a function.
 - Unlike C, etc., RETURN *does not* terminate function execution.
- DECLARE <name> <type> used to declare local variables.
- BEGIN . . . END for groups of statements.
 - Separate statements by semicolons.

Kinds of PSM Statements – (2)

- **Assignment statements:**
SET <variable> = <expression>;
 - Example: SET b = 'Bud' ;
- **Statement labels:** give a statement a label by prefixing a name and a colon.

IF Statements

- Simplest form:
IF <condition> THEN
 <statements(s)>
END IF;
- Add ELSE <statement(s)> if desired, as
IF . . . THEN . . . ELSE . . . END IF;
- Add additional cases by ELSEIF <statements(s)>:
IF ... THEN ... ELSEIF ... THEN ... ELSEIF ...
THEN ... ELSE ... END IF;

Example: IF

- Let's rate bars by how many customers they have, based on `Frequents(drinker,bar)`.
 - < 100 customers: 'unpopular'.
 - 100-199 customers: 'average'.
 - ≥ 200 customers: 'popular'.
- Function `Rate(b)` rates bar b.

Example: IF (continued)

```
CREATE FUNCTION Rate (IN b CHAR(20) )
```

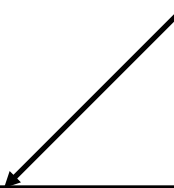
```
  RETURNS CHAR(10)
```

```
  DECLARE cust INTEGER;
```

```
  BEGIN
```

```
    SET cust = (SELECT COUNT(*) FROM Frequents  
                WHERE bar = b);
```

Number of
customers of
bar b



```
    IF cust < 100 THEN RETURN 'unpopular'  
    ELSEIF cust < 200 THEN RETURN 'average'  
    ELSE RETURN 'popular'  
    END IF;
```

Nested
IF statement



```
  END;
```

Return occurs here, not at
one of the RETURN statements



Loops

- Basic form:
 <loop name>: LOOP
 <statements>
 END LOOP;
- Exit from a loop by:
 LEAVE <loop name>;

Example: Exiting a Loop

```
loop1: LOOP
```

```
...
```

```
LEAVE loop1; ← If this statement is executed ...
```

```
...
```

```
END LOOP;
```

← Control winds up here

Other Loop Forms

- WHILE <condition>
 DO <statements>
 END WHILE;
- REPEAT <statements>
 UNTIL <condition>
 END REPEAT;

Queries

- General SELECT-FROM-WHERE queries are *not* permitted in PSM.
- There are three ways to get the effect of a query:
 1. Queries producing one value can be the expression in an assignment.
 2. Single-row SELECT . . . INTO ...
 3. Cursors

Example: Assignment/Query

- Using local variable p and `Sells(bar, beer, price)`, we can get the price Joe charges for Bud by:

```
SET p = (SELECT price FROM Sells
WHERE bar = 'Joe''s Bar'
AND beer = 'Bud' );
```

SELECT . . . INTO ...

- Another way to get the value of a query that returns one tuple is by placing **INTO <variable>** after the SELECT clause.
- **Example:**

```
SELECT price INTO p
FROM Sells
WHERE bar = 'Joe''s Bar'
      AND beer = 'Bud';
```

Cursors

- A *cursor* is essentially a tuple-variable that ranges over all tuples in the result of some query.
- Declare a cursor *c* by:
 DECLARE *c* CURSOR FOR <query>;

Opening and Closing Cursors

- To use cursor c , we must issue the command:
OPEN c ;
 - The query of c is evaluated, and c is set to point to the first tuple of the result.
- When finished with c , issue command:
CLOSE c ;

Fetching Tuples From a Cursor

- To get the next tuple from cursor *c*, issue command:

FETCH FROM *c* INTO *x*₁, *x*₂, ..., *x*_{*n*} ;

- The *x* ' s are a list of variables, one for each component of the tuples referred to by *c*.
- *c* is moved automatically to the next tuple.

Breaking Cursor Loops – (1)

- The usual way to use a cursor is to create a loop with a FETCH statement, and do something with each tuple fetched.
- A tricky point is how we get out of the loop when the cursor has no more tuples to deliver.

Breaking Cursor Loops – (2)

- Each SQL operation returns a *status*, which is a 5-digit character string.
 - For example:
 - '00000' means “Everything OK,”
 - '02000' means “Failed to find a tuple.”
- In PSM, we can get the value of the status in a variable called SQLSTATE.

Breaking Cursor Loops – (3)

- We may declare a *condition*, which is a boolean variable that is true if and only if SQLSTATE has a particular value.
- **Example:** We can declare condition NotFound to represent 02000 by:

```
DECLARE NotFound CONDITION FOR  
        SQLSTATE '02000';
```

Breaking Cursor Loops – (4)

- The structure of a cursor loop is thus:

```
cursorLoop: LOOP
```

```
...
```

```
  FETCH c INTO ... ;
```

```
  IF NotFound THEN LEAVE cursorLoop;
```

```
  END IF;
```

```
...
```

```
END LOOP;
```

Example: Cursor

- Let's write a procedure that examines `Sells(bar, beer, price)`, and raises by one dollar the price of all beers at Joe's Bar that are under three dollars.
- Yes, we could write this as a simple UPDATE, but the details are instructive anyway.

The Needed Declarations

```
CREATE PROCEDURE JoeGouge( )
```

```
    DECLARE theBeer CHAR(20);
```

```
    DECLARE thePrice REAL;
```

Used to hold
beer-price pairs
when fetching
through cursor c

```
    DECLARE NotFound CONDITION FOR
```

```
        SQLSTATE ' 02000' ;
```

```
    DECLARE c CURSOR FOR
```

Returns Joe's menu

```
        (SELECT beer, price FROM Sells
```

```
            WHERE bar = ' Joe' ' s Bar' );
```

The Procedure Body

```
BEGIN
  OPEN c;
  menuLoop: LOOP
    FETCH c INTO theBeer, thePrice;
    IF NotFound THEN LEAVE menuLoop END IF;
    IF thePrice < 3.00 THEN
      UPDATE Sells SET price = thePrice + 1.00
      WHERE bar = 'Joe' 's Bar' AND beer = theBeer;
    END IF;
  END LOOP;
  CLOSE c;
END;
```

Check if the recent
FETCH failed to
get a tuple

If Joe charges less than \$3 for
the beer, raise its price at
Joe's Bar by \$1.

The Procedure Body:

Using **CURRENT OF** Cursor

```
BEGIN
  OPEN c;
  menuLoop: LOOP
    FETCH c INTO theBeer, thePrice;
    IF NotFound THEN LEAVE menuLoop END IF;
    IF thePrice < 3.00 THEN
      UPDATE Sells SET price = thePrice + 1.00
      WHERE CURRENT OF c;
    END IF;
  END LOOP;
  CLOSE c;
END;
```

Check if the recent
FETCH failed to
get a tuple

If Joe charges less than \$3 for
the beer, raise its price at
Joe's Bar by \$1.

PL/SQL

- Oracle uses **PL/SQL**, a **variation** of SQL/PSM that helped inspire PSM.
- PL/SQL not only allows you to create and store procedures or functions, but it also can be run from Oracle's *generic query interface (SQL*Plus)*, just like any SQL statement.
- PostgreSQL: **PL/pgSQL (needed for Lab4)**
- IBM DB2: **SQL PL**
- MS SQL Server and Sybase: **Transact-SQL (T-SQL)**

Triggers and Stored Procedures

Trigger

- *Event* : typically a type of database modification, e.g., “insert on Sells”
 - *Condition* : Any SQL boolean-valued expression
 - *Action* : Any SQL statements
-
- Triggers may invoke Stored Procedures.
 - A typical trigger body (actions) may itself be thought of as an unnamed Stored Procedure.
 - In some systems, the trigger body may include many of the kinds of statements that can be in a Stored Procedure.

Approach 2: Embedded SQL

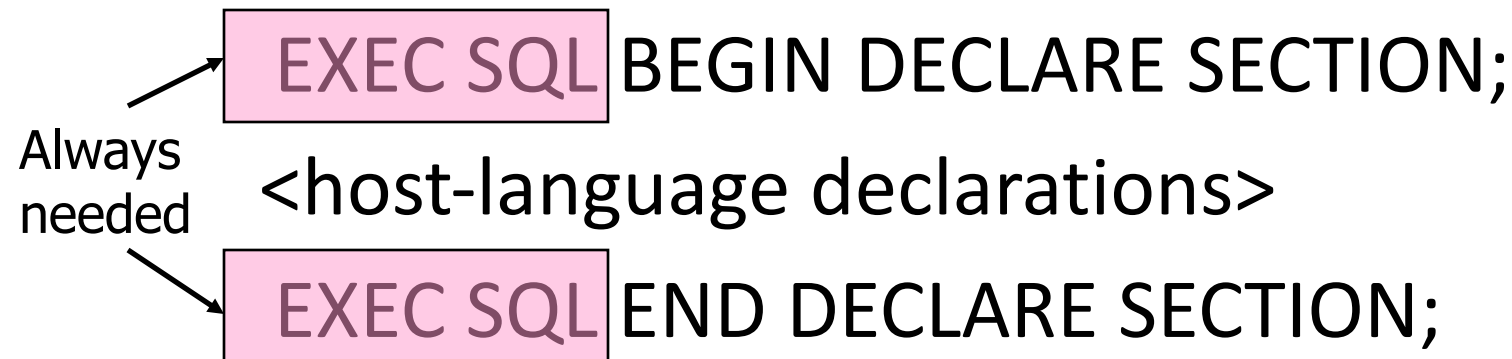
- **Key idea:** A pre-processor turns SQL statements into procedure calls that fit with the surrounding host-language code.
- All embedded SQL statements begin with EXEC SQL, so the pre-processor can find them easily.

Shared Variables

- To connect SQL and the host-language program, the two parts must share some variables.
- Declarations of shared variables are bracketed by:

Always
needed

```
EXEC SQL BEGIN DECLARE SECTION;  
<host-language declarations>  
EXEC SQL END DECLARE SECTION;
```



Use of Shared Variables

- In SQL, the shared variables must be preceded by a colon.
 - They may be used as if they were constants provided by the host-language program.
 - They may get values from SQL statements and pass those values to the host-language program.
- In the host language, shared variables behave like any other variable.

Example: Looking Up Prices

- We'll use C with embedded SQL to sketch the important parts of a function that obtains a beer and a bar, and looks up the price of that beer at that bar.
- Assumes database has the **Sells(bar, beer, price)** relation.

Example: C with SQL

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
char theBar[21], theBeer[21];
```

```
float thePrice;
```

```
EXEC SQL END DECLARE SECTION;
```

```
/* obtain values for theBar and theBeer */
```

```
EXEC SQL SELECT price INTO :thePrice
```

```
FROM Sells
```

```
WHERE bar = :theBar AND beer = :theBeer;
```

```
/* do something with thePrice */
```

Note 21-char
arrays needed
for 20 chars +
endmarker

SELECT-INTO
as in PSM

Embedded Queries

- Embedded SQL has the same limitations as PSM regarding queries:
 - SELECT-INTO for a query guaranteed to produce a single tuple.
 - Otherwise, you have to use a cursor.
 - Small syntactic differences, but the key ideas are the same.

Cursor Statements

- Declare a cursor *c* with:

```
EXEC SQL DECLARE c CURSOR FOR <query>;
```

- Open and close cursor *c* with:

```
EXEC SQL OPEN CURSOR c;
```

```
EXEC SQL CLOSE CURSOR c;
```

- Fetch from *c* by:

```
EXEC SQL FETCH c INTO <variable(s)>;
```

- You can write a macro NOT_FOUND that is true if and only if the FETCH fails to find a tuple.
- If *c* is a cursor, you may use ... **WHERE CURRENT OF *c***, just as in Stored Procedures.

Example: Print Joe's Menu

- Let's write C + SQL to print Joe's menu – the list of beer-price pairs that we find in `Sells(bar, beer, price)` with `bar = Joe's Bar`.
- A cursor will visit each Sells tuple that has `bar = Joe's Bar`.

Example: Declarations

```
EXEC SQL BEGIN DECLARE SECTION;  
    char theBeer[21]; float thePrice;  
EXEC SQL END DECLARE SECTION;
```

```
EXEC SQL DECLARE c CURSOR FOR  
    SELECT beer, price FROM Sells  
    WHERE bar = ' Joe' ' s Bar' ;
```



The cursor declaration goes outside the declare-section

Example: Executable Part

```
EXEC SQL OPEN CURSOR c;
```

```
while(1) {
```

```
    EXEC SQL FETCH c
```

```
        INTO :theBeer, :thePrice;
```

```
    if (NOT_FOUND) break;
```

```
    /* format and print theBeer and thePrice */
```

```
}
```

```
EXEC SQL CLOSE CURSOR c;
```

The C style
of breaking
loops



Need for Dynamic SQL

- Most applications use specific queries and modification statements to interact with the database.
 - The DBMS compiles EXEC SQL ... statements into specific procedure calls and produces an ordinary host-language program that uses a library.

Dynamic SQL

- Preparing a query:

```
EXEC SQL PREPARE <query-name>  
FROM <text of the query>;
```

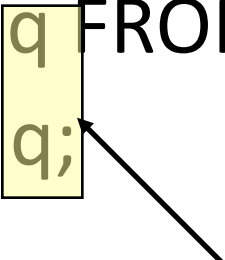
- Executing a query:

```
EXEC SQL EXECUTE <query-name>;
```

- “Prepare” means optimize query.
- Prepare once, Execute many times.

Example: A Generic Interface

```
EXEC SQL BEGIN DECLARE SECTION;  
    char query[MAX_LENGTH];  
EXEC SQL END DECLARE SECTION;  
while(1) {  
    /* issue SQL> prompt */  
    /* read user's query into array query */  
    EXEC SQL PREPARE q FROM :query;  
    EXEC SQL EXECUTE q;  
}
```



q is an SQL "query variable" representing the optimized form of whatever statement is typed into :query

Execute-Immediate

- If we are only going to execute the query once, we can combine the PREPARE and EXECUTE steps into one.

- Use:

```
EXEC SQL EXECUTE IMMEDIATE <text>;
```

Example: Generic Interface Again

```
EXEC SQL BEGIN DECLARE SECTION;  
    char query[MAX_LENGTH];  
EXEC SQL END DECLARE SECTION;  
  
while(1) {  
    /* issue SQL> prompt */  
    /* read user's query into array query */  
    EXEC SQL EXECUTE IMMEDIATE :query;  
}
```

Approach 3: Host Language/SQL Interfaces via Libraries

- The third approach to connecting databases to conventional languages is to use library calls.
 1. C + CLI
 2. Java + JDBC
 3. PHP + PEAR/DB

Three-Tier Architecture

- A common environment for using a database has three tiers of processors:
 1. *Web servers* --- talk to the user.
 2. *Application servers* --- execute business logic.
 - Often not used—logic in web tier
 3. *Database servers* --- get what the app servers (or web servers) need from the database.


Example: Amazon

- Database holds the information about products, customers, etc.
- Business logic includes things like “What do I do after someone clicks ‘checkout’ ?”
 - **Answer:** Show the “How will you pay for this?” screen.
- Presentation layer, handled on web server and web browser, handles preparation and display of web pages

Environments, Connections, Queries

- The database is, in many DB-access languages, an *environment*.
- Database servers maintain some number of *connections*, so app servers can connect to them and ask queries or perform modifications.
- The app server issues *statements* : queries and modifications, usually.

Diagram to Remember



SQL/CLI

- Instead of using a pre-processor (as in embedded SQL), we can use a library of functions.
 - The library for C is called SQL/CLI = “*Call-Level Interface*.”
 - Embedded SQL’s pre-processor will translate the EXEC SQL ... statements into CLI or similar calls, anyway.

Data Structures

- C connects to the database by structs of the following types:
 1. *Environments* : represent the DBMS installation.
 2. *Connections* : logins to the database.
 3. *Statements* : SQL statements to be passed to a connection.
 4. *Descriptions* : records about tuples from a query, or parameters of a statement.

JDBC

- *Java Database Connectivity* (JDBC) is a library similar to SQL/CLI, but with Java as the host language.
- Like CLI, but with a few differences.

Making a Connection

```
import java.sql.*;
Class.forName("com.mysql.jdbc.Driver");
Connection myCon =
    DriverManager.getConnection(...);
```

The JDBC classes

Loaded by
forName

URL of the database
your name, and password
go here.

The driver
for mySql;
others exist

Statements

- JDBC provides two classes:
 1. *Statement* is an object that can accept a string that is a SQL statement and can execute such a string.
 2. *PreparedStatement* is an object that has an associated SQL statement ready to execute.

Creating Statements

- The Connection class has methods to create Statements and PreparedStatement.

```
Statement stat1 = myCon.createStatement();
```

```
PreparedStatement stat2 =  
    myCon.prepareStatement(  
        "SELECT beer, price FROM Sells " +  
        "WHERE bar = ' Joe' ' s Bar' "  
    );
```

Executing SQL Statements

- JDBC distinguishes queries from modifications, which it calls “updates.”
- Statement and PreparedStatement each have methods `executeQuery` and `executeUpdate`.
 - For Statement: one argument: the query or modification to be executed.
 - For PreparedStatement: no argument.

Example: Update

- stat1 is a Statement.
- We can use it to insert a tuple:

```
stat1.executeUpdate(  
    "INSERT INTO Sells " +  
    "VALUES ('Brass Rail', 'Bud', 3.00)"  
);
```

Example: Query

- stat2 is a PreparedStatement holding the query "SELECT beer, price FROM Sells WHERE bar = 'Joe' 's Bar' ".
• **executeQuery** returns an object of class ResultSet; we'll examine that soon.
- The query:

```
ResultSet Menu = stat2.executeQuery();
```


Accessing the ResultSet

- An object of type ResultSet is a lot like a cursor.
- Method `next()` advances the “cursor” to the next tuple.
 - The first time `next()` is applied, it gets the first tuple.
 - If there are no more tuples, `next()` returns the value `false`.

Reminder of Example: Query

- stat2 is a PreparedStatement holding the query "SELECT beer, price FROM Sells WHERE bar = 'Joe' 's Bar' ".
• **executeQuery** returns an object of class ResultSet; we'll examine that soon.
- The query:

```
ResultSet Menu = stat2.executeQuery();
```

Accessing Components of Tuples

- When a ResultSet refers to a tuple, we can get the components of that tuple by applying certain methods to the ResultSet.
- Method `getX (i)`, where X is some type, and i is the component number, returns the value of that component.
 - The value must have type X .

Example: Accessing Components

- Menu is the ResultSet for query “SELECT beer, price FROM Sells WHERE bar = 'Joe' 's Bar' ”.
- Access beer and price from each tuple by:

```
while ( Menu.next() ) {  
    theBeer = Menu.getString(1);  
    thePrice = Menu.getFloat(2);  
    /* do something with theBeer and  
    thePrice */  
}
```

ExecuteQuery, ExecuteUpdate and Execute

- `executeQuery()`: Executes a SQL SELECT statement, and returns a `ResultSet` object.
- `executeUpdate()`: Executes a SQL UPDATE, INSERT or DELETE statement, and returns the number of affected rows.
 - May also be used with DDL, e.g., CREATE, DROP
- `execute()`: Executes either query or modification, and returns `TRUE` if query and `FALSE` if modification
 - `stat.getResultSet` for query result
 - `stat.getUpdateCount` for modification
- All methods may throw Exceptions

Executing a Stored Procedure **GoodBeers**

Assume **GoodBeers** somehow finds all the good beers that are sold at a specific bar (theBar) that sell for under a particular price (thePrice).

- We won't tell you the secret of how GoodBeers procedure works.

```
PreparedStatement stmt = mycon.prepareStatement(
    "SELECT * FROM GoodBeers(?, ?)");
stmt.setString(1,theBar);    /* first parameter */
stmt.setFloat(2,thePrice);   /* second parameter */
ResultSet result = stmt.executeQuery();
while(result.next()) {
    theBeer= result.getString(1);
    /* do something with theBeer */
}
```

Executing a Stored Function

- Executing a Stored Function is similar to executing a Stored Procedure, except that what's returned is the result of the Stored Function.
 - Can be a scalar value (as in Lab4) or a table.
 - If it's a scalar value, that can be treat that as a table with one row.
- There is another way to execute Stored Procedures and Stored Functions, using the **CallableStatement** class (instead of Statement or PreparedStatement).
- This is not described well in the PostgreSQL documentation, but it's a good approach, and you may use it if you can figure it out.

Approaches

When/Why do you use each?

1. **Stored Procedure** languages such as PSM and PL/SQL)
2. SQL statements **embedded in a host language** (e.g., C)
3. **Connection tools/libraries** such as, CLI, and JDBC