

NetApp 7<sup>th</sup> Annual University Day  
NetApp HQ, Sunnyvale, CA, USA  
17 February 2014

# Modern Database Systems: NewSQL, NoSQL and Modernized Classic Systems

C. Mohan

IBM Fellow and Former IBM India Chief Scientist  
IBM Almaden Research Center, San Jose, CA 95120, USA  
[cmohan@us.ibm.com](mailto:cmohan@us.ibm.com) <http://bit.ly/CMohan>

Talk Version of Tutorial Presented First at  
**39<sup>th</sup> International Conference on Very Large Data Bases (VLDB)**  
Riva del Garda, Italy, 29 August 2013  
[Links to Revised Tutorial Slides and Bibliography at http://bit.ly/CMnMDS](http://bit.ly/CMnMDS)

# Background + Warnings

- My personal opinions, NOT necessarily IBM's positions/opinions
- Lots of hype on NoSQL (lesser on NewSQL) in industry, academia & VC world – this **Kool-Aid** has intoxicated many ☺
- Been a details oriented, hard core DB person for > 3 decades – I go beyond hype, am not easily swayed by the latest fashion!
- Not a relational bigot – most of my results on systems topics have applied to all sorts of persistent, distributed information systems
- Have dissected many IBM **and** non-IBM systems' internals (including pre- and post-relational ones)
- Don't claim to have all the answers or make definitive assertions
- **Aim:** not exhaustive survey of modern systems or details of too many specific systems – broad brush analysis of good (sensible) and bad (nonsenSQL) of NoSQL/NewSQL, warn of pitfalls – *time to pause and introspect*. Give a few systems' interesting details.
- **NOT against NoSQL** but unhappy with choices on internals, design justifications (“anything goes”), and not learning enough from the past

## Some Quotes

- **Jack Clark, The Register, 30 August 2013:** *“The tech world is turning back toward SQL, bringing to a close a possibly misspent half-decade in which startups courted developers with promises of infinite scalability and the finest imitation-Google tools available, and companies found themselves exposed to unstable data and poor guarantees.”*
- **Google Spanner paper, October 2012:** *“We believe it is better to have application programmers deal with performance problems due to overuse of transactions as bottlenecks arise, rather than always coding around the lack of transactions.”*
- **Sean Doherty in Wired, September 2013:** *“But don’t become unnecessarily distracted by the shiny, new-fangled, NoSQL red buttons just yet. ... Relational databases may not be hot or sexy but for your important data there is no substitute.”*
- **Gartner** estimates RDBMS market at \$26B with about 9% annual growth  
**Market Research Media Ltd** expects NoSQL market to be at \$3.5B by 2018

# Introduction

- Goal: Broad Survey of Modern Database Systems (MDS)
- Drivers Behind Emergence of MDS, Benchmarks/Performance Studies
- Classes of MDS
  - ▶ Evolution of Classical DBMSs (DB2, SQL Server, Oracle, Informix, PostgreSQL)
  - ▶ Brand New Systems (NewSQL, NoSQL)
  - ▶ Hybrid Systems (Hadoop + SQL)
- Overviews of Specific Systems
  - ▶ DB2 BLU, IDAA, MS Hekaton, Oracle Exadata, Informix Warehouse Accelerator, Postgres Plus Advanced Server
  - ▶ SAP Hana, Google F1, NuoDB, VoltDB, HyPer, Calvin
  - ▶ MongoDB, DB2/Informix NoSQL, Oracle NoSQL, Aerospike, Neo4j, Facebook Tao
- Deployments or Use Case Scenarios: Facebook

Ack: Figures and tables sourced from documentation/papers/web pages!

Links to Slides and Bibliography at <http://bit.ly/CMnMDS>

# Drivers of MDS Features

- **Hardware Trends:** Multicore, Multi-socket, InfiniBand, DRAMs, SSDs, FPGAs, GPUs, Storage Class Memories (e.g., PCM)
- **Big Data** (esp. Semi Structured or Unstructured): Compression, Column Handling, Scalability (Sharding), Avoiding Data Loading, Linkage to Analytics Tools, ...
- **Wider Adoption of Database Technology:** Ease of Use, Appliances
- **Monetary Constraints:** Open Source, Commodity Processors/Disks, Local Disks, Economic Factors in General
- **Skills Shortage:** DBAs, Admins in General, SQL Knowledge, ...
- **Social Media:** Relationship Graphs, Ultra Low Response Times, Data Heterogeneity/Volume, Mobile Users
- **Internet of Things (IOT):** Smarter Planet, Sensors, Low Power, ...
- **Cloud Platforms:** Elasticity, Multi-tenancy, Security, Transborder Data Flows

# Hardware Trends

- **DRAM costs 17X SSD!** Dell \$20/GB DRAM (16 GB DIMM \$315)
- **SSD 10X HDD cost** (\$0.10/GB, 200 IOPS)

SSD: Samsung 840, 250 GB, \$180

4 drives/chassis: 1 TB, 384K read IOPS, 248K random write IOPS, \$720, 0.3 Watts

- **DBs under 100 GB: DRAM**

**100 GB – 2 TB: SSD**

**>1 PB: HDD**

- Open source Aerospike Certification Tool for SSDs for real-time use cases



# Conclusions from Workload Studies

- EPFL Studies
  - ▶ OLTP workloads in traditional DSs don't leverage aggressive micro-architectural features, wasting time in memory stalls, resulting in low instructions per cycle (IPC)
  - ▶ L1 instruction misses cause significant stall times with index probes dominating instruction & data stalls
- H-Store: significant time spent in lock/latch, buffer pool fixes/unfixes
- Sort Vs Hash join & NUMA effects revisited in multicore/multi-socket systems

QPI – QuickPath Interconnect  
IMC – Integrated Memory Controller




Figure 1: Block diagram of a typical machine. Cores communicate either through a common cache, an interconnect across socket or main memory.

# Benchmarks

- **TPC** benchmarks (e.g., TPC-C, TPC-E, TPC-H)
- **LinkBench**: Facebook's Benchmark Based on Social Graph
- **BigBench**: Benchmark for Big Data Analytics
  - Teradata, University of Toronto, InfoSizing, Oracle
  - Product Retailer Scenario
  - Addresses Variety, Velocity & Volume of Big Data
  - Structured (TPC-DS), Semistructured (user clicks), Unstructured (product reviews)
  - Queries Cover Analytics Proposed by McKinsey: Marketing (Cross Selling, Customer Micro-Segmentation, Sentiment Analysis, Enhancing Multichannel Consumer Experience), Merchandising (Assortment Optimization, Pricing Optimization), Operations (Performance Transparency, Return Analysis), Supply Chain (Inventory Management), New Business Models (Price Comparison)
  - Feasibility Illustrated via Teradata Aster Database's SQL-MR

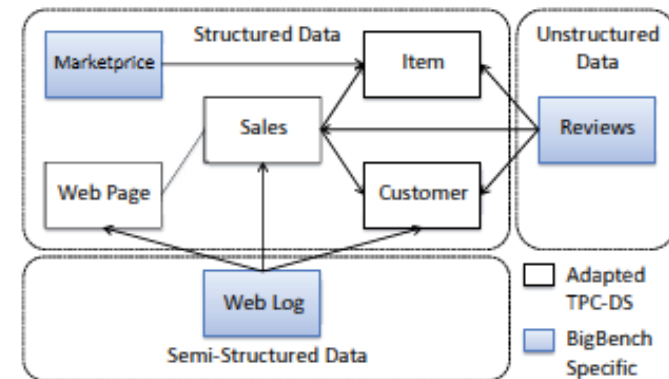

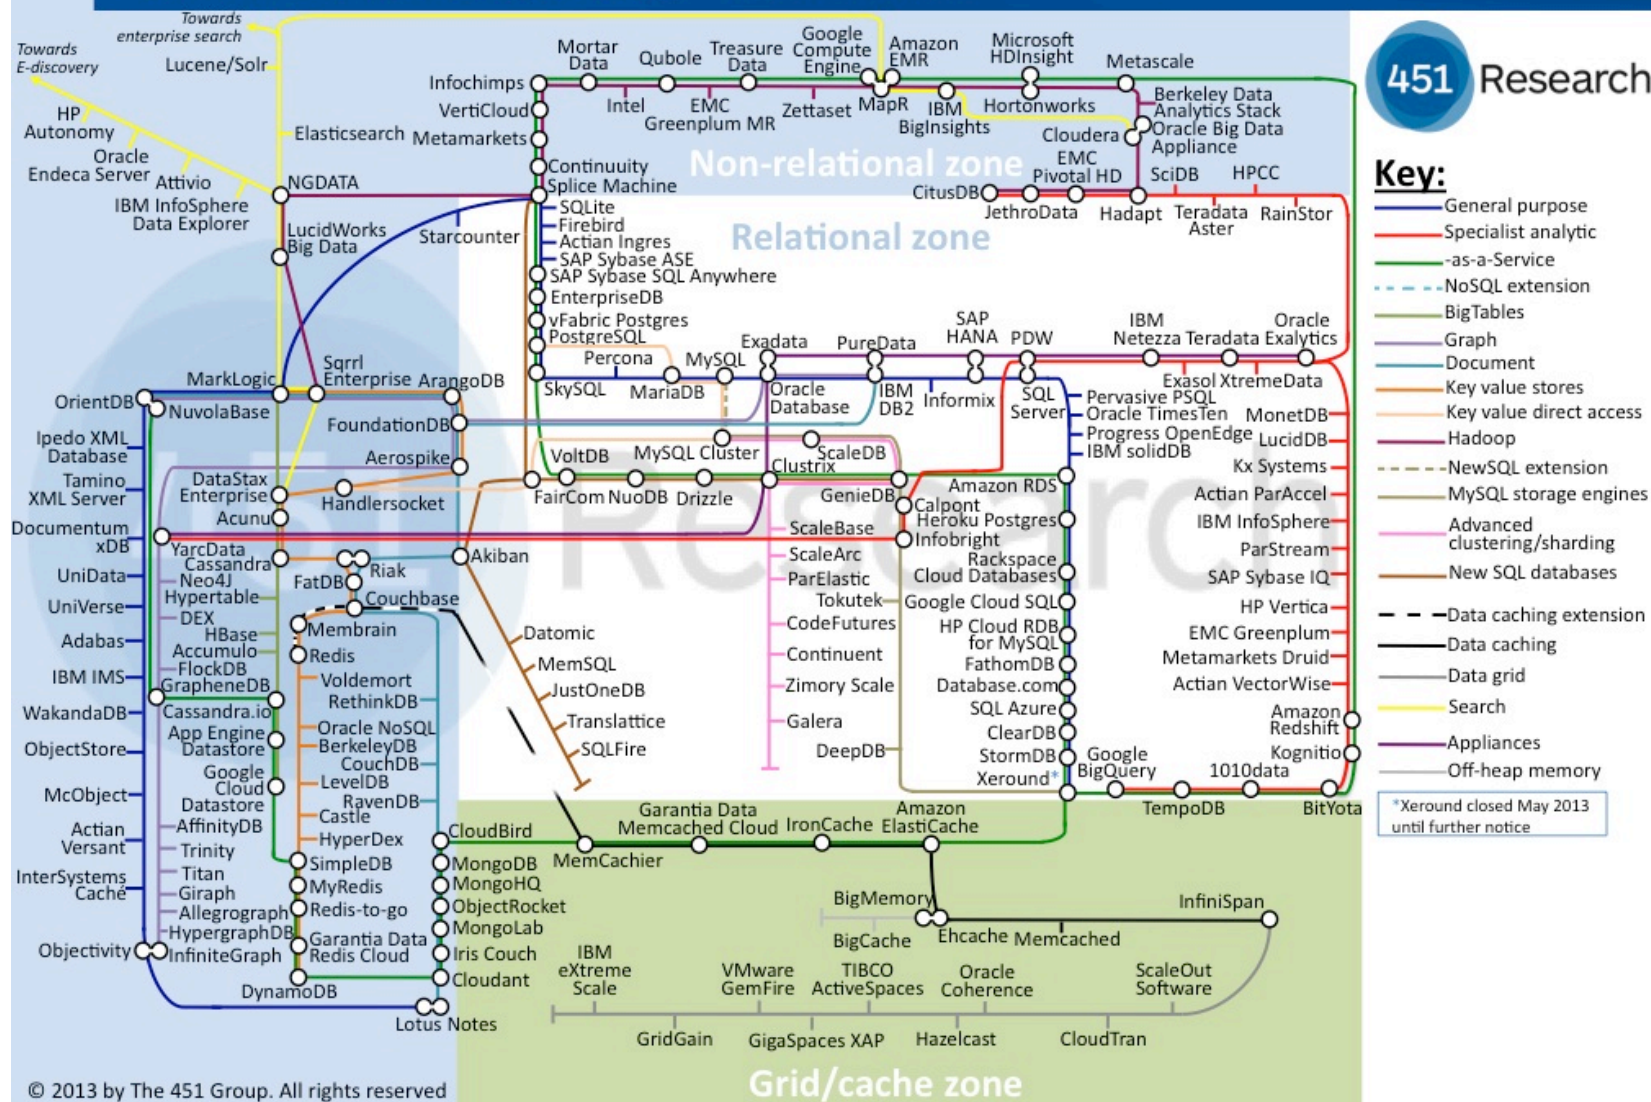


Figure 1: Big Data Benchmark Data Model

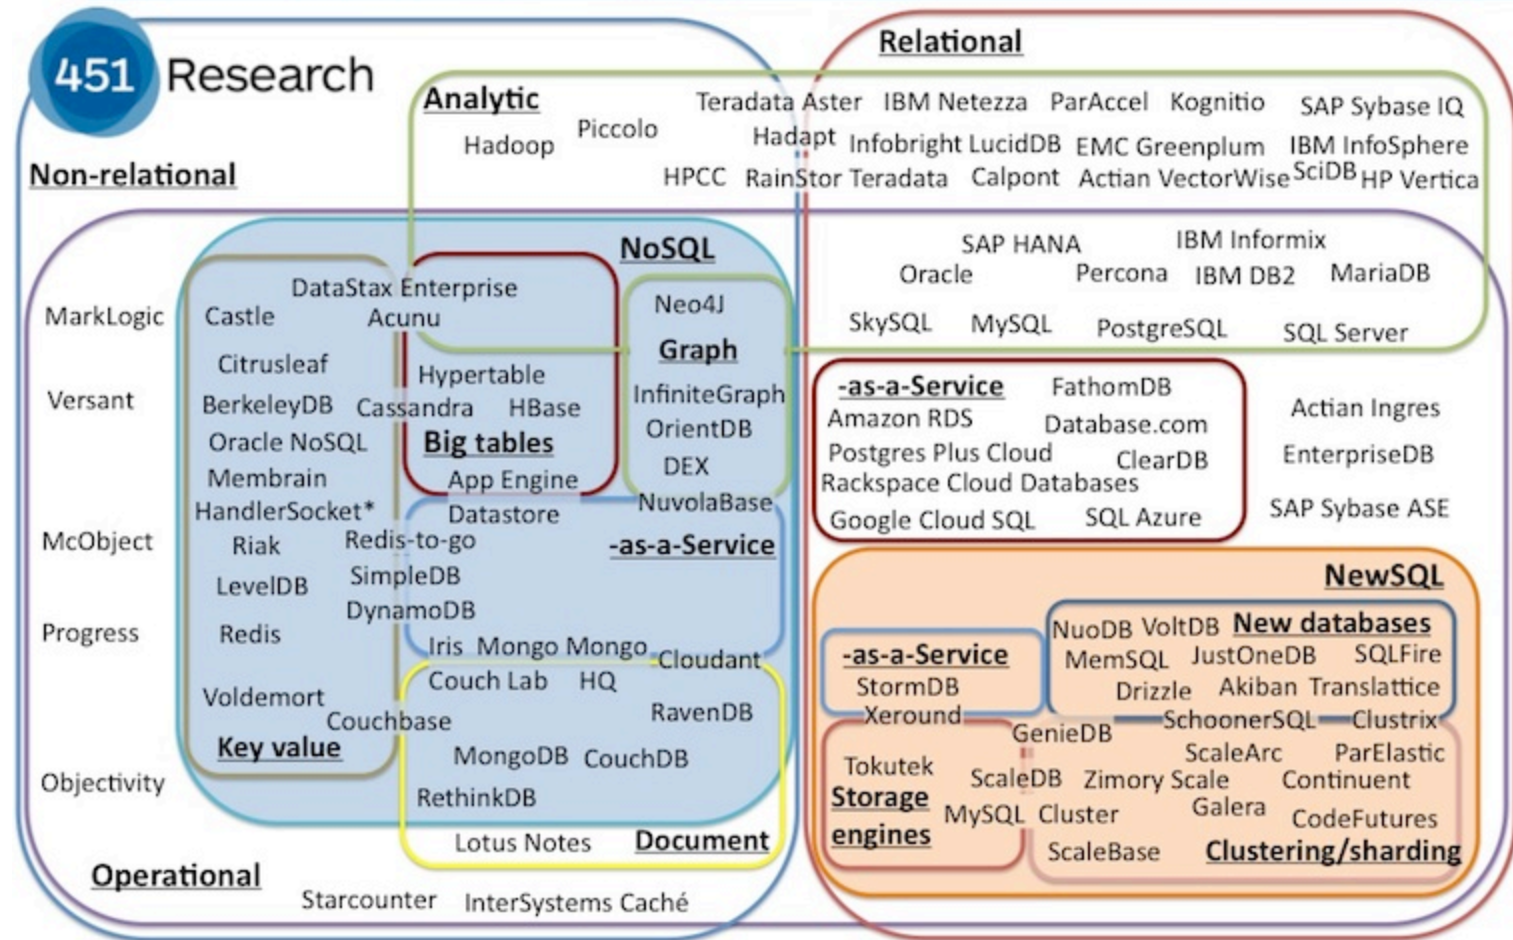


# Database Landscape Map – June 2013



# Database *Product* Landscape – Matthew Aslett

11-2012



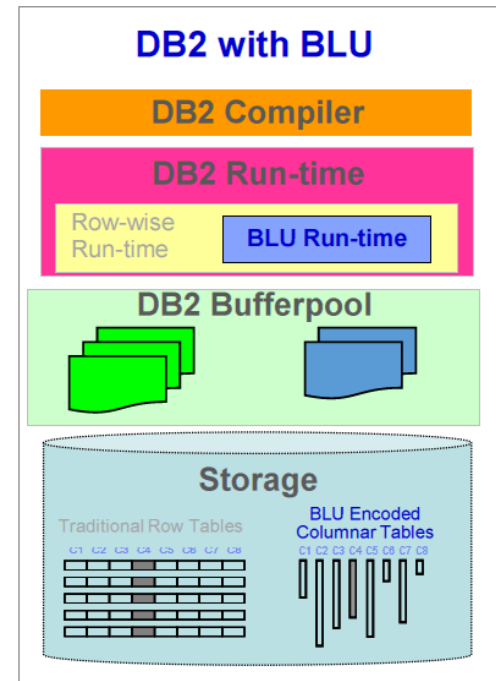
# Evolution of Classical DBMSs



# DB2 with BLU Acceleration



- New innovative technology for analytic queries
  - Columnar storage, single copy of the data
  - New run-time engine with vector (aka SIMD) processing
  - Deep multi-core optimizations and cache-aware memory management
  - “Active compression” - unique encoding for storage reduction and run-time processing without decompression
- “Revolution by Evolution”
  - Built directly into the DB2 kernel
  - BLU tables can coexist with row tables, in same schema, tablespaces, bufferpools
  - Query any combination of BLU or row data
  - Memory-optimized (not “in-memory”)
- Value : Order-of-magnitude benefits in ...
  - Performance
  - Storage savings
  - Simplicity!
- Generally available since June 2013



*Column-store (in) DB2 LUW v10.5*

V. Raman et al. *DB2 with BLU Acceleration: So Much More than Just a Column Store*.  
Proc. VLDB Endowment (PVLDB), 2013.

# IBM DB2 Analytics Accelerator (IDAA)

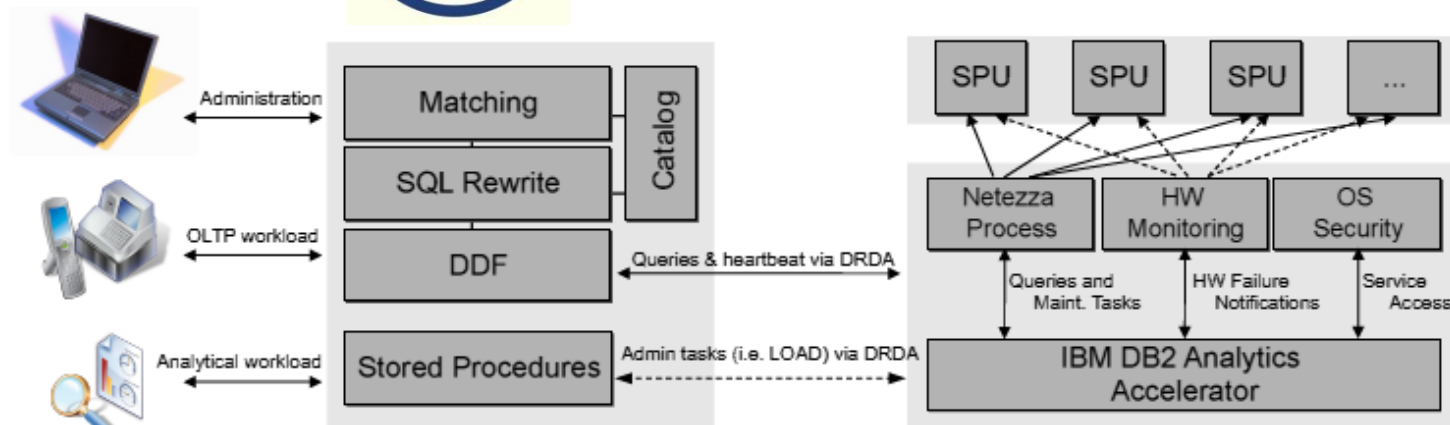


Figure 2: IDAA Architecture

- Data from DB2 z replicated continuously to Netezza
- Offload complex queries from DB2 z to Netezza for more efficient/cheaper execution
- Queries execute with snapshot semantics since Netezza supports MVCC
- Replaced IBM Smart Analytics Optimizer (ISAO) where Blink technology was used
- Evolution of Blink is what is now in DB2 BLU

Martin, D., Koeth, O., Ivanova, I., Kern, J. *Near Real-Time Analytics with IBM DB2 Analytics Accelerator*, Proc. International Conference on Extending Database Technology (EDBT), Genoa, Italy, March 2013.

# Hekaton: SQL Server's In-Memory OLTP Engine



- Alternate data manager integrated with SQL Server
- All Hekaton data in memory with lock/latch-free optimistic, multiversion concurrency control
- Declare table to be memory optimized
- T-SQL stored procs accessing only Hekaton tables compiled into machine code
- Hash and range indexes
- Optionally non-durable tables
- Index updates not logged; rebuilt during recovery from checkpoint and logs.

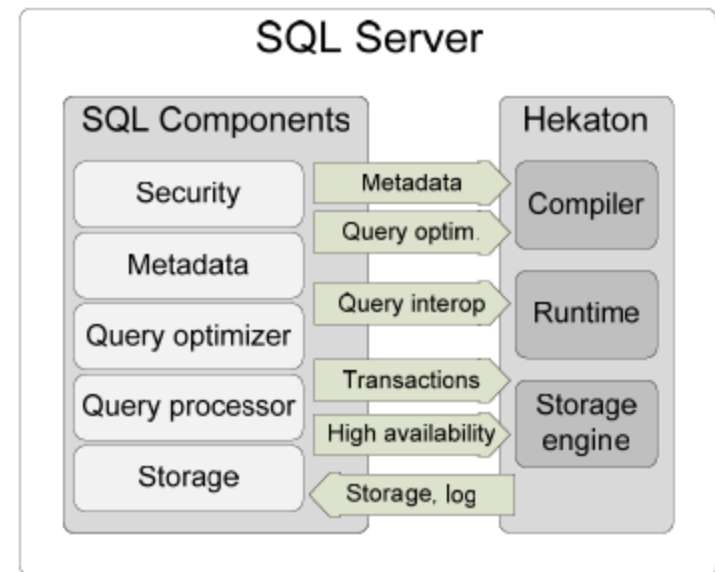


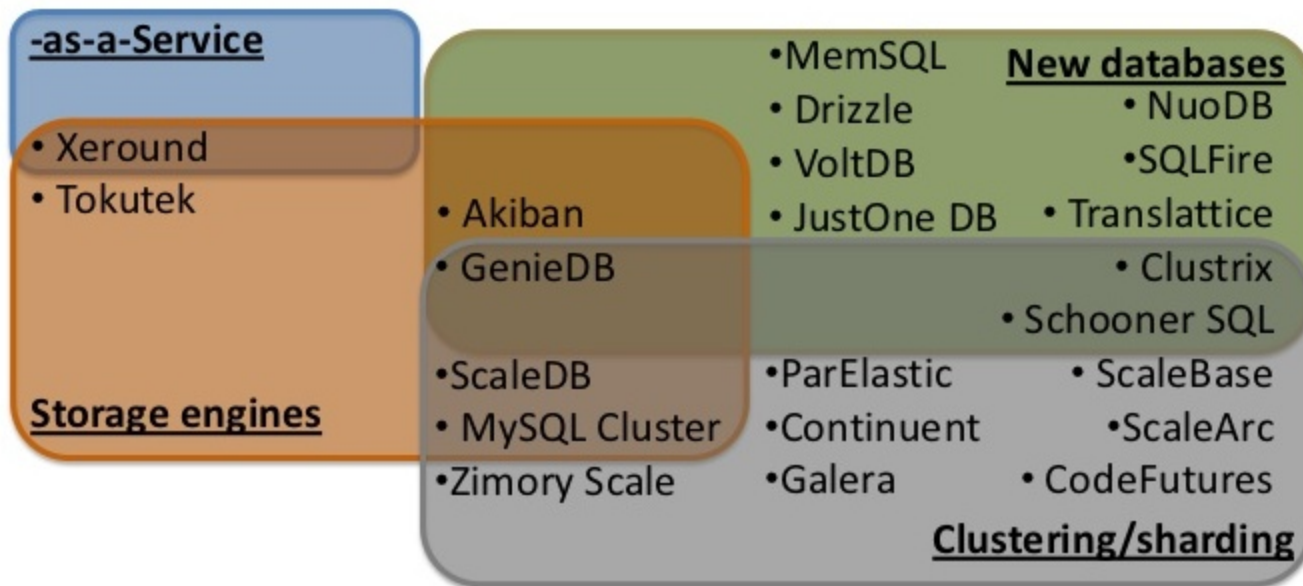
Figure 1: Hekaton's main components and integration into SQL Server.

C. Diaconu, C. Freedman, E. Ismert, P. Larson, P. Mittal, R. Stonecipher, N. Verma, M. Zwilling: *Hekaton: SQL Server's Memory-Optimized OLTP Engine*. Proc. ACM SIGMOD International Conference on Management of Data, New York, USA, June 2013.



# NewSQL Systems

# NewSQL Ecosystem - Matthew Aslett 2012



# SAP Hana (OLTP+OLAP)

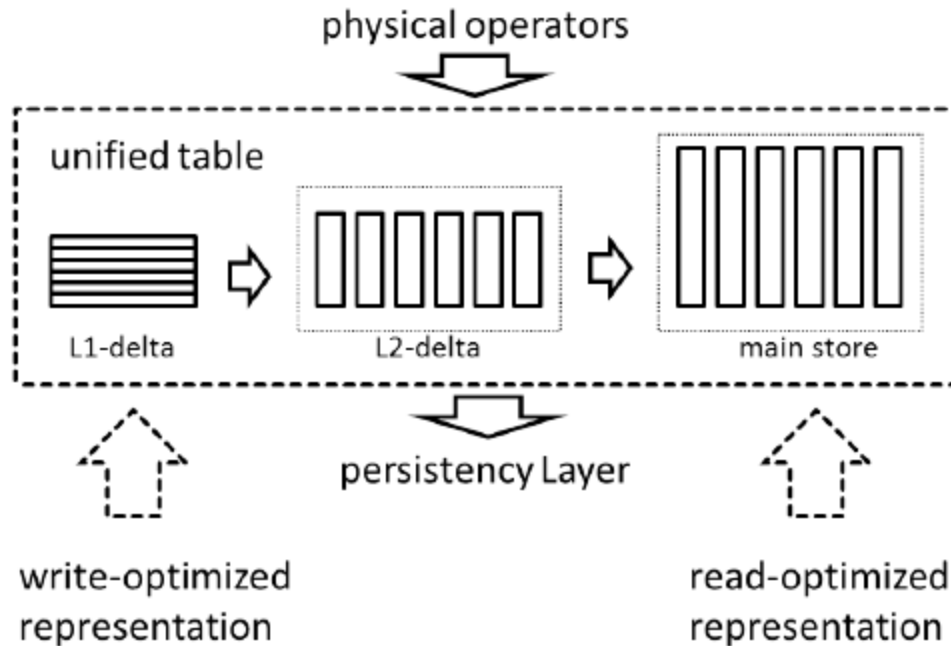


Figure 4: Overview of the unified table concept

L2 delta dictionary encoded, no compression for L1  
Main store maximally compressed with many schemes

# SAP Hana

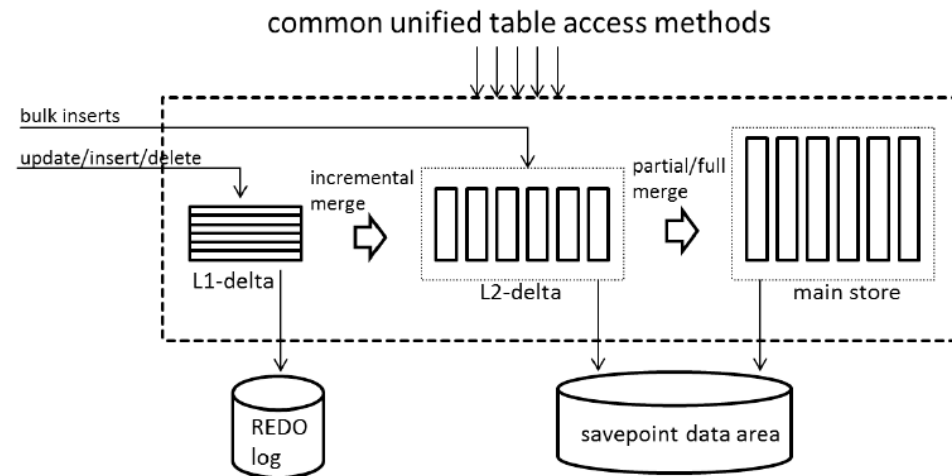


Figure 5: Overview of the persistency mechanisms of the unified table

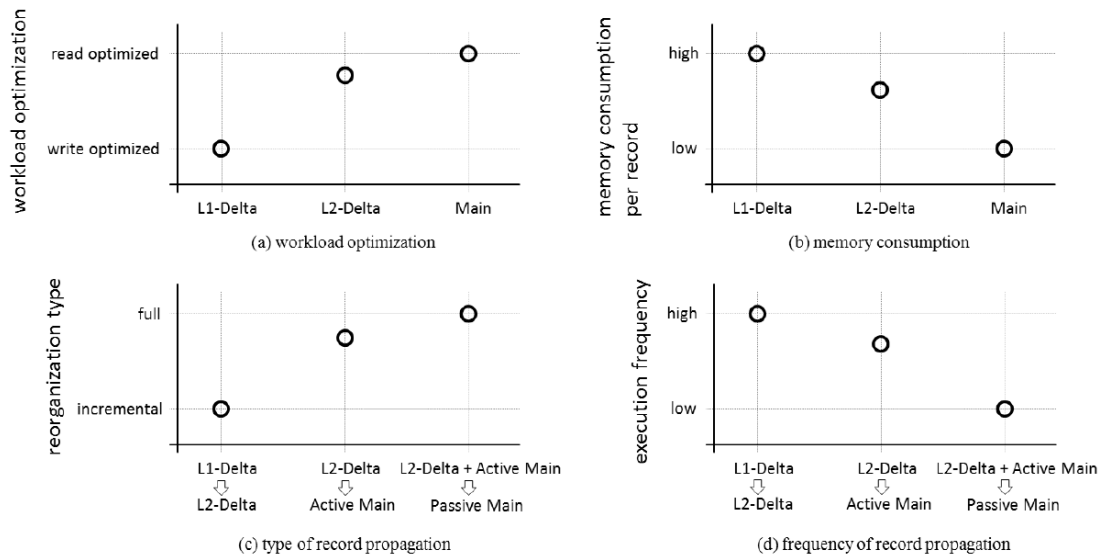


Figure 11: Characteristics of the SAP HANA database record life cycle

# Google Spanner

- Globally distributed semi-relational system, SQL-based query lang
- Transactional, relies on TrueTime and Paxos

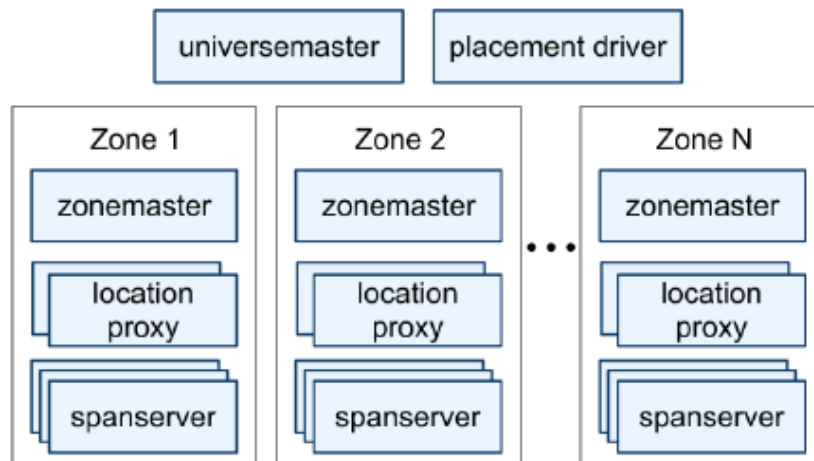


Figure 1: Spanner server organization.

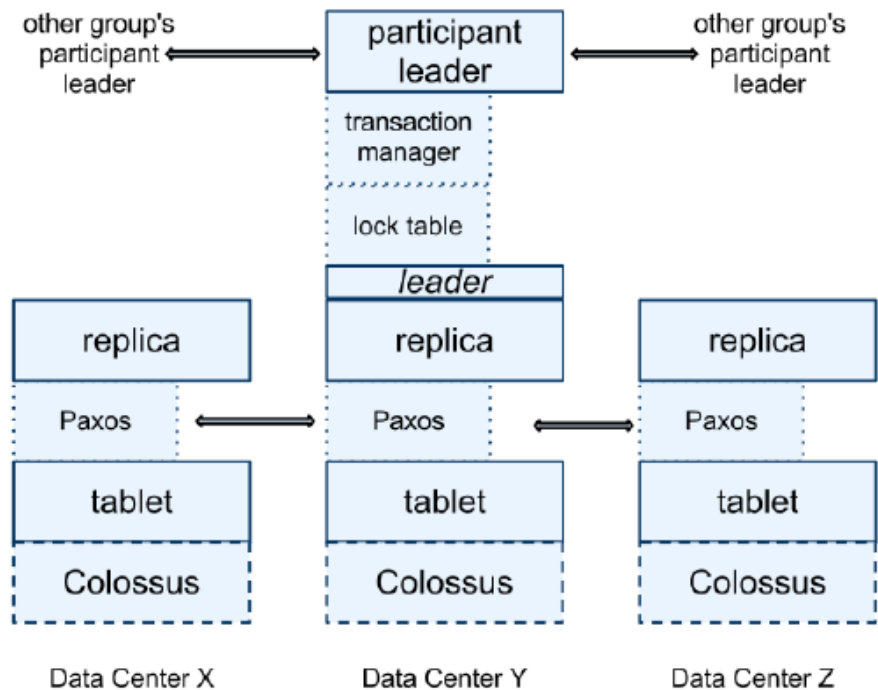


Figure 2: Spanserver software stack.

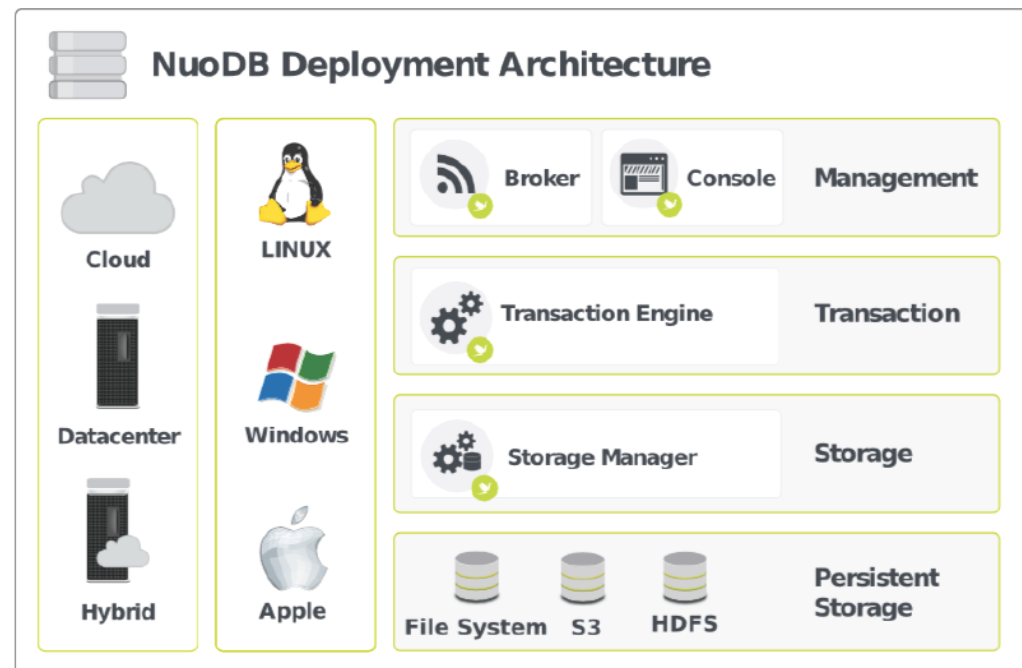
# Google F1

- Developed for use with AdWords platform replacing MySQL
- Relies on Google Spanner
- Distributed SQL queries for OLAP-style queries, using snapshot transactions
  - and two phase commit
- Optimistic concurrency control
- Transaction consistent secondary indexes
- Protocol Buffer as column data type
- Clustered hierarchical tables
- Increased read, write and commit latencies
- 5 copy replication
- Row or sub-row level locking
- Many issues found with inefficiencies of MySQL ORM layer – replaced with F1 ORM
- Multiple client consumers of a single query's results



# NuoDB

- Separates transaction management from storage management
  - Transaction Engines and Storage Managers
- TEs manage data via partial, on-demand replication
- Data resides where it is needed and used
- NuoDB Cloud on Amazon AWS
- Uses multi-version concurrency control (MVCC)




# VoltDB (H-Store product)

- H-Store origin
- In-memory, shared nothing RDBMS
- Database is sharded with serial execution of transactions within a shard – user specifies partitioning column for a table
- Synchronous replication with reexecution of transactions at replicas
- Read only remote backup for disaster recovery
- Transaction-level logging rather than ARIES-style page-oriented logging
- Periodically consistent snapshots written to disk and transaction-level log used to do recovery
- Multi-shard transactions are expensive

# VoltDB (H-Store product)

	Nodes	VoltDB	DBMSx	VoltDB Advantage
"TPC-C-like" workload (VoltDB lab)	1	53,000 TPS	1,555 TPS	45x better throughput
	12	560,000 TPS	N/A	Near-linear (.9) scaling



The **Hype** on it!


# VoltDB

## General Purpose RDBMS Processing Profile

*OLTP Through the Looking Glass, and What We Found There*

Stavros Harizopoulos, Daniel Abadi, Samuel Madden, and Michael Stonebraker

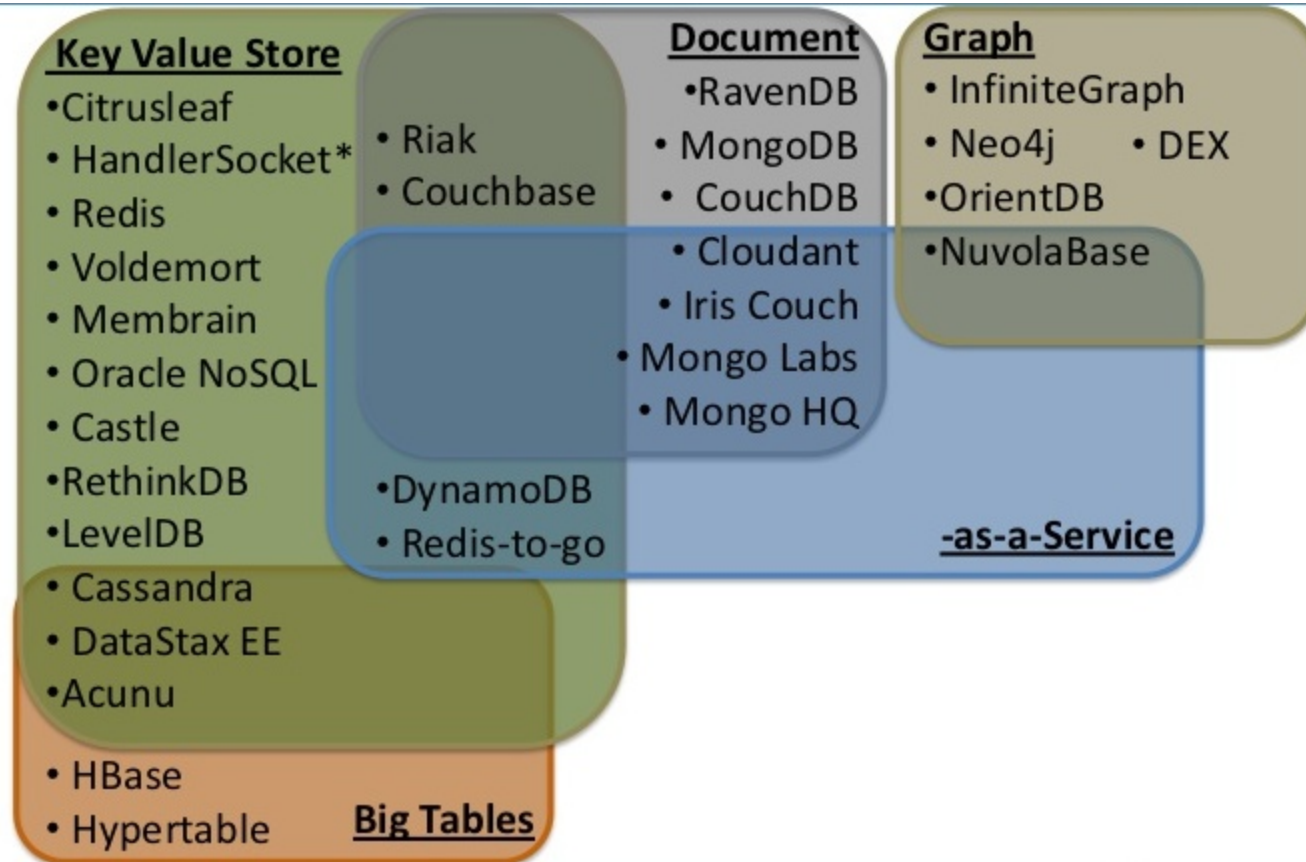
ACM SIGMOD 2008.



	VoltDB	NoSQL	Traditional RDBMS
Scale-out architecture	✓	✓	
Built-in high availability	✓	✓	
Multi-master replication	✓	✓	
ACID compliant	✓		✓
SQL data language	✓		✓
Cross-partition joins	Automatic	In app code	In app code
Cost at Web scale	\$	\$\$\$	\$\$\$\$

# NoSQL Systems

# NoSQL Ecosystem - Matthew Aslett 2012






# Motivations for NoSQL Systems

- Need for flexible data model, schema evolution, ...
- Many apps (e.g., Web2.0) need fewer features, high scale (data sizes, # of servers), very low response times
- Native support for interchange/storage formats like JSON
- Sharding RDBMS DBs rigid or unfriendly
- High-scale RDBMS too expensive for simple apps
- Need for low-latency, low-overhead API to access data
- Need to scale-out on cheap commodity nodes with locally-attached SATA disks, especially in cloud context
- Increasing use of distributed analytics
- Desire for one programming language for imperative programming and persistent data accesses (ref: POS)
- For some, access to source code for tailoring for their needs
- Flexibility with respect to data consistency

# NoSQL Systems

Focus on	Give up
<ul style="list-style-type: none"> <li>Commodity servers, networking, disks</li> <li>Easy elasticity and scalability to multiple racks (10s to 100s of servers)</li> <li>Fault-tolerance and high availability</li> </ul>	<ul style="list-style-type: none"> <li>Relational data model</li> <li>Standardized SQL APIs</li> <li>Complex queries (joins, secondary indexes)</li> <li>Stronger notion of transactions</li> </ul>

## Two Worlds



### Transactional

- Custom high-end OLTP for financial applications
- Scaleout datastores for Cloud/Web 2.0
- Examples
  - MemcacheDB, Cassandra, Dynamo, Voldemort, SimpleDB, Gigaspace, Websphere eXtreme Scale (WXS)

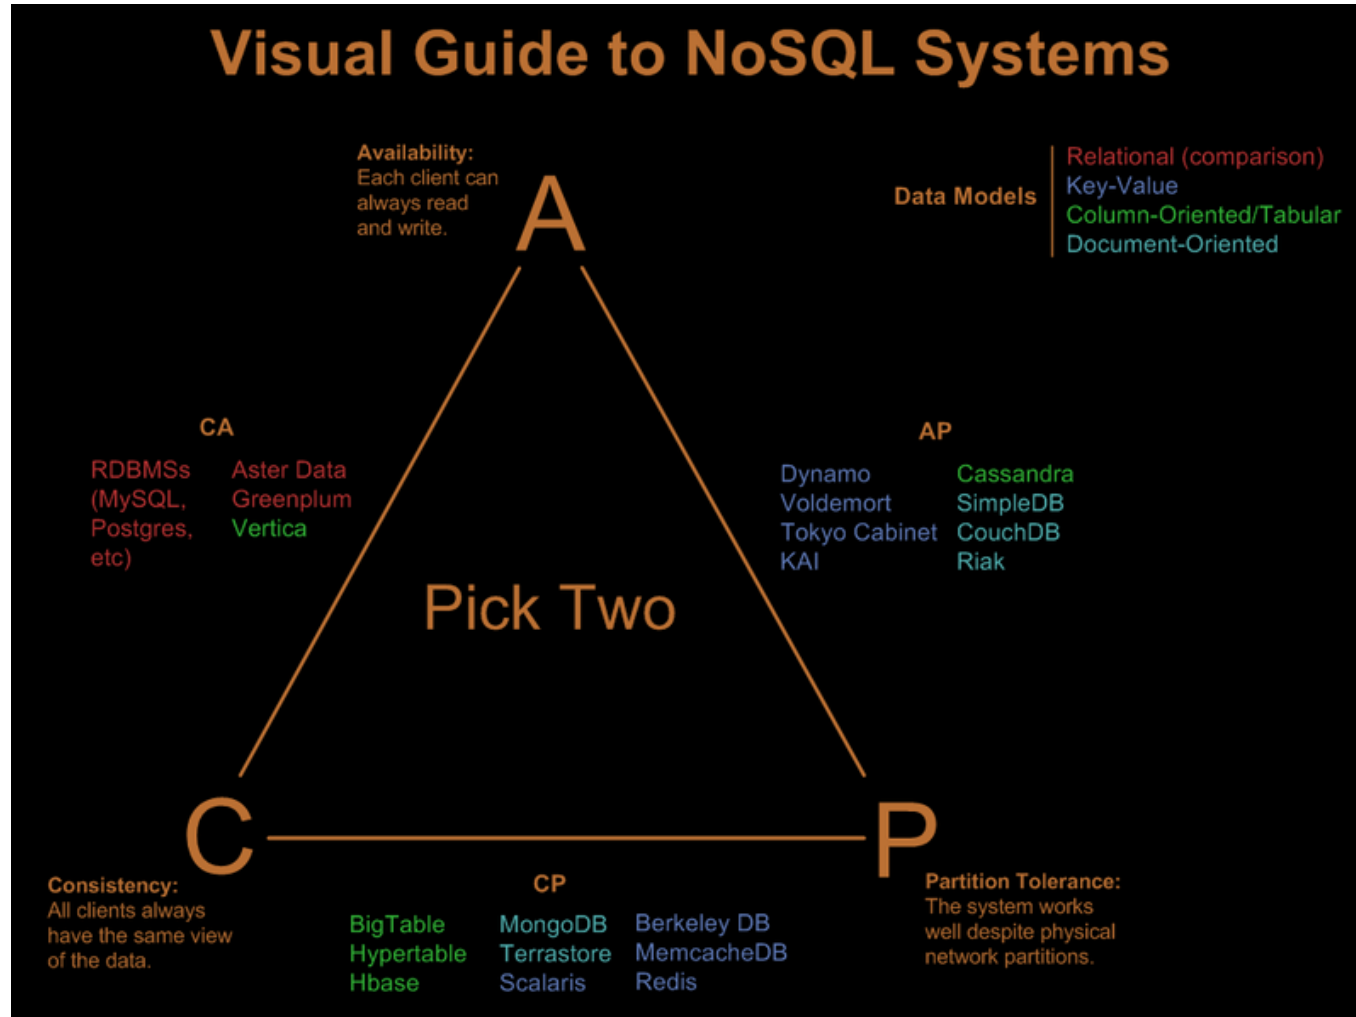
### Analytics

- Managing updates
- Support for random access and indexing
- Scaleout content store
- Examples
  - Bigtable, HBase, Hypertable

# NoSQL Stores

- Major open source projects
  - Cassandra and HBase
  - Other Projects: Voldemort, MongoDB, Hypertable
- Mushrooming of startups with tremendous interest from VCs (e.g., Aerospike)
- **“Scale-out structured storage”**: between raw files and DBMS
- Commercial flex-schema offerings based on relational technologies: DB2/Informix JSON, MonetDB

# One View of NoSQL Systems – Nathan Hurst, 3/2010?



<http://blog.nahurst.com/visual-guide-to-nosql-systems>

# Concerns/Misgivings (NonsenSQL ...)

- Throwing baby with bath water analogy comes to mind ... many very good developments of last 3 decades discarded ... but lot of them coming back ... albeit in ad hoc ways on a flaky base
- Many past systems that started simple ran into many scalability and other issues – S/38, AS/400, Lotus Notes/Domino, OODBMSs and RDBMSs with page as smallest unit of locking (e.g., mainframe DB2, Sybase/SQLServer, ObjectStore), Single level stores, ...
- Seat of the pant design in newer systems by highly inexperienced people who haven't done their homework AND who aren't documenting enough their designs for others to recognize their deficiencies
- Too many misleading claims on novelty of some design features
- Ridiculous choices like database level lock in MongoDB
- Object granularity as unit of replication locally and for disaster recovery remotely
- Myths about ACID transactions and relaxed notions of consistency

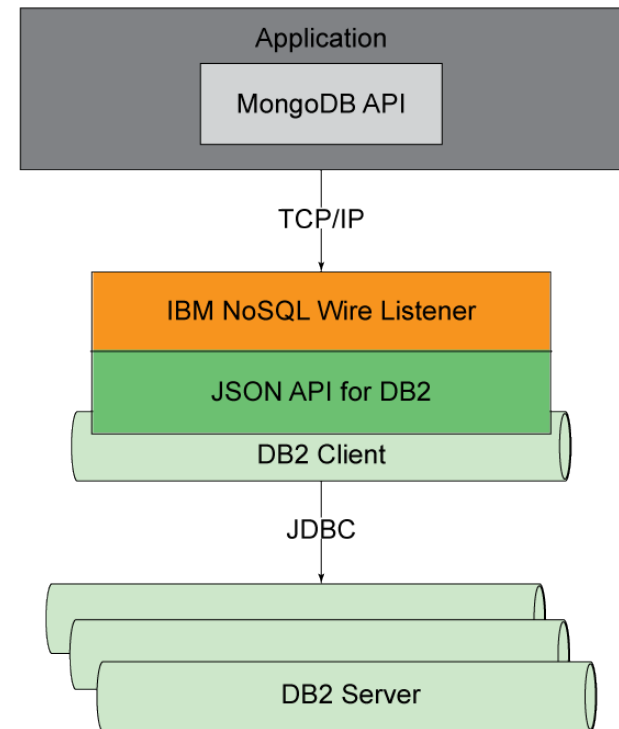
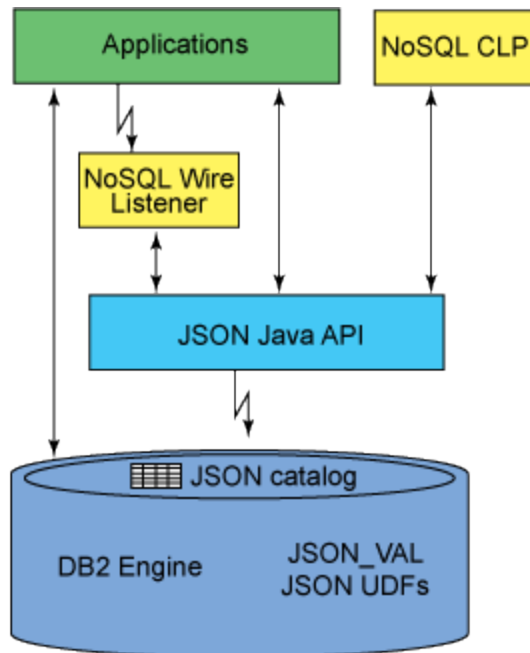
# MongoDB



- Open source document database system (similar to Lotus Notes)
- 1<sup>st</sup> release in 2009, Company behind it: MongoDB, Inc
- Native support for JSON-like docs with dynamic schemas (BSON)
- Indexing, including secondaries
- Master slave replication and high availability
- Auto sharding
- Simple query support
- MapReduce functionality could be invoked for complex aggregation
- Relies on file system for buffering
- GridFS for very large files storage
- Common use cases: operational and analytical big data, content management/delivery, mobile/social infrastructure, user data management and data hub.
- \$231M funding, 300 employees, \$1.2B evaluation!



# DB2 NoSQL JSON




In addition to support of basic features of the MongoDB query language, DB2 JSON also provides DB2 database server-specific extensions. These extensions enable applying more DB2 database features to JSON documents, for example:

- Transaction control to group multiple operations into one commit scope.
- Bi-temporal data to associate business time and system time concepts with your JSON data.

# Oracle NoSQL

## A Distributed, Scalable Key-Value Database

- Simple Data Model
  - Key-value pair (with major/minor key paradigm)
  - CRUD + iteration
- Scalability
  - Data partitioning and distribution
  - Optimized data access via intelligent driver
- High availability
  - One or more replicas
  - Resilient to partition master failures
  - No single point of failure
  - Disaster recovery through location of replicas
- Transparent load balancing
  - Reads from master or replicas
  - Driver is network topology & latency aware



Lamb, C. Oracle NoSQL Database, Presentation at International Workshop on High Performance Transaction Systems (HPTS), Asilomar, October 2011, [http://hpts.ws/papers/2011/sessions\\_2011/cwl-hpts-for-website.pdf](http://hpts.ws/papers/2011/sessions_2011/cwl-hpts-for-website.pdf)

# Oracle NoSQL Building Block

## Berkeley DB Java Edition

- Robust storage for a distributed key-value database
  - ACID transactions
  - Persistence
  - High availability
  - High throughput
  - Large capacity
  - Simple administration
- Already used in
  - Amazon Dynamo
  - Voldemort (LinkedIn)
  - GenieDB


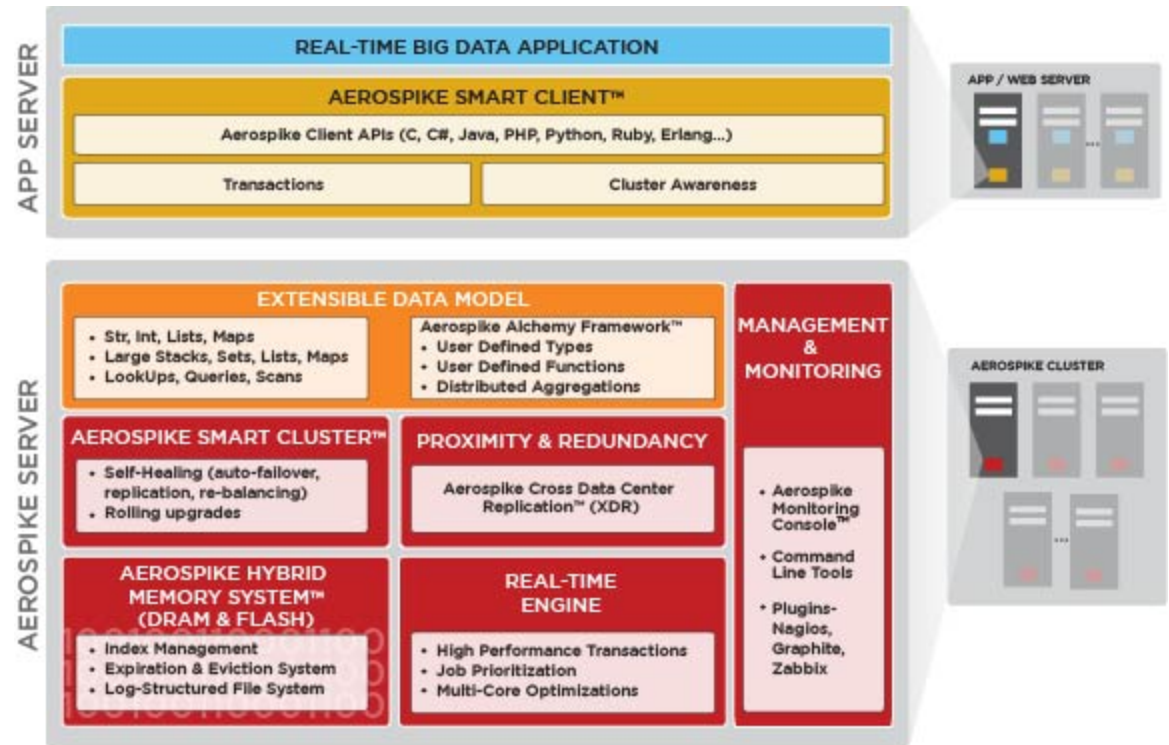


Figure 2: The Oracle NoSQL Database Flexibility Spectrum

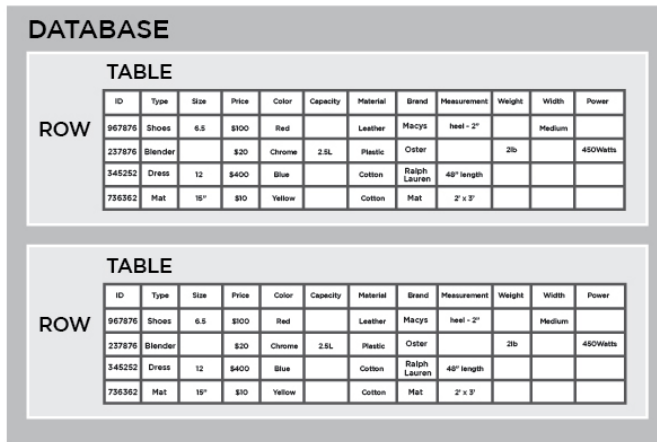
# Aerospike



- Advocates a hybrid DRAM & Flash system with indexes stored in DRAM & data in flash, disks for persistence
- Direct access to flash with own LFS, optimized via small block reads and large block writes
- Shared nothing with data distribution using randomized hashing and identical nodes
- Uses Paxos for cluster configuration, not for commit logic: all nodes have to agree, not just a majority
- Implemented in C for speed with an extensible data model
- Master/Slave and Master/Master replication supported
- Acquired alchemyDB (used Redis before) being ported onto Aerospike

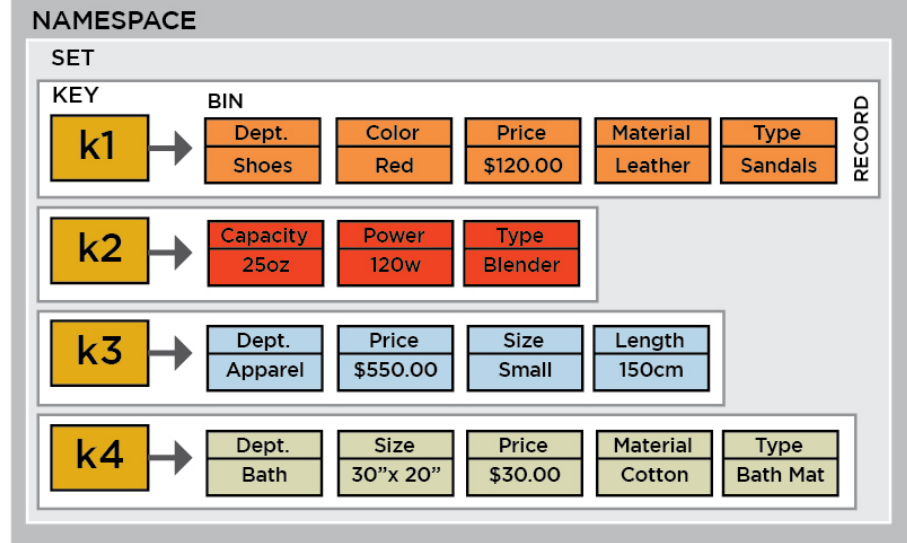
# Aerospike

## RELATIONAL DATABASE MODEL



The data is presented as relations - a collection of tables with each table consisting of a set of rows and columns

## NON-RELATIONAL DATABASE (AEROSPIKE)



The data is presented as flex-schema and provides simplicity of design, horizontal scaling and finer control over availability

- 500K TPS on \$2K server
- 1M TPS on \$5K server
- \$1 for 250 TPS
- Scaling by mitigating network-I/O-soft-interrupt overhead & avoiding unnecessary context switches – isolate NIC-Physical CPU, NUMA sensitive memory allocation, dedicated core for NIC hardware interrupt handling on > 4 core CPUs

# Facebook Case Study

- mySQL + memcache for real time access to PBs of data
- TAO for consistency across big geographies
- Haystack for managing photos (3B new ones per week)
- Apache Hadoop for mining intelligence from 100PBs of click logs
- Apache HBase for FB messages
- Data warehouse in Hive
- Scuba for real time event data handling – in-memory



# Hybrid Systems: Hadoop + SQL

# Hadoop


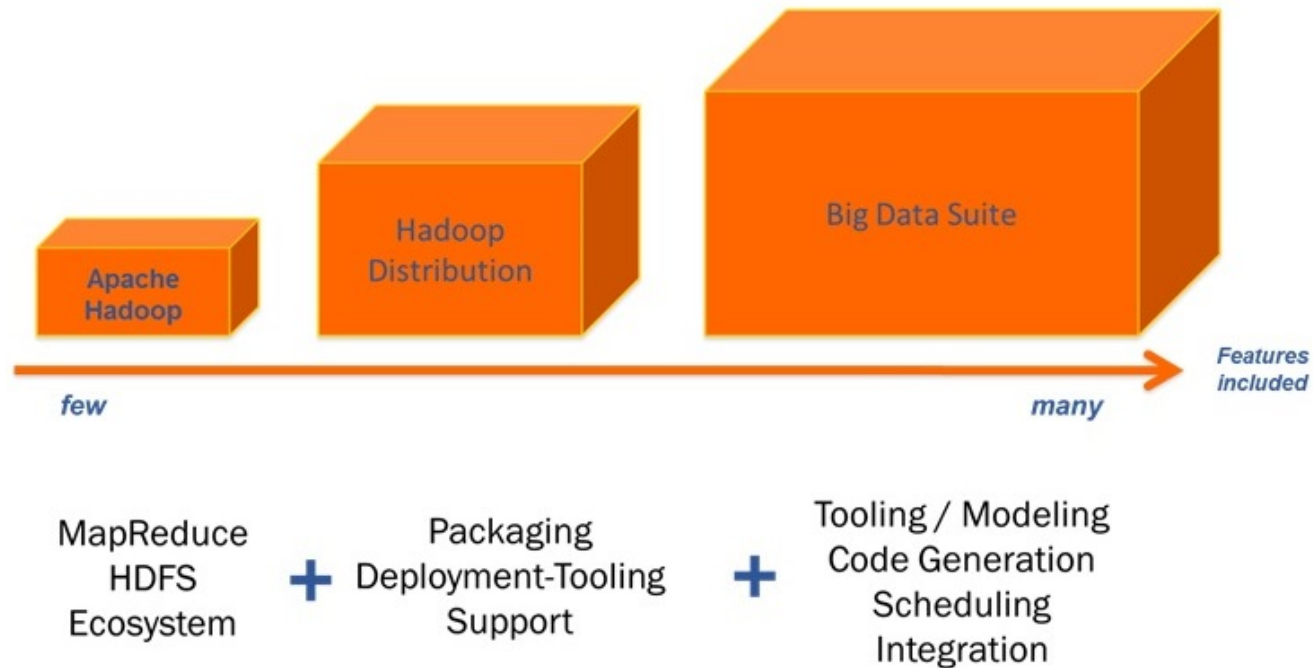


Figure 1: Hadoop Architecture



# Hadoop

- Apache Hadoop Modules: Common, Hadoop Distributed File System (HDFS), MapReduce, YARN
- Hadoop Ecosystem: Flume, HBase, Hive, Pig, Sqoop, Zookeeper
- Hadoop Distributions: Amazon Elastic MapReduce, Cloudera CDH, HortonWorks, IBM BigInsights, MapR, Pivotal HD
- Cloudera Impala: Distributed query execution engine for Apache HDFS/HBase data
- Apache Hadoop 2 (Rel 2.2.0) GAed on 15 October 2013
- Largest known production Hadoop environment: >35K nodes at Yahoo
- New: HA and Snapshots for HDFS data, YARN (aka MRv2 – ResourceManager, ApplicationMaster, NodeManager), Federation support (Multiple Independent Namenodes & Namespaces)

# Summary

- Good that many shortcomings of RDBMSs began to be addressed by user community (Web 2.0 companies)
- More rigor needed in rationalizing design choices
- Just because data model isn't relational doesn't mean internals have to be different or lessons from the past don't apply
- Let us not rediscover problems and adopt solutions with significant shortcomings
- Longer time issues not considered in addressing more immediate needs (migration, skills, tools, optimization, ...)
- Let us not allow history to repeat itself
- Shakeout and market consolidation will happen